

LevelDB 源码分析「十、其他细节」

2019.09.16 SF-Zhou

本篇为 LevelDB 源码分析的最后一篇博文，将会分析 LevelDB 中同步、原子量、单元测试和构建系统的一些细节。

1. 同步 Synchronization

LevelDB 中有大量并发访问的场景，也就需要同步的支持。LevelDB 使用的仍然是 C++ 标准库中的互斥量和条件变量，做了简单的封装 port/port_stdcxh.h：

```
#include <condition_variable> // NOLINT
#include <mutex> // NOLINT

// Thinly wraps std::mutex.
class LOCKABLE Mutex {
public:
    Mutex() = default;
    ~Mutex() = default;

    Mutex(const Mutex&) = delete;
    Mutex& operator=(const Mutex&) = delete;

    void Lock() EXCLUSIVE_LOCK_FUNCTION() { mu_.lock(); }
    void Unlock() UNLOCK_FUNCTION() { mu_.unlock(); }
    void AssertHeld() ASSERT_EXCLUSIVE_LOCK() {}
};
```

```

    void AssertMutex() ASSERT_EXCLUSIVE_LOCK() {}

private:
    friend class CondVar;
    std::mutex mu_;
};

// Thinly wraps std::condition_variable.
class CondVar {
public:
    explicit CondVar(Mutex* mu) : mu_(mu) { assert(mu != nullptr); }
    ~CondVar() = default;

    CondVar(const CondVar&) = delete;
    CondVar& operator=(const CondVar&) = delete;

    void Wait() {
        std::unique_lock<std::mutex> lock(mu_ -> mu_, std::adopt_lock);
        cv_.wait(lock);
        lock.release();
    }
    void Signal() { cv_.notify_one(); }
    void SignalAll() { cv_.notify_all(); }

private:
    std::condition_variable cv_;
    Mutex* const mu_;
};

```

其中类似 EXCLUSIVE_LOCK_FUNCTION 的宏定义于 port/thread_annotations.h ,

“EXCLUSIVE_LOCK_FUNCTION” 在 C++ 中是一个原子性操作，它用于确保在多线程环境中，对共享资源的访问是互斥的。这个宏通常用于声明一个函数，该函数在调用时会获取一个互斥锁，从而保证在函数执行期间，其他线程无法同时访问共享资源。这有助于防止数据竞争和保证程序的正确性。

作为线程安全分析的标注，在 Clang 环境下设定 `-Wthread-safety` 继而在编译期完成线程安全检查，详细资料参见文献 1。

这里分析一下 `DBImpl::Write` 中控制同步的部分：

```
class DBImpl {
    // Queue of writers.
    std::deque<Writer*> writers_ GUARDED_BY(mutex_);
    WriteBatch* tmp_batch_ GUARDED_BY(mutex_);
    ...
}

// Information kept for every waiting writer
struct DBImpl::Writer {
    explicit Writer(port::Mutex* mu)
        : batch(nullptr), sync(false), done(false), cv(mu) {}

    Status status;
    WriteBatch* batch;
    bool sync;
    bool done;
    port::CondVar cv;
};

Status DBImpl::Write(const WriteOptions& options, WriteBatch* updates) {
    Writer w(&mutex_);
    w.batch = updates;
    w.sync = options.sync;
    w.done = false;
```

```

MutexLock l(&mutex_);
writers_.push_back(&w);
while (!w.done && &w != writers_.front()) {
    w.cv.Wait();
}

if (w.done) {
    return w.status;
}

// May temporarily unlock and wait.
Status status = MakeRoomForWrite(updates == nullptr);
uint64_t last_sequence = versions_>LastSequence();
Writer* last_writer = &w;
if (status.ok() && updates != nullptr) { // nullptr batch is for compactions
    WriteBatch* updates = BuildBatchGroup(&last_writer);
    WriteBatchInternal::SetSequence(updates, last_sequence + 1);
    last_sequence += WriteBatchInternal::Count(updates);

    // Add to log and apply to memtable. We can release the lock
    // during this phase since &w is currently responsible for logging
    // and protects against concurrent loggers and concurrent writes
    // into mem_.
    {
        mutex_.Unlock();
        status = log_>AddRecord(WriteBatchInternal::Contents(updates));
        bool sync_error = false;
        if (status.ok() && options.sync) {
            status = logfile_>Sync();
            if (!status.ok()) {
                sync_error = true;
            }
        }
    }
}

```

```

    }
}
if (status.ok()) {
    status = WriteBatchInternal::InsertInto(updates, mem_);
}

mutex_.Lock();
if (sync_error) {
    // The state of the log file is indeterminate: the log record we
    // just added may or may not show up when the DB is re-opened.
    // So we force the DB into a mode where all future writes fail.
    RecordBackgroundError(status);
}
}
if (updates == tmp_batch_) tmp_batch_>Clear();

versions_>SetLastSequence(last_sequence);
}

while (true) {
    Writer* ready = writers_.front();
    writers_.pop_front();
    if (ready != &w) {
        ready->status = status;
        ready->done = true;
        ready->cv.Signal();
    }
    if (ready == last_writer) break;
}

// Notify new head of write queue

```

```

    if (!writers_.empty()) {
        writers_.front()->cv.Signal();
    }

    return status;
}

// REQUIRES: Writer list must be non-empty
// REQUIRES: First writer must have a non-null batch
WriteBatch* DBImpl::BuildBatchGroup(Writer** last_writer) {
    mutex_.AssertHeld();
    assert(!writers_.empty());
    Writer* first = writers_.front();
    WriteBatch* result = first->batch;
    assert(result != nullptr);

    size_t size = WriteBatchInternal::ByteSize(first->batch);

    // Allow the group to grow up to a maximum size, but if the
    // original write is small, limit the growth so we do not slow
    // down the small write too much.
    size_t max_size = 1 << 20;
    if (size <= (128 << 10)) {
        max_size = size + (128 << 10);
    }

    *last_writer = first;
    std::deque<Writer*>::iterator iter = writers_.begin();
    ++iter; // Advance past "first"
    for (; iter != writers_.end(); ++iter) {

```

```

Writer* w = *iter;
if (w->sync && !first->sync) {
    // Do not include a sync write into a batch handled by a non-sync write.
    break;
}

if (w->batch != nullptr) {
    size += WriteBatchInternal::ByteSize(w->batch);
    if (size > max_size) {
        // Do not make batch too big
        break;
    }

    // Append to *result
    if (result == first->batch) {
        // Switch to temporary batch instead of disturbing caller's batch
        result = tmp_batch_;
        assert(WriteBatchInternal::Count(result) == 0);
        WriteBatchInternal::Append(result, first->batch);
    }
    WriteBatchInternal::Append(result, w->batch);
}
*last_writer = w;
}
return result;
}

```

假设现在有编号为 [1, 2, 3, 4] 的四个线程基本同时调用写入操作，发生的事件如下：

1. 每个线程中各自构造 `WriteBatch`。

1. 每个线程中各自构造 `Writer` ;
2. 1 号线程较快地构造了 `MutexLock` 拿到锁, [2, 3, 4] 则阻塞在此处;
3. 1 号线程将写入请求插入双向队列中, 跳过循环, 继续向下走;
4. 1 号线程执行 `BuildBatchGroup`, 由于队列中只有自身一个请求, 不会发生合并;
5. 1 号线程执行 `mutex_.Unlock()` 释放锁, 随后执行写入操作, 完成后执行 `mutex_.Lock()` 再次获得锁;
6. 1 号线程在循环中从双向队列里将写入请求弹出, 最后通知队列顶的 2 号线程唤醒;
7. 1 号线程析构局部变量、释放锁。

在第 5 步发生释放锁的同时:

1. 2 号线程获得锁, 将写入请求插入双向队列中, 由于请求不在队列顶端, 进而进入循环、等待、释放锁;
2. 3 号线程获得锁, 将写入请求插入双向队列中, 由于请求不在队列顶端, 进而进入循环、等待、释放锁;
3. 此时 1 号线程写入完成、执行 `mutex_.Lock()` 获得锁, 4 号线程继续等待;
4. 1 号线程执行结束、释放锁, 2 号线程唤醒获得锁, 执行 `BuildBatchGroup` 将队列中的 3 号线程中的写入请求合并;
5. 2 号线程执行 `mutex_.Unlock()` 解锁, 随后执行写入操作, 完成后执行 `mutex_.Lock()` 再次获得锁。与此同时 4 号线程获得锁, 将写入请求插入双向队列中, 等待、释放锁;
6. 2 号线程在循环中从双向队列里将写入请求弹出, 将 3 号线程的写入请求标记为完成, 尝试唤醒 3 号线程;
7. 2 号线程析构局部变量、释放锁, 3 号线程唤醒、获得锁, 判断已完成, 返回、释放

1.2 与线性结构同构的叉星、伴双树。3 与线性映能、拟树树，判断二元成，返回、伴双锁；

8. 4 号线程唤醒、获得锁，正常执行。

上述合并操作依赖写入时的释放锁操作，这使得其他线程有机会加入队列、然后等待，在下次获得锁时合并队列中的其他写入请求。

2. 原子量 Atomic

由于并发访问，LevelDB 中也大量使用了原子量，且使用了两种不同的内存模型。一种是 Relaxed Ordering，其只保证原子性、不保证并发时的执行顺序，一般在计数功能中使用，例如内存池中的内存使用量：

```
size_t MemoryUsage() const {
    return memory_usage_.load(std::memory_order_relaxed);
}

char* Arena::AllocateNewBlock(size_t block_bytes) {
    char* result = new char[block_bytes];
    blocks_.push_back(result);
    memory_usage_.fetch_add(block_bytes + sizeof(char*),
                            std::memory_order_relaxed);
    return result;
}
```

另一种是 Release-Acquire Ordering，其不仅可以保证原子性，还可以保证一定程度的

执行顺序。以下摘录自参考文献 7：

If an atomic store in thread A is tagged memory_order_release and an atomic load in thread B from the same variable is tagged memory_order_acquire, all memory writes (non-atomic and relaxed atomic) that happened-before the atomic store from the point of view of thread A, become visible side-effects in thread B. That is, once the atomic load is completed, thread B is guaranteed to see everything thread A wrote to memory.

LevelDB 中的使用举例：

```
template <typename Key, class Comparator>
class SkipList {
    ...
    Node* Next(int n) {
        assert(n >= 0);
        // Use an 'acquire load' so that we observe a fully initialized
        // version of the returned Node.
        return next_[n].load(std::memory_order_acquire);
    }
    void SetNext(int n, Node* x) {
        assert(n >= 0);
        // Use a 'release store' so that anybody who reads through this
        // pointer observes a fully initialized version of the inserted node.
        next_[n].store(x, std::memory_order_release);
    }
}

template <typename Key, class Comparator>
void SkipList<Key, Comparator>::Insert(const Key& key) {
```

```

// TODO(opt): We can use a barrier-free variant of FindGreaterOrEqual()
// here since Insert() is externally synchronized.
Node* prev[kMaxHeight];
Node* x = FindGreaterOrEqual(key, prev);

// Our data structure does not allow duplicate insertion
assert(x == nullptr || !Equal(key, x->key));

int height = RandomHeight();
if (height > GetMaxHeight()) {
    for (int i = GetMaxHeight(); i < height; i++) {
        prev[i] = head_;
    }
    // It is ok to mutate max_height_ without any synchronization
    // with concurrent readers. A concurrent reader that observes
    // the new value of max_height_ will see either the old value of
    // new level pointers from head_ (nullptr), or a new value set in
    // the loop below. In the former case the reader will
    // immediately drop to the next level since nullptr sorts after all
    // keys. In the latter case the reader will use the new node.
    max_height_.store(height, std::memory_order_relaxed);
}

x = NewNode(key, height);
for (int i = 0; i < height; i++) {
    // NoBarrier_SetNext() suffices since we will add a barrier when
    // we publish a pointer to "x" in prev[i].
    x->NoBarrier_SetNext(i, prev[i]->NoBarrier_Next(i));
    prev[i]->SetNext(i, x);
}

```

```
}
```

3. 单元测试

LevelDB 中的单元测试并没有使用自家的 Google Test，而是自己实现了一套简单的测试工具，位于 util/testharness.h：

```
// An instance of Tester is allocated to hold temporary state during
// the execution of an assertion.
class Tester {
private:
    bool ok_;
    const char* fname_;
    int line_;
    std::stringstream ss_;

public:
    Tester(const char* f, int l) : ok_(true), fname_(f), line_(l) {}

    ~Tester() {
        if (!ok_) {
            fprintf(stderr, "%s:%d:%s\n", fname_, line_, ss_.str().c_str());
            exit(1);
        }
    }
}

Tester& Is(bool b, const char* msg) {
    if (!b) {
```

```

        ss_ << " Assertion failure " << msg;
        ok_ = false;
    }
    return *this;
}

```

```

Tester& IsOk(const Status& s) {
    if (!s.ok()) {
        ss_ << " " << s.ToString();
        ok_ = false;
    }
    return *this;
}

```

```

#define BINARY_OP(name, op) \
    template <class X, class Y> \
    Tester& name(const X& x, const Y& y) { \
        if (!(x op y)) { \
            ss_ << " failed: " << x << (" " #op " ") << y; \
            ok_ = false; \
        } \
        return *this; \
    }

```

```

BINARY_OP(IsEq, ==)
BINARY_OP(IsNe, !=)
BINARY_OP(IsGe, >=)
BINARY_OP(IsGt, >)
BINARY_OP(IsLe, <=)
BINARY_OP(IsLt, <)

```

```

#undef BINARY_OP

// Attach the specified value to the error message if an error has occurred
template <class V>
Tester& operator<<(const V& value) {

    if (!ok_) {
        ss_ << " " << value;
    }
    return *this;
}
};

#define ASSERT_TRUE(c) ::leveldb::test::Tester(__FILE__, __LINE__).Is((c), #c)
#define ASSERT_OK(s) ::leveldb::test::Tester(__FILE__, __LINE__).IsOk((s))
#define ASSERT_EQ(a, b) \
    ::leveldb::test::Tester(__FILE__, __LINE__).IsEq((a), (b))
#define ASSERT_NE(a, b) \
    ::leveldb::test::Tester(__FILE__, __LINE__).IsNe((a), (b))
#define ASSERT_GE(a, b) \
    ::leveldb::test::Tester(__FILE__, __LINE__).IsGe((a), (b))
#define ASSERT_GT(a, b) \
    ::leveldb::test::Tester(__FILE__, __LINE__).IsGt((a), (b))
#define ASSERT_LE(a, b) \
    ::leveldb::test::Tester(__FILE__, __LINE__).IsLe((a), (b))
#define ASSERT_LT(a, b) \
    ::leveldb::test::Tester(__FILE__, __LINE__).IsLt((a), (b))

```

BINARY_OP 宏简化了比较运算函数的定义，使用完后及时 undef 也避免了污染。继续：

```

#define TCONCAT(a, b) TCONCAT1(a, b)
#define TCONCAT1(a, b) a##b

#define TEST(base, name) \

    class TCONCAT(_Test_, name) : public base { \
    public: \
        void _Run(); \
        static void _RunIt() { \
            TCONCAT(_Test_, name) t; \
            t._Run(); \
        } \
    }; \
    bool TCONCAT(_Test_ignored_, name) = ::leveldb::test::RegisterTest( \
        #base, #name, &TCONCAT(_Test_, name)::_RunIt); \
    void TCONCAT(_Test_, name)::_Run()

// Register the specified test. Typically not used directly, but
// invoked via the macro expansion of TEST.
bool RegisterTest(const char* base, const char* name, void (*func)());

```

定义 TCONCAT1 宏是为了让 TCONCAT 像函数一样支持嵌套调用。全局变量 TCONCAT(_Test_ignored_, name) 则可以实现在 main 函数前对测试类进行注册。注册函数和执行所有测试的实现位于 util/testharness.cc :

```

namespace {
struct Test {
    const char* base;
    const char* name;

```

```

    const char* name;
    void (*func)();
};
std::vector<Test>* tests;
} // namespace

bool RegisterTest(const char* base, const char* name, void (*func)()) {
    if (tests == nullptr) {
        tests = new std::vector<Test>;
    }
    Test t;
    t.base = base;
    t.name = name;
    t.func = func;
    tests->push_back(t);
    return true;
}

int RunAllTests() {
    const char* matcher = getenv("LEVELDB_TESTS");

    int num = 0;
    if (tests != nullptr) {
        for (size_t i = 0; i < tests->size(); i++) {
            const Test& t = (*tests)[i];
            if (matcher != nullptr) {
                std::string name = t.base;
                name.push_back('.');
                name.append(t.name);
                if (strstr(name.c_str(), matcher) == nullptr) {
                    continue;
                }
            }
            num++;
        }
    }
}

```



```

    }
}
fprintf(stderr, "==== Test %s.%s\n", t.base, t.name);
(*t.func)();
++num;

}
}
fprintf(stderr, "==== PASSED %d tests\n", num);
return 0;
}

```

4. 构建系统

LevelDB 使用 CMake 作为其构建系统，搭建了跨平台、可配置的编译系统。例如使用 CMake 变量 WIN32 实现不同环境的编译：

```

if (WIN32)
    set(LEVELDB_PLATFORM_NAME LEVELDB_PLATFORM_WINDOWS)
    # TODO(cmumford): Make UNICODE configurable for Windows.
    add_definitions(-D_UNICODE -DUNICODE)
else (WIN32)
    set(LEVELDB_PLATFORM_NAME LEVELDB_PLATFORM_POSIX)
endif (WIN32)

if (WIN32)
    target_sources(leveldb
        PRIVATE
        "${PROJECT_SOURCE_DIR}/util/env_windows.cc"
        "${PROJECT_SOURCE_DIR}/util/windows_logger.h"
    )
endif (WIN32)

```

```

    ${PROJECT_SOURCE_DIR}/util/windows_logger.h
)
else (WIN32)
    target_sources(leveldb
        PRIVATE
        "${PROJECT_SOURCE_DIR}/util/env_posix.cc"
        "${PROJECT_SOURCE_DIR}/util/posix_logger.h"
    )
endif (WIN32)

```

使用 `configure_file` 和 `cmakedefine01` 将 CMake 中的变量转为代码中的宏定义：

```

check_library_exists(crc32c crc32c_value "" HAVE_CRC32C)

configure_file(
    "${PROJECT_SOURCE_DIR}/port/port_config.h.in"
    "${PROJECT_BINARY_DIR}/${LEVELDB_PORT_CONFIG_DIR}/port_config.h"
)

```

对应的 `port/port_config.h.in`：

```

// Define to 1 if you have Google CRC32C.
#if !defined(HAVE_CRC32C)
#cmakedefine01 HAVE_CRC32C
#endif // !defined(HAVE_CRC32C)

/* after compile */
// Define to 1 if you have Google CRC32C.
#if !defined(HAVE_CRC32C)

```

```
#define HAVE_CRC32C 0
#endif // !defined(HAVE_CRC32C
```

使用 `check_cxx_source_compiles` 检查编译器的特性:

```
# Test whether -Wthread-safety is available. See
# https://clang.llvm.org/docs/ThreadSafetyAnalysis.html
# -Werror is necessary because unknown attributes only generate warnings.
set(OLD_CMAKE_REQUIRED_FLAGS ${CMAKE_REQUIRED_FLAGS})
list(APPEND CMAKE_REQUIRED_FLAGS -Werror -Wthread-safety)
check_cxx_source_compiles("
struct __attribute__((lockable)) Lock {
    void Acquire() __attribute__((exclusive_lock_function()));
    void Release() __attribute__((unlock_function()));
};
struct ThreadSafeType {
    Lock lock_;
    int data_ __attribute__((guarded_by(lock_)));
};
int main() { return 0; }
" HAVE_CLANG_THREAD_SAFETY)
set(CMAKE_REQUIRED_FLAGS ${OLD_CMAKE_REQUIRED_FLAGS})

if(HAVE_CLANG_THREAD_SAFETY)
    target_compile_options(leveldb
        PUBLIC
        -Werror -Wthread-safety)
endif(HAVE_CLANG_THREAD_SAFETY)
```

References

1. "Thread Safety Analysis", *Clang Documentation*
2. "std::mutex", *C++ Reference*
3. "std::unique_lock", *C++ Reference*
4. "std::condition_variable", *C++ Reference*
5. "std::condition_variable::notify_one", *C++ Reference*
6. "std::atomic", *C++ Reference*
7. "std::memory_order", *C++ Reference*

1 comment – *powered by giscus*

Oldest

Newest



SF-Zhou Jun 17, 2020

Owner

1. Google Test 最新的框架已经移除测试框架并替换为了 Google Test

Except where otherwise noted, content on this site is licensed under a CC BY-SA 4.0 license.
Copyright©2017 SF-Zhou, All Rights Reserved. Powered by GitHub Pages and GitHub Actions.