

LevelDB 源码分析「四、高性能写操作续」

2019.08.10 SF-Zhou

本系列的上一篇文章介绍了 LevelDB 写操作中构造的 WriteBatch，以及写操作的第一步：追加日志。本篇将继续介绍写操作的第二步：插入内存数据库。

3. 内存数据库 MemTable

MemTable 即为在内存中建立的 KV 数据库，基于跳表，定义于 db/memtable.h：

```
#include <string>

#include "db/dbformat.h"
#include "db/skiplist.h"
#include "leveldb/db.h"
#include "util/arena.h"

namespace leveldb {

class InternalKeyComparator;
class MemTableIterator;

class MemTable {
public:
    // MemTables are reference counted. The initial reference count
    // is zero and the caller must call Ref() at least once.
```

```

explicit MemTable(const InternalKeyComparator& comparator);

MemTable(const MemTable&) = delete;
MemTable& operator=(const MemTable&) = delete;

// Increase reference count.
void Ref() { ++refs_; }

// Drop reference count. Delete if no more references exist.
void Unref() {
    --refs_;
    assert(refs_ >= 0);
    if (refs_ <= 0) {
        delete this;
    }
}

// Returns an estimate of the number of bytes of data in use by this
// data structure. It is safe to call when MemTable is being modified.
size_t ApproximateMemoryUsage();

// Return an iterator that yields the contents of the memtable.
//
// The caller must ensure that the underlying MemTable remains live
// while the returned iterator is live. The keys returned by this
// iterator are internal keys encoded by AppendInternalKey in the
// db/format.{h,cc} module.
Iterator* NewIterator();

// Add an entry into memtable that maps key to value at the

```

```

// specified sequence number and with the specified type.
// Typically value will be empty if type==kTypeDeletion.
void Add(SequenceNumber seq, ValueType type, const Slice& key,
         const Slice& value);

// If memtable contains a value for key, store it in *value and return true.
// If memtable contains a deletion for key, store a NotFound() error
// in *status and return true.
// Else, return false.
bool Get(const LookupKey& key, std::string* value, Status* s);

private:
    friend class MemTableIterator;
    friend class MemTableBackwardIterator;

    struct KeyComparator {
        const InternalKeyComparator comparator;
        explicit KeyComparator(const InternalKeyComparator& c) : comparator(c) {}
        int operator()(const char* a, const char* b) const;
    };

    typedef SkipList<const char*, KeyComparator> Table;

    ~MemTable(); // Private since only Unref() should be used to delete it

    KeyComparator comparator_;
    int refs_;
    Arena arena_;
    Table table_;
};

```

```
} // namespace leveldb
```

MemTable 包含比较器 `comparator_`、引用计数 `refs_`、内存池 `arena_` 和跳表 `table_` 四个成员变量。代码中首先前置声明了 `InternalKeyComparator` 类，该类的对象是 `MemTable` 的构造函数参数。该类定义为 `db/dbformat.h` 中，将在下文中介绍。随后 `MemTable` 禁止了拷贝构造和赋值操作符，不允许拷贝操作，这应该是内存池的副作用。`MemTable` 使用引用计数管理自己的生命周期，甚至其析构函数都是私有的。而接口 `ApproximateMemoryUsage` 实际上返回的是内存池中的内存使用量。接口方面提供了读写接口和迭代器接口。内部定义了结构体 `KeyComparator`，对 `InternalKeyComparator` 进行了封装。为了弄清楚 `InternalKeyComparator` 到底是什么，我们先转到 `db/dbformat.h` 查看其定义：

```
enum ValueType { kTypeDeletion = 0x0, kTypeValue = 0x1 };

static const ValueType kValueTypeForSeek = kTypeValue;

typedef uint64_t SequenceNumber;
static const SequenceNumber kMaxSequenceNumber = ((0x1ull << 56) - 1);

struct ParsedInternalKey {
    Slice user_key;
    SequenceNumber sequence;
    ValueType type;
```

```

ParsedInternalKey() {} // Intentionally left uninitialized (for speed)
ParsedInternalKey(const Slice& u, const SequenceNumber& seq, ValueType t)
    : user_key(u), sequence(seq), type(t) {}
std::string DebugString() const;
};

// Return the length of the encoding of "key".
inline size_t InternalKeyEncodingLength(const ParsedInternalKey& key) {
    return key.user_key.size() + 8;
}

static uint64_t PackSequenceAndType(uint64_t seq, ValueType t) {
    assert(seq <= kMaxSequenceNumber);
    assert(t <= kValueTypeForSeek);
    return (seq << 8) | t;
}

void AppendInternalKey(std::string* result, const ParsedInternalKey& key) {
    result->append(key.user_key.data(), key.user_key.size());
    PutFixed64(result, PackSequenceAndType(key.sequence, key.type));
}

inline bool ParseInternalKey(const Slice& internal_key,
                             ParsedInternalKey* result) {
    const size_t n = internal_key.size();
    if (n < 8) return false;
    uint64_t num = DecodeFixed64(internal_key.data() + n - 8);
    uint8_t c = num & 0xff;
    result->sequence = num >> 8;
    result->type = static_cast<ValueType>(c);
}

```

```

    result->user_key = Slice(internal_key.data(), n - 8);
    return (c <= static_cast<uint8_t>(kTypeValue));
}

inline Slice ExtractUserKey(const Slice& internal_key) {

    assert(internal_key.size() >= 8);
    return Slice(internal_key.data(), internal_key.size() - 8);
}

```

为了方便阅读这里对函数进行了重排。首先定义了两值类型 `kTypeDeletion` 和 `kTypeValue`，注意 `kTypeDeletion` 值较小。然后定义了 `SequenceNumber` 为 64 位无符号数，并且定义了其最大值为 $2^{56} - 1$ 。然后是 `ParsedInternalKey` 结构体的定义，其内部包括 `user_key`、`sequence` 和 `type`。根据类的名字就可以猜到一点了，`Internal Key` 应该是对 `User Key` 的封装，在其基础上加入了序列号 `sequence` 和值类型 `type`。函数 `PackSequenceAndType` 中可以发现，`sequence` 占用 56bit，剩下 8bit 给 `type`，一共组成 64 位无符号数，下文称之为合成序列号。所以 `InternalKeyEncodingLength` 计算的长度是 `User Key` 的长度加 8。函数 `AppendInternalKey` 可以将 `ParsedInternalKey` 转为字符串，先将 `User Key` 编码进去，再将 64 位合成序列号加入，对应的解析函数 `ParseInternalKey` 为逆过程。继续：

```

class InternalKeyComparator : public Comparator {
private:
    const Comparator* user_comparator_;

public:
    explicit InternalKeyComparator(const Comparator* c) : user_comparator_(c) {}
}

```

```

const char* Name() const override;
int Compare(const Slice& a, const Slice& b) const override;
void FindShortestSeparator(std::string* start,
                           const Slice& limit) const override;
void FindShortSuccessor(std::string* key) const override;

const Comparator* user_comparator() const { return user_comparator_; }

int Compare(const InternalKey& a, const InternalKey& b) const;
};

const char* InternalKeyComparator::Name() const {
    return "leveldb.InternalKeyComparator";
}

int InternalKeyComparator::Compare(const Slice& akey, const Slice& bkey) const {
    // Order by:
    //   increasing user key (according to user-supplied comparator)
    //   decreasing sequence number
    //   decreasing type (though sequence# should be enough to disambiguate)
    int r = user_comparator_>Compare(ExtractUserKey(akey), ExtractUserKey(bkey));
    if (r == 0) {
        const uint64_t anum = DecodeFixed64(akey.data() + akey.size() - 8);
        const uint64_t bnum = DecodeFixed64(bkey.data() + bkey.size() - 8);
        if (anum > bnum) {
            r = -1;
        } else if (anum < bnum) {
            r = +1;
        }
    }
}

```

```
    return r;  
}
```

InternalKeyComparator 继承于 LevelDB 中定义的比较器 Comparator，其源码位于 include/leveldb/comparator.h，有兴趣自己阅读。其核心函数即为 Compare，比较两个 Key 的大小。InternalKeyComparator 对 user_comparator_ 进行了封装，当比较 InternalKey 时，首先使用 user_comparator_ 对 User Key 进行比较，如果相等，则抽取合成序列号进行比较，序列号大的反而在顺序上更小，即降序排列。继续看 InternalKey 的定义：

```
class InternalKey {  
private:  
    std::string rep_;  
  
public:  
    InternalKey() {} // Leave rep_ as empty to indicate it is invalid  
    InternalKey(const Slice& user_key, SequenceNumber s, ValueType t) {  
        AppendInternalKey(&rep_, ParsedInternalKey(user_key, s, t));  
    }  
  
    bool DecodeFrom(const Slice& s) {  
        rep_.assign(s.data(), s.size());  
        return !rep_.empty();  
    }  
  
    Slice Encode() const {
```



```

        assert(!rep_.empty());
        return rep_;
    }

    Slice user_key() const { return ExtractUserKey(rep_); }

    void SetFrom(const ParsedInternalKey& p) {
        rep_.clear();
        AppendInternalKey(&rep_, p);
    }

    void Clear() { rep_.clear(); }

    std::string DebugString() const;
};

inline int InternalKeyComparator::Compare(const InternalKey& a,
                                           const InternalKey& b) const {
    return Compare(a.Encode(), b.Encode());
}

```

InternalKey 存储的信息和 ParseInternalKey 一致，只是存储形式不同，前者直接使用字符串 rep_ 存储 User Key 和合成序列号，对应的比较函数则直接使用 InternalKeyComparator 对 rep_ 进行解析后的比较。该文件中还有 InternalFilterPolicy 的定义，同样是从 Internal Key 解析出 User Key 进行操作。继续看 LookupKey：

```

// A helper class useful for DBImpl::Get()
class LookupKey {

```

```

public:
    // Initialize *this for looking up user_key at a snapshot with
    // the specified sequence number.
    LookupKey(const Slice& user_key, SequenceNumber sequence);

    LookupKey(const LookupKey&) = delete;
    LookupKey& operator=(const LookupKey&) = delete;

    ~LookupKey();

    // Return a key suitable for lookup in a MemTable.
    Slice memtable_key() const { return Slice(start_, end_ - start_); }

    // Return an internal key (suitable for passing to an internal iterator)
    Slice internal_key() const { return Slice(kstart_, end_ - kstart_); }

    // Return the user key
    Slice user_key() const { return Slice(kstart_, end_ - kstart_ - 8); }

private:
    // We construct a char array of the form:
    //      klength  varint32          <-- start_
    //      userkey  char[klength]    <-- kstart_
    //      tag      uint64
    //                                     <-- end_
    // The array is a suitable MemTable key.
    // The suffix starting with "userkey" can be used as an InternalKey.
    const char* start_;
    const char* kstart_;
    const char* end_;

```

```


    char space_[200]; // Avoid allocation for short keys
};

LookupKey::LookupKey(const Slice& user_key, SequenceNumber s) {
    size_t usize = user_key.size();

    size_t needed = usize + 13; // A conservative estimate
    char* dst;
    if (needed <= sizeof(space_)) {
        dst = space_;
    } else {
        dst = new char[needed];
    }
    start_ = dst;
    dst = EncodeVarint32(dst, usize + 8);
    kstart_ = dst;
    memcpy(dst, user_key.data(), usize);
    dst += usize;
    EncodeFixed64(dst, PackSequenceAndType(s, kValueTypeForSeek));
    dst += 8;
    end_ = dst;
}

```

在 LookupKey 中可以返回三种 Key，画图理理关系：



```
|          +-----+-----+-----+ |  
+-----+
```

回到内存数据库的定义中，KeyComparator 对 InternalKeyComparator 又进行了一次封装，来看下其具体函数定义：

```
static Slice GetLengthPrefixedSlice(const char* data) {  
    uint32_t len;  
    const char* p = data;  
    p = GetVarint32Ptr(p, p + 5, &len); // +5: we assume "p" is not corrupted  
    return Slice(p, len);  
}  
  
int MemTable::KeyComparator::operator()(const char* aptr,  
                                         const char* bptr) const {  
    // Internal keys are encoded as length-prefixed strings.  
    Slice a = GetLengthPrefixedSlice(aptr);  
    Slice b = GetLengthPrefixedSlice(bptr);  
    return comparator.Compare(a, b);  
}
```

这里的 comparator 自然是 InternalKeyComparator，从 GetLengthPrefixedSlice 推测该函数的参数为 MemTable Key。继续看 MemTable::Add 和 MemTable::Get 的实现：

```
void MemTable::Add(SequenceNumber s, ValueType type, const Slice& key,  
                  const Slice& value) {
```

```

// Format of an entry is concatenation of:
// key_size      : varint32 of internal_key.size()
// key bytes     : char[internal_key.size()]
// value_size    : varint32 of value.size()
// value bytes   : char[value.size()]

size_t key_size = key.size();
size_t val_size = value.size();
size_t internal_key_size = key_size + 8;
const size_t encoded_len = VarintLength(internal_key_size) +
                           internal_key_size + VarintLength(val_size) +
                           val_size;

char* buf = arena_.Allocate(encoded_len);
char* p = EncodeVarint32(buf, internal_key_size);
memcpy(p, key.data(), key_size);
p += key_size;
EncodeFixed64(p, (s << 8) | type);
p += 8;
p = EncodeVarint32(p, val_size);
memcpy(p, value.data(), val_size);
assert(p + val_size == buf + encoded_len);
table_.Insert(buf);
}

bool MemTable::Get(const LookupKey& key, std::string* value, Status* s) {
    Slice memkey = key.memtable_key();
    Table::Iterator iter(&table_);
    iter.Seek(memkey.data());
    if (iter.Valid()) {
        // entry format is:
        //    klength  varint32

```

```

//    userkey  char[klength]
//    tag      uint64
//    vlength  varint32
//    value    char[vlength]
// Check that it belongs to same user key.  We do not check the

// sequence number since the Seek() call above should have skipped
// all entries with overly large sequence numbers.
const char* entry = iter.key();
uint32_t key_length;
const char* key_ptr = GetVarint32Ptr(entry, entry + 5, &key_length);
if (comparator_.comparator.user_comparator()->Compare(
    Slice(key_ptr, key_length - 8), key.user_key()) == 0) {
    // Correct user key
    const uint64_t tag = DecodeFixed64(key_ptr + key_length - 8);
    switch (static_cast<ValueType>(tag & 0xff)) {
        case kTypeValue: {
            Slice v = GetLengthPrefixedSlice(key_ptr + key_length);
            value->assign(v.data(), v.size());
            return true;
        }
        case kTypeDeletion:
            *s = Status::NotFound(Slice());
            return true;
    }
}
return false;
}

```

MemTable::Add 的注释中说明了条目的格式。除了图中 key 部分的内容外，条目还

MemTable::Add 的工作与图 1 所示的格式，除了图中 Key 部分的内容外，条目还包括 Value 的存储，打包好后插入跳表中。MemTable::Get 中则根据 MemTable Key 查找对应的条目，查到后使用 user_comparator 进一步比较两者的 User Key 是否一致，如

果完全一致并且值类型不是删除，就把条目中的 Value 读出来放到结果中。最后看下迭代器的实现：

```
static const char* EncodeKey(std::string* scratch, const Slice& target) {
    scratch->clear();
    PutVarint32(scratch, target.size());
    scratch->append(target.data(), target.size());
    return scratch->data();
}

class MemTableIterator : public Iterator {
public:
    explicit MemTableIterator(MemTable::Table* table) : iter_(table) {}

    MemTableIterator(const MemTableIterator&) = delete;
    MemTableIterator& operator=(const MemTableIterator&) = delete;

    ~MemTableIterator() override = default;

    bool Valid() const override { return iter_.Valid(); }
    void Seek(const Slice& k) override { iter_.Seek(EncodeKey(&tmp_, k)); }
    void SeekToFirst() override { iter_.SeekToFirst(); }
    void SeekToLast() override { iter_.SeekToLast(); }
    void Next() override { iter_.Next(); }
```

```

void Prev() override { iter_.Prev(); }
Slice key() const override { return GetLengthPrefixedSlice(iter_.key()); }
Slice value() const override {
    Slice key_slice = GetLengthPrefixedSlice(iter_.key());
    return GetLengthPrefixedSlice(key_slice.data() + key_slice.size());
}

Status status() const override { return Status::OK(); }

private:
    MemTable::Table::Iterator iter_;
    std::string tmp_; // For passing to EncodeKey
};

Iterator* MemTable::NewIterator() { return new MemTableIterator(&table_); }

```

MemTableIterator 对跳表的迭代器进行了一层封装，跳表的条目中编码了 Key 和 Value，MemTableIterator 中对其进行了对应的解析。至此，内存数据库的代码就分析完了。

总结

上一篇和本篇分析了 LevelDB 写操作的具体过程，分析了 WriteBatch、Log 和 MemTable 的代码，理解了高性能写操作的原理。不过内存数据库的大小毕竟是有限的，总会需要把内存中的数据持久化到硬盘中，并且还需要提供高速的硬盘数据查询操作，这些又是怎样实现的呢？且听下回分解！


0 comments

Write

Preview

Aa

Sign in to comment

 Styling with Markdown is supported

Sign in with GitHub

Except where otherwise noted, content on this site is licensed under a CC BY-SA 4.0 license.
Copyright©2017 SF-Zhou, All Rights Reserved. Powered by GitHub Pages and GitHub Actions.