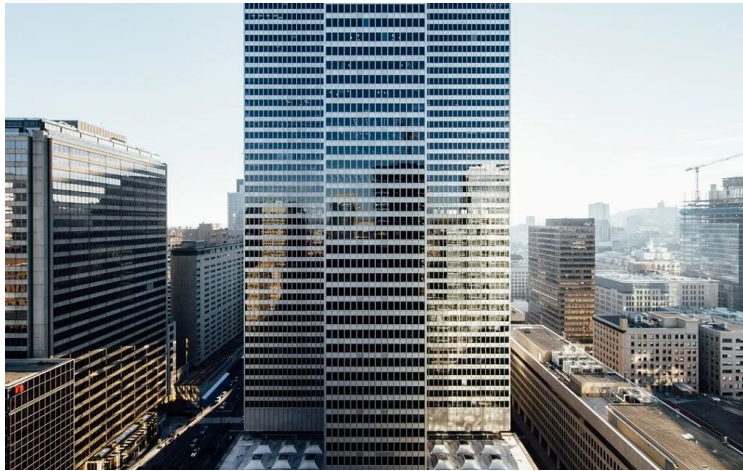


# MySQL 8.0 Server层最新架构详解

Original 道客 阿里开发者 2021-07-25 17:15

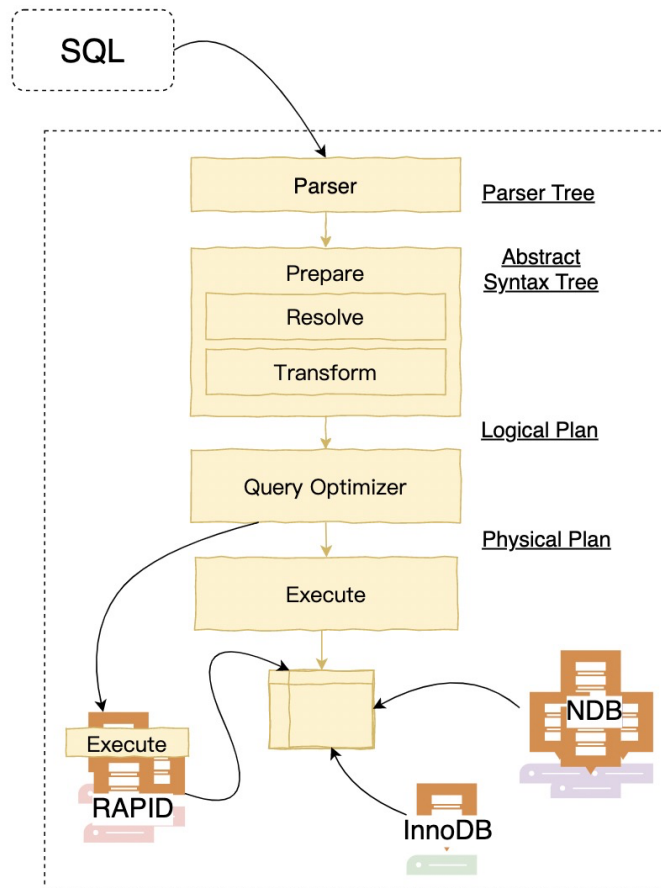
收录于合集

#mysql 5 #数据存储与数据库 16



## 一 背景和架构

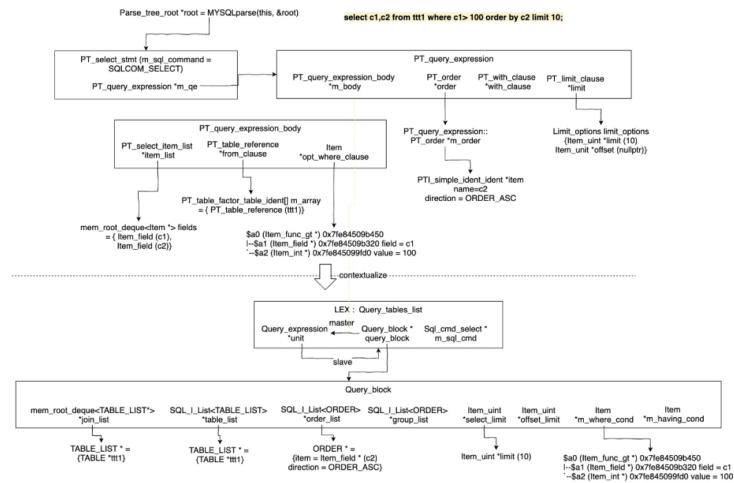
本文基于MySQL 8.0.25源码进行分析和总结。这里MySQL Server层指的是MySQL的优化器、执行器部分。我们对MySQL的理解还建立在5.6和5.7版本的理解之上，更多的是对比PostgreSQL或者传统数据库。然而从MySQL 8.0开始，持续每三个月的迭代和重构工作，使得MySQL Server层的整体架构有了质的飞越。下面来看下MySQL最新的架构。



我们可以看到最新的MySQL的分层架构和其他数据库并没有太大的区别，另外值得一提的是从图中可以看出MySQL现在更多的加强InnoDB、NDB集群和RAPID(HeatWave clusters)内存集群架构的演进。下面我们就看下具体细节，我们这次不随着官方的Feature实现和重构顺序进行理解，本文更偏向于从优化器、执行器的流程角度来演进。

## 二 MySQL 解析器Parser

首先从Parser开始，官方MySQL 8.0使用Bison进行了重写，生成Parser Tree，同时Parser Tree会contextualize生成MySQL抽象语法树（Abstract Syntax Tree）。



MySQL抽象语法树和其他数据库有些不同，是由比较让人拗口的SELECT\_LEX\_UNIT/SELECT\_LEX类交替构成的，然而这两个结构在最新的版本中已经重命名成标准的SELECT\_LEX -> Query\_block和SELECT\_LEX\_UNIT -> Query\_expression。Query\_block是代表查询块，而Query\_expression是包含多个查询块的查询表达式，包括UNION AND/OR的查询块（如SELECT \* FROM t1 union SELECT \* FROM t2）或者有多Level的ORDER BY/LIMIT（如SELECT \* FROM t1 ORDER BY a LIMIT 10）ORDER BY b LIMIT 5。

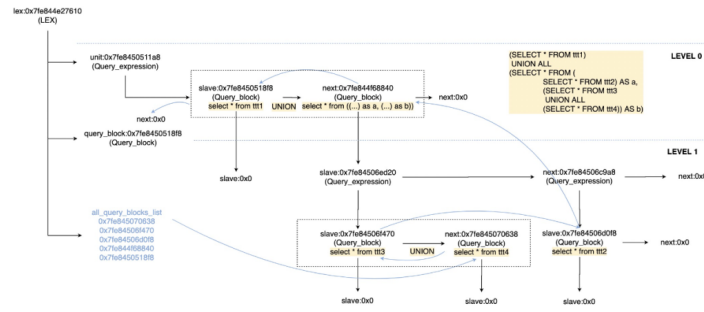
例如，来看一个复杂的嵌套查询：

```

1  (SELECT *
2    FROM ttt1)
3  UNION ALL
4  (SELECT *
5    FROM
6      (SELECT *
7        FROM ttt2) AS a,
8      (SELECT *
9        FROM ttt3
10       UNION ALL SELECT *
11       FROM ttt4) AS b)

```

在MySQL中就可以用下面方式表达：



经过解析和转换后的语法树仍然建立在Query\_block和Query\_expression的框架下，只不过有些LEVEL的query block被消除或者合并了，这里不再详细展开。

### 三 MySQL prepare/rewrite阶段

接下来我们要经过resolve和transformation过程Query\_expression::prepare->Query\_block::prepare，这个过程包括（按功能分而非完全按照执行顺序）：

#### 1 Setup and Fix

- setup\_tables: Set up table leaves in the query block based on list of tables.
- resolve\_placeholder\_tables/merge\_derived/setup\_table\_function/setup\_materialized\_derived: Resolve derived table, view or table function references in query block.
- setup\_natural\_join\_row\_types: Compute and store the row types of the top-most NATURAL/USING joins.
- setup\_wild: Expand all '\*' in list of expressions with the matching column references.
- setup\_base\_ref\_items: Set query\_block's base\_ref\_items.
- setup\_fields: Check that all given fields exists and fill struct with current data.
- setup\_conds: Resolve WHERE condition and join conditions.
- setup\_group: Resolve and set up the GROUP BY list.

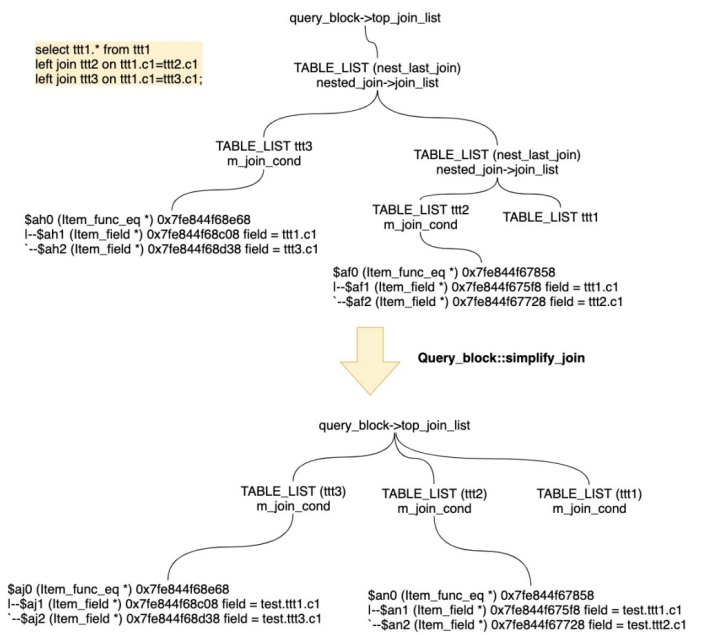
- `m_having_cond->fix_fields`: Setup the HAVING clause.
- `resolve_rollup`: Resolve items in SELECT list and ORDER BY list for rollup processing.
- `resolve_rollup_item`: Resolve an item (and its tree) for rollup processing by replacing items matching grouped expressions with `Item_rollup_group_items` and updating properties (`m_nullable`, `PROP_ROLLUP_FIELD`). Also check any GROUPING function for incorrect column.
- `setup_order`: Set up the ORDER BY clause.
- `resolve_limits`: Resolve OFFSET and LIMIT clauses.
- `Window::setup_windows1` : Set up windows after `setup_order()` and before `setup_order_final()`.
- `setup_order_final` : Do final setup of ORDER BY clause, after the query block is fully resolved.
- `setup_ftfuncs`: Setup full-text functions after resolving HAVING.
- `resolve_rollup_wfs` : Replace group by field references inside window functions with references in the presence of ROLLUP.

## 2 Transformation

- `remove_redundant_subquery_clause` : Permanently remove redundant parts from the query if
  - 1) This is a subquery
  - 2) Not normalizing a view. Removal should take place when a query involving a view is optimized, not when the view is created.
- `remove_base_options` : Remove SELECT\_DISTINCT options from a query block if can skip distinct.
- `resolve_subquery` : Resolve predicate involving subquery, perform early unconditional subquery transformations.

- Convert subquery predicate into semi-join, or
- Mark the subquery for execution using materialization, or
- Perform IN->EXISTS transformation, or
- Perform more/less ALL/ANY -> MIN/MAX rewrite
- Substitute trivial scalar-context subquery with its value
- `transform_scalar_subqueries_to_join_with_derived`: Transform eligible scalar subqueries to derived tables.
- `flatten_subqueries`: Convert semi-join subquery predicates into semi-join join nests. Convert candidate subquery predicates into semi-join join nests. This transformation is performed once in query lifetime and is irreversible.
- `apply_local_transforms`:
  - `delete_unused_merged_columns`: If query block contains one or more merged derived tables/views, walk through lists of columns in select lists and remove unused columns.
  - `simplify_joins`: Convert all outer joins to inner joins if possible
  - `prune_partitions`: Perform partition pruning for a given table and condition.
- `push_conditions_to_derived_tables`: Pushing conditions down to derived tables must be done after validity checks of grouped queries done by `apply_local_transforms()`;
- `Window::eliminate_unused_objects`: Eliminate unused window definitions, redundant sorts etc.

这里，节省篇幅，我们只举例关注下和top\_join\_list相关的simple\_joins这个函数的作用，对于Query\_block中嵌套join的简化过程。



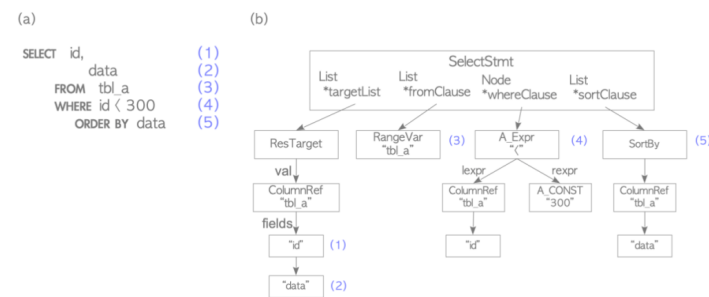
### 3 对比PostgreSQL

为了更清晰的理解标准数据库的做法，我们这里引用了PostgreSQL的这三个过程：

#### Parser

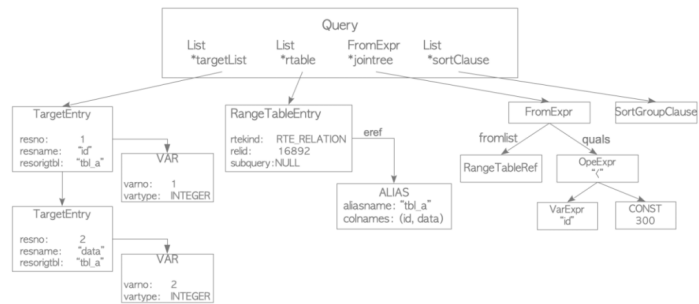
下图首先Parser把SQL语句生成parse tree。

```
1 testdb=# SELECT id, data FROM tbl_a WHERE id < 300 ORDER BY data;
```



## Analyzer/Analyser

下图展示了PostgreSQL的analyzer/analyser如何将parse tree通过语义分析后生成query tree。



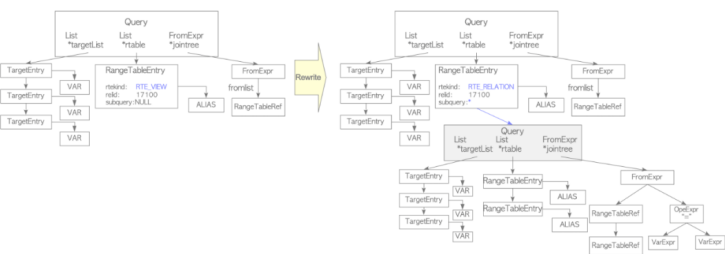
## Rewriter

Rewriter会根据规则系统中的规则把query tree进行转换改写。

```
1 sampled= CREATE VIEW employees_list
2 sampled-# AS SELECT e.id, e.name, d.name AS department
3 sampled-# FROM employees AS e, departments AS d WHERE e.depart
```

下图的例子就是一个包含views的query tree如何展开成新的query tree。

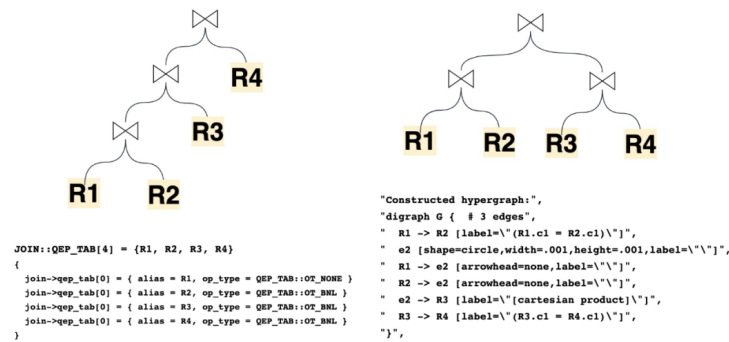
```
1 sampled= SELECT * FROM employees_list;
```



## 四 MySQL Optimize和Planning阶段



接下来我们进入了逻辑计划生成物理计划的过程，本文还是注重于结构的解析，而不去介绍生成的细节，MySQL过去在8.0.22之前，主要依赖的结构就是JOIN和QEP\_TAB。JOIN是与之对应的每个Query\_block，而QEP\_TAB对应的每个Query\_block涉及到的具体“表”的顺序、方法和执行计划。然而在8.0.22之后，新的基于Hypergraph的优化器算法成功的抛弃了QEP\_TAB结构来表达左深树的执行计划，而直接使用HyperNode/HyperEdge的图来表示执行计划。



举例可以看到数据结构表达的left deep tree和超图结构表达的bushy tree对应的不同计划展现：

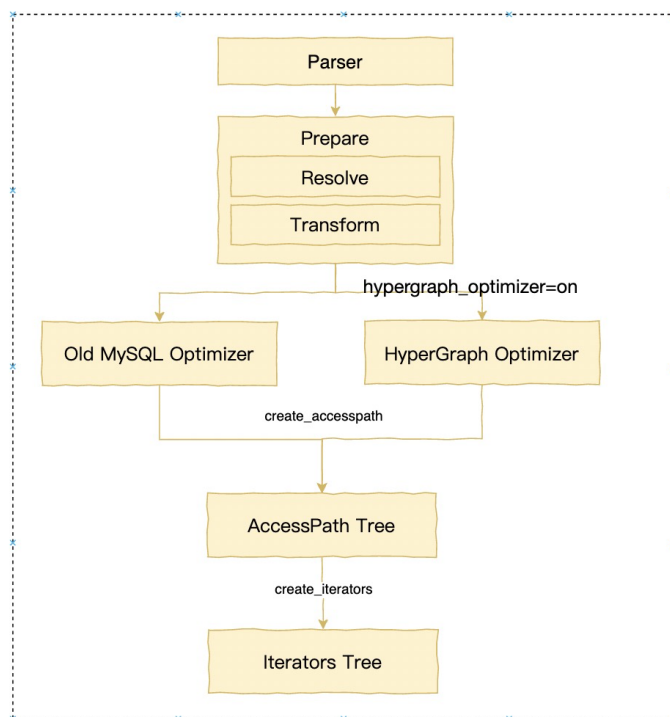
```

1 | -> Inner hash join (no condition) (cost=1.40 rows=1)
2   -> Table scan on R4 (cost=0.35 rows=1)
3   -> Hash
4     -> Inner hash join (no condition) (cost=1.05 rows=1)
5       -> Table scan on R3 (cost=0.35 rows=1)
6       -> Hash
7         -> Inner hash join (no condition) (cost=0.70 rows=1)
8           -> Table scan on R2 (cost=0.35 rows=1)
9           -> Hash
10             -> Table scan on R1 (cost=0.35 rows=1)
11
12 | -> Nested loop inner join (cost=0.55..0.55 rows=0)
13   -> Nested loop inner join (cost=0.50..0.50 rows=0)
14     -> Table scan on R4 (cost=0.25..0.25 rows=1)
15     -> Filter: (R4.c1 = R3.c1) (cost=0.35..0.35 rows=0)
16       -> Table scan on R3 (cost=0.25..0.25 rows=1)
17     -> Nested loop inner join (cost=0.50..0.50 rows=0)
18       -> Table scan on R2 (cost=0.25..0.25 rows=1)

```

```
19      -> Filter: (R2.c1 = R1.c1) (cost=0.35..0.35 rows=0)
20      -> Table scan on R1 (cost=0.25..0.25 rows=1)
```

MySQL8.0.2x为了更好的兼容两种优化器，引入了新的类AccessPath，可以认为这是MySQL为了解耦执行器和不同优化器抽象出来的Plan Tree。



## 1 老优化器的入口

老优化器仍然走JOIN::optimize来把query block转换成query execution plan (QEP)。

这个阶段仍然做一些逻辑的重写工作，这个阶段的转换可以理解为基于cost-based优化前做准备，详细步骤如下：

- Logical transformations
  - optimize\_derived : Optimize the query expression representing a derived table/view.
  - optimize\_cond : Equality/constant propagation.

- `prune_table_partitions` : Partition pruning.
- `optimize_aggregated_query` : `COUNT(*)`, `MIN()`, `MAX()` constant substitution in case of implicit grouping.
- `substitute_gc` : ORDER BY optimization, substitute all expressions in the WHERE condition and ORDER/GROUP lists that match generated columns (GC) expressions with GC fields, if any.
- Perform cost-based optimization of table order and access path selection.
  - `JOIN::make_join_plan()` : Set up join order and initial access paths.
- Post-join order optimization
  - `substitute_for_best_equal_field` : Create optimal table conditions from the where clause and the join conditions.
  - `make_join_query_block` : Inject outer-join guarding conditions.
  - Adjust data access methods after determining table condition (several times).
  - `optimize_distinct_group_order` : Optimize ORDER BY/DISTINCT.
  - `optimize_fts_query` : Perform FULLTEXT search before all regular searches.
  - `remove_eq_conds` : Removes const and eq items. Returns the new item, or nullptr if no condition.
  - `replace_index_subquery/create_access_paths_for_index_subquery` : See if this subquery can be evaluated with `subselect_indexsubquery_engine`.
  - `setup_join_buffering` : Check whether join cache could be used.

- Code generation
  - `alloc_qep(tables)` : Create QEP\_TAB array.
  - `test_skip_sort` : Try to optimize away sorting/distinct.
  - `make_join_readinfo` : Plan refinement stage: do various setup things for the executor.
  - `make_tmp_tables_info` : Setup temporary table usage for grouping and/or sorting.
  - `push_to_engines` : Push (parts of) the query execution down to the storage engines if they can provide faster execution of the query, or part of it.
  - `create_access_paths` : generated ACCESS\_PATH.

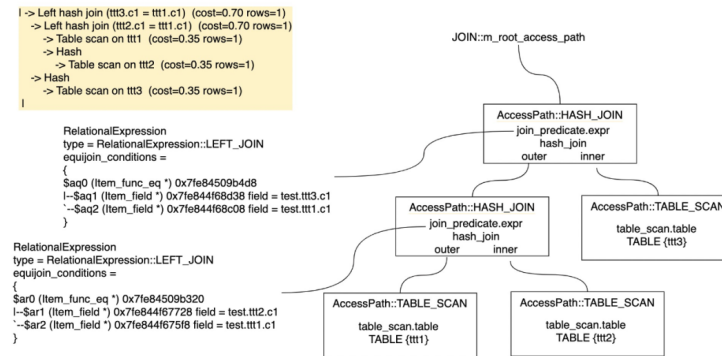
## 2 新优化器的入口

新优化器默认不打开，必须通过`set optimizer_switch="hypergraph_optimizer=on";`来打开。主要通过`FindBestQueryPlan`函数来实现，逻辑如下：

- 先判断是否属于新优化器可以支持的Query语法（`CheckSupportedQuery`），不支持的直接返回错误`ER_HYPERGRAPH_NOT_SUPPORTED_YET`。
- 转化`top_join_list`变成`JoinHypergraph`结构。由于Hypergraph是比较独立的算法层面的实现，`JoinHypergraph`结构用来更好的把数据库的结构包装到Hypergraph的edges和nodes的概念上的。
- 通过`EnumerateAllConnectedPartitions`实现论文中的DPhyp算法。
- `CostingReceiver`类包含了过去JOIN planning的主要逻辑，包括根据cost选择相应的访问路径，根据DPhyp生成的子计划进行评估，保留cost最小的子计划。

- 得到root\_path后，接下来处理group/agg/having/sort/limit的。对于Group by操作，目前Hypergraph使用sorting first + streaming aggregation的方式。

举例看下Plan (AccessPath) 和SQL的关系：



最后生成Iterator执行器框架需要的Iterator执行载体，AccessPath和Iterator是一对一的关系（Access paths are a query planning structure that correspond 1:1 to iterators）。

```

1 Query_expression::m_root_iterator = CreateIteratorFromAccessPath(.....)
2
3 unique_ptr_destroy_only<RowIterator> CreateIteratorFromAccessPath(
4     THD *thd, AccessPath *path, JOIN *join, bool eligible_for_batch_mode
5     .....
6     switch (path->type) {
7         case AccessPath::TABLE_SCAN: {
8             const auto &param = path->table_scan();
9             iterator = NewIterator<TableScanIterator>(
10                 thd, param.table, path->num_output_rows, examined_rows);
11             break;
12         }
13         case AccessPath::INDEX_SCAN: {
14             const auto &param = path->index_scan();
15             if (param.reverse) {
16                 iterator = NewIterator<IndexScanIterator<true>>(
17                 thd, param.table, param.idx, param.use_order, path->num_output_rows);

```

本文主要focus在MySQL最新版本官方的源码上，重点分析了官方的重构在多阶段和各阶段结构上的变化和联系，更多的是为了让大家了解一个全新的MySQL的发展。

## 关于我们

PolarDB 是阿里巴巴自主研发的云原生分布式关系型数据库，于2020年进入Gartner全球数据库Leader象限，并获得了2020年中国电子学会颁发的科技进步一等奖。PolarDB 基于云原生分布式数据库架构，提供大规模在线事务处理能力，兼具对复杂查询的并行处理能力，在云原生分布式数据库领域整体达到了国际领先水平，并且得到了广泛的市场认可。在阿里巴巴集团内部的最佳实践中，PolarDB还全面支撑了2020年天猫双十一，并刷新了数据库处理峰值记录，高达1.4亿TPS。欢迎有志之士加入我们，简历请投递到[daoke.wangc@alibaba-inc.com](mailto:daoke.wangc@alibaba-inc.com)，期待与您共同打造世界一流的下一代云原生分布式关系型数据库。

---

## 数据库技术图谱

数据库技术图谱由阿里云数据库专家团出品，含7个知识点，17个课程，6个体验场景，9个公开课，带你从基础到进阶玩转常见数据库，同时深入实战，学习阿里云使用各数据库的最佳实践！

点击“阅读原文”，开始学习吧~

收录于合集 [#mysql 5](#)

[上一篇 · MySQL 深潜 - 一文详解 MySQL Data Dictionary](#)

[Read more](#)

People who liked this content also liked

代码重构：面向单元测试

阿里开发者

