

Let's Build a Simple Database

Writing a sqlite clone from scratch in C

[Overview](#)

[View on GitHub \(pull requests welcome\)](#)

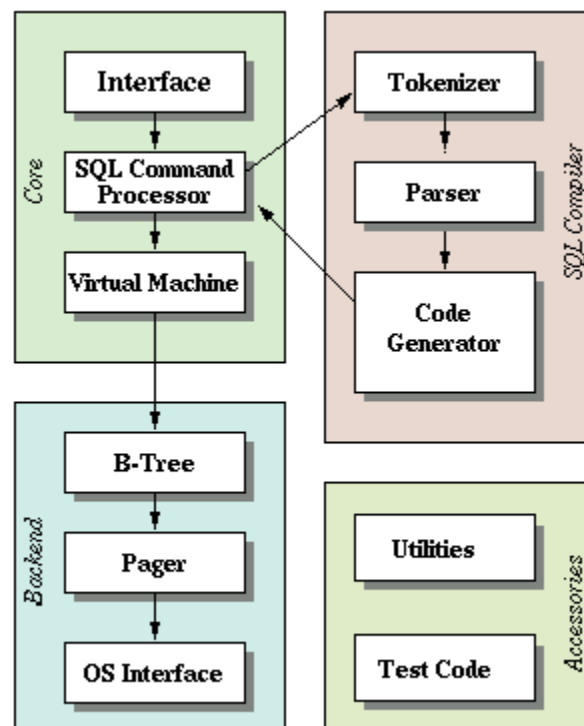
Part 2 - World's Simplest SQL Compiler and Virtual Machine

[< Part 1 - Introduction and Setting up the REPL](#)

[Part 3 - An In-Memory, Append-Only, Single-Table Database >](#)

We're making a clone of sqlite. The "front-end" of sqlite is a SQL compiler that parses a string and outputs an internal representation called bytecode.

This bytecode is passed to the virtual machine, which executes it.



SQLite Architecture (<https://www.sqlite.org/arch.html>)

Breaking things into two steps like this has a couple advantages:

- Reduces the complexity of each part (e.g. virtual machine does not worry about syntax errors)
- Allows compiling common queries once and caching the bytecode for improved performance

With this in mind, let's refactor our main function and support two new keywords in the process:

```
int main(int argc, char* argv[]) {
    InputBuffer* input_buffer = new_input_buffer();
    while (true) {
        print_prompt();
        read_input(input_buffer);

        if (strcmp(input_buffer->buffer, ".exit") == 0) {
            exit(EXIT_SUCCESS);
        } else {
            printf("Unrecognized command '%s'.\n", input_buffer->buffer);
            if (input_buffer->buffer[0] == '.') {
                switch (do_meta_command(input_buffer)) {
                    case (META_COMMAND_SUCCESS):
                        continue;
                    case (META_COMMAND_UNRECOGNIZED_COMMAND):
                        printf("Unrecognized command '%s'\n", input_buffer->buffer);
                        continue;
                }
            }

            Statement statement;
            switch (prepare_statement(input_buffer, &statement)) {
                case (PREPARE_SUCCESS):
                    break;
                case (PREPARE_UNRECOGNIZED_STATEMENT):
                    printf("Unrecognized keyword at start of '%s'.\n",
                        input_buffer->buffer);
                    continue;
            }
        }
    }
}
```

```
+     execute_statement(&statement);  
+     printf("Executed.\n");  
    }  
}
```

Non-SQL statements like `.exit` are called “meta-commands”. They all start with a dot, so we check for them and handle them in a separate function.

Next, we add a step that converts the line of input into our internal representation of a statement. This is our hacky version of the sqlite front-end.

Lastly, we pass the prepared statement to `execute_statement`. This function will eventually become our virtual machine.

Notice that two of our new functions return enums indicating success or failure:

```
typedef enum {  
    META_COMMAND_SUCCESS,  
    META_COMMAND_UNRECOGNIZED_COMMAND  
} MetaCommandResult;  
  
typedef enum { PREPARE_SUCCESS, PREPARE_UNRECOGNIZED_STATEMENT
```

“Unrecognized statement”? That seems a bit like an exception. I prefer not to use exceptions (and C doesn’t even support them), so I’m using enum result codes wherever practical. The C compiler will complain if my switch statement doesn’t handle a member of the enum, so we can feel a little more confident we handle every result of a function. Expect more result codes to be added in the future.

`do_meta_command` is just a wrapper for existing functionality that leaves room for more commands:

```
MetaCommandResult do_meta_command (InputBuffer* input_buffer) {  
    if (strcmp(input_buffer->buffer, ".exit") == 0) {  
        exit(EXIT_SUCCESS);  
    } else {  
        return META_COMMAND_UNRECOGNIZED_COMMAND;  
    }  
}
```

Our “prepared statement” right now just contains an enum with two possible values. It will contain more data as we allow parameters in statements:

```
typedef enum { STATEMENT_INSERT, STATEMENT_SELECT } StatementType;

typedef struct {
    StatementType type;
} Statement;
```

`prepare_statement` (our “SQL Compiler”) does not understand SQL right now. In fact, it only understands two words:

```
PrepareResult prepare_statement (InputBuffer* input_buffer,
                                Statement* statement) {
    if (strncmp(input_buffer->buffer, "insert", 6) == 0) {
        statement->type = STATEMENT_INSERT;
        return PREPARE_SUCCESS;
    }
    if (strcmp(input_buffer->buffer, "select") == 0) {
        statement->type = STATEMENT_SELECT;
        return PREPARE_SUCCESS;
    }

    return PREPARE_UNRECOGNIZED_STATEMENT;
}
```

Note that we use `strncmp` for “insert” since the “insert” keyword will be followed by data. (e.g. `insert 1 cstack foo@bar.com`)

Lastly, `execute_statement` contains a few stubs:

```
void execute_statement (Statement* statement) {
    switch (statement->type) {
        case (STATEMENT_INSERT):
            printf("This is where we would do an insert.\n");
            break;
        case (STATEMENT_SELECT):
            printf("This is where we would do a select.\n");
    }
```

```
        break;
    }
}
```

Note that it doesn't return any error codes because there's nothing that could go wrong yet.

With these refactors, we now recognize two new keywords!

```
~ ./db
db > insert foo bar
This is where we would do an insert.
Executed.
db > delete foo
Unrecognized keyword at start of 'delete foo'.
db > select
This is where we would do a select.
Executed.
db > .tables
Unrecognized command '.tables'
db > .exit
~
```

The skeleton of our database is taking shape... wouldn't it be nice if it stored data? In the next part, we'll implement `insert` and `select`, creating the world's worst data store. In the mean time, here's the entire diff from this part:

```
@@ -10,6 +10,23 @@ struct InputBuffer_t {
    } InputBuffer;

+typedef enum {
+  META_COMMAND_SUCCESS,
+  META_COMMAND_UNRECOGNIZED_COMMAND
+} MetaCommandResult;
+
+typedef enum { PREPARE_SUCCESS, PREPARE_UNRECOGNIZED_STATEMENT
+
+typedef enum { STATEMENT_INSERT, STATEMENT_SELECT } StatementTy
+
```

```
+typedef struct {
+  StatementType type;
+} Statement;
+
+  InputBuffer* new_input_buffer() {
+    InputBuffer* input_buffer = malloc(sizeof(InputBuffer));
+    input_buffer->buffer = NULL;
@@ -40,17 +57,67 @@ void close_input_buffer(InputBuffer* input_b
+    free(input_buffer);
+  }

+MetaCommandResult do_meta_command(InputBuffer* input_buffer) {
+  if (strcmp(input_buffer->buffer, ".exit") == 0) {
+    close_input_buffer(input_buffer);
+    exit(EXIT_SUCCESS);
+  } else {
+    return META_COMMAND_UNRECOGNIZED_COMMAND;
+  }
+}
+
+PrepareResult prepare_statement(InputBuffer* input_buffer,
+                                Statement* statement) {
+  if (strncmp(input_buffer->buffer, "insert", 6) == 0) {
+    statement->type = STATEMENT_INSERT;
+    return PREPARE_SUCCESS;
+  }
+  if (strcmp(input_buffer->buffer, "select") == 0) {
+    statement->type = STATEMENT_SELECT;
+    return PREPARE_SUCCESS;
+  }
+  return PREPARE_UNRECOGNIZED_STATEMENT;
+}
+
+void execute_statement(Statement* statement) {
+  switch (statement->type) {
+    case (STATEMENT_INSERT):
+      printf("This is where we would do an insert.\n");
+      break;
+    case (STATEMENT_SELECT):
```

```
+     printf("This is where we would do a select.\n");
+     break;
+ }
+}
+
+ int main(int argc, char* argv[]) {
+     InputBuffer* input_buffer = new_input_buffer();
+     while (true) {
+         print_prompt();
+         read_input(input_buffer);
+
+         if (strcmp(input_buffer->buffer, ".exit") == 0) {
+             close_input_buffer(input_buffer);
+             exit(EXIT_SUCCESS);
+         } else {
+             printf("Unrecognized command '%s'.\n", input_buffer->buffer);
+             if (input_buffer->buffer[0] == '.') {
+                 switch (do_meta_command(input_buffer)) {
+                     case (META_COMMAND_SUCCESS):
+                         continue;
+                     case (META_COMMAND_UNRECOGNIZED_COMMAND):
+                         printf("Unrecognized command '%s'\n", input_buffer->buffer);
+                         continue;
+                 }
+             }
+
+             Statement statement;
+             switch (prepare_statement(input_buffer, &statement)) {
+                 case (PREPARE_SUCCESS):
+                     break;
+                 case (PREPARE_UNRECOGNIZED_STATEMENT):
+                     printf("Unrecognized keyword at start of '%s'.\n",
+                             input_buffer->buffer);
+                     continue;
+             }
+
+             execute_statement(&statement);
+             printf("Executed.\n");
+         }
+     }
+ }
```

[< Part 1 - Introduction and Setting up the REPL](#)

[Part 3 - An In-Memory, Append-Only, Single-Table Database >](#)

[rss](#) | [subscribe by email](#)

This project is maintained by [cstack](#)

Hosted on GitHub Pages — Theme by [orderedlist](#)