

# Let's Build a Simple Database

Writing a sqlite clone from scratch in C

[Overview](#)

[View on GitHub \(pull requests welcome\)](#)

---

## Part 4 - Our First Tests (and Bugs)

[< Part 3 - An In-Memory, Append-Only, Single-Table Database](#)

[Part 5 - Persistence to Disk >](#)

We've got the ability to insert rows into our database and to print out all rows. Let's take a moment to test what we've got so far.

I'm going to use [RSpec](#) to write my tests because I'm familiar with it, and the syntax is fairly readable.

I'll define a short helper to send a list of commands to our database program then make assertions about the output:

```
describe 'database' do
  def run_script(commands)
    raw_output = nil
    IO.popen("./db", "r+") do |pipe|
      commands.each do |command|
        pipe.puts command
      end

      pipe.close_write

      # Read entire output
      raw_output = pipe.gets(nil)
    end
    raw_output.split("\n")
  end
end
```

```
end

it 'inserts and retrieves a row' do
  result = run_script([
    "insert 1 user1 person1@example.com",
    "select",
    ".exit",
  ])
  expect(result).to match_array([
    "db > Executed.",
    "db > (1, user1, person1@example.com)",
    "Executed.",
    "db > ",
  ])
end
end
```

This simple test makes sure we get back what we put in. And indeed it passes:

```
bundle exec rspec
```

```
.
```

```
Finished in 0.00871 seconds (files took 0.09506 seconds to load)
1 example, 0 failures
```

Now it's feasible to test inserting a large number of rows into the database:

```
it 'prints error message when table is full' do
  script = (1..1401).map do |i|
    "insert #{i} user#{i} person#{i}@example.com"
  end
  script << ".exit"
  result = run_script(script)
  expect(result[-2]).to eq('db > Error: Table full.')
end
```

Running tests again...

```
bundle exec rspec
```

```
..
```

```
Finished in 0.01553 seconds (files took 0.08156 seconds to load)
2 examples, 0 failures
```

Sweet, it works! Our db can hold 1400 rows right now because we set the maximum number of pages to 100, and 14 rows can fit in a page.

Reading through the code we have so far, I realized we might not handle storing text fields correctly. Easy to test with this example:

```
it 'allows inserting strings that are the maximum length' do
  long_username = "a"*32
  long_email = "a"*255
  script = [
    "insert 1 #{long_username} #{long_email}",
    "select",
    ".exit",
  ]
  result = run_script(script)
  expect(result).to match_array([
    "db > Executed.",
    "db > (1, #{long_username}, #{long_email})",
    "Executed.",
    "db > ",
  ])
end
```

And the test fails!

#### Failures:

1) database allows inserting strings that are the maximum length  
Failure/Error: raw\_output.split("\n")

#### ArgumentError:

invalid byte sequence in UTF-8  
# ./spec/main\_spec.rb:14:in `split'

```
# ./spec/main_spec.rb:14:in `run_script'  
# ./spec/main_spec.rb:48:in `block (2 levels) in <top (req
```

If we try it ourselves, we'll see that there's some weird characters when we try to print out the row. (I'm abbreviating the long strings):

```
db > insert 1 aaaaa... aaaaa...  
Executed.  
db > select  
(1, aaaaa...aaa\0, aaaaa...aaa\0)  
Executed.  
db >
```

What's going on? If you take a look at our definition of a Row, we allocate exactly 32 bytes for username and exactly 255 bytes for email. But [C strings](#) are supposed to end with a null character, which we didn't allocate space for. The solution is to allocate one additional byte:

```
const uint32_t COLUMN_EMAIL_SIZE = 255;  
typedef struct {  
    uint32_t id;  
    - char username[COLUMN_USERNAME_SIZE];  
    - char email[COLUMN_EMAIL_SIZE];  
    + char username[COLUMN_USERNAME_SIZE + 1];  
    + char email[COLUMN_EMAIL_SIZE + 1];  
} Row;
```

And indeed that fixes it:

```
bundle exec rspec  
...  
  
Finished in 0.0188 seconds (files took 0.08516 seconds to load)  
3 examples, 0 failures
```

We should not allow inserting usernames or emails that are longer than column size. The spec for that looks like this:

```

it 'prints error message if strings are too long' do
  long_username = "a"*33
  long_email = "a"*256
  script = [
    "insert 1 #{long_username} #{long_email}",
    "select",
    ".exit",
  ]
  result = run_script(script)
  expect(result).to match_array([
    "db > String is too long.",
    "db > Executed.",
    "db > ",
  ])
end

```

In order to do this we need to upgrade our parser. As a reminder, we're currently using `scanf()`:

```

if (strcmp(input_buffer->buffer, "insert", 6) == 0) {
  statement->type = STATEMENT_INSERT;
  int args_assigned = sscanf(
    input_buffer->buffer, "insert %d %s %s", &(statement->row_to_insert.id,
    statement->row_to_insert.username, statement->row_to_insert.email);
  if (args_assigned < 3) {
    return PREPARE_SYNTAX_ERROR;
  }
  return PREPARE_SUCCESS;
}

```

But `scanf` has some disadvantages. If the string it's reading is larger than the buffer it's reading into, it will cause a buffer overflow and start writing into unexpected places. We want to check the length of each string before we copy it into a Row structure. And to do that, we need to divide the input by spaces.

I'm going to use `strtok()` to do that. I think it's easiest to understand if you see it in action:

```

+PrepareResult prepare_insert(InputBuffer* input_buffer, Statement* statement) {
+    statement->type = STATEMENT_INSERT;
+
+    char* keyword = strtok(input_buffer->buffer, " ");
+    char* id_string = strtok(NULL, " ");
+    char* username = strtok(NULL, " ");
+    char* email = strtok(NULL, " ");
+
+    if (id_string == NULL || username == NULL || email == NULL) {
+        return PREPARE_SYNTAX_ERROR;
+    }
+
+    int id = atoi(id_string);
+    if (strlen(username) > COLUMN_USERNAME_SIZE) {
+        return PREPARE_STRING_TOO_LONG;
+    }
+    if (strlen(email) > COLUMN_EMAIL_SIZE) {
+        return PREPARE_STRING_TOO_LONG;
+    }
+
+    statement->row_to_insert.id = id;
+    strcpy(statement->row_to_insert.username, username);
+    strcpy(statement->row_to_insert.email, email);
+
+    return PREPARE_SUCCESS;
+}
+
+PrepareResult prepare_statement(InputBuffer* input_buffer,
+                                Statement* statement) {
+    if (strncmp(input_buffer->buffer, "insert", 6) == 0) {
+        return prepare_insert(input_buffer, statement);
+    }
+    statement->type = STATEMENT_INSERT;
+    int args_assigned = sscanf(
+        input_buffer->buffer, "insert %d %s %s", &(statement->row_to_insert.id),
+        statement->row_to_insert.username, statement->row_to_insert.email);
+    if (args_assigned < 3) {
+        return PREPARE_SYNTAX_ERROR;
+    }
+    return PREPARE_SUCCESS;
+}

```

```
}
```

Calling `strtok` successively on the input buffer breaks it into substrings by inserting a null character whenever it reaches a delimiter (space, in our case). It returns a pointer to the start of the substring.

We can call `strlen()` on each text value to see if it's too long.

We can handle the error like we do any other error code:

```
enum PrepareResult_t {  
    PREPARE_SUCCESS,  
+   PREPARE_STRING_TOO_LONG,  
    PREPARE_SYNTAX_ERROR,  
    PREPARE_UNRECOGNIZED_STATEMENT  
};
```

```
switch (prepare_statement(input_buffer, &statement)) {  
    case (PREPARE_SUCCESS):  
        break;  
+   case (PREPARE_STRING_TOO_LONG):  
+       printf("String is too long.\n");  
+       continue;  
    case (PREPARE_SYNTAX_ERROR):  
        printf("Syntax error. Could not parse statement.\n");  
        continue;
```

Which makes our test pass

```
bundle exec rspec  
....  
  
Finished in 0.02284 seconds (files took 0.116 seconds to load)  
4 examples, 0 failures
```

While we're here, we might as well handle one more error case:

```
it 'prints an error message if id is negative' do
```

```

script = [
    "insert -1 cstack foo@bar.com",
    "select",
    ".exit",
]
result = run_script(script)
expect(result).to match_array([
    "db > ID must be positive.",
    "db > Executed.",
    "db > ",
])
end

```

```

enum PrepareResult_t {
    PREPARE_SUCCESS,
+   PREPARE_NEGATIVE_ID,
    PREPARE_STRING_TOO_LONG,
    PREPARE_SYNTAX_ERROR,
    PREPARE_UNRECOGNIZED_STATEMENT
@@ -148,9 +147,6 @@ PrepareResult prepare_insert(InputBuffer* in
    }

    int id = atoi(id_string);
+   if (id < 0) {
+       return PREPARE_NEGATIVE_ID;
+   }
    if (strlen(username) > COLUMN_USERNAME_SIZE) {
        return PREPARE_STRING_TOO_LONG;
    }
@@ -230,9 +226,6 @@ int main(int argc, char* argv[]) {
    switch (prepare_statement(input_buffer, &statement)) {
        case (PREPARE_SUCCESS):
            break;
+       case (PREPARE_NEGATIVE_ID):
+           printf("ID must be positive.\n");
+           continue;
        case (PREPARE_STRING_TOO_LONG):
            printf("String is too long.\n");
            continue;

```



Alright, that's enough testing for now. Next is a very important feature: persistence! We're going to save our database to a file and read it back out again.

It's gonna be great.

Here's the complete diff for this part:

```
@@ -22,6 +22,8 @@

enum PrepareResult_t {
    PREPARE_SUCCESS,
+   PREPARE_NEGATIVE_ID,
+   PREPARE_STRING_TOO_LONG,
    PREPARE_SYNTAX_ERROR,
    PREPARE_UNRECOGNIZED_STATEMENT
};
@@ -34,8 +36,8 @@
#define COLUMN_EMAIL_SIZE 255
typedef struct {
    uint32_t id;
-   char username[COLUMN_USERNAME_SIZE];
-   char email[COLUMN_EMAIL_SIZE];
+   char username[COLUMN_USERNAME_SIZE + 1];
+   char email[COLUMN_EMAIL_SIZE + 1];
} Row;

@@ -150,18 +152,40 @@ MetaCommandResult do_meta_command(InputBuffer*
    }
}

-PrepareResult prepare_statement(InputBuffer* input_buffer,
-                                Statement* statement) {
-   if (strncmp(input_buffer->buffer, "insert", 6) == 0) {
+PrepareResult prepare_insert(InputBuffer* input_buffer, Statement* statement) {
    statement->type = STATEMENT_INSERT;
-   int args_assigned = sscanf(
-       input_buffer->buffer, "insert %d %s %s", &(statement->row_to_insert.id),
-       statement->row_to_insert.username, statement->row_to_insert.email);
-   if (args_assigned < 3) {
```

```
+
+ char* keyword = strtok(input_buffer->buffer, " ");
+ char* id_string = strtok(NULL, " ");
+ char* username = strtok(NULL, " ");
+ char* email = strtok(NULL, " ");
+
+ if (id_string == NULL || username == NULL || email == NULL) {
+     return PREPARE_SYNTAX_ERROR;
+ }
+
+ int id = atoi(id_string);
+ if (id < 0) {
+     return PREPARE_NEGATIVE_ID;
+ }
+ if (strlen(username) > COLUMN_USERNAME_SIZE) {
+     return PREPARE_STRING_TOO_LONG;
+ }
+ if (strlen(email) > COLUMN_EMAIL_SIZE) {
+     return PREPARE_STRING_TOO_LONG;
+ }
+
+ statement->row_to_insert.id = id;
+ strcpy(statement->row_to_insert.username, username);
+ strcpy(statement->row_to_insert.email, email);
+
+ return PREPARE_SUCCESS;
+
+}
+PrepareResult prepare_statement(InputBuffer* input_buffer,
+                                Statement* statement) {
+    if (strncmp(input_buffer->buffer, "insert", 6) == 0) {
+        return prepare_insert(input_buffer, statement);
+    }
+    if (strcmp(input_buffer->buffer, "select") == 0) {
+        statement->type = STATEMENT_SELECT;
+
+@@ -223,6 +247,12 @@ int main(int argc, char* argv[]) {
+    switch (prepare_statement(input_buffer, &statement)) {
+        case (PREPARE_SUCCESS):
+            break;
+
+        case (PREPARE_NEGATIVE_ID):
```

```
+     printf("ID must be positive.\n");
+     continue;
+   case (PREPARE_STRING_TOO_LONG):
+     printf("String is too long.\n");
+     continue;
+   case (PREPARE_SYNTAX_ERROR):
+     printf("Syntax error. Could not parse statement.\n");
+     continue;
```

And we added tests:

```
+describe 'database' do
+  def run_script(commands)
+    raw_output = nil
+    IO.popen("./db", "r+") do |pipe|
+      commands.each do |command|
+        pipe.puts command
+      end
+
+      pipe.close_write
+
+      # Read entire output
+      raw_output = pipe.gets(nil)
+    end
+    raw_output.split("\n")
+  end
+
+  it 'inserts and retrieves a row' do
+    result = run_script([
+      "insert 1 user1 person1@example.com",
+      "select",
+      ".exit",
+    ])
+    expect(result).to match_array([
+      "db > Executed.",
+      "db > (1, user1, person1@example.com)",
+      "Executed.",
+      "db > ",
+    ])
+  end
end
```

```
+ end
+
+ it 'prints error message when table is full' do
+   script = (1..1401).map do |i|
+     "insert #{i} user#{i} person#{i}@example.com"
+   end
+   script << ".exit"
+   result = run_script(script)
+   expect(result[-2]).to eq('db > Error: Table full.')
+ end
+
+ it 'allows inserting strings that are the maximum length' do
+   long_username = "a"*32
+   long_email = "a"*255
+   script = [
+     "insert 1 #{long_username} #{long_email}",
+     "select",
+     ".exit",
+   ]
+   result = run_script(script)
+   expect(result).to match_array([
+     "db > Executed.",
+     "db > (1, #{long_username}, #{long_email})",
+     "Executed.",
+     "db > ",
+   ])
+ end
+
+ it 'prints error message if strings are too long' do
+   long_username = "a"*33
+   long_email = "a"*256
+   script = [
+     "insert 1 #{long_username} #{long_email}",
+     "select",
+     ".exit",
+   ]
+   result = run_script(script)
+   expect(result).to match_array([
+     "db > String is too long.",
+     "db > Executed.",
+   ])
```

```
+      "db > ",
+    ])
+  end
+
+  it 'prints an error message if id is negative' do
+    script = [
+      "insert -1 cstack foo@bar.com",
+      "select",
+      ".exit",
+    ]
+    result = run_script(script)
+    expect(result).to match_array([
+      "db > ID must be positive.",
+      "db > Executed.",
+      "db > ",
+    ])
+  end
+end
```

[< Part 3 - An In-Memory, Append-Only, Single-Table Database](#)

[Part 5 - Persistence to Disk >](#)

---

[rss](#) | [subscribe by email](#)

This project is maintained by [cstack](#)

Hosted on GitHub Pages — Theme by [orderedlist](#)