

7.2. Common Instructions

In this section, we discuss several common assembly instructions. Table 1 lists the most foundational instructions in x86 (and thus x64) assembly.

Table 1. Most Common Instructions

Instruction	Translation
<code>mov S, D</code>	$S \rightarrow D$ (copies value of <i>S</i> into <i>D</i>)
<code>add S, D</code>	$S + D \rightarrow D$ (adds <i>S</i> to <i>D</i> and stores result in <i>D</i>)
<code>sub S, D</code>	$D - S \rightarrow D$ (subtracts <i>S</i> from <i>D</i> and stores result in <i>D</i>)

Therefore, the sequence of instructions

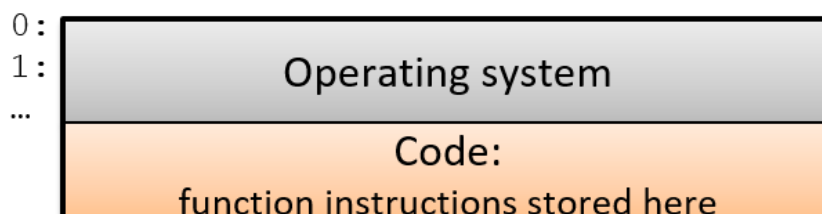
```
mov    -0x4(%rbp), %eax
add    $0x2, %eax
```

translates to:

- Copy the value at location `%rbp + -0x4` in *memory* ($M[\%rbp - 0x4]$) to register `%eax`.
- Add the value `0x2` to register `%eax`, and store the result in register `%eax`.

The three instructions shown in Table 1 also form the building blocks for instructions that maintain the organization of the program stack (i.e., the **call stack**). Recall that registers `%rbp` and `%rsp` refer to the *frame pointer* and *stack pointer*, respectively, and are reserved by the compiler for call stack management. Recall from our earlier discussion on [program memory](#) that the call stack typically stores local variables and parameters and helps the program track its own execution (see Figure 1). On x86-64 systems, the execution stack grows toward *lower* addresses. Like all stack data structures, operations occur at the "top" of the stack.

Parts of Program Memory



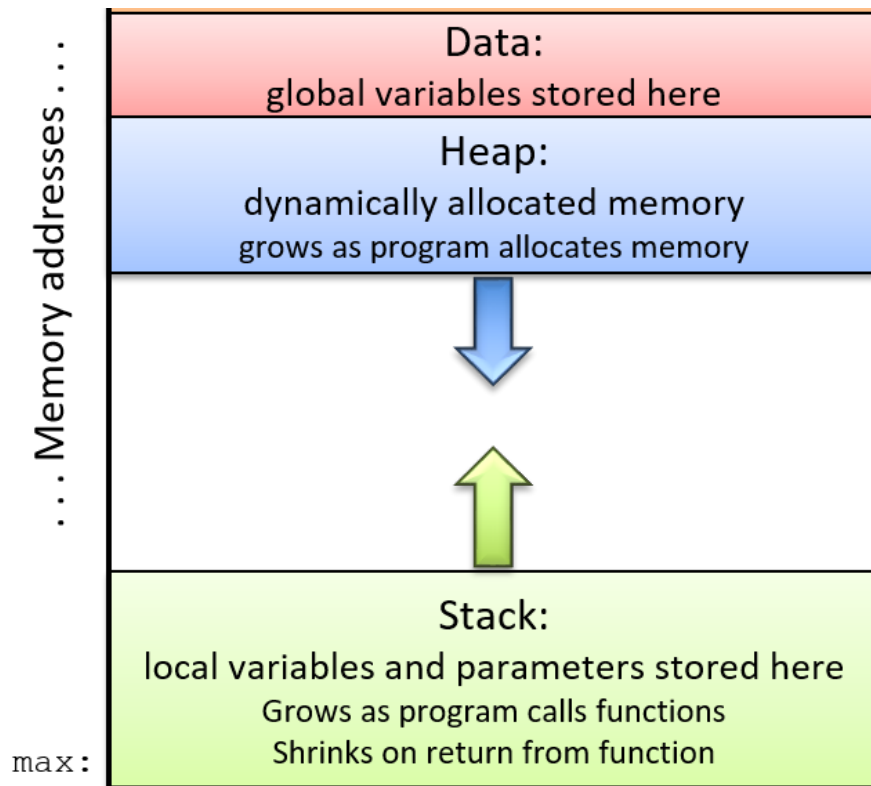


Figure 1. The parts of a program's address space

The x86-64 ISA provides two instructions (Table 2) to simplify call stack management.

Table 2. Stack Management Instructions

Instruction	Translation
push S	<p>Pushes a copy of S onto the top of the stack. Equivalent to:</p> <pre>sub \$0x8, %rsp mov S, (%rsp)</pre>
pop D	<p>Pops the top element off the stack and places it in location D. Equivalent to:</p> <pre>mov (%rsp), D add \$0x8, %rsp</pre>

Notice that while the three instructions in Table 1 require two operands, the push and pop instructions in Table 2 require only one operand apiece.

7.2.1. Putting It All Together: A More Concrete Example

Let's take a closer look at the `adder2` function:

```
//adds two to an integer and returns the result  
int adder2(int a) {  
    return a + 2;  
}
```

C

and its corresponding assembly code:

```
0000000000400526 <adder2>:  
400526:      55                push    %rbp  
400527:      48 89 e5          mov     %rsp,%rbp  
40052a:      89 7d fc          mov     %edi,-0x4(%rbp)  
40052d:      8b 45 fc          mov     -0x4(%rbp),%eax  
400530:      83 c0 02          add     $0x2,%eax  
400533:      5d                pop     %rbp  
400534:      c3                retq
```

The assembly code consists of a `push` instruction, followed by three `mov` instructions, an `add` instruction, a `pop` instruction, and finally a `retq` instruction. To understand how the CPU executes this set of instructions, we need to revisit the structure of [program memory](#). Recall that every time a program executes, the operating system allocates the new program's address space (also known as **virtual memory**). [Virtual memory](#) and the related concept of [processes](#) are covered in greater detail in Chapter 13; for now, it suffices to think of a process as the abstraction of a running program and virtual memory as the memory that is allocated to a single process. Every process has its own region of memory called the **call stack**. Keep in mind that the call stack is located in process/virtual memory, unlike registers (which are located on the CPU).

Figure 2 depicts a sample state of the call stack and registers prior to the execution of the `adder2` function.

```

0x526  push  %rbp
0x527  mov   %rsp, %rbp
0x52a  mov   %edi, -0x4(%rbp)
0x52d  mov   -0x4(%rbp), %eax
0x530  add   $0x2, %eax
0x533  pop   %rbp
0x534  retq

```

Registers	
%eax	0x123
%edi	0x28
%rsp	0xd28
%rbp	0xd40
%rip	0x526




Figure 2. Execution stack prior to execution

Notice that the stack grows toward *lower* addresses. Register `%eax` contains a junk value. The single parameter to the `adder2` function (`a`) is stored in register `%rdi` by convention. Since `a` is of type `int`, it is stored in component register `%edi`, which is shown in Figure 2. Likewise, since the `adder2` function returns an `int`, component register `%eax` is used for the return value instead of `%rax`.

The addresses associated with the instructions in the code segment of program memory (0x400526-0x400534) have been shortened to (0x526-0x534) to improve figure readability. Likewise, the addresses associated with the call stack segment of program memory have been shortened to 0xd28-0xd1c from 0x7fffffffdd28 - 0x7fffffffdd1c. In truth, call stack addresses occur at much higher addresses in program memory than code segment addresses.


Pay close attention to the initial values of registers `%rsp` and `%rbp`: they are 0xd28 and 0xd40, respectively. The red (upper-left) arrow in the following figures visually indicates the currently executing instruction. The `%rip` register (or instruction pointer) shows the next instruction to execute. Initially, `%rip` contains address 0x526, which corresponds to the first instruction in the `adder2` function.

```

➔ 0x526  push  %rbp
    0x527  mov   %rsp, %rbp
    0x52a  mov   %edi, -0x4(%rbp)
    0x52d  mov   -0x4(%rbp), %eax
    0x530  add   $0x2, %eax
    0x533  pop   %rbp
    0x534  retq

```

Registers	
%eax	0x123
%edi	0x28
%rsp	0xd20
%rbp	0xd40
%rip	0x527



The first instruction (`push %rbp`) places a copy of the value in `%rbp` (or `0xd40`) on top of the stack. After it executes, the `%rip` register advances to the address of the next instruction to execute (`0x527`). The `push` instruction decrements the stack pointer by 8 ("growing" the stack by 8 bytes), resulting in a new `%rsp` value of `0xd20`. Recall that the `push %rbp` instruction is equivalent to:

```

sub $8, %rsp
mov %rbp, (%rsp)

```


In other words, subtract 8 from the stack pointer and place a copy of the contents of `%rbp` in the location pointed to by the dereferenced stack pointer, `(%rsp)` .

```

0x526  push  %rbp
→ 0x527  mov   %rsp, %rbp
0x52a  mov   %edi, -0x4(%rbp)
0x52d  mov   -0x4(%rbp), %eax
0x530  add   $0x2, %eax
0x533  pop   %rbp
0x534  retq

```

Registers	
%eax	0x123
%edi	0x28
%rsp	0xd20
%rbp	0xd20
%rip	0x52a



call stack


Recall that the structure of the `mov` instruction is `mov S, D`, where `S` is the source location, and `D` is the destination. Thus, the next instruction (`mov %rsp, %rbp`) updates the value of `%rbp` to `0xd20`. The register `%rip` advances to the address of the next instruction to execute, or `0x52a`.

```

0x526  push  %rbp
0x527  mov   %rsp, %rbp
→ 0x52a  mov   %edi, -0x4(%rbp)
0x52d  mov   -0x4(%rbp), %eax
0x530  add   $0x2, %eax
0x533  pop   %rbp
0x534  retq

```

Registers	
%eax	0x123
%edi	0x28
%rsp	0xd20
%rbp	0xd20
%rip	0x52d



call stack


Next, `mov %edi, -0x4(%rbp)` is executed. This is a bit more complicated than the last `mov` instruction. Let's parse it piece by piece. First, recall that the first parameter to any function is stored in register `%rdi`. Since `a` is of type `int`, the compiler stores the first parameter in component register `%edi`. Next, the operand `-0x4(%rbp)` translates to `M[%rbp - 0x4]`. Since `%rbp` contains the value `0xd20`, subtracting 4 from it yields `0xd1c`. Therefore, the `mov` instruction copies the value of register `%edi` (or

0x28) to location 0xd1c on the stack. The instruction pointer advances to address 0x52d, the next address to be executed.

Note that storing the value 0x28 does not affect the stack pointer (%rsp). Therefore, as far as the program is concerned, the "top" of this stack is still address 0xd20.

```
0x526  push  %rbp
0x527  mov   %rsp, %rbp
0x52a  mov   %edi, -0x4(%rbp)
➔ 0x52d  mov   -0x4(%rbp), %eax
0x530  add   $0x2, %eax
0x533  pop   %rbp
0x534  retq
```

Registers	
%eax	0x28
%edi	0x28
%rsp	0xd20
%rbp	0xd20
%rip	0x530




call stack

The next mov instruction (mov -0x4(%rbp), %eax) copies the value at stack location 0xd1c (i.e., M[%rbp - 0x4] or 0x28) and stores it in register %eax . Register %rip advances to the next instruction to be executed, or 0x530.

```
0x526  push  %rbp
0x527  mov   %rsp, %rbp
0x52a  mov   %edi, -0x4(%rbp)
➔ 0x52d  mov   -0x4(%rbp), %eax
0x530  add   $0x2, %eax
0x533  pop   %rbp
0x534  retq
```

Registers	
%eax	0x2A
%edi	0x28
%rsp	0xd20
%rbp	0xd20
%rip	0x533



call stack


Next, `add $0x2, %eax` is executed. Recall that the `add` instruction has the form `add S, D` and places the quantity `S + D` in the destination `D`. So, `add $0x2, %eax` adds the constant value `0x2` to the value stored in `%eax` (or `0x28`), resulting in the value `0x2A` being stored in register `%eax`. Register `%rip` advances to point to the next instruction to be executed, or `0x533`.

```

0x526  push  %rbp
0x527  mov   %rsp, %rbp
0x52a  mov   %edi, -0x4(%rbp)
0x52d  mov   -0x4(%rbp), %eax
0x530  add   $0x2, %eax
➔ 0x533  pop   %rbp
0x534  retq

```

Registers	
%eax	0x2A
%edi	0x28
%rsp	0xd28
%rbp	0xd40
%rip	0x534



call stack

The next instruction that executes is `pop %rbp`. This instruction "pops" the value off the top of the call stack and places it in destination register `%rbp`. Recall that this instruction is equivalent to the following sequence of two instructions:

```

mov (%rsp), %rbp
add $8, %rsp

```

Recall that the top of the stack is `0xd20`, since that is the value stored in `%rsp`. Therefore, once this instruction executes, the value `(%rsp)` (i.e., `M[0xd20]`) is copied into register `%rbp`. Thus, `%rbp` now contains the value `0xd40`. The stack pointer *increments* by 8, since the stack grows toward lower addresses (and consequently *shrinks* toward higher ones). The new value of `%rsp` is `0xd28`, and `%rip` now points to the address of the last instruction to execute (i.e., `0x534`).

The last instruction executed is `retq`. We will talk more about what happens with `retq` in future sections when we discuss function calls, but for now it suffices to know that it prepares the call stack for returning from a function. By convention, the register `%rax` always contains the return value (if one exists). In this case, because `adder2` is of type `int`, the return value is stored in component register `%eax`, and the function returns the value `0x2A`, or 42.

Before we continue, note that the final values in registers `%rsp` and `%rbp` are `0xd28` and `0xd40`, respectively, which are the *same values as when the function started executing!* This is normal and expected behavior with the call stack. The purpose of the call stack is to store the temporary variables and data of each function as it executes in the context of a program. Once a function completes executing, the stack returns to the state it was in prior to the function call. As a result, it is common to see the following two instructions at the beginning of a function:

```
push %rbp
mov %rsp, %rbp
```

and the following two instructions at the end of a function:

```
pop %rbp
retq
```

Contents

7.2. Common Instructions

7.2.1. Putting It All Together: A More Concrete Example

Copyright (C) 2020 Dive into Systems, LLC.

Dive into Systems, is licensed under the Creative Commons [Attribution-NonCommercial-NoDerivatives 4.0 International](https://creativecommons.org/licenses/by-nc-nd/4.0/) (CC BY-NC-ND 4.0).