

## 7.9. structs in Assembly

A `struct` is another way to create a collection of data types in C. Unlike arrays, structs enable different data types to be grouped together. C stores a `struct` like a single-dimension array, where the data elements (fields) are stored contiguously.

Let's revisit `struct studentT` from Chapter 1:

```
struct studentT {  
    char name[64];  
    int age;  
    int grad_yr;  
    float gpa;  
};  
  
struct studentT student;
```

C

Figure 1 shows how `student` is laid out in memory. Each  $x_i$  denotes the address of a particular field.




Figure 1. The memory layout of a struct `studentT`

The fields are stored contiguously next to one another in memory in the order in which they are declared. In Figure 1, the `age` field is allocated at the memory location directly after the `name` field (at byte offset  $x_{64}$ ) and is followed by the `grad_yr` (byte offset  $x_{68}$ ) and `gpa` (byte offset  $x_{72}$ ) fields. This organization enables memory-efficient access to the fields.

To understand how the compiler generates assembly code to work with a `struct`, consider the function `initStudent`:

```
void initStudent(struct studentT *s, char *nm, int ag, int gr, float g) {  
    strncpy(s->name, nm, 64);  
    s->grad_yr = gr;  
    s->age = ag;  
    s->gpa = g;  
}
```

C

The `initStudent` function uses the base address of a `struct studentT` as its first parameter, and the desired values for each field as its remaining parameters. The following listing depicts this function in assembly:

Dump of assembler code for function `initStudent`:

Address	Disassembly	Comment
0x4006aa	<+0>: push %rbp	#save rbp
0x4006ab	<+1>: mov %rsp,%rbp	#update rbp (new stack frame)
0x4006ae	<+4>: sub \$0x20,%rsp	#add 32 bytes to stack frame
0x4006b2	<+8>: mov %rdi,-0x8(%rbp)	#copy 1st param to %rbp-0x8
(s)		
0x4006b6	<+12>: mov %rsi,-0x10(%rbp)	#copy 2nd param to %rbp-0x10
(nm)		
0x4006ba	<+16>: mov %edx,-0x14(%rbp)	#copy 3rd param to %rbp-0x14
(ag)		
0x4006bd	<+19>: mov %ecx,-0x18(%rbp)	#copy 4th param to %rbp-0x18
(gr)		
0x4006c0	<+22>: movss %xmm0,-0x1c(%rbp)	#copy 5th param to %rbp-0x1c
(g)		
0x4006c5	<+27>: mov -0x8(%rbp),%rax	#copy s to %rax
0x4006c9	<+31>: mov -0x10(%rbp),%rcx	#copy nm to %rcx
0x4006cd	<+35>: mov \$0x40,%edx	#copy 0x40 (or 64) to %edx
0x4006d2	<+40>: mov %rcx,%rsi	#copy nm to %rsi
0x4006d5	<+43>: mov %rax,%rdi	#copy s to %rdi
0x4006d8	<+46>: callq 0x400460 <strncpy@plt>	#call strncpy(s->name, nm, 64)
0x4006dd	<+51>: mov -0x8(%rbp),%rax	#copy s to %rax
0x4006e1	<+55>: mov -0x18(%rbp),%edx	#copy gr to %edx
0x4006e4	<+58>: mov %edx,0x44(%rax)	#copy gr to %rax+0x44 (s->grad_yr)
0x4006e7	<+61>: mov -0x8(%rbp),%rax	#copy s to %rax
0x4006eb	<+65>: mov -0x14(%rbp),%edx	#copy ag to %edx
0x4006ee	<+68>: mov %edx,0x40(%rax)	#copy ag to %rax+0x40 (s->age)
0x4006f1	<+71>: mov -0x8(%rbp),%rax	#copy s to %rax
0x4006f5	<+75>: movss -0x1c(%rbp),%xmm0	#copy g to %xmm0
0x4006fa	<+80>: movss %xmm0,0x48(%rax)	#copy g to %rax+0x48
0x400700	<+86>: leaveq	#prepare stack to exit
function		
0x400701	<+87>: retq	#return (void func, %rax ignored)

Being mindful of the byte offsets of each field is key to understanding this code. Here are a few things to keep in mind.

- The `strncpy` call takes the base address of the `name` field of `s`, the address of array `nm`, and a length specifier as its three arguments. Recall that because `name` is the first field in the `struct studentT`, the address of `s` is synonymous with the address of `s->name`.

```
0x4006b2 <+8>: mov    %rdi,-0x8(%rbp)      #copy 1st param to %rbp-0x8
(s)
0x4006b6 <+12>: mov    %rsi,-0x10(%rbp)     #copy 2nd param to %rpb-0x10
(nm)
0x4006ba <+16>: mov    %edx,-0x14(%rbp)            #copy 3rd param to %rbp-0x14
(ag)
0x4006bd <+19>: mov    %ecx,-0x18(%rbp)            #copy 4th param to %rbp-0x18
(gr)
0x4006c0 <+22>: movss  %xmm0,-0x1c(%rbp)           #copy 5th param to %rbp-0x1c
(g)
0x4006c5 <+27>: mov    -0x8(%rbp),%rax             #copy s to %rax
0x4006c9 <+31>: mov    -0x10(%rbp),%rcx            #copy nm to %rcx
0x4006cd <+35>: mov    $0x40,%edx                  #copy 0x40 (or 64) to %edx
0x4006d2 <+40>: mov    %rcx,%rsi                   #copy nm to %rsi
0x4006d5 <+43>: mov    %rax,%rdi                   #copy s to %rdi
0x4006d8 <+46>: callq 0x400460 <strncpy@plt> #call strncpy(s->name, nm,
64)
```

- This code snippet contains the previously undiscussed register (`%xmm0`) and instruction (`movss`). The `%xmm0` register is an example of a register reserved for floating-point values. The `movss` instruction indicates that the data being moved onto the call stack is of type single-precision floating point.
- The next part of the code (instructions `<initStudent+51>` thru `<initStudent+58>`) places the value of the `gr` parameter at an offset of `0x44` (or `68`) from the start of `s`. Revisiting the memory layout in Figure 1 shows that this address corresponds to `s->grad_yr`:

```
0x4006dd <+51>: mov    -0x8(%rbp),%rax             #copy s to %rax
0x4006e1 <+55>: mov    -0x18(%rbp),%edx            #copy gr to %edx
0x4006e4 <+58>: mov    %edx,0x44(%rax)             #copy gr to %rax+0x44 (s-
>grad_yr)
```

- The next section of code (instructions `<initStudent+61>` thru `<initStudent+68>`) copies the `ag` parameter to the `s->age` field of the `struct`, which is located at an offset of `0x40` (or `64`) bytes from the address of `s`:

```

0x4006e7 <+61>: mov    -0x8(%rbp),%rax    #copy s to %rax
0x4006eb <+65>: mov    -0x14(%rbp),%edx    #copy ag to %edx
0x4006ee <+68>: mov    %edx,0x40(%rax)    #copy ag to %rax+0x40 (s-
>age)

```

- Lastly, the `g` parameter value is copied to the `s->gpa` field (byte offset 72 or 0x48) of the `struct`. Notice the use of the `%xmm0` register since the data contained at location `%rbp-0x1c` is single-precision floating point:

```

0x4006f1 <+71>: mov    -0x8(%rbp),%rax    #copy s to %rax
0x4006f5 <+75>: movss -0x1c(%rbp),%xmm0    #copy g to %xmm0
0x4006fa <+80>: movss %xmm0,0x48(%rax)    #copy g to %rax+0x48

```

### 7.9.1. Data Alignment and structs

Consider the following modified declaration of `struct studentT`:

```


struct studentTM {
    char name[63]; //updated to 63 instead of 64
    int  age;
    int  grad_yr;
    float gpa;
};

struct studentTM student2;

```


C

The size of the `name` field is modified to be 63 bytes, instead of the original 64. Consider how this affects the way the `struct` is laid out in memory. It may be tempting to visualize it as in Figure 2.



*Figure 2. An incorrect memory layout for the updated struct studentTM. Note that the struct's "name" field is reduced from 64 to 63 bytes.*

In this depiction, the `age` field occurs in the byte immediately following the `name` field. But this is incorrect. Figure 3 depicts the actual layout in memory.



*Figure 3. The correct memory layout for the updated struct `studentTM`. Byte  $x_{63}$  is added by the compiler to satisfy memory alignment constraints, but it doesn't correspond to any of the fields.*

x64's alignment policy requires that two-byte data types (i.e., `short`) reside at a two-byte-aligned address, four-byte data types (i.e., `int`, `float`, and `unsigned`) reside at four-byte-aligned addresses, whereas larger data types (`long`, `double`, and pointer data) reside at eight-byte-aligned addresses. For a `struct`, the compiler adds empty bytes as **padding** between fields to ensure that each field satisfies its alignment requirements. For example, in the `struct` declared in Figure 3, the compiler adds a byte of padding at byte  $x_{63}$  to ensure that the `age` field starts at an address that is at a multiple of four. Values aligned properly in memory can be read or written in a single operation, enabling greater efficiency.

Consider what happens when a `struct` is defined as follows:

```
struct studentTM {  
    int age;  
    int grad_yr;  
    float gpa;  
    char name[63];  
};  
  
struct studentTM student3;
```

C

Moving the `name` array to the end ensures that `age`, `grad_yr`, and `gpa` are four-byte aligned. Most compilers will remove the filler byte at the end of the `struct`. However, if the `struct` is ever used in the context of an array (e.g., `struct studentTM courseSection[20];`) the compiler will once again add the filler byte as padding between each `struct` in the array to ensure that alignment requirements are properly met.

## Contents

### 7.9. structs in Assembly

#### 7.9.1. Data Alignment and structs

Copyright (C) 2020 Dive into Systems, LLC.

*Dive into Systems*, is licensed under the Creative Commons [Attribution-NonCommercial-NoDerivatives 4.0 International](https://creativecommons.org/licenses/by-nc-nd/4.0/) (CC BY-NC-ND 4.0).