

7.1. Diving into Assembly: Basics

For a first look at x64 assembly, we modify the `adder` function from [Chapter 6](#) to simplify its behavior. The modified function (`adder2`) is shown below:

```
#include <stdio.h>

//adds two to an integer and returns the result
int adder2(int a) {
    return a + 2;
}

int main(void){
    int x = 40;
    x = adder2(x);
    printf("x is: %d\n", x);
    return 0;
}
```

C

To compile this code, use the following command:

```
$ gcc -o adder adder.c
```

Next, let's view the corresponding assembly of this code by using the `objdump` command:

```
$ objdump -d adder > output
$ less output
```

Search for the code snippet associated with `adder2` by typing `/adder2` while examining the file `output` using `less`. The section associated with `adder2` should look similar to the following:

Assembly output for the `adder2` function

```
0000000000400526 <adder2>:
  400526:    55                push    %rbp
  400527:    48 89 e5          mov     %rsp,%rbp
  40052a:    89 7d fc          mov     %edi,-0x4(%rbp)
  40052d:    8b 45 fc          mov     -0x4(%rbp),%eax
  400530:    83 c0 02          add     $0x2,%eax
```

```
400533 :      5d                pop    %rbp
400534 :      c3                retq
```

Don't worry if you don't understand what's going on just yet. We will cover assembly in greater detail in later sections. For now, let's study the structure of these individual instructions.

Each line in the preceding example contains an instruction's 64-bit address in program memory, the bytes corresponding to the instruction, and the plaintext representation of the instruction itself. For example, 55 is the machine code representation of the instruction `push %rbp`, and the instruction occurs at address 0x400526 in program memory. Note that 0x400526 is an abbreviation of the full 64-bit address associated with the `push %rbp` instruction; the leading zeroes are ignored for readability.

It is important to note that a single line of C code often translates to multiple instructions in assembly. The operation `a + 2` is represented by the two instructions `mov -0x4(%rbp), %eax` and `add $0x2, %eax`.

WARNING

Your assembly may look different!

If you are compiling your code along with us, you may notice that some of your assembly examples look different from what is shown in this book. The precise assembly instructions that are output by any compiler depend on that compiler's version and the underlying operating system. Most of the assembly examples in this book were generated on systems running Ubuntu or Red Hat Enterprise Linux (RHEL).

In the examples that follow, we do not use any optimization flags. For example, we compile any example file (`example.c`) using the command `gcc -o example example.c`. Consequently, there are many seemingly redundant instructions in the examples that follow. Remember that the compiler is not "smart" — it simply follows a series of rules to translate human-readable code into machine language. During this translation process, it is not uncommon for some redundancy to occur. Optimizing compilers remove many of these redundancies during optimization, which is covered in a [later chapter](#).

7.1.1. Registers

Recall that a **register** is a word-sized storage unit located directly on the CPU. There may be separate registers for data, instructions, and addresses. For example, the Intel CPU has a total of 16 registers for storing 64-bit data:

`%rax`, `%rbx`, `%rcx`, `%rdx`, `%rdi`, `%rsi`, `%rsp`, `%rbp`, and `%r8 - %r15`. All the registers save for `%rsp` and `%rbp` hold general-purpose 64-bit data. While a program may interpret a register's contents

as, say, an integer or an address, the register itself makes no distinction. Programs can read from or write to all sixteen registers.

The registers `%rsp` and `%rbp` are known as the **stack pointer** and the **frame pointer** (or **base pointer**), respectively. The compiler reserves these registers for operations that maintain the layout of the program stack. For example, register `%rsp` always points to the top of the stack. In earlier x86 systems (e.g., IA32), the frame pointer commonly tracked the base of the active stack frame and helped to reference parameters. However, the base pointer is less frequently used in x86-64 systems. Compilers typically store the first six parameters in registers `%rdi`, `%rsi`, `%rdx`, `%rcx`, `%r8` and `%r9`, respectively. Register `%rax` stores the return value from a function.

The last register worth mentioning is `%rip` or the **instruction pointer**, sometimes called the **program counter** (PC). It points to the next instruction to be executed by the CPU. Unlike the 16 registers mentioned previously, programs cannot write directly to register `%rip`.

7.1.2. Advanced Register Notation

Since x86-64 is an extension of the 32-bit x86 architecture (which itself was an extension of an earlier 16-bit version), the ISA provides mechanisms to access the lower 32 bits, 16 bits, and lower bytes of each register. Table 1 lists each of the 16 registers and the ISA notations to access their component bytes.

Table 1. x86-64 Registers and Mechanisms for Accessing Lower Bytes

64-bit Register	32-bit Register	Lower 16 Bits	Lower 8 Bits
<code>%rax</code>	<code>%eax</code>	<code>%ax</code>	<code>%al</code>
<code>%rbx</code>	<code>%ebx</code>	<code>%bx</code>	<code>%bl</code>
<code>%rcx</code>	<code>%ecx</code>	<code>%cx</code>	<code>%cl</code>
<code>%rdx</code>	<code>%edx</code>	<code>%dx</code>	<code>%dl</code>
<code>%rdi</code>	<code>%edi</code>	<code>%di</code>	<code>%dil</code>
<code>%rsi</code>	<code>%esi</code>	<code>%si</code>	<code>%sil</code>
<code>%rsp</code>	<code>%esp</code>	<code>%sp</code>	<code>%spl</code>
<code>%rbp</code>	<code>%ebp</code>	<code>%bp</code>	<code>%bpl</code>
<code>%r8</code>	<code>%r8d</code>	<code>%r8w</code>	<code>%r8b</code>
<code>%r9</code>	<code>%r9d</code>	<code>%r9w</code>	<code>%r9b</code>

64-bit Register	32-bit Register	Lower 16 Bits	Lower 8 Bits
%r10	%r10d	%r10w	%r10b
%r11	%r11d	%r11w	%r11b
%r12	%r12d	%r12w	%r12b
%r13	%r13d	%r13w	%r13b
%r14	%r14d	%r14w	%r14b
%r15	%r15d	%r15w	%r15b

The first eight registers (`%rax` , `%rbx` , `%rcx` , `%rdx` , `%rdi` , `%rsi` , `%rsp` , and `%rbp`) are 64-bit extensions of 32-bit registers in x86 and have a common mechanism for accessing their lower 32 bits, lower 16 bits, and least-significant byte. To access the lower 32 bits of the first eight registers, simply replace the `r` in the register name with `e` . Thus, the register corresponding to the lower 32 bits of register `%rax` is register `%eax` . To access the lower 16 bits of each of these eight registers, reference the last two letters of the register's name. So, the mechanism to access the lower two bytes of register `%rax` is `%ax` .




Figure 1. The names that refer to subsets of register %rax.

The ISA provides a separate mechanism to access the eight-bit components within the lower 16 bits of the first four listed registers. Figure 1 depicts the access mechanisms for register `%rax` . The *higher* and *lower* bytes within the lower 16 bits of the first four listed registers can be accessed by taking the last two letters of the register name and replacing the last letter with either an `h` (for *higher*) or an `l` (for *lower*) depending on which byte is desired. For example, `%al` references the lower eight-bits of register `%ax` , whereas `%ah` references the higher eight-bits of register `%ax` . These eight-bit registers are commonly used for storing single-byte values for certain operations, such as bitwise shifts (a 32-bit register cannot be shifted more than 32 places, and the number 32 requires only a single byte of storage).

WARNING

Compiler may choose component registers depending on type

When reading assembly code, keep in mind that the compiler typically uses the 64-bit registers when dealing with 64-bit values (e.g., pointers or `long` types) and the 32-bit component registers when dealing with 32-bit types (e.g., `int`). In x86-64, it is very common to see 32-bit component registers intermixed with the full 64-bit registers. For example, in the `adder2` function shown in the previous example, the compiler references component register `%eax` instead of `%rax` since `int` types typically take up 32 bits (four bytes) of space on 64-bit systems. If the `adder2` function had a `long` parameter instead of a `int`, the compiler would store `a` in register `%rax` instead of register `%eax`.

The last eight registers (`%r8 - %r15`) were not part of the IA32 ISA. However, they also have mechanisms to access their different byte components. To access the lower 32 bits, 16 bits, or byte of the last eight registers, append the letter `d`, `w`, or `b`, respectively, to the end of the register's name. Thus, `%r9d` accesses the lower 32 bits of register `%r9`, whereas `%r9w` accesses the lower 16 bits, and `%r9b` accesses the lowest byte of register `%r9`.

7.1.3. Instruction Structure

Each instruction consists of an operation code (or **opcode**) that specifies what it does, and one or more **operands** that tell the instruction how to do it. For example, the instruction `add $0x2, %eax` has the opcode `add` and the operands `$0x2` and `%eax`.

Each operand corresponds to a source or destination location for a specific operation. Two operand instructions typically follow the source, destination (`S, D`) format, where the first operand specifies a source register, and the second operand specifies the destination.

There are multiple types of operands:

- **Constant (literal)** values are preceded by the `$` sign. For example, in the instruction `add $0x2, %eax`, `$0x2` is a literal value that corresponds to the hexadecimal value `0x2`.
- **Register** forms refer to individual registers. The instruction `mov %rsp, %rbp` specifies that the value in the source register (`%rsp`) should be copied to the destination location (register `%rbp`).
- **Memory** forms correspond to some value inside main memory (RAM) and are commonly used for address lookups. Memory address forms can contain a combination of registers and constant values. For example, in the instruction `mov -0x4(%rbp), %eax`, the operand `-0x4(%rbp)` is an example of a memory form. It loosely translates to "add `-0x4` to the value in register `%rbp` (i.e., subtract `0x4` from `%rbp`), and then perform a memory lookup." If this sounds like a pointer dereference, that's because it is!

7.1.4. An Example with Operands

The best way to explain operands in detail is to present a quick example. Suppose that memory contains the following values:

Address	Value
0x804	0xCA
0x808	0xFD
0x80c	0x12
0x810	0x1E

Let's also assume that the following registers contain the values shown:

Register	Value
%rax	0x804
%rbx	0x10
%rcx	0x4
%rdx	0x1

Then the operands in Table 2 evaluate to the values shown there. Each row of the table matches an operand with its form (e.g., constant, register, memory), how it is translated, and its value. Note that the notation $M[x]$ in this context denotes the value at the memory location specified by address x .

Table 2. Example operands

Operand	Form	Translation	Value
%rcx	Register	%rcx	0x4
(%rax)	Memory	$M[\%rax]$ or $M[0x804]$	0xCA
\$0x808	Constant	0x808	0x808
0x808	Memory	$M[0x808]$	0xFD
0x8(%rax)	Memory	$M[\%rax + 8]$ or $M[0x80c]$	0x12

Operand	Form	Translation	Value
(%rax, %rcx)	Memory	M[%rax + %rcx] or M[0x808]	0xFD
0x4(%rax, %rcx)	Memory	M[%rax + %rcx + 4] or M[0x80c]	0x12
0x800(,%rdx,4)	Memory	M[0x800 + %rdx*4] or M[0x804]	0xCA
(%rax, %rdx, 8)	Memory	M[%rax + %rdx*8] or M[0x80c]	0x12

In Table 2, the notation `%rcx` indicates the value stored in register `%rcx`. In contrast, `M[%rax]` indicates that the value inside `%rax` should be treated as an address, and to dereference (look up) the value at that address. Therefore, the operand `(%rax)` corresponds to `M[0x804]` which corresponds to the value `0xCA`.

A few important notes before continuing. Although Table 2 shows many valid operand forms, not all forms can be used interchangeably in all circumstances. Specifically:

- Constant forms cannot serve as destination operands.
- Memory forms cannot serve as *both* the source and destination operand in a single instruction.
- In cases of scaling operations (see the last two operands in Table 2), the scaling factor is a third parameter in the parentheses. Scaling factors can be one of 1, 2, 4, or 8.

Table 2 is provided as a reference; however, understanding key operand forms will help improve the reader's speed in parsing assembly language.

7.1.5. Instruction Suffixes

In several cases in upcoming examples, common and arithmetic instructions have a suffix that indicates the size (associated with the *type*) of the data being operated on at the code level. The compiler automatically translates code to instructions with the appropriate suffix. Table 3 shows the common suffixes for x86-64 instructions.

Table 3. Example Instruction Suffixes

Suffix	C Type	Size (bytes)
b	char	1

Suffix	C Type	Size (bytes)
w	short	2
l	int or unsigned	4
s	float	4
q	long , unsigned long , all pointers	8
d	double	8

Note that instructions involved with conditional execution have different suffixes based on the evaluated condition. We cover instructions associated with conditional execution in a [later section](#).

Contents

- 7.1. Diving into Assembly: Basics
 - 7.1.1. Registers
 - 7.1.2. Advanced Register Notation
 - 7.1.3. Instruction Structure
 - 7.1.4. An Example with Operands
 - 7.1.5. Instruction Suffixes

Copyright (C) 2020 Dive into Systems, LLC.

Dive into Systems, is licensed under the Creative Commons [Attribution-NonCommercial-NoDerivatives 4.0 International](#) (CC BY-NC-ND 4.0).