

C++ Lock-free Atomic Shared Pointer

2020.06.14 SF-Zhou

1. Shared Pointer

使用原子量引用计数实现一个简易的共享指针：

```
#include <atomic>
#include <cassert>
#include <iostream>
#include <memory>

template <class T>
class ReferenceCount {
public:
    ReferenceCount(std::unique_ptr<T> ptr) : ptr_(std::move(ptr)), cnt_(1) {}

    T *Ptr() const { return ptr_.get(); }

    ReferenceCount *Ref() {
        ++cnt_;
        return this;
    }

    void Deref() {
        if (--cnt_ == 0) {
```

```

        delete this;
    }
}

private:
    std::unique_ptr<T> ptr_;
    std::atomic_uint32_t cnt_;
};

template <class T>
class SharedPtr {
public:
    SharedPtr(std::unique_ptr<T> ptr = nullptr) noexcept
        : rc_(new ReferenceCount<T>(std::move(ptr))) {}

    ~SharedPtr() { rc_>Deref(); }

    T *Load() const { return rc_>Ptr(); }

    void Store(const SharedPtr &other) {
        auto old = rc_;
        rc_ = other.Copy();
        old->Deref();
    }

private:
    ReferenceCount<T> *Copy() const { return rc_>Ref(); }

private:
    ReferenceCount<T> *rc_;

```

```

};

class A {
public:
    A(int value) : value_(value) {
        std::cout << "A(" << value_ << ")" << std::endl;
    }
    ~A() { std::cout << "~A(" << value_ << ")" << std::endl; }

    int Value() const { return value_; }

private:
    int value_ = 0;
};

int main() {
    SharedPtr<A> a;
    assert(a.Load() == nullptr);

    SharedPtr<A> b(std::make_unique<A>(7));
    assert(b.Load() != nullptr);
    A &r = *b.Load();
    assert(r.Value() == 7);

    a.Store(b);
    assert(a.Load());
    assert(a.Load()->Value() == 7);

    b.Store(SharedPtr<A>(std::make_unique<A>(9)));
    assert(b.Load());

```

```
assert(b.Load()->Value() == 9);

a.Store(SharedPtr<A>());
assert(a.Load() == nullptr);
}
```

标准库中的 `std::shared_ptr` 的实现与之类似。仔细观察可以发现 `ReferenceCount` 是线程安全的，无论是 `Ref` 还是 `Unref`，使用原子量都可以保证计数准确，并且有且仅有一次析构。但 `SharedPtr` 中对 `rc_` 的操作并不是线程安全的，例如两个线程同时执行 `Store`，可能会对同一个 `rc_` 对象重复执行两次 `Deref`。所以只能支持单线程写或多线程读。

2. Atomic Shared Pointer

如果希望线程安全，最简单的方案自然是加锁。可以在 `SharedPtr` 的 `Load` / `Store` / `Copy` 函数中加自旋锁或互斥锁，标准库也是这样实现的，但显然锁的开销有点大。

仔细分析这里的 `Store` 的过程，一来需要将原先的计数 -1，二来需要从新计数中复制指针并 +1 计数，如果可以原子化的实现这一步骤，就可以实现无锁的共享指针。直觉地写出如下的代码：

```
template <class T>
class SharedPtr {
public:
    SharedPtr(std::unique_ptr<T> ptr = nullptr) noexcept {
        rc_.store(new ReferenceCount<T>(std::move(ptr)));
    }
}
```

```

~SharedPtr() { rc_.load()->Deref(); }

T *Load() const { return rc_.load()->Ptr(); }

void Store(const SharedPtr &other) {
    auto copy = other.rc_.load()->Ref();
    auto old = rc_.exchange(copy);
    old->Deref();
}

private:
    std::atomic<ReferenceCount<T> *> rc_;
};

```

然而随意地多线程跑下 Store，会发现这段代码是不靠谱的。仔细分析 Store 的过程，可以发现 other.rc_.load()->Ref() 并不是原子的。当一个线程完成 other.rc_.load() 后，可能另一个线程执行 old->Deref()，此时引用计数对象已经完成析构，也就没法再执行后面的 Ref 操作。这里的 Load 也是如此。换句话说，这里需要保证计数对象存活地情况下执行 Ref()。

仔细思考下，这里无法基于可能被析构的 rc_ 做引用计数的原子加。一个可行的方案是增加本地引用计数。除了 rc_ 指向的全局引用计数外，再增加一个本地引用计数变量。在 Load 时首先原子地增加本地引用计数，并在 Release 时删去。那么如何使得 rc_ 也能感知到本地引用计数的存在、不至于提前“自杀”呢？一个简单粗暴的方法是预支。首先在 rc_ 指向的全局引用计数上增加一个大计数，用来表示共享指针提前预支的引用计数量，

保证它不会因为外界的原因先析构掉。后面每次 Load 的时候，从预支的计数中取出，CAS 更新本地计数剩余量，最后 Release 时再减去剩下的预支计数量。

本地引用计数如果使用独立的变量存储，就需要使用 128 位的 CAS 操作了，但这个操作是很低效的。好在 x64 平台上，指针的高 16 位是全 0 的，刚好可以用来放本地引用计数，也就可以直接使用 64 位的 CAS 操作了。这也就是 folly 中的无锁共享指针的实现方法，简化实现如下：

```
#include <atomic>
#include <iostream>
#include <memory>
#include <thread>
#include <vector>

template <class T>
class ReferenceCount {
public:
    ReferenceCount(std::unique_ptr<T> ptr) : ptr_(std::move(ptr)), cnt_(1) {}

    T *Ptr() const { return ptr_.get(); }

    ReferenceCount *Ref(uint32_t cnt = 1) {
        cnt_.fetch_add(cnt);
        return this;
    }

    void Deref(uint32_t cnt = 1) {
        if (cnt_.fetch_sub(cnt) == cnt) {
```

```

        delete this;
    }
}

private:
    std::unique_ptr<T> ptr_;
    std::atomic_uint32_t cnt_;
};

template <class T>
class AtomicSharedPtr;

template <class T>
class SharedPtr {
public:
    SharedPtr(std::unique_ptr<T> ptr = nullptr) noexcept
        : rc_(new ReferenceCount<T>(std::move(ptr))) {}

    ~SharedPtr() { rc_>Deref(); }

    T *Load() const { return rc_>Ptr(); }

    void Store(const SharedPtr &other) {
        auto old = rc_;
        rc_ = other.Copy();
        old->Deref();
    }

private:
    SharedPtr(ReferenceCount<T> *rc) : rc_(rc) {}

```

```

friend class AtomicSharedPtr<T>;

ReferenceCount<T> *Copy() const { return rc_->Ref(); }

private:
    ReferenceCount<T> *rc_;
};

template <class T>
class AtomicSharedPtr {
public:
    ~AtomicSharedPtr() { Release(rc_.load()); }

    SharedPtr<T> Load() { return SharedPtr<T>(Acquire()); }

    void Store(SharedPtr<T> &ptr) {
        ptr.rc_->Ref(kCnt);
        auto old = rc_.exchange((uint64_t)ptr.rc_ | (kCnt << 48));
        Release(old);
    }

private:
    ReferenceCount<T> *Acquire() {
        uint64_t local = 0;
        do {
            local = rc_.load();
        } while (!rc_.compare_exchange_weak(local, local - (1ull << 48)));
        return reinterpret_cast<ReferenceCount<T>*>(local & (-1ull >> 16));
    }
}

```



```

static void Release(uint64_t local) {
    if (local == 0) {
        return;
    }
    uint32_t local_cnt = (local >> 48);
    reinterpret_cast<ReferenceCount<T> *>(local & (-1ull >> 16))
        ->Deref(local_cnt);
}

private:
    static constexpr uint64_t kCnt = 0x2000;
    std::atomic<uint64_t> rc_{0};
};

std::atomic<int32_t> cnt{0};
class A {
public:
    A(int value) : value_(value) { ++cnt; }

    ~A() { --cnt; }

    int Value() const { return value_; }

private:
    int value_ = 0;
};

int main() {
    constexpr uint32_t N = 1000000;
    constexpr uint32_t T = 4;

```

```

AtomicSharedPtr<A> x;
AtomicSharedPtr<A> y;
std::vector<std::thread> threads;
for (uint32_t t = 0; t < T; ++t) {
    threads.emplace_back([&] {
        for (uint32_t i = 0; i < N; ++i) {
            SharedPtr<A> a(std::make_unique<A>(t * N + i));
            x.Store(a);
            SharedPtr<A> b = x.Load();
            y.Store(b);
        }
    });
}

for (auto &thread : threads) {
    thread.join();
}

return cnt.load() == 1 ? 0 : -1;
}

```

点击[此处](#)查看线上运行结果。

如果本地引用计数不足了怎么办？继续预支一笔就好。设定一个阈值，小于阈值时就预支一笔，并 CAS 更新本地引用计数，使其始终保持足够的余额，可参考文献 3 中 folly 的实现。

以下代码截取自 folly，删除了 folly 本身的一些调用，可直接在 C++ 11 中使用：

```
#include <atomic>
#include <cassert>
#include <climits>
#include <memory>
#include <thread>
#include <type_traits>
#include <vector>

// copy from
// https://github.com/facebook/folly/blob/master/folly/concurrency/detail/AtomicSharedPtr
// https://github.com/facebook/folly/blob/master/folly/PackedSyncPtr.h
// https://github.com/facebook/folly/blob/master/folly/concurrency/AtomicSharedPtr.h

#if !__x86_64__
#error "PackedSyncPtr is x64 specific code."
#endif

namespace std {
namespace detail {

// This implementation is specific to libstdc++, now accepting
// diffs for other libraries.

// Specifically, this adds support for two things:
// 1) incrementing/decrementing the shared count by more than 1 at a time
// 2) Getting the thing the shared_ptr points to, which may be different from
//    the aliased pointer.

class shared_ptr_internals {
public:
```

```

template <typename T, typename... Args>
static std::shared_ptr<T> make_ptr(Args &&... args) {
    return std::make_shared<T>(std::forward<Args...>(args...));
}
typedef std::__shared_count<std::_S_atomic> shared_count;
typedef std::_Sp_counted_base<std::_S_atomic> counted_base;
template <typename T>
using CountedPtr = std::shared_ptr<T>;

template <typename T>
static counted_base *get_counted_base(const std::shared_ptr<T> &bar);

static void inc_shared_count(counted_base *base, long count);

template <typename T>
static void release_shared(counted_base *base, long count);

template <typename T>
static T *get_shared_ptr(counted_base *base);

template <typename T>
static T *release_ptr(std::shared_ptr<T> &p);

template <typename T>
static std::shared_ptr<T> get_shared_ptr_from_counted_base(counted_base *base,
                                                            bool inc = true);

private:
/* Accessors for private members using explicit template instantiation */
struct access_shared_ptr {

```

```
typedef shared_count std::__shared_ptr<const void, std::_S_atomic>::*type;
friend type fieldPtr(access_shared_ptr);
};

struct access_base {
    typedef counted_base *shared_count::*type;
    friend type fieldPtr(access_base);
};

struct access_use_count {
    typedef _Atomic_word counted_base::*type;
    friend type fieldPtr(access_use_count);
};

struct access_weak_count {
    typedef _Atomic_word counted_base::*type;
    friend type fieldPtr(access_weak_count);
};

struct access_counted_ptr_ptr {
    typedef const void
        *std::_Sp_counted_ptr<const void *, std::_S_atomic>::*type;
    friend type fieldPtr(access_counted_ptr_ptr);
};

struct access_shared_ptr_ptr {
    typedef const void *std::__shared_ptr<const void, std::_S_atomic>::*type;
    friend type fieldPtr(access_shared_ptr_ptr);
};
```

```

struct access_refcount {
    typedef shared_count std::__shared_ptr<const void, std::_S_atomic>::*type;
    friend type fieldPtr(access_refcount);
};

template <typename Tag, typename Tag::type M>
struct Rob {
    friend typename Tag::type fieldPtr(Tag) { return M; }
};
};

template struct shared_ptr_internals::Rob<
    shared_ptr_internals::access_shared_ptr,
    &std::__shared_ptr<const void, std::_S_atomic>::_M_refcount>;
template struct shared_ptr_internals::Rob<
    shared_ptr_internals::access_base,
    &shared_ptr_internals::shared_count::_M_pi>;
template struct shared_ptr_internals::Rob<
    shared_ptr_internals::access_use_count,
    &shared_ptr_internals::counted_base::_M_use_count>;
template struct shared_ptr_internals::Rob<
    shared_ptr_internals::access_weak_count,
    &shared_ptr_internals::counted_base::_M_weak_count>;
template struct shared_ptr_internals::Rob<
    shared_ptr_internals::access_counted_ptr_ptr,
    &std::_Sp_counted_ptr<const void *, std::_S_atomic>::_M_ptr>;
template struct shared_ptr_internals::Rob<
    shared_ptr_internals::access_shared_ptr_ptr,
    &std::__shared_ptr<const void, std::_S_atomic>::_M_ptr>;
template struct shared_ptr_internals::Rob<

```

```

        shared_ptr_internals::access_refcount,
        &std::__shared_ptr<const void, std::_S_atomic>::_M_refcount>;

template <typename T>
inline shared_ptr_internals::counted_base *
shared_ptr_internals::get_counted_base(const std::shared_ptr<T> &bar) {
    // reinterpret_pointer_cast<const void>
    // Not quite C++ legal, but explicit template instantiation access to
    // private members requires full type name (i.e. shared_ptr<const void>, not
    // shared_ptr<T>)
    const std::shared_ptr<const void> &ptr(
        reinterpret_cast<const std::shared_ptr<const void> &>(bar));
    return (ptr.*fieldPtr(access_shared_ptr{})).*fieldPtr(access_base{});
}

inline void shared_ptr_internals::inc_shared_count(counted_base *base,
                                                    long count) {
    // Check that we don't exceed the maximum number of atomic_shared_ptrs.
    // Consider setting EXTERNAL_COUNT lower if this CHECK is hit.
    assert(base->_M_get_use_count() + count < INT_MAX);
    __gnu_cxx::__atomic_add_dispatch(&(base->*fieldPtr(access_use_count{})),
                                    count);
}

template <typename T>
inline void shared_ptr_internals::release_shared(counted_base *base,
                                                  long count) {
    // If count == 1, this is equivalent to base->_M_release()
    auto &a = base->*fieldPtr(access_use_count{});
    if (__gnu_cxx::__exchange_and_add_dispatch(&a, -count) == count) {

```

```

    base->_M_dispose();

    auto &b = base->*fieldPtr(access_weak_count{});
    if (__gnu_cxx::__exchange_and_add_dispatch(&b, -1) == 1) {
        base->_M_destroy();
    }
}
}

template <typename T>
inline T *shared_ptr_internals::get_shared_ptr(counted_base *base) {
    // See if this was a make_shared allocation
    auto inplace = base->_M_get_deleter(typeid(std::_Sp_make_shared_tag));
    if (inplace) {
        return (T *)inplace;
    }
    // Could also be a _Sp_counted_deleter, but the layout is the same
    using derived_type = std::_Sp_counted_ptr<const void *, std::_S_atomic>;
    auto ptr = reinterpret_cast<derived_type *>(base);
    return (T *)(ptr->*fieldPtr(access_counted_ptr_ptr{}));
}

template <typename T>
inline T *shared_ptr_internals::release_ptr(std::shared_ptr<T> &p) {
    auto res = p.get();
    std::shared_ptr<const void> &ptr(
        reinterpret_cast<std::shared_ptr<const void> &>(p));
    ptr.*fieldPtr(access_shared_ptr_ptr{}) = nullptr;
    (ptr.*fieldPtr(access_refcount{})).*fieldPtr(access_base{}) = nullptr;
    return res;
}

```



```

}

template <typename T>
inline std::shared_ptr<T>
shared_ptr_internals::get_shared_ptr_from_counted_base(counted_base *base,
                                                         bool inc) {

    if (!base) {
        return nullptr;
    }
    std::shared_ptr<const void> newp;
    if (inc) {
        inc_shared_count(base, 1);
    }
    newp.*fieldPtr(access_shared_ptr_ptr{}) =
        get_shared_ptr<const void>(base); // _M_ptr
    (newp.*fieldPtr(access_refcount{})).*fieldPtr(access_base{}) = base;
    // reinterpret_pointer_cast<T>
    auto res = reinterpret_cast<std::shared_ptr<T> *>(&newp);
    return std::move(*res);
}

template <class T>
class PackedSyncPtr {
    // This just allows using this class even with T=void. Attempting
    // to use the operator* or operator[] on a PackedSyncPtr<void> will
    // still properly result in a compile error.
    typedef typename std::add_lvalue_reference<T>::type reference;

public:
    /*

```

```

* If you default construct one of these, you must call this init()
* function before using it.
*
* (We are avoiding a constructor to ensure gcc allows us to put
* this class in packed structures.)
*/
void init(T *initialPtr = nullptr, uint16_t initialExtra = 0) {
    auto intPtr = reinterpret_cast<uintptr_t>(initialPtr);
    assert(!(intPtr >> 48));
    data_ = intPtr;
    setExtra(initialExtra);
}

/*
* Sets a new pointer. You must hold the lock when calling this
* function, or else be able to guarantee no other threads could be
* using this PackedSyncPtr<>.
*/
void set(T *t) {
    auto intPtr = reinterpret_cast<uintptr_t>(t);
    auto shiftedExtra = uintptr_t(extra()) << 48;
    assert(!(intPtr >> 48));
    data_ = (intPtr | shiftedExtra);
}

/*
* Get the pointer.
*
* You can call any of these without holding the lock, with the
* normal types of behavior you'll get on x64 from reading a pointer

```

```

    * without locking.
    */
T *get() const { return reinterpret_cast<T *>(data_ & (-1ull >> 16)); }
T *operator->() const { return get(); }
reference operator*() const { return *get(); }
reference operator[](std::ptrdiff_t i) const { return get()[i]; }

/*
 * Access extra data stored in unused bytes of the pointer.
 *
 * It is ok to call this without holding the lock.
 */
uint16_t extra() const { return data_ >> 48; }

/*
 * Don't try to put anything into this that has the high bit set:
 * that's what we're using for the mutex.
 *
 * Don't call this without holding the lock.
 */
void setExtra(uint16_t extra) {
    assert(!(extra & 0x8000));
    auto ptr = data_ & (-1ull >> 16);
    data_ = ((uintptr_t(extra) << 48) | ptr);
}

private:
    uintptr_t data_;
};

```

```

static_assert(std::is_pod<PackedSyncPtr<void>>::value,
              "PackedSyncPtr must be kept a POD type.");
static_assert(sizeof(PackedSyncPtr<void>) == 8,
              "PackedSyncPtr should be only 8 bytes---something is "
              "messed up");

} // namespace detail

template <typename T, typename CountedDetail = detail::shared_ptr_internals>
class atomic_shared_ptr {
    using SharedPtr = typename CountedDetail::template CountedPtr<T>;
    using BasePtr = typename CountedDetail::counted_base;
    using PackedPtr = detail::PackedSyncPtr<BasePtr>;

public:
    atomic_shared_ptr() noexcept { init(); }
    explicit atomic_shared_ptr(SharedPtr foo) /* noexcept */
        : atomic_shared_ptr() {
        store(std::move(foo));
    }
    atomic_shared_ptr(const atomic_shared_ptr<T> &) = delete;

    ~atomic_shared_ptr() { store(SharedPtr(nullptr)); }
    void operator=(SharedPtr desired) /* noexcept */ {
        store(std::move(desired));
    }
    void operator=(const atomic_shared_ptr<T> &) = delete;

    bool is_lock_free() const noexcept {
        // lock free unless more than EXTERNAL_OFFSET threads are

```

```

    // contending and they all get unlucky and scheduled out during
    // load().
    //
    // TODO: Could use a lock-free external map to fix this
    // corner case.
    return true;
}

SharedPtr load(
    std::memory_order order = std::memory_order_seq_cst) const noexcept {
    auto local = takeOwnedBase(order);
    return get_shared_ptr(local, false);
}

/* implicit */ operator SharedPtr() const { return load(); }

void store(SharedPtr n, std::memory_order order =
    std::memory_order_seq_cst) /* noexcept */ {
    auto newptr = get_newptr(std::move(n));
    auto old = ptr_.exchange(newptr, order);
    release_external(old);
}

SharedPtr exchange(
    SharedPtr n,
    std::memory_order order = std::memory_order_seq_cst) /* noexcept */ {
    auto newptr = get_newptr(std::move(n));
    auto old = ptr_.exchange(newptr, order);

    SharedPtr old_ptr;

```

```

    if (old.get()) {
        old_ptr = get_shared_ptr(old);
        release_external(old);
    }

    return old_ptr;
}

bool compare_exchange_weak(
    SharedPtr &expected, const SharedPtr &n,
    std::memory_order mo = std::memory_order_seq_cst) noexcept {
    return compare_exchange_weak(expected, n, mo, mo);
}

bool compare_exchange_weak(SharedPtr &expected, const SharedPtr &n,
                           std::memory_order success,
                           std::memory_order failure) /* noexcept */ {
    auto newptr = get_newptr(n);
    PackedPtr oldptr, expectedptr;

    oldptr = takeOwnedBase(success);
    if (!owners_eq(oldptr, CountedDetail::get_counted_base(expected))) {
        expected = get_shared_ptr(oldptr, false);
        release_external(newptr);
        return false;
    }
    expectedptr = oldptr; // Need oldptr to release if failed
    if (ptr_.compare_exchange_weak(expectedptr, newptr, success, failure)) {
        if (oldptr.get()) {
            release_external(oldptr, -1);
        }
    }
}

```

```

    }
    return true;
} else {
    if (oldptr.get()) {
        expected = get_shared_ptr(oldptr, false);
    } else {
        expected = SharedPtr(nullptr);
    }
    release_external(newptr);
    return false;
}
}

bool compare_exchange_weak(
    SharedPtr &expected, SharedPtr &&desired,
    std::memory_order mo = std::memory_order_seq_cst) noexcept {
    return compare_exchange_weak(expected, desired, mo, mo);
}

bool compare_exchange_weak(SharedPtr &expected, SharedPtr &&desired,
    std::memory_order success,
    std::memory_order failure) /* noexcept */ {
    return compare_exchange_weak(expected, desired, success, failure);
}

bool compare_exchange_strong(
    SharedPtr &expected, const SharedPtr &n,
    std::memory_order mo = std::memory_order_seq_cst) noexcept {
    return compare_exchange_strong(expected, n, mo, mo);
}

bool compare_exchange_strong(SharedPtr &expected, const SharedPtr &n,
    std::memory_order success,
    std::memory_order failure) /* noexcept */ {

```

```

    auto local_expected = expected;
    do {
        if (compare_exchange_weak(expected, n, success, failure)) {
            return true;
        }
    } while (local_expected == expected);

    return false;
}

bool compare_exchange_strong(
    SharedPtr &expected, SharedPtr &&desired,
    std::memory_order mo = std::memory_order_seq_cst) noexcept {
    return compare_exchange_strong(expected, desired, mo, mo);
}

bool compare_exchange_strong(SharedPtr &expected, SharedPtr &&desired,
                             std::memory_order success,
                             std::memory_order failure) /* noexcept */ {
    return compare_exchange_strong(expected, desired, success, failure);
}

private:
    // Matches packed_sync_pointer. Must be > max number of local
    // counts. This is the max number of threads that can access this
    // atomic_shared_ptr at once before we start blocking.
    static constexpr unsigned EXTERNAL_OFFSET{0x2000};
    // Bit signifying aliased constructor
    static constexpr unsigned ALIASED_PTR{0x4000};

    mutable std::atomic<PackedPtr> ptr_;

```



```

void add_external(BasePtr *res, int64_t c = 0) const {
    assert(res);
    CountedDetail::inc_shared_count(res, EXTERNAL_OFFSET + c);
}

void release_external(PackedPtr &res, int64_t c = 0) const {
    if (!res.get()) {
        return;
    }
    int64_t count = get_local_count(res) + c;
    int64_t diff = EXTERNAL_OFFSET - count;
    assert(diff >= 0);
    CountedDetail::template release_shared<T>(res.get(), diff);
}

PackedPtr get_newptr(const SharedPtr &n) const {
    BasePtr *newval;
    unsigned count = 0;
    if (!n) {
        newval = nullptr;
    } else {
        newval = CountedDetail::get_counted_base(n);
        if (n.get() != CountedDetail::template get_shared_ptr<T>(newval)) {
            // This is an aliased sharedptr. Make an un-aliased one
            // by wrapping in *another* shared_ptr.
            auto data = CountedDetail::template make_ptr<SharedPtr>(n);
            newval = CountedDetail::get_counted_base(data);
            count = ALIASED_PTR;
            // (add external must happen before data goes out of scope)
            add_external(newval);
        } else {
            add_external(newval);
        }
    }
}

```

```
    }  
}
```

```
PackedPtr newptr;  
newptr.init(newval, count);
```

```
    return newptr;  
}
```

```
PackedPtr get_newptr(SharedPtr &&n) const {  
    BasePtr *newval;  
    unsigned count = 0;  
    if (!n) {  
        newval = nullptr;  
    } else {  
        newval = CountedDetail::get_counted_base(n);  
        if (n.get() != CountedDetail::template get_shared_ptr<T>(newval)) {  
            // This is an aliased sharedptr. Make an un-aliased one  
            // by wrapping in *another* shared_ptr.  
            auto data = CountedDetail::template make_ptr<SharedPtr>(std::move(n));  
            newval = CountedDetail::get_counted_base(data);  
            count = ALIASED_PTR;  
            CountedDetail::release_ptr(data);  
            add_external(newval, -1);  
        } else {  
            CountedDetail::release_ptr(n);  
            add_external(newval, -1);  
        }  
    }  
}
```

```
PackedPtr newptr;
```

```

    newptr.init(newval, count);

    return newptr;
}
void init() {
    PackedPtr data;
    data.init();
    ptr_.store(data);
}

unsigned int get_local_count(const PackedPtr &p) const {
    return p.extra() & ~ALIASED_PTR;
}

// Check pointer equality considering wrapped aliased pointers.
bool owners_eq(PackedPtr &p1, BasePtr *p2) {
    bool aliased1 = p1.extra() & ALIASED_PTR;
    if (aliased1) {
        auto p1a = CountedDetail::template get_shared_ptr_from_counted_base<T>(
            p1.get(), false);
        return CountedDetail::get_counted_base(p1a) == p2;
    }
    return p1.get() == p2;
}

SharedPtr get_shared_ptr(const PackedPtr &p, bool inc = true) const {
    bool aliased = p.extra() & ALIASED_PTR;

    auto res = CountedDetail::template get_shared_ptr_from_counted_base<T>(
        p.get(), inc);

```

```

    if (aliased) {
        auto aliasedp =
            CountedDetail::template get_shared_ptr_from_counted_base<SharedPtr>(
                p.get());
        res = *aliasedp;
    }
    return res;
}

```

```

/* Get a reference to the pointer, either from the local batch or
 * from the global count.
 *
 * return is the base ptr, and the previous local count, if it is
 * needed for compare_and_swap later.
 */

```

```

PackedPtr takeOwnedBase(std::memory_order order) const noexcept {
    PackedPtr local, newlocal;
    local = ptr_.load(std::memory_order_acquire);
    while (true) {
        if (!local.get()) {
            return local;
        }
        newlocal = local;
        if (get_local_count(newlocal) + 1 > EXTERNAL_OFFSET) {
            // spinlock in the rare case we have more than
            // EXTERNAL_OFFSET threads trying to access at once.
            std::this_thread::yield();
            // Force DeterministicSchedule to choose a different thread
            local = ptr_.load(std::memory_order_acquire);
        } else {

```

```

        newlocal.setExtra(newlocal.extra() + 1);
        assert(get_local_count(newlocal) > 0);
        if (ptr_.compare_exchange_weak(local, newlocal, order)) {
            break;
        }
    }
}

// Check if we need to push a batch from local -> global
auto batchcount = EXTERNAL_OFFSET / 2;
if (get_local_count(newlocal) > batchcount) {
    CountedDetail::inc_shared_count(newlocal.get(), batchcount);
    putOwnedBase(newlocal.get(), batchcount, order);
}

return newlocal;
}

void putOwnedBase(BasePtr *p, unsigned int count,
                  std::memory_order mo) const noexcept {
    PackedPtr local = ptr_.load(std::memory_order_acquire);
    while (true) {
        if (local.get() != p) {
            break;
        }
        auto newlocal = local;
        if (get_local_count(local) > count) {
            newlocal.setExtra(local.extra() - count);
        } else {
            // Otherwise it may be the same pointer, but someone else won

```

```

        // the compare_exchange below, local count was already made
        // global. We decrement the global count directly instead of
        // the local one.
        break;
    }
    if (ptr_.compare_exchange_weak(local, newlocal, mo)) {
        return;
    }
}

CountedDetail::template release_shared<T>(p, count);
}
};

} // namespace std

// example
std::atomic<int32_t> cnt{0};
class A {
public:
    A(int value) : value_(value) { ++cnt; }

    ~A() { --cnt; }

    int Value() const { return value_; }

private:
    int value_ = 0;
};

```

```
int main() {
    constexpr uint32_t N = 1000000;
    constexpr uint32_t T = 4;

    std::atomic_shared_ptr<A> x;
    std::atomic_shared_ptr<A> y;
    std::vector<std::thread> threads;
    for (uint32_t t = 0; t < T; ++t) {
        threads.emplace_back([&] {
            for (uint32_t i = 0; i < N; ++i) {
                auto a = std::make_shared<A>(t * N + i);
                x.store(a);
                auto b = x.load();
                y.store(b);
            }
        });
    }

    for (auto &thread : threads) {
        thread.join();
    }

    return cnt.load() == 1 ? 0 : -1;
}
```

References

1. `std::shared_ptr`, *C++ Reference*
2. `std::atomic(std::shared_ptr)`, *C++ Reference*

3. folly::atomic_shared_ptr, Facebook

3 comments – powered by giscus

Oldest

Newest



uestc-lfs Sep 11, 2020

```
ReferenceCount *Ref() {  
    ++cnt_; // 1  
    return this;  
}  
  
void Deref() {  
    if (--cnt_ == 0) { // 2  
        delete this; // 3  
    }  
}
```

这个不是线程安全的吧, 可能的执行循序有 2 -> 1 -> 3

↑ 1



0 replies



SF-Zhou Sep 11, 2020

Owner

@uestc-lfs

```
ReferenceCount *Ref() {
```



```
referencecount Ref() {  
    ++cnt_; // 1  
    return this;  
}  
  
void Deref() {  
    if (--cnt_ == 0) { // 2  
        delete this; // 3  
    }  
}
```

这个不是线程安全的吧, 可能的执行循序有 2 -> 1 -> 3

你说的这种情况确实会导致线程不安全, 但在 SharedPtr 单线程写或多线程读场景下并不会出现。

↑ 1



0 replies



uestc-lfs Sep 11, 2020

哦哦, 是 获取 Outlook for Android<<https://aka.ms/ghei36>>

...

↑ 1



0 replies

Write

Preview

Aa

Sign in to comment

Sign in to comment

 Styling with Markdown is supported

Sign in with GitHub

Except where otherwise noted, content on this site is licensed under a CC BY-SA 4.0 license.
Copyright©2017 SF-Zhou, All Rights Reserved. Powered by GitHub Pages and GitHub Actions.