


Lecture 6 More on the LLVM Compiler

Jonathan Burket

Special thanks to Deby Katz,
Luke Zarko, and Gabe Weisz for their slides

Carnegie Mellon

Visualizing the LLVM Compiler System



The LLVM Optimizer is a series of “passes”

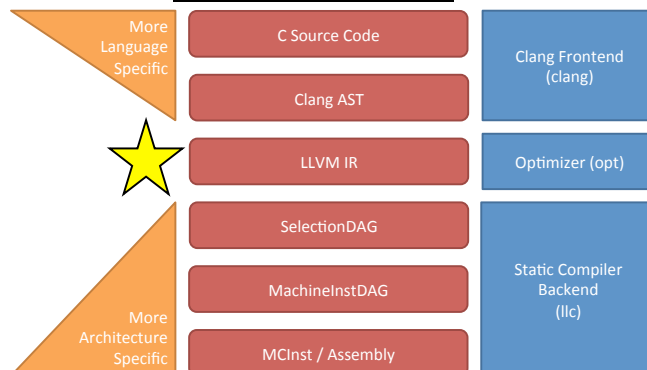
- Analysis and optimization passes, run one after another
- Analysis passes do not change code, optimization passes do

LLVM Intermediate Form is a **Virtual Instruction Set**

- Language- and target-independent form
- Used to perform the same passes for all source and target languages
- Internal Representation (IR) and external (persistent) representation

Carnegie Mellon

LLVM: From Source to Binary



Carnegie Mellon

LLVM IR

In-Memory Data Structure

Bitcode (.bc files)

```

42 43 C0 DE 21 0C 00 00
06 10 32 39 92 01 84 0C
0A 32 44 24 48 0A 90 21
18 00 00 00 98 00 00 00
E6 C6 21 1D E6 A1 1C DA
...
  
```

Text Format (.ll files)

```

define i32 @main() #0 {
entry:
    %retval = alloca i32, align 4
    %a = alloca i32, align 4
    ...
  
```

llvm-dis

llvm-asm

Bitcode files and LLVM IR text files are **lossless serialization formats!**
We can pause optimization and come back later.

Carnegie Mellon

Doing Things The LLVM Way - Strings

- LLVM does not use either `char *` or `std::string` to represent strings internally.
 - `char *`s use null terminators to represent strings -> bad for byte strings
 - `std::strings` are ok, but not recommended for performance reasons
- If it takes a `StringRef`, it can take a `std::string` or string literal as well:

```
getFunction("myfunc")           getFunction(std::string("myfunc"))

getFunction(StringRef("myfunc",6))    All Equivalent!
```

- Use `.str()` on a `StringRef` to get a `std::string`.
- Should avoid using C-style strings altogether

Carnegie Mellon

Doing Things The LLVM Way - Strings

- Forget `std::cout` and `std::cerr`! You want `outs()`, `errs()`, and `null()`.
 - Oddly, there's not equivalent of `std::endl` in LLVM

Printing the Name of a Function:

```
std::cout << F->getName().str() << std::endl;
outs() << F->getName() << "\n";
```

Printing an Instruction:

```
Instruction *I;
I->dump();           or           outs() << *I << "\n";
```

Printing an Entire Basic Block:

```
BB->dump()           or           outs() << *BB << "\n";
```

Carnegie Mellon

Doing Things the LLVM Way – Data Structures

- LLVM provides lots of data structures:
 - `BitVector`, `DenseMap`, `DenseSet`, `ImmutableList`, `ImmutableMap`, `ImmutableSet`, `IntervalMap`, `IndexedMap`, `MapVector`, `PriorityQueue`, `SetVector`, `ScopedHashTable`, `SmallBitVector`, `SmallPtrSet`, `SmallSet`, `SmallString`, `SmallVector`, `SparseBitVector`, `SparseSet`, `StringMap`, `StringRef`, `StringSet`, `Triple`, `TinyPtrVector`, `PackedVector`, `FoldingSet`, `UniqueVector`, `ValueMap`
- Provide better performance through *specialization*
- STL data structures work fine as well
- Only use these data structures in HW if you really want to

Carnegie Mellon

LLVM Instructions <--> Values

```
int main() {
  int x;
  int y = 2;
  int z = 3;
  x = y+z;
  y = x+z;
  z = x+y;
}

; Function Attrs: nounwind
define i32 @main() #0 {
entry:
  %add = add nsw i32 2, 3
  %add1 = add nsw i32 %add, 3
  %add2 = add nsw i32 %add, %add1
  ret i32 0
}
```

clang + mem2reg

Instruction I: `%add1 = add nsw i32 %add, 3`

You can't "get" `%add1` from Instruction I.
Instruction serves as the Value `%add1`.

Operand 1 Operand 2

Carnegie Mellon

LLVM Instructions <--> Values

```

int main() {
  int x = 1;
  int y = 2;
  int z = 3;
  x = y+z;
  y = x+z;
  z = x+y;
}

; Function Attrs: nounwind
define i32 @main() #0 {
entry:
  %add = add nsw i32 2, 3
  %add1 = add nsw i32 %add, 3
  %add2 = add nsw i32 %add, %add1
  ret i32 0
}

```

clang + mem2reg

Instruction I: %add1 = add nsw i32 %add, 3

outs() << *(I.getOperand(0)); → "%add = add nsw i32 2, 3"

outs() << *(I.getOperand(0)->getOperand(0)); → "2"

Only makes sense for an SSA Compiler [More on this Later]

Carnegie Mellon

Doing Things The LLVM Way – Casting and Type Introspection

Given a Value *v, what kind of Value is it?

isa<Argument>(v)

Is v an instance of the Argument class?


Argument *v = cast<Argument>(v)

I *know* v is an Argument, perform the cast.
Causes assertion failure if you are wrong.

Argument *v = dyn_cast<Argument>(v)

Cast v to an Argument if it is an argument,
otherwise return NULL. Combines both *isa*
and *cast* in one command.

dyn_cast is not to be confused with the C++
dynamic_cast operator!



Carnegie Mellon

Doing Things The LLVM Way – Casting and Type Introspection

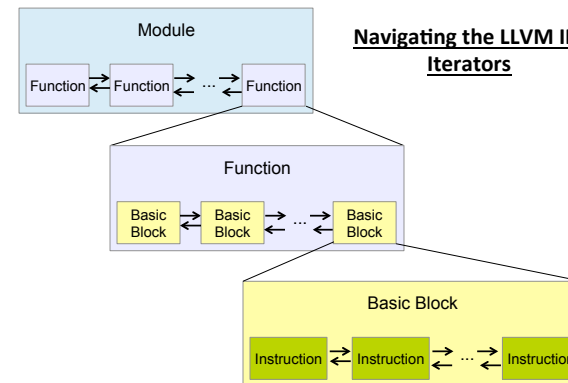
```

void analyzeInstruction(Instruction * I) {
  if (CallInst * CI = dyn_cast<CallInst>(I)) {
    outs() << "I'm a Call Instruction!\n";
  }
  if (UnaryInstruction * UI = dyn_cast<UnaryInstruction>(I)) {
    outs() << "I'm a Unary Instruction!\n";
  }
  if (CastInstruction * CI = dyn_cast<CastInstruction>(I)) {
    outs() << "I'm a Cast Instruction!\n";
  }
  ...
}

```

Carnegie Mellon

Navigating the LLVM IR: Iterators



Carnegie Mellon

Navigating the LLVM IR - Iterators

- **Module::iterator**
 - Modules are the large units of analysis
 - Iterates through the functions in the module
- **Function::iterator**
 - Iterates through a function's basic blocks
- **BasicBlock::iterator**
 - Iterates through the instructions in a basic block
- **Value::use_iterator**
 - Iterates through *uses* of a value
 - Recall that instructions are treated as values
- **User::op_iterator**
 - Iterates over the operands of an instruction (the "user" is the instruction)
 - Prefer to use convenient accessors defined by many instruction classes

Carnegie Mellon

Navigating the LLVM IR - Iterators

Iterate through every instruction in a function:

```
for (Function::iterator FI = func->begin(), FE = func->end(); FI != FE; ++FI) {
    for (BasicBlock::iterator BBI = FI->begin(), BBE = FI->end(); BBI != BBE; ++BBI) {
        outs() << "Instruction: " << *BBI << "\n";
    }
}
```

Using *InstIterator* (Provided by "llvm/IR/InstIterator.h"):

```
for (inst_iterator I = inst_begin(F), E = inst_end(F); I != E; ++I)
    outs() << *I << "\n";
```

Carnegie Mellon

Navigating the LLVM IR - Hints on Using Iterators

- Be very careful about modifying any part of the object iterated over during iteration
 - Can cause unexpected behavior
- Use ++i rather than i++ and pre-compute the end
 - Avoid problems with iterators doing unexpected things while you are iterating
 - Especially for fancier iterators
- There is overlap among iterators
 - E.g., InstIterator walks over all instructions in a function, overlapping range of Function::iterator and BasicBlock::iterator

Carnegie Mellon

Navigating the LLVM IR – More Iterators

Finding a Basic Block's predecessors/successors:

```
#include "llvm/Support/CFG.h"
BasicBlock *BB = ...;

for (pred_iterator PI = pred_begin(BB), E = pred_end(BB); PI != E; ++PI) {
    BasicBlock *Pred = *PI;
    // ...
}
```

Many useful iterators are defined outside of Function, BasicBlock, etc.

Carnegie Mellon

Navigating the LLVM IR – Casting and Iterators

```
for (Function::iterator FI = func->begin(), FE = func->end(); FI != FE; ++FI) {
    for (BasicBlock::iterator BBI = FI->begin(), BBE = FI->end(); BBI != BBE; ++BBI) {
        Instruction *I = BBI;

        if (CallInst *CI = dyn_cast<CallInst>(I)) {
            outs() << "I'm a Call Instruction!\n";
        }
        if (UnaryInstruction *UI = dyn_cast<UnaryInstruction>(I)) {
            outs() << "I'm a Unary Instruction!\n";
        }
        if (CastInstruction *CI = dyn_cast<CastInstruction>(I)) {
            outs() << "I'm a Cast Instruction!\n";
        }
        ...
    }
}
```

Very Common Code Pattern!

Carnegie Mellon

Navigating the LLVM IR – Visitor Pattern

```
struct MyVisitor : public InstVisitor<MyVisitor> {
    void visitCallInst(CallInst &CI) {
        outs() << "I'm a Call Instruction!\n";
    }
    void visitUnaryInstruction(UnaryInstruction &UI) {
        outs() << "I'm a Unary Instruction!\n";
    }
    void visitCastInst(CastInst &CI) {
        outs() << "I'm a Cast Instruction!\n";
    }
    void visitMul(BinaryOperator &I) {
        outs() << "I'm a multiplication Instruction!\n";
    }
}
MyVisitor MV;
MV.visit(F);
```

No need for iterators or casting!

A given instruction only triggers one method: a CastInst will not call visitUnaryInstruction if visitCastInst is defined.

You can case out on operators too, even if there isn't a specific class for them!

Carnegie Mellon

Writing Passes - Changing the LLVM IR

- Getting Rid of Instructions:
 - eraseFromParent()
 - Remove from basic block, drop all references, delete
 - removeFromParent()
 - Remove from basic block
 - Use if you will re-attach the instruction
 - Does not drop references (or clear the use list), so if you don't re-attach it Bad Things will happen
- moveBefore/insertBefore/insertAfter are also available
- replaceInstWithValue and replaceInstWithInst are also useful to have

Carnegie Mellon

Writing Passes – Adding New Instructions

```
define i32 @main() #0 {
entry:
    %add = add nsw i32 2, 2
    %add1 = add nsw i32 %add, 2
    %mul = mul nsw i32 %add, %add1
    %sub = sub nsw i32 %add1, %mul
    %add2 = add nsw i32 %mul, 5
    ret i32 %sub
}

define i32 @main() #0 {
entry:
    %add = add nsw i32 2, 2
    %add1 = add nsw i32 %add, 2
    %0 = add i32 %add, 0
    %mul = mul nsw i32 %add, %add1
    %sub = sub nsw i32 %add1, %mul
    %add2 = add nsw i32 %mul, 5
    ret i32 %sub
}
```

Let I = The "%mul = mul nsw i32 %add, %add1" instruction

```
Instruction *newInst = BinaryOperator::Create(Instruction::Add, I.getOperand(0),
ConstantInt::get(I.getOperand(0)->getType(), 0));
I.getParent()->getInstList().insert(&I, newInst)
```

Carnegie Mellon

Writing Passes - Correctness

- When you modify code, be careful not to change the meaning!
- You can break module invariants while in your pass, but you should repair them before you finish
 - Some module invariant examples:
 - Types of binary operator parameters are the same
 - Terminator instructions only at the end of BasicBlocks
 - Functions are called with correct argument types
 - Instructions belong to Basic blocks
 - Entry node has no predecessor
- opt automatically runs a pass (-verify) to check module invariants
 - But it doesn't check correctness in general!**
- Think about multi-threaded correctness

Carnegie Mellon


LLVM Pass Manager

- Compiler is organized as a series of "passes":
 - Each pass is one analysis or transformation
- Seven types of passes:
 - ImmutablePass: doesn't do much
 - LoopPass: process loops
 - RegionPass: process single-entry, single-exit portions of code
 - ModulePass: general interprocedural pass
 - CallGraphSCCPass: bottom-up on the call graph
 - FunctionPass: process a function at a time
 - BasicBlockPass: process a basic block at a time
- Constraints imposed (e.g. FunctionPass):
 - FunctionPass can only look at "current function"
 - Cannot maintain state across functions

Carnegie Mellon

LLVM Pass Manager

- Given a set of passes, the PassManager tries to *optimize the execution time* of the set of passes
 - Share information between passes
 - Pipeline execution of passes
- PassManager must understand how passes interact
 - Passes may require *information* from other passes
 - Passes may require *transformations* applied by other passes
 - Passes may *invalidate information or transformations* applied by other passes



Carnegie Mellon

LLVM Pass Manager

- The *getAnalysisUsage* function defines how a pass interacts with other passes
- Given *getAnalysisUsage(AnalysisUsage &AU)* for PassX:
 - AU.addRequired<PassY>() → PassY must be executed first
 - AU.addPreserves<PassY>() → PassY is still preserved by running PassX
 - AU.setPreservesAll() → PassX preserves all previous passes
 - AU.setPreservesCFG() → PassX might make some changes, but not to the CFG
 - If nothing is specified, it is assumed all previous passes are invalidated

```

DeadStoreElimination Pass:
void getAnalysisUsage(AnalysisUsage &AU) const override {
    AU.setPreservesCFG();
    AU.addRequired<DominatorTreeWrapperPass>();
    AU.addRequired<AliasAnalysis>();
    AU.addRequired<MemoryDependenceAnalysis>();
    AU.addPreserved<AliasAnalysis>();
    AU.addPreserved<DominatorTreeWrapperPass>();
    AU.addPreserved<MemoryDependenceAnalysis>();
}
  
```

Carnegie Mellon

opt tool: LLVM modular optimizer

Invoke arbitrary sequence of passes:

Completely control PassManager from command line

Supports loading passes as plugins from .so files

```
opt -load ./foo.so -pass1 -pass2 -pass3 x.bc -o y.bc
```

Passes “register” themselves:

When you write a pass, you must write the registration

```
RegisterPass<FunctionInfo> X("function-info",
    "15745: Function Information");
```

Carnegie Mellon

Using Passes

- For homework assignments, do not use passes provided by LLVM unless instructed to
 - We want you to implement the passes yourself to understand how they really work
- For projects, you can use whatever you want
 - Your own passes or LLVM's passes

Carnegie Mellon

Useful LLVM Passes: mem2reg

```
define i32 @main() #0 {
entry:
  %retval = alloca i32, align 4
  %a = alloca i32, align 4
  %b = alloca i32, align 4
  store i32 0, i32* %retval
  store i32 5, i32* %a, align 4
  store i32 3, i32* %b, align 4
  %0 = load i32* %a, align 4
  %1 = load i32* %b, align 4
  %sub = sub nsw i32 %0, %1
  ret i32 %sub
}
```



```
define i32 @main() #0 {
entry:
  %sub = sub nsw i32 5, 3
  ret i32 %sub
}
```

Not always possible:
Sometimes stack operations
are too complex

Carnegie Mellon

```
int main(int argc,
    char * argv[]) {
    int x;
    if (argc > 5) {
        x = 2;
    }
    else {
        x = 3;
    }
    return x;
}
```

clang + mem2reg

```
define i32 @main(i32 %argc,
    i8** %argv) #0 {
entry:
  %cmp = icmp sgt i32 %argc, 5
  br i1 %cmp, label %if.then,
    label %if.else

if.then:
  br label %if.end

if.else:
  br label %if.end

if.end:
  %x.0 = phi i32
    [ 2, %if.then ],
    [ 3, %if.else ]
  ret i32 %x.0
}
```

%x.0 = 2 if we go through %if.then
%x.0 = 3 if we go through %if.else

Carnegie Mellon

```
int main(int argc, char * argv[]) {
    int vals[4] = {2,4,8,16};
    int x = 0;
    vals[1] = 3;
    x += vals[0];
    x += vals[1];
    x += vals[2];
    return x;
}
```

Carnegie Mellon

```
@main.vals = private unnamed_addr constant [4 x i32]
    [i32 2, i32 4, i32 8, i32 16], align 4
define i32 @main(i32 %argc, i8** %argv) #0 {
entry:
    %vals = alloca [4 x i32], align 4
    %0 = bitcast [4 x i32]* %vals to i8*
    call void @llvm.memcpy.p0i8.p0i8.i32(i8* %0, i8* bitcast ([4 x i32]*
@main.vals to i8*), i32 16, i32 4, i1 false)
    %arrayidx = getelementptr inbounds [4 x i32]* %vals, i32 0, i32 1
    store i32 3, i32* %arrayidx, align 4
    %arrayidx1 = getelementptr inbounds [4 x i32]* %vals, i32 0, i32 0
    %i1 = load i32* %arrayidx1, align 4
    %add = add nsw i32 0, %i1
    %arrayidx2 = getelementptr inbounds [4 x i32]* %vals, i32 0, i32 1
    %i2 = load i32* %arrayidx2, align 4
    %add3 = add nsw i32 %add, %i2
    %arrayidx4 = getelementptr inbounds [4 x i32]* %vals, i32 0, i32 2
    %i3 = load i32* %arrayidx4, align 4
    %add5 = add nsw i32 %add3, %i3
    ret i32 %add5
}
```

mem2reg does not solve all of our pointer problems ☹

Carnegie Mellon

```
int main(int argc, char * arv[]) {
    int x = 0;
    for (int i=0; i<argc; i++) {
        x += i;
    }
    return x;
}
```

Carnegie Mellon

```
define i32 @main(i32 %argc, i8** %arv) #0 {
entry:
    br label %for.cond

for.cond:
    %x.0 = phi i32 [ 0, %entry ], [ %add, %for.inc ]
    %i.0 = phi i32 [ 0, %entry ], [ %inc, %for.inc ]
    %cmp = icmp slt i32 %i.0, %argc
    br i1 %cmp, label %for.body, label %for.end

for.body:
    %add = add nsw i32 %x.0, %i.0
    br label %for.inc

for.inc:
    %inc = add nsw i32 %i.0, 1
    br label %for.cond

for.end:
    ret i32 %x.0
}
```

Names are just strings, they should not be used to identify loops

There are no "Loop" Instructions: Loops are implemented with conditions and jumps

Carnegie Mellon


Useful LLVM Passes – Loop Information (-loops)

- llvm/Analysis/LoopInfo.h
- Reveals:
 - The basic blocks in a loop
 - Headers and pre-headers
 - Exit and exiting blocks
 - Back edges
 - “Canonical induction variable”
 - Loop Count

Carnegie Mellon

Useful LLVM Passes - Simplify CFG (-simplifycfg)

- Performs basic cleanup
 - Removes unnecessary basic blocks by merging unconditional branches if the second block has only one predecessor



- Removes basic blocks with no predecessors
- Eliminates phi nodes for basic blocks with a single predecessor
- Removes unreachable blocks

Carnegie Mellon

Useful LLVM Passes

- Scalar Evolution (-scalar-evolution)
 - Tracks changes to variables through nested loops
- Target Data (-targetdata)
 - Gives information about data layout on the target machine
 - Useful for generalizing target-specific optimizations
- Alias Analyses
 - Several different passes give information about aliases
 - If you know that different names refer to different locations, you have more freedom to reorder code, etc.

Carnegie Mellon

Other Useful Passes

- **Liveness-based dead code elimination**
 - Assumes code is dead unless proven otherwise
- **Sparse conditional constant propagation**
 - Aggressively search for constants
- **Correlated propagation**
 - Replace select instructions that select on a constant
- **Loop invariant code motion**
 - Move code out of loops where possible
- **Dead global elimination**
- **Canonicalize induction variables**
 - All loops start from 0
- **Canonicalize loops**
 - Put loop structures in standard form

Carnegie Mellon

Some Useful LLVM Documentation

LLVM Programmer's Manual

<http://llvm.org/docs/ProgrammersManual.html>

LLVM Language Reference Manual

<http://llvm.org/docs/LangRef.html>

Writing an LLVM Pass

<http://llvm.org/docs/WritingAnLLVMPass.html>

LLVM's Analysis and Transform Passes

<http://llvm.org/docs/Passes.html>

LLVM Internal Documentation

<http://llvm.org/docs/doxygen/html/>

Remember: We're using LLVM 3.5, but the documentation is always for the most up to date code (i.e. for the upcoming LLVM 3.6)

Carnegie Mellon