

15-213

"The course that gives CMU its Zip!"

Memory Management I: Dynamic Storage Allocation March 2, 2000

Topics

- Explicit memory allocation
- Data structures
- Mechanisms

class14.ppt

class14.ppt

- 2 -

CS 213 S'00

Harsh Reality #3 *Memory Matters*

Memory is not unbounded

- It must be allocated and managed
- Many applications are memory dominated
 - Especially those based on complex, graph algorithms

Memory referencing bugs especially pernicious

- Effects are distant in both time and space

Memory performance is not uniform


- Cache and virtual memory effects can greatly affect program performance
- Adapting program to characteristics of memory system can lead to major speed improvements

class14.ppt

- 2 -

CS 213 S'00

Dynamic Storage Allocation



Explicit vs. Implicit Storage Allocator

- Explicit: application allocates and frees space
 - E.g., malloc and free in C
- Implicit: application allocates, but does not free space
 - E.g. garbage collection in Java, ML or Lisp

Allocation

- In both cases the storage allocator provides an abstraction of memory as a set of blocks
- Does out free memory blocks to application

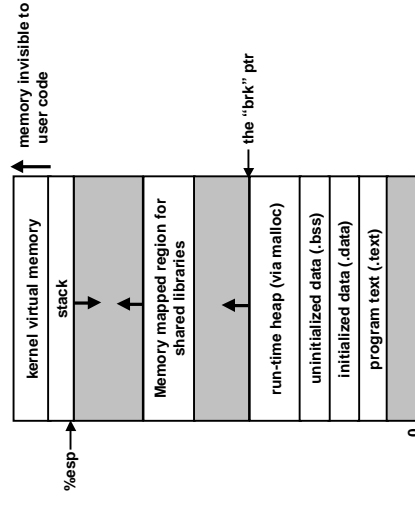
Will discuss explicit storage allocation today

class14.ppt

- 3 -

CS 213 S'00

Process memory image



class14.ppt

- 4 -


CS 213 S'00

Malloc package

- ```
void *malloc(int size)
```
- if successful:
    - returns a pointer to a memory block of at least **size** bytes
    - if **size==0**, returns NULL
  - if unsuccessful: returns NULL
- ```
void free(void *p)
```
- returns the block pointed at by **p** to pool of available memory
 - **p** must come from a previous call to **malloc()**.

Assumptions made in this lecture

- memory is word addressed (each word can hold a pointer)



class14.ppt

– 5 –

CS 213 S'00

Constraints

Applications:

- Can issue arbitrary sequence of allocation and free requests
- Free requests must correspond to an allocated block

Allocators


- Can't control number or size of allocated blocks
- Must respond immediately to all allocation requests
 - i.e., can't reorder or buffer requests
- Must allocate blocks from free memory
- Must align blocks so they satisfy all alignment requirements
 - usually 8 byte alignment
- Can only manipulate and modify free memory
- Can't move the allocated blocks once they are allocated
 - i.e., compaction is not allowed

class14.ppt

– 7 –

CS 213 S'00

Allocation example



class14.ppt

– 6 –

CS 213 S'00

Goals of good malloc/free

Primary goals

- Good time performance for **malloc** and **free**
 - Ideally should take constant time (not always possible)
 - Should certainly not take linear time in the number of blocks
- Good space usage
 - User allocated structures should be large fraction of operating-system allocated pages
 - Need to avoid fragmentation

Some other goals

- Good locality properties
 - structures allocated close in time should be close in space
 - “similar” objects should be allocated close in space
- Robust
 - can check that **free(p1)** is on a valid allocated object **p1**
 - can check that memory references are to allocated space


class14.ppt

– 8 –

CS 213 S'00

Fragmentation

Tendency for free blocks to become smaller over time leading to wasted space



No general solution assuming we cannot move blocks
We will consider several heuristics

class14.ppt


- 9 -

CS 213 S'00

Knowing how much to free

Standard method

- keep the length of a structure in the word preceeding the structure
 - This word is often called the *header field*
- requires an extra word for every allocated structure




class14.ppt

- 11 -

CS 213 S'00

Implementation issues

- How do we know how much memory to free just given a pointer?
- How do we keep track of the free blocks?
- What do we do with the extra space when allocating a structure that is smaller than the free block it is placed in?
- How do we pick a block to use for allocation -- many might fit?
- How do we reinsert freed block?



class14.ppt

- 10 -

CS 213 S'00

Keeping track of free blocks

- Method 1:** implicit list using lengths -- links all blocks



- Method 2:** explicit list among the free blocks using pointers within the free blocks



- Method 3:** segregated free lists

- Different free lists for different size classes

- Method 4:** blocks sorted by size

- Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

class14.ppt


- 12 -

CS 213 S'00

Method 1: implicit list

Need to identify whether each block is free or allocated

- Can use extra bit
- Bit can be put in the same word as the size if block sizes are always multiples of two (mask out low order bit when reading size).



class14.ppt


- 13 -

CS 213 S'00


Implicit list: allocating in a free block

Allocating in a free block - *splitting*

- Since allocated space might be smaller than free space, we need to split the block



```
void addblock(ptr p, int l) {
    int newsize = ((l + 1) >> 1) << 1; // add 1 and round up
    int oldsize = *p & -2;             // mask out low bit
    *p = newsize | 1;                 // set new length
    if (newsize < oldsize)
        *(p+newsize) = oldsize - newsize; // set length in remaining
                                           // part of block
    addblock(p, 2);
}
```



class14.ppt

- 15 -

CS 213 S'00

Implicit list: finding a free block

First fit:

- Search list from beginning, choose first free block that fits
- ```
p = start;
while ((p < end) ||
 (*p & 1) ||
 (*p <= len)); // not passed end
 // already allocated
 // too small
```
- Can take linear time in total number of blocks (allocated and free)
  - In practice it can cause "splinters" at beginning of list

### Next fit:

- Like first-fit, but search list from location of end of previous search
- Does a better job of spreading out the free blocks

### Best fit:

- Search the list, choose the free block with the closest size that fits
- Keeps fragments small --- usually helps fragmentation
- Will typically run slower than first-fit

class14.ppt


- 14 -

CS 213 S'00

## Implicit list: freeing a block

### Simplest implementation:

- Only need to clear allocated flag
- ```
void free_block(ptr p) { *p = *p & -2; }
```
- But can lead to "false fragmentation"



There is enough free space, but the allocator won't be able to find it

class14.ppt

- 16 -

CS 213 S'00


Implicit list: coalescing

Join with next and/or previous block if they are free

- Coalescing with next block

```
void free_block(ptr p) {
    *p = *p - 2;           // clear allocated flag
    next = p + *p;         // find next block
    if ((*next & 1) == 0)
        *p = *p + *next;   // add to this block if
                            // not allocated
}
```

Constant time coalescing (case 2)




class14.ppt

- 21 -

CS 213 S'00

Constant time coalescing (case 3)




class14.ppt

- 22 -

CS 213 S'00

Constant time coalescing (case 4)



class14.ppt

- 23 -

CS 213 S'00

Implicit lists: Summary

- **Implementation:** very simple
 - **Allocate:** linear time worst case
 - **Free:** constant time worst case -- even with coalescing
 - **Memory usage:** will depend on placement policy
 - First fit, next fit or best fit
- Not used in practice for malloc/free because of linear time allocate. Used in many special purpose applications.

class14.ppt

- 24 -

CS 213 S'00

Keeping track of free blocks

- **Method 1:** implicit list using lengths -- links all blocks



- **Method 2:** explicit list among the free blocks using pointers within the free blocks



- **Method 3:** segregated free lists
 - Different free lists for different size classes
- **Method 4:** blocks sorted by size
 - Can use a balanced tree (e.g. Red-Black tree) with pointers within each free block, and the length used as a key

class14.ppt

- 25 -

CS 213 S'00

Linked list of free blocks

Allocation

- Splice block out of the free list
- Split the block
- If remaining space, put space back onto the free list

Free


- Determine if coalescing with neighboring block
 - If not coalescing, add block to free list
 - If coalescing with next block, need to splice next block out of the free list, and add self into it
 - If coalescing with previous block, only need to modify lengths of previous block
 - If coalescing with both previous and next, then need to splice the next block out of the free list (but not add self)

class14.ppt

- 27 -


CS 213 S'00

Linked list of free blocks



Use data space for link pointers

- Typically doubly linked
- Still need header and footer for coalescing



- It is important to realize that links are not necessarily in the same order as the blocks

class14.ppt

- 26 -

CS 213 S'00

Linked list of free blocks

Comparison to implicit list:

- Allocate is linear time in number of free blocks instead of total blocks -- much faster allocates when most of the memory is full
- Slightly more complicated allocate and free since needs to splice blocks in and out of the list
- Some extra space for the links (4 words needed for each block)

Main use of linked lists is in conjunction with segregated free lists

- Keep multiple linked lists of different size classes, or possibly for different types of objects

class14.ppt

- 28 -

CS 213 S'00

For more information

D. Knuth, "The Art of Computer Programming, Second Edition", Addison Wesley, 1973

- the classic reference on dynamic storage allocation

Wilson et al, "Dynamic Storage Allocation: A Survey and Critical Review", Proc. 1995 Int'l Workshop on Memory Management, Kinross, Scotland, Sept, 1995.

- comprehensive survey
- `$classdir/doc/dsa.ps`

class14.ppt

— 29 —

CS 213 S00