

How Computational Graphs are Constructed in PyTorch



by Preferred Networks

In the previous [post](#) we went over the theoretical foundations of automatic differentiation and reviewed the implementation in PyTorch. In this post, we will be showing the parts of PyTorch involved in creating the graph and executing it. In order to understand the following contents, please read @ezyang's wonderful [blog post](#) about PyTorch internals.

Autograd components

First of all, let's look at where the different components of autograd live:

[tools/autograd](#): Here we can find the definition of the derivatives as we saw in the previous post [derivatives.yaml](#), several python scripts and a folder called [templates](#). These scripts and the templates are used at building time to generate the C++ code for the derivatives as specified in the yaml file. Also, the scripts here generate wrappers for the regular ATen functions so that the computational graph can be constructed.

[torch/autograd](#): This folder is where the autograd components that can be used directly from python are located. In [function.py](#) we find the actual definition of `torch.autograd.Function`, a class used by users to write their own differentiable functions in python as per the documentation. [functional.py](#) holds components for functionally computing the jacobian vector product, hessian, and other gradient related computations of a given function. The rest of the files have additional components such as gradient checkers, anomaly detection, and the autograd profiler.

[torch/csrc/autograd](#): This is where the graph creation and execution-related code lives. All this code is written in C++, since it is a critical part that is required to be extremely performant. Here we have several files that implement the engine, metadata storage, and all the needed components. Alongside this, we have several files whose names start with `_python_`, and their main responsibility is to allow python objects to be used in the autograd engine.

Graph Creation

[Previously](#), we described the creation of a computational graph. Now, we will see how PyTorch creates these graphs with references to the actual codebase.

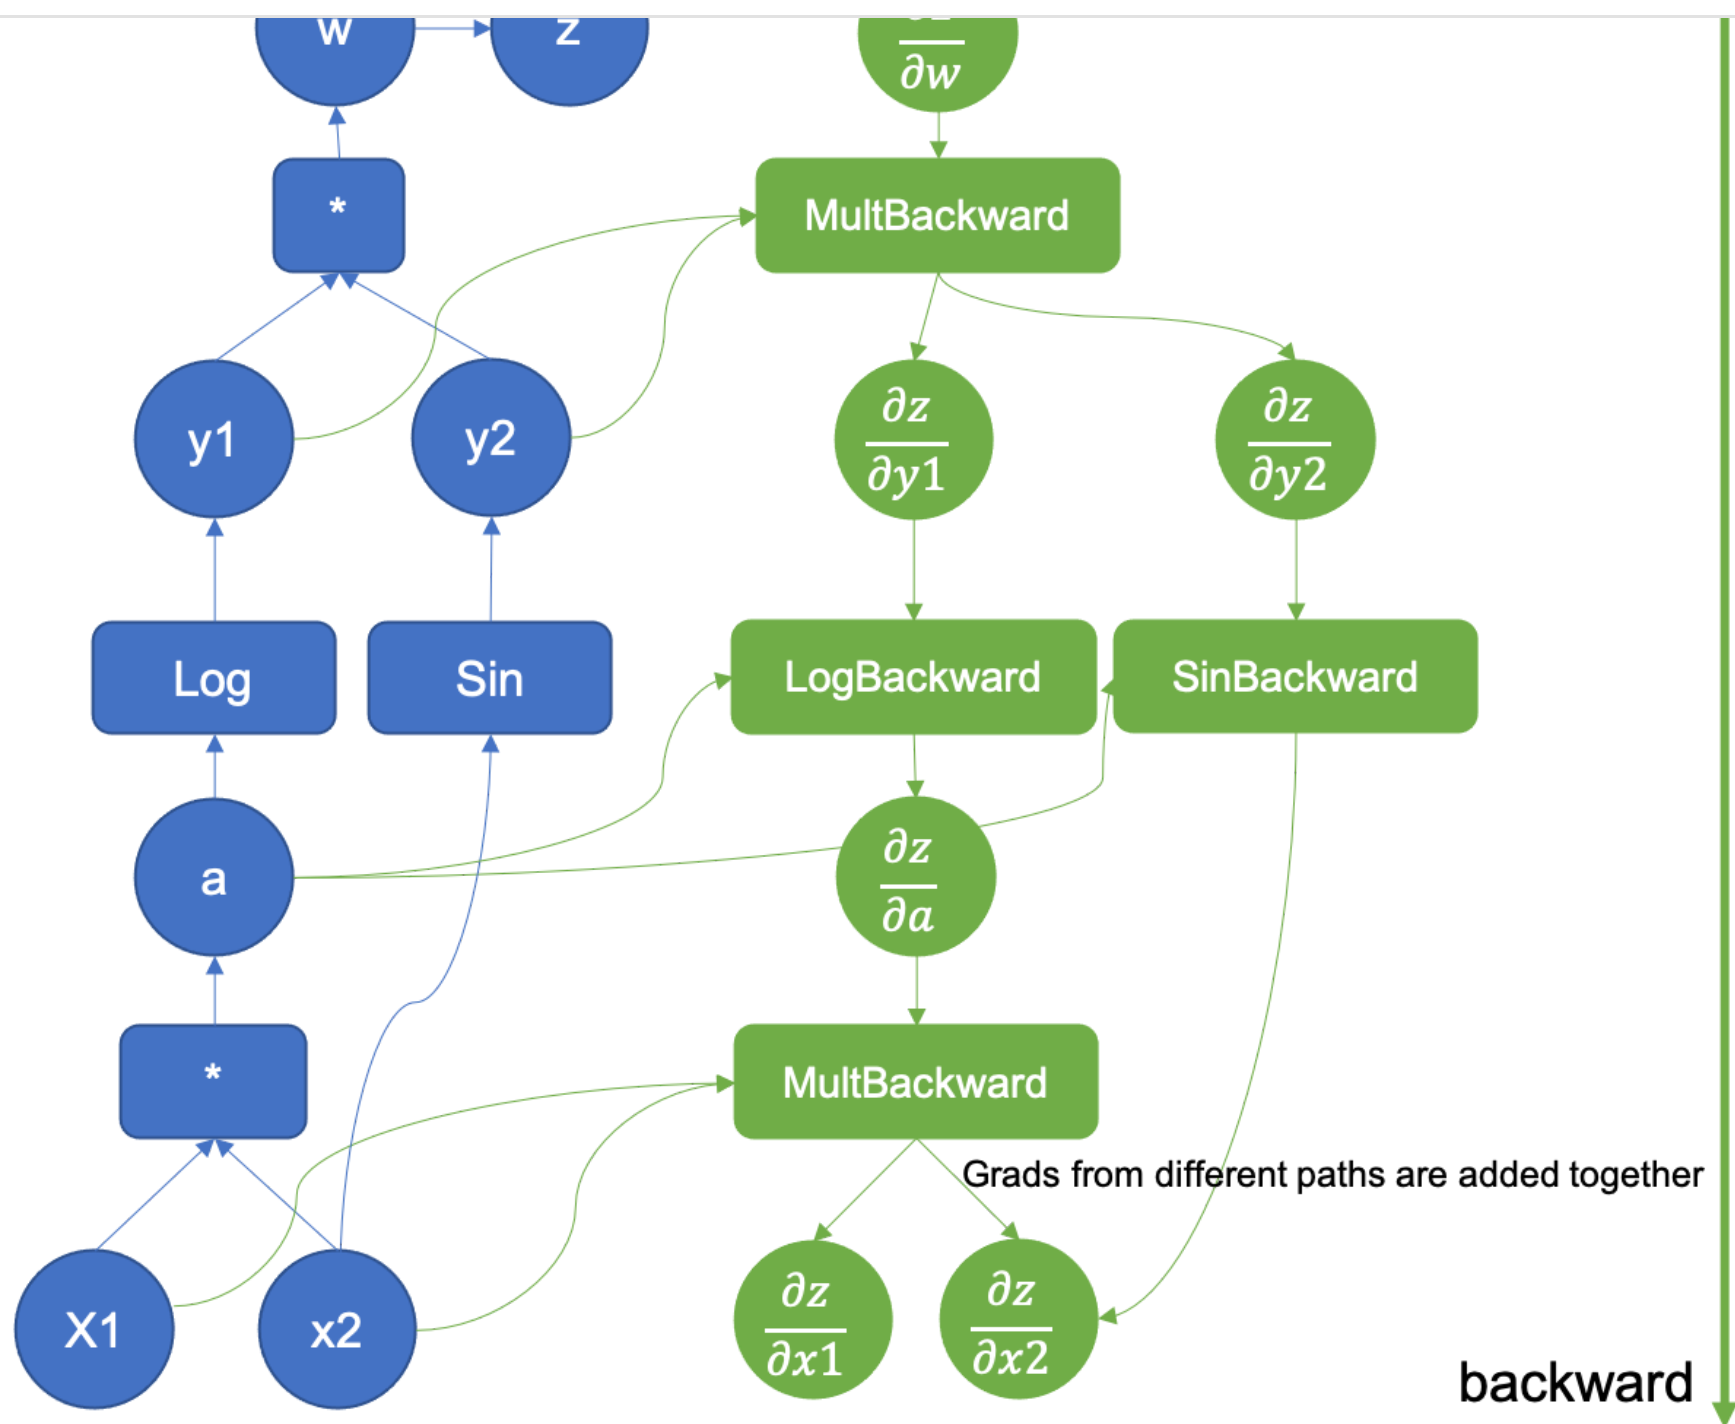


Figure 1: Example of an augmented computational graph

It all starts when in our python code, where we request a tensor to require the gradient.

```
>>> x = torch.tensor([0.5, 0.75], requires_grad=True)
```

When the `requires_grad` flag is set in tensor creation, c10 will **allocate** an `AutogradMeta` object that is used to hold the graph information.

```
void TensorImpl::set_requires_grad(bool requires_grad) {
    ...
    if (!autograd_meta_)
        autograd_meta_ = impl::GetAutogradMetaFactory()->make();
    autograd_meta_->set_requires_grad(requires_grad, this);
}
```

The `AutogradMeta` object is defined in `torch/csrc/autograd/variable.h` as follows:

```
struct TORCH_API AutogradMeta : public c10::AutogradMetaInterface {
    std::string name_;

    Variable grad_;
    std::shared_ptr<Node> grad_fn_;
    std::weak_ptr<Node> grad_accumulator_;
    // other fields and methods
    ...
};
```

The most important fields in this structure are the computed gradient in `grad_` and a pointer to the function `grad_fn` that will be called by the engine to produce the actual gradient. Also, there is a gradient accumulator object that is used to add together all the different gradients where this tensor is involved as we will see in the graph execution.

function that is implemented in ATen. Let it be the multiplication as in our previous blog post example. The resulting tensor has a field called `grad_fn` that is essentially a pointer to the function that will be used to compute the gradient of that operation.

```
>>> x = torch.tensor([0.5, 0.75], requires_grad=True)
>>> v = x[0] * x[1]
>>> v
tensor(0.3750, grad_fn=<MulBackward0>)
```

Here we see that the tensors' `grad_fn` has a `MulBackward0` value. This function is the same that was written in the [derivatives.yaml](#) file, and its C++ code was generated automatically by all the scripts in `tools/autograd`. It's auto-generated source code can be seen in `torch/csrc/autograd/generated/Functions.cpp`.

```
variable_list MulBackward0::apply(variable_list&& grads) {
    std::lock_guard<std::mutex> lock(mutex_);

    IndexRangeGenerator gen;
    auto self_ix = gen.range(1);
    auto other_ix = gen.range(1);
    variable_list grad_inputs(gen.size());
    auto& grad = grads[0];
    auto self = self_.unpack();
    auto other = other_.unpack();
    bool any_grad_defined = any_variable_defined(grads);
    if (should_compute_output({ other_ix })) {
        auto grad_result = any_grad_defined ? (mul_tensor_backward(grad, self, other_scalar_type)) : Tensor();
        copy_range(grad_inputs, other_ix, grad_result);
    }
    if (should_compute_output({ self_ix })) {
        auto grad_result = any_grad_defined ? (mul_tensor_backward(grad, other, self_scalar_type)) : Tensor();
        copy_range(grad_inputs, self_ix, grad_result);
    }
    return grad_inputs;
}
```

The `grad_fn` objects inherit from the `TraceableFunction` class, a descendant of `Node` with just a property set to enable tracing for debugging and optimization purposes. A graph by definition has nodes and edges, so these functions are indeed the nodes of the computational graph that are linked together by using `Edge` objects to enable the graph traversal later on.

The `Node` definition can be found in the [torch/csrc/autograd/function.h](#) file.

```
struct TORCH_API Node : std::enable_shared_from_this<Node> {
    ...
    /// Evaluates the function on the given inputs and returns the result of the
    /// function call.
    variable_list operator()(variable_list&& inputs) {
        ...
    }

protected:
    /// Performs the 'Node's actual operation.
    virtual variable_list apply(variable_list&& inputs) = 0;
    ...
    edge_list next_edges_;
}
```

Essentially we see that it has an override of the `operator()` that performs the call to the actual function, and a pure virtual function called `apply`. The automatically generated functions override this `apply` method as we saw in the `MulBackward0` example above. Finally, the node also has a list of edges to enable graph connectivity.

The `Edge` object is used to link `Node`s together and its implementation is straightforward.

```
struct Edge {
    ...
    /// The function this 'Edge' points to.
    std::shared_ptr<Node> function;
    /// The identifier of a particular input to the function.
    uint32_t input_nr;
};
```

It only requires a function pointer (the actual `grad_fn` objects that the edges link together), and an input number that acts as an id for the edge.

Linking nodes together

When we invoke the product operation of two tensors, we enter into the realm of autogenerated code. All the scripts that we saw in `tools/autograd` fill a series of templates that wrap the differentiable functions in ATen. These functions have code to construct the backward graph during the forward pass.

Let's take a look at how the tensor multiplication generated function looks like. The code has been simplified, but it can be found in the `torch/csrc/autograd/generated/VariableType_4.cpp` file when compiling pytorch from source.

```
at::Tensor mul_Tensor(c10::DispatchKeySet ks, const at::Tensor & self, const at::Tensor & other) {
    ...
    auto _any_requires_grad = compute_requires_grad( self, other );
    std::shared_ptr<MulBackward0> grad_fn;
    if (_any_requires_grad) {
        // Creates the link to the actual grad_fn and links the graph for backward traversal
        grad_fn = std::shared_ptr<MulBackward0>(new MulBackward0(), deleteNode);
        grad_fn->set_next_edges(collect_next_edges( self, other ));
        ...
    }
    ...
    // Does the actual function call to ATen
    auto _tmp = ([&]() {
        at::AutoDispatchBelowADInplaceOrView guard;
        return at::redispatch::mul(ks & c10::after_autograd_keyset, self_, other_);
    })();

    auto result = std::move(_tmp);
    if (grad_fn) {
        // Connects the result to the graph
        set_history(flatten_tensor_args( result ), grad_fn);
    }
    ...
    return result;
}
```

Let's walk through the most important lines of this code. First of all, the `grad_fn` object is created with: ``grad_fn = std::shared_ptr(new MulBackward0(), deleteNode);``.

After the `grad_fn` object is created, the edges used to link the nodes together are created by using the `grad_fn->set_next_edges(collect_next_edges(self, other));` calls.

```
struct MakeNextFunctionList : IterArgs<MakeNextFunctionList> {
    edge_list next_edges;
    using IterArgs<MakeNextFunctionList>::operator();
    void operator()(const Variable& variable) {
        if (variable.defined()) {
            next_edges.push_back(impl::gradient_edge(variable));
        } else {
            next_edges.emplace_back();
        }
    }
    void operator()(const c10::optional<Variable>& variable) {
        if (variable.has_value() && variable->defined()) {
            next_edges.push_back(impl::gradient_edge(*variable));
        } else {
            next_edges.emplace_back();
        }
    }
};

template <typename... Variables>
edge_list collect_next_edges(Variables&&... variables) {
    detail::MakeNextFunctionList make;
    make.apply(std::forward<Variables>(variables)...);
    return std::move(make.next_edges);
}
```

Given an input variable (it's just a regular tensor), `collect_next_edges` will create an `Edge` object by calling `impl::gradient_edge`

```
Edge gradient_edge(const Variable& self) {
    // If grad_fn is null (as is the case for a leaf node), we instead
    // interpret the gradient function to be a gradient accumulator, which will
    // accumulate its inputs into the grad property of the variable. These
    // nodes get suppressed in some situations, see "suppress gradient
    // accumulation" below. Note that only variables which have `requires_grad =
    // True` can have gradient accumulators.
    if (const auto& gradient = self.grad_fn()) {
        return Edge(gradient, self.output_nr());
    } else {
        return Edge(grad_accumulator(self), 0);
    }
}
```

To understand how edges work, let's assume that an early executed function produced two output tensors, both with their `grad_fn` set, each tensor also has an `output_nr` property with the order in which they were returned. When creating the edges for the current `grad_fn`, an `Edge` object per input variable will be created. The edges will point to the variable's `grad_fn` and will also track the `output_nr` to establish ids used when traversing the graph. In the case that the input variables are

After the edges are created, the `grad_fn` graph Node object that is being currently created will hold them using the `set_next_edges` function. This is what connects `grad_fn` s together, producing the computational graph.

```
void set_next_edges(edge_list&& next_edges) {
    next_edges_ = std::move(next_edges);
    for(const auto& next_edge : next_edges_) {
        update_topological_nr(next_edge);
    }
}
```


Now, the forward pass of the function will execute, and after the execution `set_history` will connect the output tensors to the `grad_fn` Node.

```
inline void set_history(
    at::Tensor& variable,
    const std::shared_ptr<Node>& grad_fn) {
    AT_ASSERT(grad_fn);
    if (variable.defined()) {
        // If the codegen triggers this, you most likely want to add your newly added function
        // to the DONT_REQUIRE_DERIVATIVE list in tools/autograd/gen_variable_type.py
        TORCH_INTERNAL_ASSERT(isDifferentiableType(variable.scalar_type()));
        auto output_nr =
            grad_fn->add_input_metadata(variable);
        impl::set_gradient_edge(variable, {grad_fn, output_nr});
    } else {
        grad_fn->add_input_metadata(Node::undefined_input());
    }
}
```

`set_history` calls `set_gradient_edge`, which just copies the `grad_fn` and the `output_nr` to the `AutogradMeta` object that the tensor has.

```
void set_gradient_edge(const Variable& self, Edge edge) {
    auto* meta = materialize_autograd_meta(self);
    meta->grad_fn_ = std::move(edge.function);
    meta->output_nr_ = edge.input_nr;
    // For views, make sure this new grad_fn_ is not overwritten unless it is necessary
    // in the VariableHooks::grad_fn below.
    // This logic is only relevant for custom autograd Functions for which multiple
    // operations can happen on a given Tensor before its gradient edge is set when
    // exiting the custom Function.
    auto diff_view_meta = get_view_autograd_meta(self);
    if (diff_view_meta && diff_view_meta->has_bw_view()) {
        diff_view_meta->set_attr_version(self._version());
    }
}
```

This tensor now will be the input to another function and the above steps will be all repeated. Check the animation below to see how the graph is created.



1. We start with a tensor

With `requires_grad` set that is the input to a function

Figure 2: Animation that shows the graph creation

the graph!

An autograd python defined function looks like the following:

```
class Exp(torch.autograd.Function):
    @staticmethod
    def forward(ctx, i):
        result = i.exp()
        ctx.save_for_backward(result)
        return result

    @staticmethod
    def backward(ctx, grad_output):
        result, = ctx.saved_tensors
        return grad_output * result

# Call the function
Exp.apply(torch.tensor(0.5, requires_grad=True))
# Outputs: tensor(1.6487, grad_fn=<ExpBackward>)
```

In the above snippet autograd detected our python function when creating the graph. All of this is possible thanks to the `Function` class. Let's take a look at what happens when we call `apply`.

`apply` is defined in the `torch._C._FunctionBase` class, but this class is not present in the python source. `_FunctionBase` is defined in C++ by using the python C API to hook C functions together into a single python class. We are looking for a function named `THPFunction_apply`.

```
PyObject *THPFunction_apply(PyObject *cls, PyObject *inputs)
{
    // Generates the graph node
    THPObjectPtr backward_cls(PyObject_GetAttrString(cls, "_backward_cls"));
    if (!backward_cls) return nullptr;
    THPObjectPtr ctx_obj(PyObject_CallFunctionObjArgs(backward_cls, nullptr));
    if (!ctx_obj) return nullptr;
    THPFunction* ctx = (THPFunction*)ctx_obj.get();

    auto cdata = std::shared_ptr<PyNode>(new PyNode(std::move(ctx_obj)), deleteNode);
    ctx->cdata = cdata;

    // Prepare inputs and allocate context (grad fn)
    // Unpack inputs will collect the edges
    auto info_pair = unpack_input<false>(inputs);
    UnpackedInput& unpacked_input = info_pair.first;
    InputFlags& input_info = info_pair.second;

    // Initialize backward function (and ctx)
    bool is_executable = input_info.is_executable;
    cdata->set_next_edges(std::move(input_info.next_edges));
    ctx->needs_input_grad = input_info.needs_input_grad.release();
    ctx->is_variable_input = std::move(input_info.is_variable_input);

    // Prepend ctx to input_tuple, in preparation for static method call
    auto num_args = PyTuple_GET_SIZE(inputs);
    THPObjectPtr ctx_input_tuple(PyTuple_New(num_args + 1));
    if (!ctx_input_tuple) return nullptr;
    Py_INCREF(ctx);
    PyTuple_SET_ITEM(ctx_input_tuple.get(), 0, (PyObject*)ctx);
    for (int i = 0; i < num_args; ++i) {
        PyObject *arg = PyTuple_GET_ITEM(unpacked_input.input_tuple.get(), i);
        Py_INCREF(arg);
        PyTuple_SET_ITEM(ctx_input_tuple.get(), i + 1, arg);
    }

    // Call forward
    THPObjectPtr tensor_outputs;
    {
        AutoGradMode grad_mode(false);
        THPObjectPtr forward_fn(PyObject_GetAttrString(cls, "forward"));
        if (!forward_fn) return nullptr;
        tensor_outputs = PyObject_CallObject(forward_fn, ctx_input_tuple);
        if (!tensor_outputs) return nullptr;
    }

    // Here is where the outputs gets the tensors tracked
    return process_outputs(cls, cdata, ctx, unpacked_input, inputs, std::move(tensor_outputs),
                          is_executable, node);

    END_HANDLE_TH_ERRORS
}
```

Although this code is hard to read at first due to all the python API calls, it essentially does the same thing as the auto-generated forward functions that we saw for ATen:

The `grad_fn` object is created in:

```
// Generates the graph node
THPObjectPtr backward_cls(PyObject_GetAttrString(cls, "_backward_cls"));
if (!backward_cls) return nullptr;
THPObjectPtr ctx_obj(PyObject_CallFunctionObjArgs(backward_cls, nullptr));
if (!ctx_obj) return nullptr;
THPFunction* ctx = (THPFunction*)ctx_obj.get();

auto cdata = std::shared_ptr<PyNode>(new PyNode(std::move(ctx_obj)), deleteNode);
ctx->cdata = cdata;
```

Basically, it asks the python API to get a pointer to the Python object that can execute the user-written function. Then it wraps it into a `PyNode` object that is a specialized `Node` object that calls the python interpreter with the provided python function when `apply` is executed during the forward pass. Note that in the code `cdata` is the actual `Node` object that is part of the graph. `ctx` is the object that is passed to the python `forward` / `backward` functions and it is used to store autograd related information by both, the user's function and PyTorch.

As in the regular C++ functions we also call `collect_next_edges` to track the inputs `grad_fn` objects, but this is done in `unpack_input` :

```
template<bool enforce_variables>
std::pair<UnpackedInput, InputFlags> unpack_input(PyObject *args) {
    ...
    flags.next_edges = (flags.is_executable ? collect_next_edges(unpacked.input_vars) : edge_list());
    return std::make_pair(std::move(unpacked), std::move(flags));
}
```

After this, the edges are assigned to the `grad_fn` by just doing `cdata->set_next_edges(std::move(input_info.next_edges));` and the forward function is called through the python interpreter C API.

Once the output tensors are returned from the forward pass, they are processed and converted to variables inside the `process_outputs` function.

```
PyObject* process_outputs(PyObject *op_obj, const std::shared_ptr<PyNode>& cdata,
                          THPFunction* grad_fn, const UnpackedInput& unpacked,
                          PyObject *inputs, THPObjectPtr&& raw_output, bool is_executable,
                          torch::jit::Node* node) {
    ...
    _wrap_outputs(cdata, grad_fn, unpacked.input_vars, raw_output, outputs, is_executable);
    _trace_post_record(node, op_obj, unpacked.input_vars, outputs, is_inplace, unpack_output);
    if (is_executable) {
        _save_variables(cdata, grad_fn);
    } ...
    return outputs.release();
}
```

Here, `_wrap_outputs` is in charge of setting the forward outputs `grad_fn` to the newly created one. For this, it calls another `_wrap_outputs` function defined in a different [file](#), so the process here gets a little confusing.

```
static void _wrap_outputs(const std::shared_ptr<PyNode>& cdata, THPFunction *self,
                          const variable_list &input_vars, PyObject *raw_output, PyObject *outputs, bool is_executable)
{
    auto cdata_if_executable = is_executable ? cdata : nullptr;
    ...

    // Wrap only the tensor outputs.
    // This calls csrc/autograd/custom_function.cpp
    auto wrapped_outputs = _wrap_outputs(input_vars, non_differentiable, dirty_inputs, raw_output_vars, cdata_if_executable);
    ...
}
```

The called `_wrap_outputs` is the one in charge of setting the autograd metadata in the output tensors:

```
const at::ArrayRef<at::optional<variable>>> raw_outputs,
const std::shared_ptr<Node> &cdata) {

std::unordered_set<at::TensorImpl*> inputs;
...
// Sets the grad_fn and output_nr of an output Variable.
auto set_history = [&](Variable& var, uint32_t output_nr, bool is_input, bool is_modified,
                    bool is_differentiable) {
    // Lots of checks
    if (!is_differentiable) {
        ...
    } else if (is_input) {
        // An input has been returned, but it wasn't modified. Return it as a view
        // so that we can attach a new grad_fn to the Variable.
        // Run in no_grad mode to mimic the behavior of the forward.
        {
            AutoGradMode grad_mode(false);
            var = var.view_as(var);
        }
        impl::set_gradient_edge(var, {cdata, output_nr});
    } else if (cdata) {
        impl::set_gradient_edge(var, {cdata, output_nr});
    }
};
```

And this is where `set_gradient_edge` was called and this is how a user-written python function gets included in the computational graph with its associated backward function!

Closing remarks

This blog post is intended to be a code overview on how PyTorch constructs the actual computational graphs that we discussed in the previous post. The next entry will deal with how the autograd engine executes these graphs.

Docs

Access comprehensive developer documentation for PyTorch

[View Docs](#)

Tutorials

Get in-depth tutorials for beginners and advanced developers

[View Tutorials](#)

Resources

Find development resources and get your questions answered

[View Resources](#)

PyTorch

Get Started

Features

Ecosystem

Blog

Contributing

Security

Resources

Tutorials

Docs

Discuss

GitHub Issues

Brand Guidelines

Stay up to date

Facebook

Twitter

YouTube

LinkedIn

Mastodon

PyTorch Podcasts

Spotify

Apple

Google

Amazon