

PaxosStore 源码分析「三、共识协议」

2020.02.17 SF-Zhou

本系列的前两篇分析了 PaxosStore 中网络通信和消息传递的实现，本篇将正式介绍 Paxos 算法，并分析 PaxosStore 中共识协议的实现。

1. Paxos 算法

Paxos 算法是 Leslie Lamport 于 1990 年提出的一种基于消息传递且具有高度容错特性的共识 (consensus) 算法。其论文于 1998 年 TOCS 会议上首次公开发表，中间的八年显然是有故事的。Paxos 算法已经问世 30 年了，至今依然折磨着学习分布式系统的同学们。入门学习的话，建议阅读作者 2001 年重新描述的论文 "Paxos Made Simple"。下面笔者说说自己对 Paxos 的理解。



达成共识 from "黑金"

首先 Paxos 是一个共识算法，即最终目标是达成共识，至于共识是什么、对不对，在这里并不重要，重要的是达成共识后，共识不可修改；其次，一个经典 Paxos 实例只能达

成一个共识，或者说确定 (chosen) 一个值，这一点很重要；最后，算法执行的环境是异步通信环境，使用非拜占庭模型，允许消息延迟、丢失、乱序，但不允许数据损坏 (corruption) 。

Paxos 算法中一共有三种角色：proposers, acceptors 和 learners，这里先忽略 learners。算法的步骤描述如下（摘录自论文原文 2.2 节）：

Phase 1.

1. A proposer selects a proposal number n and sends a prepare request with number n to a majority of acceptors.
2. If an acceptor receives a prepare request with number n greater than that of any prepare request to which it has already responded, then it responds to the request with a promise not to accept any more proposals numbered less than n and with the highest-numbered proposal (if any) that it has accepted.

Phase 2.

1. If the proposer receives a response to its prepare requests (numbered n) from a majority of acceptors, then it sends an accept request to each of those acceptors for a proposal numbered n with a value v , where v is the value of the highest-numbered proposal among the responses, or is any value if the responses reported no proposals.

2. If an acceptor receives an accept request for a proposal numbered n , it accepts the proposal unless it has already responded to a prepare request having a number greater than n .

因为消息本身的传递是不可靠的，所以可以从各个角色的响应来思考 Paxos 的流程，这会比面向过程的思考容易些。可以认为每个角色都是一个独立的线程，当收到特定的消息时做出对应的响应。这里先定义下消息类和角色自身的状态类：

```
struct Prepare {
    int n;
};

struct Accept {
    int n;
    void *value;
};

struct Promised {
    int n;
    Accept *proposal = nullptr;
};

struct Accepted {
    int n;
};

class Acceptor {
    int last_n = 0;
```

```
Accept *proposal = nullptr;  
};
```

对于 Acceptor, 其会对两种请求做出响应:

1. Prepare: 如果 $\text{prepare.n} > \text{acceptor.last_n}$, 则更新 acceptor.last_n , 并返回 $\text{Promised}(\text{n}, \text{acceptor.proposal})$;
2. Accept: 如果 $\text{accept.n} \geq \text{acceptor.last_n}$, 则接受该提案, 更新 acceptor.last_n 和 acceptor.proposal 并返回 $\text{Accepted}(\text{n})$ 。

对于 Proposer, 除主动发送 Prepare 请求外, 同时会接收两种响应:

1. Promised: 当 Promised 数量足够组成多数派时, 进入算法的 Phase 2、广播 Accept 请求;
2. Accepted: 当 Accepted 数量足够组成多数派时, Proposer 确认提议通过, 反之 Proposer 不确定, 但提案可能还是通过的。

Paxos 算法里, 对确定 (chosen) 的理解至关重要。当多数派通过提案的一瞬间, 共识即已达成且不可推翻, 理解了这一点就理解了 Paxos。

对于单个 Acceptor, 个体的 Accepted 和多数派的 Chosen 是有本质区别的, 个体的 Accepted 是可以被覆盖的。而当 Chosen 发生后, 之后的提案返回的 Promised 中得到的多数派里一定会返回群体 Chosen 的值, 这保证了提案的值不会被推翻。关于这点, 可以参考文献 2 中的例子。

Paxos 算法的证明网上可以找到很多，然后 MIT 6.824 2015 年前的课程里有对应的实验，参见文献 3。对应的代码笔者自己找了一份，放在 GitHub 上，可以自己完成以加深理解。2016 之后他们改用 Raft 了，有兴趣的话可以继续完成 Raft 的学习。

2. 协议实现

PaxosStore 的实现中弃用了原版的 Paxos 消息协议，并提出了使用**半对称消息**传递来实现 Paxos 过程，这里参考其论文中的描述。首先定义提案 Proposal：

$$\mathcal{P} = (n, v)$$

其中 n 表示提案号， v 表示提案的值。定义状态 State：

$$\mathcal{S} = (m, \mathcal{P})$$

其中 m 表示承诺过不会拒绝的最小提案号， \mathcal{P} 表示已经接受的提案。定义 \mathcal{S}_X 为机器 X 的状态，定义 \mathcal{S}_Y^X 为机器 X 已知的机器 Y 的状态，称之为视图状态。显然 $\mathcal{S}_Y^X (X \neq Y)$ 与 \mathcal{S}_Y 可能是不同的，不同机器上的状态依赖 Paxos 过程完成同步。定义机器 X 发送给机器 Y 的消息 Message：

$$\mathbb{M}_{X \rightarrow Y} = \{\mathcal{S}_X^X, \mathcal{S}_Y^X\}$$

简单来说，该消息中包含了 X 自己的实际状态，以及自己已知的对方的视图状态。PaxosStore 全局仅使用这一种消息来实现 Paxos 协议，具体的协议过程如下（摘录自原论

文) :

Algorithm 1: Paxos implementation in PaxosStore.

Input: intention proposal number m_i

```
1: Procedure Issue( $m_i$ ): /* invoked at  $N_A$  */
2:    $\mathcal{S}_A^A \leftarrow$  actual state of  $r_A$ 
3:   if  $\mathcal{S}_A^A.m < m_i$  then
4:      $\mathcal{S}_A^A.m \leftarrow m_i$ 
5:     write  $\mathcal{S}_A^A$  to the PaxosLog entry of  $r_A$ 
6:     foreach remote replica node  $N_X$  do
7:        $\lfloor$  send  $\mathbb{M}_{A \rightarrow X}$ 
```

Input: proposal \mathcal{P}_i with $\mathcal{P}_i.n = m_i$

```
8: Procedure Issue( $\mathcal{P}_i$ ): /* invoked at  $N_A$  */
9:    $\mathcal{S}_A^A \leftarrow$  actual state of  $r_A$ 
10:   $\mathbb{S}^A \leftarrow$  all the states of  $r$  maintained at  $N_A$ 
11:  if  $\left| \left\{ \forall \mathcal{S}_X^A \in \mathbb{S}^A \mid \mathcal{S}_X^A.m = \mathcal{P}_i.n \right\} \right| \times 2 > \left| \mathbb{S}^A \right|$  then
12:    if  $\left| \left\{ \forall \mathcal{S}_X^A \in \mathbb{S}^A \mid \mathcal{S}_X^A.\mathcal{P}.v \neq null \right\} \right| > 0$  then
13:       $\mathcal{P}' \leftarrow$  the proposal with maximum  $\mathcal{P}.n$  in  $\mathbb{S}^A$ 
14:       $\mathcal{S}_A^A.\mathcal{P} \leftarrow (\mathcal{P}_i.n, \mathcal{P}'.v)$ 
15:    else
16:       $\lfloor \mathcal{S}_A^A.\mathcal{P} \leftarrow \mathcal{P}_i$ 
17:      write  $\mathcal{S}_A^A$  to the PaxosLog entry of  $r_A$ 
18:      foreach remote replica node  $N_X$  do
19:         $\lfloor$  send  $\mathbb{M}_{A \rightarrow X}$ 
```

Input: message $\mathbb{M}_{X \rightarrow Y}$ sent from N_X to N_Y

```

20: Procedure OnMessage( $\mathbb{M}_{X \rightarrow Y}$ ): /* invoked at  $N_Y$  */
21:    $\mathcal{S}_X^X, \mathcal{S}_Y^X \leftarrow \mathbb{M}_{X \rightarrow Y}$ 
22:   UpdateStates( $Y, \mathcal{S}_X^X$ )
23:   if  $\mathcal{S}_Y^Y$  is changed then
24:     write  $\mathcal{S}_Y^Y$  to the PaxosLog entry of  $r_Y$ 
25:     if IsValueChosen( $Y$ ) is true then commit
26:   if  $\mathcal{S}_Y^X.m < \mathcal{S}_Y^Y.m$  or  $\mathcal{S}_Y^X.\mathcal{P}.n < \mathcal{S}_Y^Y.\mathcal{P}.n$  then
27:     send  $\mathbb{M}_{Y \rightarrow X}$ 

```

Input: node ID Y , actual state \mathcal{S}_X^X of r_X

```

28: Function UpdateStates( $Y, \mathcal{S}_X^X$ ): /* invoked at  $N_Y$  */
29:    $\mathcal{S}_X^Y \leftarrow$  view state of  $r_X$  stored in  $N_Y$ 
30:    $\mathcal{S}_Y^Y \leftarrow$  actual state of  $r_Y$ 
31:   if  $\mathcal{S}_X^Y.m < \mathcal{S}_X^X.m$  then  $\mathcal{S}_X^Y.m \leftarrow \mathcal{S}_X^X.m$ 
32:   if  $\mathcal{S}_X^Y.\mathcal{P}.n < \mathcal{S}_X^X.\mathcal{P}.n$  then  $\mathcal{S}_X^Y.\mathcal{P} \leftarrow \mathcal{S}_X^X.\mathcal{P}$ 
33:   if  $\mathcal{S}_Y^Y.m < \mathcal{S}_X^X.m$  then  $\mathcal{S}_Y^Y.m \leftarrow \mathcal{S}_X^X.m$ 
34:   if  $\mathcal{S}_Y^Y.m \leq \mathcal{S}_X^X.\mathcal{P}.n$  then  $\mathcal{S}_Y^Y.\mathcal{P} \leftarrow \mathcal{S}_X^X.\mathcal{P}$ 

```

Input: node ID Y

Output: whether the proposals in N_Y form a majority

```

35: Function IsValueChosen( $Y$ ): /* invoked at  $N_Y$  */
36:    $\mathbb{S}^Y \leftarrow$  all the states of  $r$  maintained at  $N_Y$ 
37:    $n' \leftarrow$  occurrence count of the most frequent  $\mathcal{P}.n$  in  $\mathbb{S}^Y$ 
38:   return  $n' \times 2 > |\mathbb{S}^Y|$ 

```


当节点 A 收到一个写请求时, 其会触发 **Issue**(m_i), 发送 Prepare 消息给所有节点; 节点 X 收到消息后触发 **OnMessage**($\mathbb{M}_{A \rightarrow X}$), 更新自身的状态, 若自身状态更新则会发送消息回节点 A ; 节点 A 收到回复的消息同样会执行 **OnMessage**($\mathbb{M}_{X \rightarrow A}$), 判断 Prepare 是否被多数派接受, 若是则触发 **Issue**(\mathcal{P}_i) 发送 Accept 消息给所有节点; 而后是一轮类似的消息处理, 节点 X 收到消息并更新, 回复消息, 节点 A 收到消息判断 **IsValueChosen**.

可以看到, PaxosStore 使用了统一的消息格式, 同时使用了统一的消息处理函数 **OnMessage**, 以此驱动整个 Paxos 协议过程。这大大简化了代码的实现, 按照论文描述的, 核心算法实现只使用了约 800 行 C++ 代码。这种实现的正确性已经在微信的大规模应用上得到验证, 另外也通过了基于 TLA+ 的形式化规约和验证, 见参考文献 5。

3. 代码实现

PaxosStore 中基本是按照论文中描述的方案来实现的。首先状态 \mathcal{S} 对应 src/Common.h 中的 EntryRecord_t:

```
struct PaxosValue_t {
    uint64_t iValueID;
    vector<uint64_t> vecValueUUID;

    bool bHasValue;
    string strValue;
```

```

PaxosValue_t() : iValueID(0), bHasValue(false) {}
PaxosValue_t(uint64_t iValueID_, const vector<uint64_t> vecValueUUID_, bool bHasValue_,
              const string &strValue_)
    : iValueID(iValueID_),
      vecValueUUID(vecValueUUID_),
      bHasValue(bHasValue_),
      strValue(strValue_) {}

bool operator==(const PaxosValue_t &tOther) const {
    if (iValueID == tOther.iValueID && vecValueUUID == tOther.vecValueUUID &&
        bHasValue == tOther.bHasValue && strValue == tOther.strValue) {
        return true;
    }

    return false;
}
};

struct EntryRecord_t {
    uint32_t iPreparedNum;
    uint32_t iPromisedNum;
    uint32_t iAcceptedNum;

    PaxosValue_t tValue;
    bool bChosen;

    bool bCheckedEmpty; // For Read Opt.

    uint64_t iStoredValueID; // For PutValue Opt.

```

```

bool operator==(const EntryRecord_t &tOther) const {
    if (iPreparedNum == tOther.iPreparedNum && iPromisedNum == tOther.iPromisedNum &&
        iAcceptedNum == tOther.iAcceptedNum && tValue == tOther.tValue &&
        bChosen == tOther.bChosen) {
        return true;
    }

    return false;
}
};

```

$S.m$ 对应 iPromisedNum, $S.P.n$ 对应 iAcceptedNum, $S.P.v$ 对应 tValue。每个节点会存储自己的状态, 以及其他节点的视图状态, 这些状态被存储于 clsEntryStateMachine 中, 定义于 src/EntryState.h:

```

enum enumEntryState {
    kEntryStateNormal = 0,
    kEntryStatePromiseLocal,
    kEntryStatePromiseRemote,
    kEntryStateMajorityPromise,
    kEntryStateAcceptRemote,
    kEntryStateAcceptLocal,
    kEntryStateChosen
};

class clsEntryStateMachine {
public:

```

```
static uint32_t s_iAcceptorNum;
static uint32_t s_iMajorityNum;
static uint32_t GetAcceptorID(uint64_t iValueID);

private:
    int m_iEntryState;

    uint32_t m_iMaxPreparedNum;

    uint32_t m_iMostAcceptedNum;
    uint32_t m_iMostAcceptedNumCnt;

    std::vector<EntryRecord_t> m_atRecord;

    uint32_t CountAcceptedNum(uint32_t iAcceptedNum);
    uint32_t CountPromisedNum(uint32_t iPromisedNum);

    int CalcEntryState(uint32_t iLocalAcceptorID);

    int MakeRealRecord(EntryRecord_t &tRecord);

    void UpdateMostAcceptedNum(const EntryRecord_t &tRecord);
    bool GetValueByAcceptedNum(uint32_t iAcceptedNum, PaxosValue_t &tValue);

public:
    static int Init(clsConfigure *poConf);

    clsEntryStateMachine() {
        m_iEntryState = kEntryStateNormal;
    }
}
```

```

    m_iMaxPreparedNum = 0;

    m_iMostAcceptedNum = 0;
    m_iMostAcceptedNumCnt = 0;

    m_atRecord.resize(s_iAcceptorNum);
    for (uint32_t i = 0; i < s_iAcceptorNum; ++i) {
        InitEntryRecord(&m_atRecord[i]);
    }
}

~clsEntryStateMachine() {}

int GetEntryState() { return m_iEntryState; }

uint32_t GetNextPreparedNum(uint32_t iLocalAcceptorID);

const EntryRecord_t &GetRecord(uint32_t iAcceptorID);

int Update(uint64_t iEntityID, uint64_t iEntry, uint32_t iLocalAcceptorID, uint32_t iAcceptorNum,
           const EntryRecord_t &tRecordMayWithValueIDOnly);

int AcceptOnMajorityPromise(uint32_t iLocalAcceptorID, const PaxosValue_t &tValue,
                           bool &bAcceptPreparedValue);

void SetStoredValueID(uint32_t iLocalAcceptorID);

// For readonly cmd.
void ResetAllCheckedEmpty();
void SetCheckedEmpty(uint32_t iAcceptorID);

```

```
bool IsLocalEmpty();
bool IsReadOK();

bool IsRemoteCompeting();

string ToString();

uint32_t CalcSize();
};
```

m_atRecord 中存储了自身的状态和其他节点的视图状态，可以通过 iAcceptorID 访问。整个 clsEntryStateMachine 构成一个状态机，接收到消息时通过执行 Update（对应算法描述中的 **UpdateStates**）更新自身的实际状态和其他节点的视图状态，进而推进 Paxos 的流程。当前 Paxos 执行的阶段使用 m_iEntryState 表示，初始化阶段是 kEntryStateNormal，最终 Chosen 的阶段是 kEntryStateChosen。下面是状态机代码的实现：

```
#include "EntryState.h"

namespace Certain {

uint32_t clsEntryStateMachine::s_iAcceptorNum = 0;
uint32_t clsEntryStateMachine::s_iMajorityNum = 0;

// 从 ValueID 中解析 AcceptorID
uint32_t clsEntryStateMachine::GetAcceptorID(uint64_t iValueID) {
    uint32_t iProposalNum = iValueID & 0xffffffff;
```

```

    AssertLess(0, iProposalNum);
    return (iProposalNum - 1) % s_iAcceptorNum;
}

// 初始化, 设定 Acceptor 和 Majority 的数量
int clsEntryStateMachine::Init(clsConfigure *poConf) {
    s_iAcceptorNum = poConf->GetAcceptorNum();
    s_iMajorityNum = (s_iAcceptorNum >> 1) + 1;

    return 0;
}

// 获取 Acceptor[i].EntryRecord
const EntryRecord_t &clsEntryStateMachine::GetRecord(uint32_t iAcceptorID) {
    return m_atRecord[iAcceptorID];
}

// 获取下一个 PreparedNum
uint32_t clsEntryStateMachine::GetNextPreparedNum(uint32_t iLocalAcceptorID) {
    if (m_iMaxPreparedNum == 0) {
        m_iMaxPreparedNum = iLocalAcceptorID + 1;
    } else {
        m_iMaxPreparedNum += s_iAcceptorNum;
    }

    return m_iMaxPreparedNum;
}

// 清空所有 CheckedEmpty 状态
void clsEntryStateMachine::ResetAllCheckedEmpty() {

```

```

    for (uint32_t i = 0; i < s_iAcceptorNum; ++i) {
        m_atRecord[i].bCheckedEmpty = false;
    }
}

// 设定 CheckedEmpty 状态
void clsEntryStateMachine::SetCheckedEmpty(uint32_t iAcceptorID) {
    m_atRecord[iAcceptorID].bCheckedEmpty = true;
}

// 检查本地是否是 Normal 状态
bool clsEntryStateMachine::IsLocalEmpty() {
    // (TODO)rock: use tla to check
    return m_iEntryState == kEntryStateNormal;
}

// 设定本地 Record 的 ValueID
void clsEntryStateMachine::SetStoredValueID(uint32_t iLocalAcceptorID) {
    EntryRecord_t &tLocalRecord = m_atRecord[iLocalAcceptorID];
    if (tLocalRecord.tValue.iValueID > 0) {
        tLocalRecord.iStoredValueID = tLocalRecord.tValue.iValueID;
    }
}

bool clsEntryStateMachine::IsReadOK() {
    uint32_t iCount = 0;

    for (uint32_t i = 0; i < s_iAcceptorNum; ++i) {
        if (m_atRecord[i].bCheckedEmpty && m_atRecord[i].iPromisedNum == 0) {
            iCount++;
        }
    }
}

```



```
    }  
}  
  
    CertainLogDebug("iCount %u", iCount);  
    return iCount >= s_iMajorityNum;  
}  
  
// 统计 Accepted 数量  
uint32_t clsEntryStateMachine::CountAcceptedNum(uint32_t iAcceptedNum) {  
    uint32_t iCount = 0;  
  
    for (uint32_t i = 0; i < s_iAcceptorNum; ++i) {  
        if (m_atRecord[i].iAcceptedNum == iAcceptedNum) {  
            iCount++;  
        }  
    }  
    return iCount;  
}  
  
// 统计 PromisedNum 数量  
uint32_t clsEntryStateMachine::CountPromisedNum(uint32_t iPromisedNum) {  
    uint32_t iCount = 0;  
  
    for (uint32_t i = 0; i < s_iAcceptorNum; ++i) {  
        if (m_atRecord[i].iPromisedNum == iPromisedNum) {  
            iCount++;  
        }  
    }  
    return iCount;  
}
```

// 根据 AcceptedNum 获取 Value

```
bool clsEntryStateMachine::GetValueByAcceptedNum(uint32_t iAcceptedNum, PaxosValue_t &tVa
    for (uint32_t i = 0; i < s_iAcceptorNum; ++i) {
        if (m_atRecord[i].iAcceptedNum == iAcceptedNum) {
            tValue = m_atRecord[i].tValue;
            return true;
        }
    }

    return false;
}
```

// 更新最多 Accepted 的 Num

```
void clsEntryStateMachine::UpdateMostAcceptedNum(const EntryRecord_t &tRecord) {
    if (m_iMostAcceptedNum == tRecord.iAcceptedNum) {
        m_iMostAcceptedNumCnt++;
        return;
    }
}
```

```
uint32_t iCount = CountAcceptedNum(tRecord.iAcceptedNum);
```

```
if (m_iMostAcceptedNumCnt < iCount) {
    m_iMostAcceptedNum = tRecord.iAcceptedNum;
    m_iMostAcceptedNumCnt = iCount;
}
}
```

// 计算当前的状态

```
int clsEntryStateMachine::CalcEntryState(uint32_t iLocalAcceptorID) {
```

```
EntryRecord_t &tLocalRecord = m_atRecord[iLocalAcceptorID];

m_iEntryState = kEntryStateNormal;

if (tLocalRecord.bChosen) {
    m_iEntryState = kEntryStateChosen;
    return 0;
}

if (m_iMostAcceptedNumCnt >= s_iMajorityNum) {
    m_iEntryState = kEntryStateChosen;

    if (tLocalRecord.iAcceptedNum != m_iMostAcceptedNum) {
        if (!GetValueByAcceptedNum(m_iMostAcceptedNum, tLocalRecord.tValue)) {
            return -1;
        }
        tLocalRecord.iAcceptedNum = m_iMostAcceptedNum;
    }

    tLocalRecord.bChosen = true;
    return 0;
}

if (tLocalRecord.iPromisedNum > tLocalRecord.iPreparedNum) {
    m_iEntryState = kEntryStatePromiseRemote;
    if (tLocalRecord.iAcceptedNum >= tLocalRecord.iPromisedNum) {
        m_iEntryState = kEntryStateAcceptRemote;
    }
    return 0;
}
```

```

if (tLocalRecord.iPromisedNum != tLocalRecord.iPreparedNum) {
    return -2;
}

// iPromisedNum == 0 means null.
if (tLocalRecord.iPromisedNum > 0) {
    m_iEntryState = kEntryStatePromiseLocal;

    uint32_t iLocalPromisedNum = tLocalRecord.iPromisedNum;
    uint32_t iPromisedNumCnt = CountPromisedNum(iLocalPromisedNum);

    if (iPromisedNumCnt >= s_iMajorityNum) {
        m_iEntryState = kEntryStateMajorityPromise;
    }

    if (tLocalRecord.iAcceptedNum == tLocalRecord.iPromisedNum) {
        m_iEntryState = kEntryStateAcceptLocal;
    }
}

if (tLocalRecord.iAcceptedNum > tLocalRecord.iPromisedNum) {
    m_iEntryState = kEntryStateAcceptRemote;
}

return 0;
}

// 通过 ValueID 获取 tRecord.tValue
int clsEntryStateMachine::MakeRealRecord(EntryRecord_t &tRecord) {

```

```

for (uint32_t i = 0; i < s_iAcceptorNum; ++i) {
    EntryRecord_t &tRealRecord = m_atRecord[i];

    if (tRecord.tValue.iValueID != tRealRecord.tValue.iValueID) {
        continue;
    }

    #if CERTAIN_DEBUG
    // For check only.
    if (tRecord.tValue.strValue.size() > 0) {
        if (tRecord.tValue.strValue != tRealRecord.tValue.strValue ||
            (tRecord.tValue.vecValueUUID.size() > 0 && tRealRecord.tValue.vecValueUUID.size() > 0 &&
             tRecord.tValue.vecValueUUID != tRealRecord.tValue.vecValueUUID)) {
            CertainLogFatal("CRC32(%u, %u) BUG record: %s lrecord[%u]: %s",
                           CRC32(tRecord.tValue.strValue), CRC32(tRealRecord.tValue.strValue),
                           EntryRecordToString(tRecord).c_str(), i,
                           EntryRecordToString(tRealRecord).c_str());

            Assert(false);
        }
    }
    #endif
    tRecord.tValue = tRealRecord.tValue;
    break;
}

if (tRecord.iAcceptedNum > 0) {
    if (tRecord.tValue.iValueID > 0) {
        return 0;
    }
}

```

```

        return -1;
    }

    return 1;
}

// 转为字符串
string clsEntryStateMachine::ToString() {
    string strState;
    for (uint32_t i = 0; i < s_iAcceptorNum; ++i) {
        if (i > 0) {
            strState += " ";
        }
        EntryRecord_t &tRecord = m_atRecord[i];
        strState += EntryRecordToString(tRecord);
    }
    return strState;
}

int clsEntryStateMachine::Update(uint64_t iEntityID, uint64_t iEntry, uint32_t iLocalAcce
                                uint32_t iAcceptorID,
                                const EntryRecord_t &tRecordMayWithValueIDOnly) {

#ifdef CERTAIN_DEBUG
    RETURN_RANDOM_ERROR_WHEN_IN_DEBUG_MODE();
#endif

    int iRet;

    if (m_iEntryState == kEntryStateChosen) {

```

```

    return -1;
}

if (iLocalAcceptorID >= s_iAcceptorNum || iAcceptorID >= s_iAcceptorNum) {
    return -2;
}

m_iEntryState = kEntryStateNormal;
EntryRecord_t &tLocalRecord = m_atRecord[iLocalAcceptorID];
EntryRecord_t &tRemoteRecord = m_atRecord[iAcceptorID];

// Make record into real when it comes in machine state.
EntryRecord_t tRecord = tRecordMayWithValueIDOnly;
iRet = CheckEntryRecordMayWithVIDOnly(tRecord);
if (iRet != 0) {
    CertainLogFatal("CheckEntryRecordMayWithVIDOnly ret %d", iRet);
    return -3;
}

iRet = MakeRealRecord(tRecord);
if (iRet < 0) {
    CertainLogFatal("MakeRealRecord ret %d", iRet);
    return -4;
}

if (!tRecord.tValue.bHasValue && tRecord.tValue.iValueID > 0) {
    if (tLocalRecord.iAcceptedNum > tRecord.iAcceptedNum) {
        // The remote acceptor supposed that this acceptor know the V.
        // But the V stored in tLocalRecord has been overridden,

```

```

        // Ignore the accepted message.
        tRecord.iAcceptedNum = 0;
        tRecord.tValue.iValueID = 0;
    } else {
        return -5;
    }
}

iRet = CheckEntryRecord(tRecord);
if (iRet != 0) {
    CertainLogFatal("CheckEntryRecord ret %d", iRet);
    return -6;
}

if (tRecord.bChosen) {
    // For check only.
    if (tRemoteRecord.iAcceptedNum >= tRecord.iAcceptedNum) {
        if (tRemoteRecord.tValue.iValueID != tRecord.tValue.iValueID) {
            return -7;
        }
    }
}

tLocalRecord.iAcceptedNum = tRecord.iAcceptedNum;
tLocalRecord.tValue = tRecord.tValue;
tLocalRecord.bChosen = true;

tRemoteRecord.iAcceptedNum = tRecord.iAcceptedNum;
tRemoteRecord.tValue = tRecord.tValue;
tRemoteRecord.bChosen = true;

```



```

    m_iEntryState = kEntryStateChosen;

    return m_iEntryState;
}

// 1. update m_iMaxPreparedNum
uint32_t iGlobalMaxPreparedNum = max(tRecord.iPreparedNum, tRecord.iPromisedNum);
if (m_iMaxPreparedNum < iGlobalMaxPreparedNum) {
    uint32_t iNextPreparedNum = m_iMaxPreparedNum;

    while (iNextPreparedNum <= iGlobalMaxPreparedNum) {
        m_iMaxPreparedNum = iNextPreparedNum;

        if (iNextPreparedNum == 0) {
            iNextPreparedNum = iLocalAcceptorID + 1;
        } else {
            iNextPreparedNum += s_iAcceptorNum;
        }
    }
}

// 2. update iPreparedNum
if (tRemoteRecord.iPreparedNum < tRecord.iPreparedNum) {
    tRemoteRecord.iPreparedNum = tRecord.iPreparedNum;
}

// 3. update old remote record
if (iAcceptorID != iLocalAcceptorID && tRecord.iAcceptedNum > tRemoteRecord.iAcceptedNum) {
    tRemoteRecord.iAcceptedNum = tRecord.iAcceptedNum;
    if (tRemoteRecord.tValue.iValueID != tRecord.tValue.iValueID) {

```

```

    tRemoteRecord.tValue = tRecord.tValue;
}

UpdateMostAcceptedNum(tRecord);
}

// 4. update value for remote accept first when use PreAuth
if (tRecord.iAcceptedNum == 0) {
    if (tRecord.tValue.iValueID > 0) {
        if (0 == tRecord.iPromisedNum) {
            return -8;
        }

        if (iLocalAcceptorID == iAcceptorID) {
            if (tLocalRecord.iAcceptedNum != 0 || tLocalRecord.tValue.iValueID != 0) {
                return -9;
            }
            tLocalRecord.tValue = tRecord.tValue;
        } else if (tRecord.iPromisedNum <= s_iAcceptorNum) {
            if (tRecord.iPromisedNum == 0 || tRecord.iPreparedNum != tRecord.iPromisedNum) {
                return -10;
            }
            tRecord.iAcceptedNum = tRecord.iPromisedNum;
        }
    }
}

// store the newest status of remote record
if (iLocalAcceptorID != iAcceptorID && tRemoteRecord.iAcceptedNum == 0) {
    if (tRemoteRecord.tValue.iValueID > 0 && tRecord.tValue.iValueID > 0 &&
        (tRemoteRecord.tValue.iValueID != tRecord.tValue.iValueID ||

```

```

        tRemoteRecord.tValue.strValue != tRecord.tValue.strValue)) {
    return -11;
}
tRemoteRecord.tValue = tRecord.tValue;
}
}

// 5. update old local record
if (tRecord.iAcceptedNum > tLocalRecord.iAcceptedNum &&
    tRecord.iAcceptedNum >= tLocalRecord.iPromisedNum) {
    tLocalRecord.iAcceptedNum = tRecord.iAcceptedNum;
    if (tLocalRecord.tValue.iValueID != tRecord.tValue.iValueID) {
        tLocalRecord.tValue = tRecord.tValue;
    }

    UpdateMostAcceptedNum(tRecord);
}

// 6. update iPromisedNum
if (tRemoteRecord.iPromisedNum < tRecord.iPromisedNum) {
    tRemoteRecord.iPromisedNum = tRecord.iPromisedNum;
}
if (tLocalRecord.iPromisedNum < tRecord.iPromisedNum) {
    tLocalRecord.iPromisedNum = tRecord.iPromisedNum;
}

iRet = CalcEntryState(iLocalAcceptorID);
if (iRet < 0) {
    CertainLogFatal("CalcEntryState ret %d", iRet);
    return -12;
}

```

[illegible]

```

bool &bAcceptPreparedValue) {

#ifdef CERTAIN_DEBUG
    RETURN_RANDOM_ERROR_WHEN_IN_DEBUG_MODE();
#endif

    int iRet;

    bAcceptPreparedValue = false;
    if (m_iEntryState != kEntryStateMajorityPromise) {
        return -1;
    }
    EntryRecord_t &tLocalRecord = m_atRecord[iLocalAcceptorID];
    if (tLocalRecord.iPreparedNum != tLocalRecord.iPromisedNum ||
        tLocalRecord.iAcceptedNum >= tLocalRecord.iPromisedNum) {
        return -2;
    }

    // 取得最大的 AcceptedNum 及其对应的 Value
    for (uint32_t i = 0; i < s_iAcceptorNum; ++i) {
        if (tLocalRecord.iAcceptedNum < m_atRecord[i].iAcceptedNum) {
            tLocalRecord.iAcceptedNum = m_atRecord[i].iAcceptedNum;
            tLocalRecord.tValue = m_atRecord[i].tValue;
        }
    }

    if (tLocalRecord.iAcceptedNum == 0) {
        tLocalRecord.tValue = tValue;
        bAcceptPreparedValue = true;
    }

    if (tLocalRecord.iAcceptedNum > tLocalRecord.iPromisedNum) {

```

```

        return -3;
    }
    tLocalRecord.iAcceptedNum = tLocalRecord.iPromisedNum;

    iRet = CheckEntryRecord(tLocalRecord);
    if (iRet != 0) {
        CertainLogFatal("CheckEntryRecord ret %d", iRet);
        return -4;
    }

    iRet = CalcEntryState(iLocalAcceptorID);
    if (iRet < 0) {
        CertainLogFatal("CalcEntryState ret %d", iRet);
        return -5;
    }

    if (m_iEntryState != kEntryStateAcceptLocal) {
        return -6;
    }

    if ((tLocalRecord.iPromisedNum - 1) % s_iAcceptorNum != iLocalAcceptorID) {
        return -7;
    }

    return 0;
}

bool clsEntryStateMachine::IsRemoteCompeting() {
    if (m_iEntryState == kEntryStatePromiseRemote || m_iEntryState == kEntryStateAcceptRemo
        return true;

```

```
}

return false;
}

// 统计自身数据大小
uint32_t clsEntryStateMachine::CalcSize() {
    uint32_t iSize = sizeof(clsEntryStateMachine);

    for (uint32_t i = 0; i < s_iAcceptorNum; ++i) {
        if (m_atRecord[i].tValue.strValue.size() == 0) {
            continue;
        }

        bool bToAdd = true;
        for (uint32_t j = 0; j < i; ++j) {
            if (m_atRecord[i].tValue.iValueID == m_atRecord[j].tValue.iValueID) {
                bToAdd = false;
                break;
            }
        }

        if (bToAdd) {
            iSize += m_atRecord[i].tValue.strValue.size();
        }
    }

    return iSize;
}
```

```
} // namespace Certain
```

算法描述中其他部分则主要实现于 src/EntityWorker.cpp。该文件还实现了 CatchUp / Recovery 等功能，代码非常长，这里摘录出 Paxos 流程的核心代码：

```
// A 节点发起 Paxos 流程
int clsEntityWorker::DoWithClientCmd(clsClientCmd *poCmd) {
    uint64_t iEntityID = poCmd->GetEntityID();
    EntryInfo_t *ptInfo = m_poEntryMng->FindEntryInfo(iEntityID, iEntry);
    uint32_t iLocalAcceptorID = ptEntityInfo->iLocalAcceptorID;

    if (ptInfo == NULL) {
        // 创建 Entry, 包含了 clsEntryStateMachine
        ptInfo = m_poEntryMng->CreateEntryInfo(ptEntityInfo, iEntry);
    }

    clsEntryStateMachine *poMachine = ptInfo->poMachine;
    int iEntryState = poMachine->GetEntryState();

    // 生成可用的 ProposalNum
    uint32_t iProposalNum = poMachine->GetNextPreparedNum(iLocalAcceptorID);

    // 创建一个新的状态
    EntryRecord_t tTempRecord;
    InitEntryRecord(&tTempRecord);

    // 自己当然直接 Promise
```



```

tTempRecord.iPreparedNum = iProposalNum;
tTempRecord.iPromisedNum = iProposalNum;

// 更新状态机
iEntryState =
    poMachine->Update(iEntityID, iEntry, iLocalAcceptorID, iLocalAcceptorID, tTempRecord);

// 获取自身状态
const EntryRecord_t &tUpdatedRecord = poMachine->GetRecord(iLocalAcceptorID);

// 构造消息
clsPaxosCmd *poPaxosCmd =
    new clsPaxosCmd(iLocalAcceptorID, iEntityID, iEntry, &tUpdatedRecord, NULL);

// 加入 PLog Req 队列进行持久化, 可以暂时忽略
iRet = clsPLogWorker::EnterPLogReqQueue(poPaxosCmd);
return eRetCodePtrReuse;
}

// A 节点经过 PLogWorker 持久化后, 会广播 Message
void clsEntityWorker::BroadcastToRemote(EntryInfo_t *ptInfo, clsEntryStateMachine *poMachine,
                                         clsClientCmd *poCmd) {
    EntityInfo_t *ptEntityInfo = ptInfo->ptEntityInfo;

    uint64_t iEntityID = ptEntityInfo->iEntityID;
    uint64_t iEntry = ptInfo->iEntry;

    uint32_t iLocalAcceptorID = ptEntityInfo->iLocalAcceptorID;

    // 遍历所有节点

```

```

for (uint32_t i = 0; i < m_iAcceptorNum; ++i) {
    // 忽略自己
    if (i == iLocalAcceptorID) {
        continue;
    }
    // 获取自身状态和目标节点的视图状态
    const EntryRecord_t &tSrc = poMachine->GetRecord(iLocalAcceptorID);
    const EntryRecord_t &tDest = poMachine->GetRecord(i);
    // 构造消息
    clsPaxosCmd *po = new clsPaxosCmd(iLocalAcceptorID, iEntityID, iEntry, &tSrc, &tDest)
    // 设定发送目标
    po->SetDestAcceptorID(i);
    // 使用 IO Worker 发送消息
    m_poIOWorkerRouter->GoAndDeleteIfFailed(po);
}
}

// x 节点通过网络接收到消息后
int clsEntityWorker::UpdateRecord(clsPaxosCmd *poPaxosCmd) {
    uint32_t iAcceptorID = poPaxosCmd->GetSrcAcceptorID();
    uint64_t iEntityID = poPaxosCmd->GetEntityID();
    uint64_t iEntry = poPaxosCmd->GetEntry();

    EntityInfo_t *ptEntityInfo = m_poEntityMng->FindEntityInfo(iEntityID);

    uint32_t iLocalAcceptorID = ptEntityInfo->iLocalAcceptorID;

    EntryInfo_t *ptInfo = m_poEntryMng->FindEntryInfo(iEntityID, iEntry);

    // 获取状态机

```

```

clsEntryStateMachine *poMachine = ptInfo->poMachine;
// 读取当前自己的状态
const EntryRecord_t tOldRecord = poMachine->GetRecord(iLocalAcceptorID);

// 获取消息中节点 A 发送的实际状态 S_A^A 和视图状态 S_X^A
const EntryRecord_t &tSrcRecord = poPaxosCmd->GetSrcRecord();
const EntryRecord_t &tDestRecord = poPaxosCmd->GetDestRecord();

// 更新 X 自身的状态机
int iRet = poMachine->Update(iEntityID, iEntry, iLocalAcceptorID, iAcceptorID, tSrcRecord);

// 获取实际状态 S_X
const EntryRecord_t &tNewRecord = poMachine->GetRecord(iLocalAcceptorID);
// 获取视图状态 S_A^X
const EntryRecord_t &tRemoteRecord = poMachine->GetRecord(iAcceptorID);

// 判断视图状态 S_X^A 与 S_X 是否一致
bool bRemoteUpdated = IsEntryRecordUpdated(tDestRecord, tNewRecord);

// 判断实际状态 S_X 是否存在更新
bool bLocalUpdated = IsEntryRecordUpdated(tOldRecord, tNewRecord);

// 构造消息
clsPaxosCmd *po =
    new clsPaxosCmd(iLocalAcceptorID, iEntityID, iEntry, &tNewRecord, &tRemoteRecord);
po->SetDestAcceptorID(iAcceptorID);
// 如果视图状态 S_X^A 与 S_X 不一致, 那么说明 A 节点需要更新对 X 节点的视图状态, 也就需要发送
ptInfo->bRemoteUpdated = bRemoteUpdated;

if (bLocalUpdated) {

```

```

        // 如果自身的实际状态存在更新, 则需要使用 PLog 将其持久化
        iRet = clsPLogWorker::EnterPLogReqQueue(po);
    } else {
        // 自身状态不存在更新, 直接进入回复阶段
        SyncEntryRecord(ptInfo, po->GetDestAcceptorID(), po->GetUUID());
    }

    return 0;
}

// x 节点持久化完成后, 发送回复消息
void clsEntityWorker::SyncEntryRecord(EntryInfo_t *ptInfo, uint32_t iDestAcceptorID,
                                       uint64_t iUUID) {
    EntityInfo_t *ptEntityInfo = ptInfo->ptEntityInfo;

    uint64_t iEntityID = ptEntityInfo->iEntityID;
    uint64_t iEntry = ptInfo->iEntry;

    uint32_t iLocalAcceptorID = ptEntityInfo->iLocalAcceptorID;

    clsEntryStateMachine *poMachine = ptInfo->poMachine;
    const EntryRecord_t &tSrcRecord = poMachine->GetRecord(iLocalAcceptorID);

    // 如果需要发送回复
    if (ptInfo->bRemoteUpdated) {
        const EntryRecord_t &tDestRecord = poMachine->GetRecord(iDestAcceptorID);

        if (!tDestRecord.bChosen) {
            // 构造发送回去的消息
            clsPaxosCmd *po =

```

```

        new clsPaxosCmd(iLocalAcceptorID, iEntityID, iEntry, &tSrcRecord, &tDestRecord)

// 设定发送目标
po->SetDestAcceptorID(iDestAcceptorID);

po->SetMaxChosenEntry(uint64_t(ptEntityInfo->iMaxChosenEntry));
// 通过 IO Worker 发送回 A 节点
m_poIOWorkerRouter->GoAndDeleteIfFailed(po);
    }
}
}

// 节点 A 收到回复消息后, 同样会执行 UpdateRecord
int clsEntityWorker::UpdateRecord(clsPaxosCmd *poPaxosCmd) {
    ...

    // 更新 A 自身的状态机
    int iRet = poMachine->Update(iEntityID, iEntry, iLocalAcceptorID, iAcceptorID, tSrcRecord, tDestRecord);

    // 判断 A 是否已经获得多数派的 Promised
    if (poMachine->GetEntryState() == kEntryStateMajorityPromise) {
        // 构造 Value
        PaxosValue_t tValue;
        tValue.bHasValue = true;
        tValue.iValueID = ptEntityInfo->poClientCmd->GetWriteBatchID();
        tValue.strValue = ptEntityInfo->poClientCmd->GetWriteBatch();
        tValue.vecValueUUID = ptEntityInfo->poClientCmd->GetWBUUID();

        bool bAcceptPreparedValue = false;
        // 尝试在 Accept 请求中使用请求 Value
    }
}

```

```
iRet = poMachine->AcceptOnMajorityPromise(iLocalAcceptorID, tValue, bAcceptPreparedVa
// bAcceptPreparedValue 会返回是否使用该 Value
// 如果 x 节点的 Promised 回复中已有 Value, 则该操作会返回失败
}

// 继续进行 PLog 持久化 / 发送回复消息
...
}
```

注意这里为了让流程清晰，删减了大量代码。上述代码大概是 Paxos 完整流程的一半，后面还有另一半发起 Accept 请求、确定值的过程，过程和上方基本一致，就不重复了。

4. 总结

PaxosStore 使用自己提出的半对称消息来实现 Paxos 过程，从实现来看确实简化了代码的复杂度，每个节点收到消息时走的都是同一套流程。当然目前只分析了最简单的 Paxos 流程，代码中还有错误处理、数据持久化、预授权优化等内容，这些将会在后续的博文中逐步分析。

References

1. "Paxos Made Simple", *Leslie Lamport*
2. "可靠分布式系统基础：paxos的直观解释", *drdrxp*
3. "6.824: Distributed Systems", *MIT*
4. "PaxosStore: High-availability Storage Made Practical in WeChat", *Jianjun Zheng*

5. "腾讯PaxosStore中共识算法的验证", 黄宇

0 comments

Write

Preview

Aa

Except where otherwise noted, content on this site is licensed under a CC BY-SA 4.0 license.
Copyright©2017 SF-Zhou, All Rights Reserved. Powered by GitHub Pages and GitHub Actions.