

[Click here to download the source code to this post](#)



## Automatic Differentiation Part 2: Implementation Using Micrograd

8:19

### Table of Contents

- [Automatic Differentiation Part 2: Implementation Using Micrograd](#)
  - [Introduction](#)
    - [What Is a Neural Network?](#)
  - [Having Problems Configuring Your Development Environment?](#)
  - [About micrograd](#)
  - [Imports and Setup](#)
  - [The Value Class](#)
  - [Addition](#)
    - [Compute Gradient](#)
  - [Multiplication](#)
    - [Compute Gradient](#)

Power

- [Power](#)
  - [Compute Gradient](#) [Click here to download the source code to this post](#)
- [Negation](#)
- [Subtraction](#)
- [Division](#)
- [Rectified Linear Unit](#)
- [The Global Backward](#)
- [Build a Multilayer Perceptron with micrograd](#)
  - [Module](#)
  - [Neuron](#)
  - [Layer](#)
  - [Multilayer Perceptron](#)
  - [Train the MLP](#)
- [Summary](#)
  - [Citation Information](#)

## [Automatic Differentiation Part 2: Implementation Using Micrograd](#)

In this tutorial, you will learn how automatic differentiation works with the help of a Python package named `micrograd`.



## Automatic Differentiation Part 2:

~~Click here to download the source code to this post~~  
Implementation Using Micrograd



(<https://pyimagesearch.com/wp-content/uploads/2022/12/autodiff-2-featured.png>)

This lesson is the last of a 2-part series on **Autodiff 101 — Understanding Automatic Differentiation from Scratch**:

- 1 ***Automatic Differentiation Part 1: Understanding the Math*** (<https://pyimg.co/pyxml>)
- 2 ***Automatic Differentiation Part 2: Implementation Using Micrograd*** (<https://pyimg.co/ra6ow>) (today's tutorial)

To learn how to implement automatic differentiation using Python, *just keep reading*.

# Automatic Differentiation Part 2: Implementation Using Micrograd

## Introduction

### What Is a Neural Network?

A Neural Network is a mathematical abstraction of our brain (at least, that is how it all started). The system consists of many learnable knobs (weights and biases) and a simple operation

(dot product). The Neural Network takes in inputs and uses an objective function that we need to optimize by turning the knobs. The best way to tune the knobs is to use the gradient of the objective function with respect to all the individual knobs as a signal.

~~Click here to download the source code to this post~~

It will take a long time if you sit down and try to calculate the gradient by hand. So, to bypass this process, we use the concept of automatic differentiation.

In the **previous tutorial (<https://pyimg.co/pyxml>)**, we deeply studied the mathematics of automatic differentiation. This tutorial will apply the concepts and work our way into understanding an automatic differentiation Python package from scratch.

The package that we will talk about today is called **[micrograd](https://github.com/karpathy/micrograd)** (<https://github.com/karpathy/micrograd>). This is an open-source Python package created by **Andrej Karpathy (<https://karpathy.ai/>)**. We have studied the **video lecture (<https://youtu.be/VMj-3S1tku0>)**, where Andrej built the package from scratch. Here, we break down the video lecture into a blog where we add our thoughts to enrich the content.

## **Having Problems Configuring Your Development Environment?**

---



[\(https://pyimagesearch.com/pyimagesearch-university/\)](https://pyimagesearch.com/pyimagesearch-university/)

Having trouble configuring your dev environment? Want access to pre-configured Jupyter Notebooks running on Google Colab? Be sure to join **PyImageSearch University** (<https://pyimagesearch.com/pyimagesearch-university/>) — you'll be up and running with this tutorial in a matter of minutes.

All that said, are you:

- Short on time?
- Learning on your employer's administratively locked system?
- Wanting to skip the hassle of fighting with the command line, package managers, and virtual environments?
- **Ready to run the code *right now* on your Windows, macOS, or Linux system?**

Then join **PyImageSearch University** (<https://pyimagesearch.com/pyimagesearch-university/>) today!

**Gain access to Jupyter Notebooks for this tutorial and other PyImageSearch guides that are *pre-configured* to run on Google Colab's ecosystem right in your web browser! No installation required.**

And best of all, these Jupyter Notebooks will run on Windows, macOS, and Linux!

[Click here to download the source code to this post](#)

## About micrograd

`micrograd` is a Python package built to understand how the reverse accumulation (backpropagation) process works in a modern deep learning package like PyTorch or Jax. It is a simple automatic differentiation package that works with **scalars** only.

## Imports and Setup

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | import math
2. | import random
3. | from typing import List, Tuple, Union
4. | from matplotlib import pyplot as plt
```

## The Value Class

We start things off by defining the `Value` class. To work on tracing and backpropagation later, it becomes essential to wrap raw scalar values into the `Value` class.

When wrapped inside the `Value` class, the scalar value is considered a **Node** of a Graph. When we use `Value`s and build an equation, the equation is considered a **Directed Acyclic Graph** ([https://en.wikipedia.org/wiki/Directed\\_acyclic\\_graph](https://en.wikipedia.org/wiki/Directed_acyclic_graph)) (DAG). With the help of *calculus* and *graph traversal*, we compute the gradients of the nodes automatically (autodiff) and backpropagate through them.

The `Value` class has the following attributes:

- `data` : The raw float data that needs to be wrapped inside the `Value` class.

• `grad` : This will hold the calculated derivatives of the nodes. The calculated derivative is the gradient

- `grad` : This will hold the **global derivative** of the node. The global derivative is the partial derivative of the root node (final node) with respect to the current node.  
[Click here to download the source code to this post](#)
- `_backward` : This is a private method that computes the global derivative of the children of the current node.
- `_prev` : The children of the current node.

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | class Value(object):
2. |     """
3. |         We need to wrap the raw data into a class that will store the
4. |         metadata to help in automatic differentiation.
5. |
6. |         Attributes:
7. |             data (float): The data for the Value node.
8. |             _children (Tuple): The children of the current node.
9. |     """
10. |
11. | def __init__(self, data: float, _children: Tuple = ()):
12. |     # The raw data for the Value node.
13. |     self.data = data
14. |
15. |     # The partial gradient of the last node with respect to this
16. |     # node. This is also termed as the global gradient.
17. |     # Gradient 0.0 means that there is no effect of the change
18. |     # of the last node with respect to this node. On
19. |     # initialization it is assumed that all the variables have no
20. |     # effect on the entire architecture.
21. |     self.grad = 0.0
22. |
23. |     # The function that derives the gradient of the children nodes
24. |     # of the current node. It is easier this way, because each node
25. |     # is built from children nodes and an operation. Upon back-propagation
26. |     # the current node can easily fill in the gradients of the children.
27. |     # Note: The global gradient is the multiplication of the local gradient
28. |     # and the flowing gradient from the parent.
29. |     self._backward = lambda: None
30. |
31. |     # Define the children of this node.
32. |     self._prev = set(_children)
33. |
34. | def __repr__(self):
35. |     # This is the string representation of the Value node.
36. |     return f"Value(data={self.data}, grad={self.grad})"
```

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Build a Value node
2. | raw_data = 5.0
```

```

3. | print(f"Raw Data(data={raw_data}, type={type(raw_data)})")
4. | value_node = Value(data=raw_data)
5. | Click here to download the source code to this post
6. | # Calling the `__repr__` function here
7. | print(value_node)

```

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

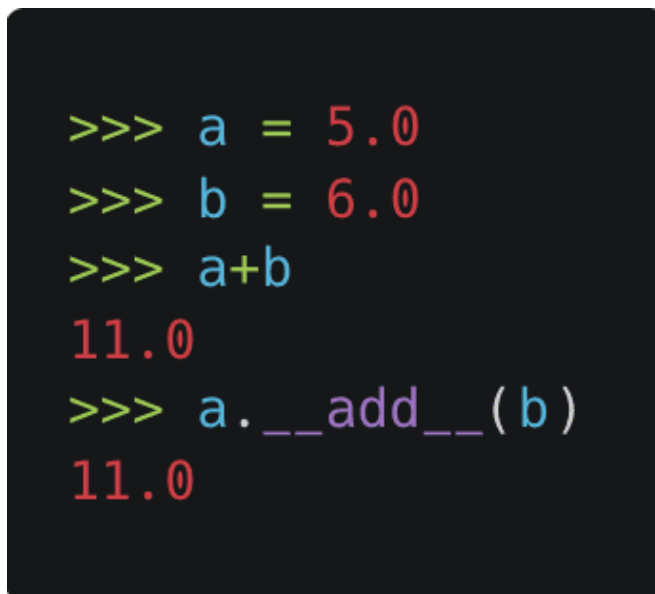
1. | >>> Raw Data(data=5.0, type=<class 'float'>)
2. | >>> Value(data=5.0, grad=0.0)

```

## Addition

Now that we have built our `Value` class, we need to define the primitive operations and their `_backward` functions. This will help trace each node's operations and back-propagate the gradients through the DAG expression.

In this section, we deal with the **addition** operation. This will help in two values being added together. Python classes have a special method `__add__` called when we use the `+` operator, as shown in **Figure 1**.



```

>>> a = 5.0
>>> b = 6.0
>>> a+b
11.0
>>> a.__add__(b)
11.0

```

(<https://pyimagesearch.com/wp-content/uploads/2022/12/add-code.png>)

**Figure 1:** `__add__` dunder function (source: image by the authors).



Here we create the `custom_addition` function that is later assigned to the `__add__` method of the `Value` class. This is done for us to focus on the addition method and discard everything that is not important to the addition operation.

[Click here to download the source code to this post](#)

The addition operation is as simple as it gets:

- 1 The `self` and the `other` nodes as an argument to the call. We then take their `data` and apply addition.
- 2 The result is then wrapped inside the `Value` class.
- 3 The `out` node is initialized, where we mention that `self` and `other` are its children.

## Compute Gradient

We will have this section for every primitive operation that we define. For example, to compute the global gradient of the children nodes, we need to define the local gradient of the `addition` operation.

Let us consider a node  $c$  that is built by adding two children nodes  $a$  and  $b$ . Then, the partial derivatives of  $c$  are derived in **Figure 2**.

$$c = a + b$$

$$\frac{\partial c}{\partial a} = 1$$

$$\frac{\partial c}{\partial b} = 1$$

[Click here to download the source code to this post](#)

(<https://pyimagesearch.com/wp-content/uploads/2022/12/math-1.png>)

**Figure 2:** Local gradient of addition operation (source: image by the authors).

Now think of backpropagation. The partial derivative of the loss (objective) function  $l$  is already deduced for  $c$ . This means we have  $(\partial l)/(\partial c)$ . This gradient needs to flow to the child nodes  $a$  and  $b$ , respectively.

Applying the chain rule, we get the global gradient for  $a$  and  $b$ , as shown in **Figure 3**.

$$\frac{\partial l}{\partial a} = \frac{\partial l}{\partial c} \frac{\partial c}{\partial a} = 1$$

Click here to download the source code to this post

$$\frac{\partial l}{\partial a} = \frac{\partial l}{\partial c} \cdot \frac{\partial c}{\partial a} = \frac{\partial l}{\partial c} \cdot 1$$

$$\frac{\partial l}{\partial b} = \frac{\partial l}{\partial c} \cdot \frac{\partial c}{\partial b} = \frac{\partial l}{\partial c} \cdot 1$$

(<https://pyimagesearch.com/wp-content/uploads/2022/12/math-2.png>)

**Figure 3:** Global derivative of addition operation (source: image by the authors).

The addition operation acts like a **router** to the gradients flowing in. It routes the gradients to all the children.

► **Note:** In the `_backward` functions that we define, we accumulate the gradients of the children with the `+=` operation. This is done to bypass a unique case. Suppose we have  $c = a + a$ . Here we know that the expression can be simplified to  $c = 2a$ , but our `_backward` for `__add__` does not know how to do this. The `__backward__` in `__add__` treats one  $a$  as `self` and the other  $a$  as `other`. If the gradients are not accumulated, we will see a discrepancy with the gradients.

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. def custom_addition(self, other: Union["Value", float]) -> "Value":
2.     """
3.     The addition operation for the Value class.
4.     Args:
5.         other (Union["Value", float]): The other value to add to this one.
6.     Usage:
7.         >>> x = Value(2)
8.         >>> y = Value(3)
9.         >>> z = x + y
10.        >>> z.data
11.
12.        5
13.    """
14.    # If the other value is not a Value, then we need to wrap it.
15.    other = other if isinstance(other, Value) else Value(other)
16.    # Create a new Value node that will be the output of the addition.
17.    out = Value(data=self.data + other.data, children=(self, other))
```

[Click here to download the source code to this post](#)

```

17. out = value(data=self.data + other.data, _children=(self, other))
18.
19. def _backward():
20.     # Local gradient.
21.     # x = a + b
22.     # dx/da = 1
23.     # dx/db = 1
24.     # Global gradient with chain rule:
25.     # dy/da = dy/dx . dx/da = dy/dx . 1
26.     # dy/db = dy/dx . dx/db = dy/dx . 1
27.     self.grad += out.grad * 1.0
28.     other.grad += out.grad * 1.0
29.
30.     # Set the backward function on the output node.
31.     out._backward = _backward
32.     return out
33.
34. def custom_reverse_addition(self, other):
35.     """
36.     Reverse addition operation for the Value class.
37.     Args:
38.         other (float): The other value to add to this one.
39.     Usage:
40.         >>> x = Value(2)
41.         >>> y = Value(3)
42.         >>> z = y + x
43.         >>> z.data
44.         5
45.     """
46.     # This is the same as adding. We can reuse the __add__ method.
47.     return self + other
48.
49.
50. Value.__add__ = custom_addition
51. Value.__radd__ = custom_reverse_addition

```

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. # Build a and b
2. a = Value(data=5.0)
3. b = Value(data=6.0)
4.
5. # Print the addition
6. print(f"{a} + {b} => {a+b}")

```

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. >>> Value(data=5.0, grad=0.0) + Value(data=6.0, grad=0.0) => Value(data=11.0, grad=0.0)

```

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. # Add a and b
2. c = a + b

```

```
3. |  
4. | # Assign a global gradient to c  
5. | c.grad = 11.0 Click here to download the source code to this post  
6. | print(f"c => {c}")  
7. |  
8. | # Now apply `_backward` to c  
9. | c._backward()  
10. | print(f"a => {a}")  
11. | print(f"b => {b}")
```

→ [Launch Jupyter Notebook on Google Colab](#)

---

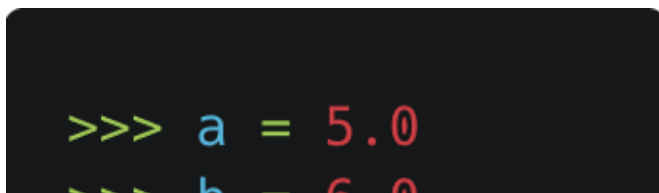
Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | >>> c => Value(data=11.0, grad=11.0)  
2. | >>> a => Value(data=5.0, grad=11.0)  
3. | >>> b => Value(data=6.0, grad=11.0)
```

► **Note:** The global gradient of *c* is routed to *a* and *b*.

## Multiplication

In this section, we deal with the **multiplication** operation. Python classes have a special method `__mul__` called when we use the `*` operator, as shown in **Figure 4**.



```
>>> a = 5.0  
>>> b = 6.0
```

```
>>> b = 0.0
>>> a*b
30.0
>>> a.__mul__(b)
30.0
```

<https://pyimagesearch.com/wp-content/uploads/2022/12/mul-code.png>

**Figure 4:** `__mul__` dunder method (source: image by the authors).

We get the `self` and the `other` nodes as an argument to the call. We then take their `data` and apply multiplication. The result is then wrapped inside the `Value` class. Finally, the `out` node is initialized, where we mention that `self` and `other` are its children.

## Compute Gradient

Let us consider a node  $c$  that is built by multiplying two children nodes  $a$  and  $b$ . Then, the partial derivatives of  $c$  are shown in **Figure 5**.

$$c = a \times b$$

[Click here to download the source code to this post](#)

$$\frac{\partial c}{\partial a} = b$$

$$\frac{\partial c}{\partial b} = a$$

(<https://pyimagesearch.com/wp-content/uploads/2022/12/math-3.png>)

**Figure 5:** Local gradient of multiplication operation (source: image by the authors).

Now think of backpropagation. The partial derivative of the loss (objective) function  $l$  is already deduced for  $c$ . This means we have  $(\partial l)/(\partial c)$ . This gradient needs to flow to the children nodes  $a$  and  $b$ , respectively.

Applying the chain rule, we get the global gradient for  $a$  and  $b$ , as shown in **Figure 6**.

$$\frac{\partial l}{\partial a} = \frac{\partial l}{\partial c} \frac{\partial c}{\partial a}$$

$$\frac{\partial l}{\partial b} = \frac{\partial l}{\partial c} \frac{\partial c}{\partial b}$$

$$\frac{\partial}{\partial a} \frac{\partial}{\partial l} = \frac{\partial}{\partial c} \cdot \frac{\partial}{\partial a} = \frac{\partial}{\partial c} \cdot 1$$

$$\frac{\partial}{\partial b} = \frac{\partial}{\partial c} \cdot \frac{\partial}{\partial b} = \frac{\partial}{\partial c} \cdot a$$

[Click here to download the source code to this post](https://pyimagesearch.com/wp-content/uploads/2022/12/math-4.png)

(<https://pyimagesearch.com/wp-content/uploads/2022/12/math-4.png>)

**Figure 6:** Global gradient of multiplication operation (source: image by the authors).

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. def custom_multiplication(self, other: Union["Value", float]) -> "Value":
2.     """
3.     The multiplication operation for the Value class.
4.     Args:
5.         other (float): The other value to multiply to this one.
6.     Usage:
7.         >>> x = Value(2)
8.         >>> y = Value(3)
9.         >>> z = x * y
10.        >>> z.data
11.        6
12.     """
13.     # If the other value is not a Value, then we need to wrap it.
14.     other = other if isinstance(other, Value) else Value(other)
15.
16.     # Create a new Value node that will be the output of
17.     # the multiplication.
18.     out = Value(data=self.data * other.data, _children=(self, other))
19.
20.     def _backward():
21.         # Local gradient:
22.         # x = a * b
23.         # dx/da = b
24.         # dx/db = a
25.         # Global gradient with chain rule:
26.         # dy/da = dy/dx . dx/da = dy/dx . b
27.         # dy/db = dy/dx . dx/db = dy/dx . a
28.
29.         self.grad += out.grad * other.data
30.         other.grad += out.grad * self.data
31.
32.     # Set the backward function on the output node.
33.     out._backward = _backward
34.     return out

```



34. |  
 35. | def custom\_reverse\_multiplication(self, other):  
 36. | """  
 37. | [Click here to download the source code to this post](#)  
 38. | Reverse multiplication operation for the Value class.  
 39. | Args:  
 40. | other (float): The other value to multiply to this one.  
 41. | Usage:  
 42. | >>> x = Value(2)  
 43. | >>> y = Value(3)  
 44. | >>> z = y \* x  
 45. | >>> z.data  
 46. | 6  
 47. | """  
 48. | # This is the same as multiplying. We can reuse the \_\_mul\_\_ method.  
 49. | return self \* other  
 50. |  
 51. | Value.\_\_mul\_\_ = custom\_multiplication  
 52. | Value.\_\_rmul\_\_ = custom\_reverse\_multiplication

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | # Build a and b
2. | a = Value(data=5.0)
3. | b = Value(data=6.0)
4. |
5. | # Print the multiplication
6. | print(f"{a} * {b} => {a*b}")
```

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | >>> Value(data=5.0, grad=0.0) * Value(data=6.0, grad=0.0) => Value(data=30.0, grad=0.0)
```

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | # Multiply a and b
2. | c = a * b
3. |
4. | # Assign a global gradient to c
5. | c.grad = 11.0
6. | print(f"c => {c}")
7. |
8. | # Now apply `_backward` to c
9. | c._backward()
10. | print(f"a => {a}")
11. | print(f"b => {b}")
```

→ [Launch Jupyter Notebook on Google Colab](#)

---

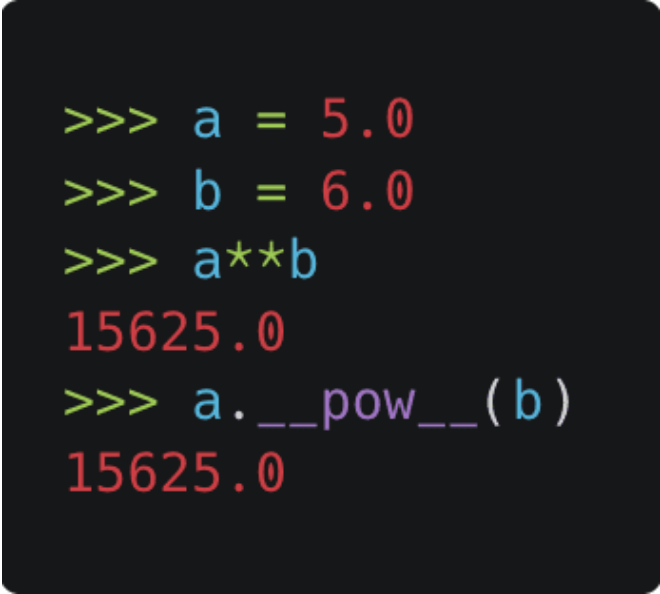
Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | >>> c => Value(data=30.0, grad=11.0)
2. | >>> a => Value(data=5.0, grad=66.0)
3. | >>> b => Value(data=5.0, grad=55.0)
```

[Click here to download the source code to this post](#)

## Power

In this section, we deal with the **power** operation. Python classes have a special method `__pow__` that is called when we use the `**` operator, as shown in **Figure 7**.



```
>>> a = 5.0
>>> b = 6.0
>>> a**b
15625.0
>>> a.__pow__(b)
15625.0
```

(<https://pyimagesearch.com/wp-content/uploads/2022/12/pow-code.png>)

**Figure 7:** `__pow__` dunder function (source: image by the authors).

After obtaining the `self` and the `other` nodes as an argument to the call, we take their `data` and apply the power operation.

## Compute Gradient

Let us consider a node  $c$  that is built by multiplying two children nodes  $a$  and  $b$ . Then, the partial derivatives of  $c$  are derived in **Figure 8**.

$$c = a^b$$

$$\frac{\partial c}{\partial a} = ba^{b-1}$$

[Click here to download the source code to this post](#)

(<https://pyimagesearch.com/wp-content/uploads/2022/12/math-5.png>)

**Figure 8:** Local gradient of power operation (source: image by the authors).

Now think of backpropagation. The partial derivative of the loss (objective) function  $l$  is already deduced for  $c$ . This means we have  $(\partial l)/(\partial c)$ . This gradient needs to flow to the child node  $a$ .

Applying the chain rule, we get the global gradient for  $a$  and  $b$ , as shown in **Figure 9**.

$$\frac{\partial l}{\partial a} = \frac{\partial l}{\partial c} \cdot \frac{\partial c}{\partial a} = \frac{\partial l}{\partial c} \cdot ba^{b-1}$$

(<https://pyimagesearch.com/wp-content/uploads/2022/12/math-6.png>)

**Figure 9:** Global gradient of power operation (source: image by the authors).

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | def custom_power(self, other):
2. |     """
3. |     The power operation for the Value class.
4. |     Args:
5. |         other (float): The other value to raise this one to.
6. |
7. |     Usage:
8. |         >>> x = Value(2)
9. |         >>> z = x ** 2.0
10. |         >>> z.data
11. |         4
12. |     """
```

```

12. |         assert isinstance(
13. |             other, (int, float)
14. |         ), "only supporting int/float powers for now"
15. |
16. |         # Create a new Value node that will be the output of the power.
17. |         out = Value(data=self.data ** other, _children=(self,))
18. |
19. |         def _backward():
20. |             # Local gradient:
21. |             # x = a ** b
22. |             # dx/da = b * a ** (b - 1)
23. |             # Global gradient:
24. |             # dy/da = dy/dx . dx/da = dy/dx . b * a ** (b - 1)
25. |             self.grad += out.grad * (other * self.data ** (other - 1))
26. |
27. |         # Set the backward function on the output node.
28. |         out._backward = _backward
29. |         return out
30. |
31. |
32. | Value.__pow__ = custom_power

```

**[Click here to download the source code to this post](#)**

→ **[Launch Jupyter Notebook on Google Colab](#)**

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Build a
2. | a = Value(data=5.0)
3. | # For power operation we will use
4. | # the raw data and not wrap it into
5. | # a node. This is done for simplicity.
6. | b = 2.0
7. |
8. | # Print the power operation
9. | print(f"{a} ** {b} => {a**b}")

```

→ **[Launch Jupyter Notebook on Google Colab](#)**

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | >>> Value(data=5.0, grad=0.0) ** 2.0 => Value(data=25.0, grad=0.0)

```

→ **[Launch Jupyter Notebook on Google Colab](#)**

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Raise a to the power of b
2. | c = a ** b
3. |
4. | # Assign a global gradient to c
5. | c.grad = 11.0
6. | print(f"c => {c}")
7. |
8. | # Now apply `_backward` to c
9. | c._backward()
10. | print(f"a => {a}")
11. | print(f"b => {b}")

```

→ [Launch Jupyter Notebook on Google Colab](#)

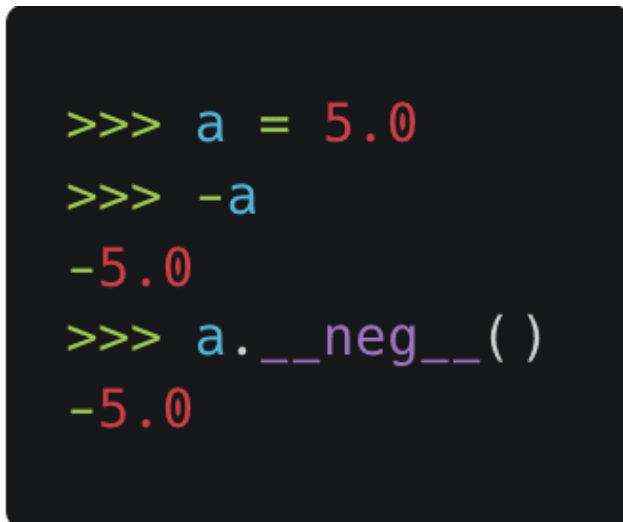
[Click here to download the source code to this post](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | >>> c => Value(data=25.0, grad=11.0)
2. | >>> a => Value(data=5.0, grad=110.0)
3. | >>> b => 2.0
```

## Negation

For the **negation** operation, we will be using the `__mul__` operation defined above. In addition, Python classes have a special method `__neg__` called when we use the unary `-` operator, as shown in **Figure 10**.



```
>>> a = 5.0
>>> -a
-5.0
>>> a.__neg__()
-5.0
```

(<https://pyimagesearch.com/wp-content/uploads/2022/12/neg-code.png>)

**Figure 10:** `__neg__` dunder function (source: image by the authors).

This means the `_backward` of negation will be taken care of, and we would not have to define it explicitly.

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | def custom_negation(self):
2. |     """
3. |     Negation operation for the Value class.
4. |     Usage:
5. |         >>> x = Value(2)
6. |         >>> z = -x
7. |         >>> z.data
8. |         2
```

```

8. |         -2
9. |         """
10. |         # This is the same as multiplying by -1. We can reuse the
11. |         # __mul__ method.
12. |         return self * -1
13. |
14. | Value.__neg__ = custom_negation

```

[Click here to download the source code to this post](#)

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Build `a`
2. | a = Value(data=5.0)
3. |
4. | # Print the negation
5. | print(f"Negation of {a} => {(-a)}")

```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | >>> Negation of Value(data=5.0, grad=0.0) => Value(data=-5.0, grad=0.0)

```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Negate a
2. | c = -a
3. |
4. | # Assign a global gradient to c
5. | c.grad = 11.0
6. | print(f"c => {c}")
7. |
8. | # Now apply `_backward` to c
9. | c._backward()
10. | print(f"a => {a}")

```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | >>> c => Value(data=-5.0, grad=11.0)
2. | >>> a => Value(data=5.0, grad=-11.0)

```

## Subtraction

The **subtraction** operation can be handled with `__add__` and `__neg__`. In addition, Python classes have a special method `sub` called when we use the `-` operator. as shown in

Figure 11.

[Click here to download the source code to this post](#)

```
>>> a = 5.0
>>> b = 4.0
>>> a - b
1.0
>>> a.__sub__(b)
1.0
```

<https://pyimagesearch.com/wp-content/uploads/2022/12/sub-code.png>

Figure 11: `__sub__` dunder function (source: image by the authors).

This will help us delegate the `_backward` subtraction operation to the addition and negation operations.

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. def custom_subtraction(self, other):
2.     """
3.     Subtraction operation for the Value class.
4.     Args:
5.         other (float): The other value to subtract to this one.
6.     Usage:
7.         >>> x = Value(2)
8.         >>> y = Value(3)
9.         >>> z = x - y
10.        >>> z.data
11.        -1
12.     """
13.     # This is the same as adding the negative of the other value.
14.     # We can reuse the __add__ and the __neg__ methods.
15.     return self + (-other)

16.
17. def custom_reverse_subtraction(self, other):
18.     """
19.     Reverse subtraction operation for the Value class.
20.     Args:
21.         other (float): The other value to subtract to this one.
22.     """
```

```

22. |         usage:
23. |         >>> x = Value(2)
24. |         >>> y = Value(3)
25. |         >>> z = y - x
26. |         >>> z.data
27. |         1
28. |         """
29. |         # This is the same as subtracting. We can reuse the __sub__ method.
30. |         return other + (-self)
31. |
32. |
33. | Value.__sub__ = custom_subtraction
34. | Value.__rsub__ = custom_reverse_subtraction

```

[Click here to download the source code to this post](#)

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Build a and b
2. | a = Value(data=5.0)
3. | b = Value(data=4.0)
4. |
5. | # Print the negation
6. | print(f"{a} - {b} => {(a-b)}")

```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | >>> Value(data=5.0, grad=0.0) - Value(data=4.0, grad=0.0) => Value(data=1.0, grad=0.0)

```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Subtract b from a
2. | c = a - b
3. |
4. | # Assign a global gradient to c
5. | c.grad = 11.0
6. | print(f"c => {c}")
7. |
8. | # Now apply `_backward` to c
9. | c._backward()
10. | print(f"a => {a}")
11. | print(f"b => {b}")

```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | >>> c => Value(data=1.0, grad=11.0)
2. | >>> a => Value(data=5.0, grad=11.0)
3. | >>> b => Value(data=4.0, grad=0.0)

```

► **Note:** The gradients did not flow as they were supposed to on paper. Why? Can you figure



out the answer to this?

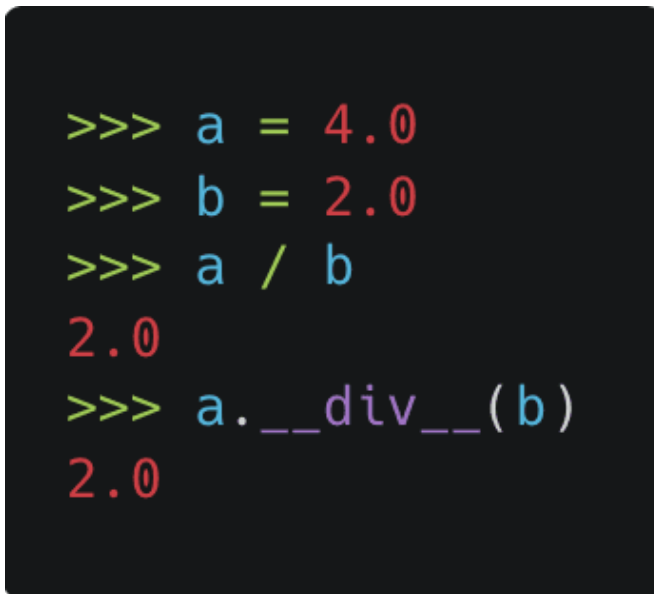
**[Click here to download the source code to this post](#)**

► **Hint:** The subtraction operation consists of more than one primitive operation: negation and addition.

We will discuss this later in the tutorial.

## Division

The **division** operation can be handled with `__mul__` and `__pow__`. In addition, Python classes have a special method `__div__` called when we use the `/` operator, as shown in **Figure 12**.



```
>>> a = 4.0
>>> b = 2.0
>>> a / b
2.0
>>> a.__div__(b)
2.0
```

<https://pyimagesearch.com/wp-content/uploads/2022/12/div-code.png>

**Figure 12:** `__div__` dunder function (source: image by the authors).

This will help us delegate the `_backward` division operation to the power operation.

→ **[Launch Jupyter Notebook on Google Colab](#)**

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | def custom_division(self, other):
2. |     """
3. |     Division operation for the Value class.
4. |     Args:
5. |         other (float): The other value to divide to this one.
```

```

6. |         Usage:
7. |             >>> x = Value(10)
8. |             >>> y = Value(5)
9. |             >>> z = x / y
10. |             >>> z.data
11. |             2
12. |         """
13. |         # Use the __pow__ method to implement division.
14. |         return self * other ** -1
15. |
16. | def custom_reverse_division(self, other):
17. |     """
18. |     Reverse division operation for the Value class.
19. |     Args:
20. |         other (float): The other value to divide to this one.
21. |     Usage:
22. |         >>> x = Value(10)
23. |         >>> y = Value(5)
24. |         >>> z = y / x
25. |         >>> z.data
26. |         0.5
27. |     """
28. |     # Use the __pow__ method to implement division.
29. |     return other * self ** -1
30. |
31. |
32. | Value.__truediv__ = custom_division
33. | Value.__rtruediv__ = custom_reverse_division

```

[Click here to download the source code to this post](#)

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Build a and b
2. | a = Value(data=6.0)
3. | b = Value(data=3.0)
4. |
5. | # Print the negation
6. | print(f"{a} / {b} => {(a/b)}")

```

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | >>> Value(data=6.0, grad=0.0) / Value(data=3.0, grad=0.0) => Value(data=2.0, grad=0.0)

```

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Divide a with b
2. | c = a / b
3. |
4. | # Assign a global gradient to c
5. | c.grad = 11.0
6. | print(f"c => {c}")
7. |
8. | # Now compute the backward pass

```

```

8. | # Now apply _backward to c
9. | c._backward()
10. | print(f"a => {a}")
11. | print(f"b => {b}")

```

[Click here to download the source code to this post](#)

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | >>> c => Value(data=2.0, grad=11.0)
2. | >>> a => Value(data=6.0, grad=3.6666666666666665)
3. | >>> b => Value(data=3.0, grad=0.0)

```

► With division, we see the same problem with gradient flow as we had seen with subtraction. Have you figured out the problem yet? 🤖

## Rectified Linear Unit

In this section, we introduce nonlinearity. ReLU is **not** a primitive function; we would need to build the function and also the `_backward` function for it.

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | def relu(self):
2. |     """
3. |     The ReLU activation function.
4. |     Usage:
5. |         >>> x = Value(-2)
6. |         >>> y = x.relu()
7. |         >>> y.data
8. |         0
9. |     """
10. |     out = Value(data=0 if self.data < 0 else self.data, _children=(self,))
11. |
12. |     def _backward():
13. |         # Local gradient:
14. |         # x = relu(a)
15. |         # dx/da = 0 if a < 0 else 1
16. |         # Global gradient:
17. |         # dy/da = dy/dx . dx/da = dy/dx . (0 if a < 0 else 1)
18. |         self.grad += out.grad * (out.data > 0)
19. |
20. |     # Set the backward function on the output node.
21. |     out._backward = _backward
22. |     return out
23. |
24. |
25. | Value.relu = relu

```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd [Click here to download the source code to this post](#)

```
1. | # Build a
2. | a = Value(data=6.0)
3. |
4. | # Print a and the negation
5. | print(f"ReLU ({a}) => {(a.relu())}")
6. | print(f"ReLU (-{a}) => {((-a).relu())}")
```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | >>> ReLU (Value(data=6.0, grad=0.0)) => Value(data=6.0, grad=0.0)
2. | >>> ReLU (-Value(data=6.0, grad=0.0)) => Value(data=0, grad=0.0)
```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | # Build a and b
2. | a = Value(3.0)
3. | b = Value(-3.0)
4. |
5. | # Apply relu on both the nodes
6. | relu_a = a.relu()
7. | relu_b = b.relu()
8. |
9. | # Assign a global gradients
10. | relu_a.grad = 11.0
11. | relu_b.grad = 11.0
12. |
13. | # Now apply `_backward`
14. | relu_a._backward()
15. | print(f"a => {a}")
16. | relu_b._backward()
17. | print(f"b => {b}")
```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | >>> a => Value(data=3.0, grad=11.0)
2. | >>> b => Value(data=-3.0, grad=0.0)
```

## The Global Backward

Until now, we have devised primitive and non-primitive (ReLU) functions with their individual `backward` methods. Each primitive can back-prop the flowing gradients to its children only.

We now have to devise a method to iterate over all such primitive nodes in a DAG (built equation) and back-propagate the gradient over the entire expression.

To make that happen, the `Value` call needs a global `backward` method. We apply the `backward` function on the last (final) node of the DAG. The function performs the following operations:

- Sorts the DAG in a topological order
- Sets the `grad` of the last node as 1.0
- Iterates over the topologically sorted graph and applies the `_backward` method of each primitive.

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | def backward(self):
2. |     """
3. |     The backward pass of the backward propagation algorithm.
4. |     Usage:
5. |         >>> x = Value(2)
6. |         >>> y = Value(3)
7. |         >>> z = x * y
8. |         >>> z.backward()
9. |         >>> x.grad
10. |         3
11. |         >>> y.grad
12. |         2
13. |     """
14. |     # Build an empty list which will hold the
15. |     # topologically sorted graph
16. |     topo = []
17. |
18. |     # Build a set of all the visited nodes
19. |     visited = set()
20. |
21. |     # A closure to help build the topologically sorted graph
22. |     def build_topo(node: "Value"):
23. |         if node not in visited:
24. |             # If node is not visited add the node to the
25. |             # visited set.
26. |
27. |             visited.add(node)
28. |
29. |             # Iterate over the children of the node that
30. |             # is being visited
31. |             for child in node._prev:
32. |                 # Apply recursion to build the topologically sorted
33. |                 # graph of the children
```

```

32.         # graph of the children
33.         build_topo(child)
34.
35.         # Only append node to the topologically sorted list
36.         # if all its children are visited.
37.         topo.append(node)
38.
39.         # Call the `build_topo` method on self
40.         build_topo(self)
41.
42.         # Go one node at a time and apply the chain rule
43.         # to get its gradient
44.         self.grad = 1.0
45.         for node in reversed(topo):
46.             node._backward()
47.
48. Value.backward = backward

```

**[Click here to download the source code to this post](#)**

→ **[Launch Jupyter Notebook on Google Colab](#)**

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. # Now create an expression that uses a lot of
2. # primitive operations
3. a = Value(2.0)
4. b = Value(3.0)
5. c = a+b
6. d = 4.0
7. e = c**d
8. f = Value(6.0)
9. g = e/f
10.
11.
12. print("BEFORE backward")
13. for element in [a, b, c, d, e, f, g]:
14.     print(element)
15.
16. # Backward on the final node will backprop
17. # the gradients through the entire DAG
18. g.backward()
19.
20.
21. print("AFTER backward")
22. for element in [a, b, c, d, e, f, g]:
23.     print(element)

```

→ **[Launch Jupyter Notebook on Google Colab](#)**

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. >>> BEFORE backward
2. >>> Value(data=2.0, grad=0.0)

3. >>> Value(data=3.0, grad=0.0)
4. >>> Value(data=5.0, grad=0.0)
5. >>> 4.0
6. >>> Value(data=625.0, grad=0.0)
7. >>> Value(data=6.0, grad=0.0)
8. >>> Value(data=104.16666666666666, grad=0.0)
~

```

```

9. |
10. | >>> AFTER backward
11. | >>> Value(data=2.0, grad=83.33333333333333)
12. | >>> Value(data=3.0, grad=83.33333333333333)
13. | >>> Value(data=5.0, grad=83.33333333333333)
14. | >>> 4.0
15. | >>> Value(data=625.0, grad=0.16666666666666666)
16. | >>> Value(data=6.0, grad=-17.361111111111111)
17. | >>> Value(data=104.16666666666666, grad=1.0)

```

[Click here to download the source code to this post](#)

Remember the problem we had with `__sub__` and `__div__` ? The gradients did not backpropagate according to the rules of calculus. There is nothing wrong with implementing the `_backward` function.

However, the two operations ( `__sub__` and `__div__` ) are built with more than one primitive operation ( `__neg__` and `__add__` for `__sub__` ; `__mul__` and `__pow__` for `__div__` ).

This creates an intermediate node that prohibits the gradients from flowing to the children properly (remember, `_backward` is not supposed to backpropagate the gradients through the entire DAG).

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Solve the problem with subtraction
2. | a = Value(data=6.0)
3. | b = Value(data=3.0)
4. |
5. | c = a - b
6. | c.backward()
7. | print(f"c => {c}")
8. | print(f"a => {a}")
9. | print(f"b => {b}")

```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | c => Value(data=3.0, grad=1.0)
2. | a => Value(data=6.0, grad=1.0)
3. | b => Value(data=3.0, grad=-1.0)

```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Solve the problem with division
2. | a = Value(data=6.0)

```

```

3. | b = Value(data=3.0)
4. |
5. | c = a / b
6. | c.backward()
7. | print(f"c => {c}")
8. | print(f"a => {a}")
9. | print(f"b => {b}")

```

[Click here to download the source code to this post](#)

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | >>> c => Value(data=2.0, grad=1.0)
2. | >>> a => Value(data=6.0, grad=0.3333333333333333)
3. | >>> b => Value(data=3.0, grad=-0.6666666666666666)

```

## Build a Multilayer Perceptron with micrograd

What good does it do if we just build the `Value` class and not build a Neural Network with it?

In this section, we build a very simple Neural Network (a Multilayer Perceptron) and use it to model a simple dataset.

### Module

This is the parent class. The `Module` class has two methods:

- `zero_grad` : This is used to zero out all the gradients of the parameters.
- `parameters` : This function is built to be overwritten. This would eventually get us the parameters of the **neurons**, **layers**, and the **mlp**.

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | class Module(object):
2. |     """
3. |     The parent class for all neural network modules.
4. |     """
5. |
6. |     def zero_grad(self):
7. |         # Zero out the gradients of all parameters.
8. |         for p in self.parameters():

```



```

8. |         for p in self.parameters():
9. |             p.grad = 0
10. |
11. |     def parameters(self):
12. |         # Initialize a parameters function that all the children will
13. |         # override and return a list of parameters.
14. |         return []

```

[Click here to download the source code to this post](#)

## Neuron

This serves as the unit of our Neural Network upon which the entire architecture is built. It has a list of weights and a bias. The function of a Neuron is shown in **Figure 13**.

$$\text{Neuron}(x) = \sum_{i=1}^{\text{number of inputs}} x_i \times w_i + b$$

(<https://lh3.googleusercontent.com/72gIWKOj7BBnpzaBltJUJyDOatKNp-ejGXxuFTAKplzHeUrCIN1dJMCDAPuJqxPdf3ueGS0zxS4dbSS2mZrD4QWTqO8JKZKgiSJOBHDO1fSVIHTuYTgMrN-V4vnxBNSHNzmb7MIECVZTJcD9CV9TJKLR91ztWcZr5HhvwzBVQupl9AMT8D-DYXKN1ZQ>)

**Figure 13:** Anatomy of a neuron (source: image by the authors).

→ [Launch Jupyter Notebook on Google Colab](#)

```

Automatic Differentiation Part 2: Implementation Using Micrograd
1. | class Neuron(Module):
2. |     """
3. |     A single neuron.
4. |     Parameters:
5. |         number_inputs (int): number of inputs
6. |         is_nonlinear (bool): whether to apply ReLU nonlinearity
7. |         name (int): the index of neuron
8. |     """
9. |
10. |
11. |     def __init__(self, number_inputs: int, name, is_nonlinear: bool = True):
12. |         # Create weights for the neuron. The weights are initialized
13. |         # from a random uniform distribution.
14. |         self.weights = [Value(data=random.uniform(-1, 1)) for _ in range(number_inputs)]
15. |
16. |         # Create bias for the neuron.
17. |         self.bias = Value(data=0.0)

```

```

16. |         self.bias = value(data=0.0)
17. |         self.is_nonlinear = is_nonlinear
18. |
19. |         self.name = name
20. |
21. |     def __call__(self, x: List["Value"]) -> "Value":
22. |         # Compute the dot product of the input and the weights. Add the
23. |         # bias to the dot product.
24. |         act = sum(
25. |             ((wi * xi) for wi, xi in zip(self.weights, x)),
26. |             self.bias
27. |         )
28. |
29. |         # If activation is mentioned, apply ReLU to it.
30. |         return act.relu() if self.is_nonlinear else act
31. |
32. |     def parameters(self):
33. |         # Get the parameters of the neuron. The parameters of a neuron
34. |         # is its weights and bias.
35. |         return self.weights + [self.bias]
36. |
37. |     def __repr__(self):
38. |         # Print a better representation of the neuron.
39. |         return f"Neuron {self.name} (Number={len(self.weights)}, Non-Linearity={'ReLU' if
self.is_nonlinear else 'None'})"

```

**[Click here to download the source code to this post](#)**

→ **[Launch Jupyter Notebook on Google Colab](#)**

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | x = [2.0, 3.0]
2. | neuron = Neuron(number_inputs=2, name=1)
3. | print(neuron)
4. | out = neuron(x)
5. | print(f"Output => {out}")

```

→ **[Launch Jupyter Notebook on Google Colab](#)**

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | >>> Neuron 1 (Number=2, Non-Linearity=ReLU)
2. | >>> Output => Value(data=2.3063230206881347, grad=0.0)

```

## Layer

A layer is built of a number of Neurons.

→ **[Launch Jupyter Notebook on Google Colab](#)**

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | class Layer(Module):
2. |     """

```

[Click here to download the source code to this post](#)

```
3. | A layer of neurons.
4. | Parameters:
5. |     number_inputs (int): number of inputs
6. |     number_outputs (int): number of outputs
7. |     name (int): index of the layer
8. | """
9. |
10. |
11. | def __init__(self, number_inputs: int, number_outputs: int, name: int, **kwargs):
12. |     # A layer is a list of neurons.
13. |     self.neurons = [
14. |         Neuron(number_inputs=number_inputs, name=idx, **kwargs) for idx in
15. |         range(number_outputs)
16. |     ]
17. |     self.name = name
18. |     self.number_outputs = number_outputs
19. |
20. | def __call__(self, x: List["Value"]) -> Union[List["Value"], "Value"]:
21. |     # Iterate over all the neurons and compute the output of each.
22. |     out = [n(x) for n in self.neurons]
23. |     return out if self.number_outputs != 1 else out[0]
24. |
25. | def parameters(self):
26. |     # The parameters of a layer is the parameters of all the neurons.
27. |     return [p for n in self.neurons for p in n.parameters()]
28. |
29. | def __repr__(self):
30. |     # Print a better representation of the layer.
31. |     layer_str = "\n".join(f'    - {str(n)}' for n in self.neurons)
32. |     return f"Layer {self.name} \n{layer_str}\n"
```

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | x = [2.0, 3.0]
2. | layer = Layer(number_inputs=2, number_outputs=3, name=1)
3. | print(layer)
4. | out = layer(x)
5. | print(f"Output => {out}")
```

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | >>> Layer 1
2. | >>>   - Neuron 0 (Number=2, Non-Linearity=ReLU)
3. | >>>   - Neuron 1 (Number=2, Non-Linearity=ReLU)
4. | >>>   - Neuron 2 (Number=2, Non-Linearity=ReLU)
5. |
6. | >>> Output => [Value(data=0, grad=0.0), Value(data=1.1705131190055296, grad=0.0),
7. |               Value(data=3.0608608028649344, grad=0.0)]
```

→ [Launch Jupyter Notebook on Google Colab](#)

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | x = [2.0, 3.0]
2. | layer = Layer(number_inputs=2, number_outputs=1, name=1)
```

```

3. | print(layer)
4. | out = layer(x)
5. | print(f"Output {out}")

```

[Click here to download the source code to this post](#)

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | >>> Layer 1
2. | >>>      - Neuron 0 (Number=2, Non-Linearity=ReLU)
3. |
4. | >>> Output => Value(data=2.3123867684232247, grad=0.0)

```

## Multilayer Perceptron

A Multilayer Perceptron (MLP) is built of a number of Layers.

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | class MLP(Module):
2. |     """
3. |     The Multi-Layer Perceptron (MLP) class.
4. |     Parameters:
5. |         number_inputs (int): number of inputs.
6. |         list_number_outputs (List[int]): number of outputs in each layer.
7. |     """
8. |
9. |     def __init__(self, number_inputs: int, list_number_outputs: List[int]):
10. |         # Get the number of inputs and all the number of outputs in
11. |         # a single list.
12. |         total_size = [number_inputs] + list_number_outputs
13. |
14. |         # Build layers by connecting each layer to the previous one.
15. |         self.layers = [
16. |             # Do not use non linearity in the last layer.
17. |             Layer(
18. |                 number_inputs=total_size[i],
19. |                 number_outputs=total_size[i + 1],
20. |                 name=i,
21. |                 is_nonlinear=i != len(list_number_outputs) - 1
22. |             )
23. |             for i in range(len(list_number_outputs))
24. |         ]
25. |
26. |     def __call__(self, x: List["Value"]) -> List["Value"]:
27. |         # Iterate over the layers and compute the output of
28. |         # each sequentially.
29. |         for layer in self.layers:
30. |             x = layer(x)
31. |         return x
32. |
33. |     def parameters(self):
34. |         # Get the parameters of the MLP

```

```

34. |         # Get the parameters of the MLP
35. |         return [p for layer in self.layers for p in layer.parameters()]
36. |
37. |     def __repr__(self):
38. |         # Print a better representation of the MLP.
39. |         mlp_str = "\n".join(f' - {str(layer)}' for layer in self.layers)
40. |         return f"MLP of \n{mlp_str}"

```

**[Click here to download the source code to this post](#)**

→ **[Launch Jupyter Notebook on Google Colab](#)**

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | x = [2.0, 3.0]
2. | mlp = MLP(number_inputs=2, list_number_outputs=[3, 3, 1])
3. | print(mlp)
4. | out = mlp(x)
5. | print(f"Output => {out}")

```

→ **[Launch Jupyter Notebook on Google Colab](#)**

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | >>> MLP of
2. | >>>   - Layer 0
3. | >>>     - Neuron 0 (Number=2, Non-Linearity=ReLU)
4. | >>>     - Neuron 1 (Number=2, Non-Linearity=ReLU)
5. | >>>     - Neuron 2 (Number=2, Non-Linearity=ReLU)
6. |
7. | >>>   - Layer 1
8. | >>>     - Neuron 0 (Number=3, Non-Linearity=ReLU)
9. | >>>     - Neuron 1 (Number=3, Non-Linearity=ReLU)
10. | >>>     - Neuron 2 (Number=3, Non-Linearity=ReLU)
11. |
12. | >>>   - Layer 2
13. | >>>     - Neuron 0 (Number=3, Non-Linearity=None)
14. |
15. | >>> Output => Value(data=-0.3211612402687316, grad=0.0)

```

## **Train the MLP**

In this section, we will create a small dataset and try to understand how to model the dataset with our MLP.

→ **[Launch Jupyter Notebook on Google Colab](#)**

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Build a dataset
2. | xs = [
3. |     [0.5, 0.5, 0.70],
4. |     [0.4, -0.1, 0.5],
5. |     [-0.2, -0.75, 1.0],
6. | ]

```

```
7. | ys = [0.0, 1.0, 0.0]
```

**[Click here to download the source code to this post](#)**  
**→ [Launch Jupyter Notebook on Google Colab](#)**

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | # Build an MLP
2. | mlp = MLP(number_inputs=3, list_number_outputs=[3, 3, 1])
```

In the following code snippet, we define three functions:

- `forward` : The forward function takes the `mlp` and the inputs. The inputs are forwarded through the `mlp`, and we obtain the predictions from the `mlp`.
- `compute_loss` : We have ground truth and predictions. This function computes the loss between the two. We will optimize our `mlp` to make the loss go to zero.
- `update_mlp` : In this function, we update the parameters (weights and biases) of our `mlp` with the gradient information.

**→ [Launch Jupyter Notebook on Google Colab](#)**

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | def forward(mlp: "MLP", xs: List[List[float]]) -> List["Value"]:
2. |     # Get the predictions upon forwarding the input data through
3. |     # the mlp
4. |     ypred = [mlp(x) for x in xs]
5. |     return ypred
```

**→ [Launch Jupyter Notebook on Google Colab](#)**

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | def compute_loss(ys: List[int], ypred: List["Value"]) -> "Value":
2. |     # Obtain the L2 distance of the prediction and ground truths
3. |     loss = sum(
4. |         [(ygt - yout)**2 for ygt, yout in zip(ys, ypred)]
5. |     )
6. |     return loss
```

**→ [Launch Jupyter Notebook on Google Colab](#)**

---

Automatic Differentiation Part 2: Implementation Using Micrograd

```
1. | def update_mlp(mlp: "MLP"):
2. |     # Iterate over all the layers of the MLP
3. |     for layer in mlp.layers:
4. |         # Iterate over all the neurons of each layer
5. |         for neuron in layer.neurons:
```

```

6. |         # Iterate over all the weights of each neuron
7. |         for weight in neuron.weights:
8. |             # Update the data of the weight with the
9. |             # gradient information.
10. |             weight.data -= (1e-2 * weight.grad)
11. |         # Update the data of the bias with the
12. |         # gradient information.
13. |         neuron.bias.data -= (1e-2 * neuron.bias.grad)

```

**[Click here to download the source code to this post](#)**

→ **[Launch Jupyter Notebook on Google Colab](#)**

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Define the epochs for which we want to run the training process.
2. | epochs = 50
3. |
4. | # Define a loss list to help log the loss.
5. | loss_list = []
6. |
7. | # Iterate each epoch and train the model.
8. | for idx in range(epochs):
9. |     # Step 1: Forward the inputs to the mlp and get the predictions
10. |     ypred = forward(mlp, xs)
11. |     # Step 2: Compute Loss between the predictions and the ground truths
12. |     loss = compute_loss(ys, ypred)
13. |     # Step 3: Ground the gradients. These accumulate which is not desired.
14. |     mlp.zero_grad()
15. |     # Step 4: Backpropagate the gradients through the entire architecture
16. |     loss.backward()
17. |     # Step 5: Update the mlp
18. |     update_mlp(mlp)
19. |     # Step 6: Log the loss
20. |     loss_list.append(loss.data)
21. |     print(f"Epoch {idx}: Loss {loss.data: 0.2f}")

```

→ **[Launch Jupyter Notebook on Google Colab](#)**

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | Epoch 0: Loss  0.95
2. | Epoch 1: Loss  0.89
3. | Epoch 2: Loss  0.81
4. | Epoch 3: Loss  0.74
5. | Epoch 4: Loss  0.68
6. | Epoch 5: Loss  0.63
7. | Epoch 6: Loss  0.59
8. | .
9. | .
10. | Epoch 47: Loss  0.24
11. | Epoch 48: Loss  0.23
12. | Epoch 49: Loss  0.22

```

→ **[Launch Jupyter Notebook on Google Colab](#)**

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Plot the loss
2. | plt.plot(loss_list)

```


```

3. | plt.grid()
4. | plt.ylabel("Loss")
5. | plt.xlabel("Epoch")
6. | plt.show()

```

[Click here to download the source code to this post](#)

The loss plot is shown in **Figure 14**.



<https://pyimagesearch.com/wp-content/uploads/2022/12/plot-loss.png>

**Figure 14:** Loss plot (source: image by the authors).

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | # Inference
2. | pred = mlp(xs[0])
3. | ygt = ys[0]
4. |
5. | print(f"Prediction => {pred.data: 0.2f}")
6. | print(f"Ground Truth => {ygt: 0.2f}")

```

→ [Launch Jupyter Notebook on Google Colab](#)

Automatic Differentiation Part 2: Implementation Using Micrograd

```

1. | >>> Prediction => 0.14
2. | >>> Ground Truth => 0.00

```

**What's next? We recommend [PyImageSearch](#) University**