

# sketch2sky

What I Cannot Create, I Do Not Understand —Richard Feynman And I



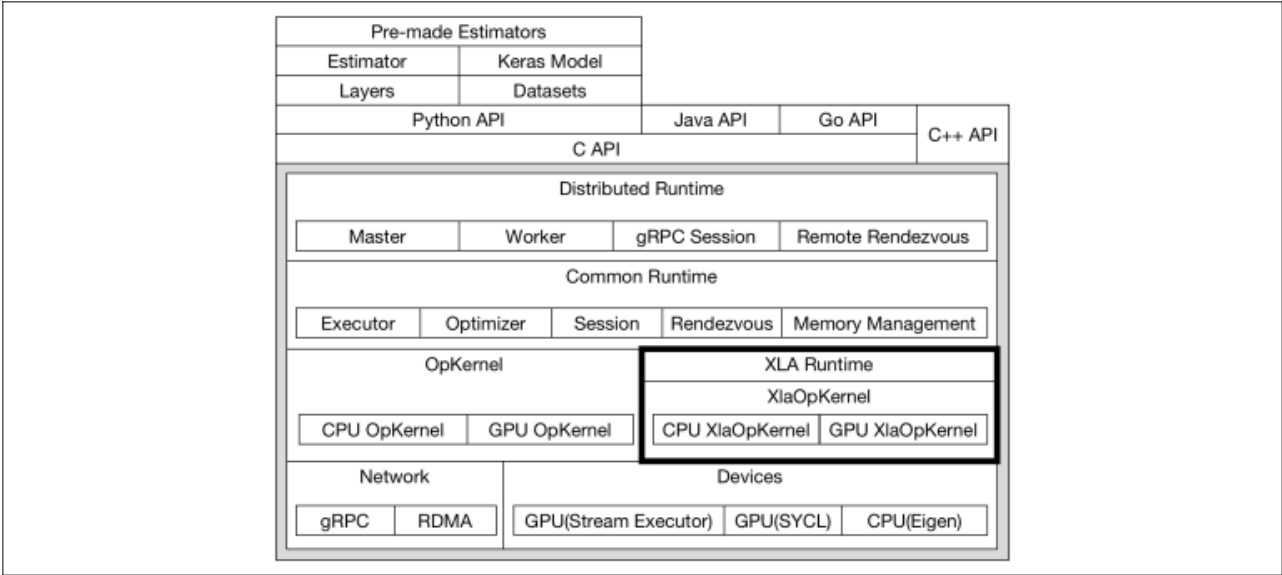
Primary Menu

## Tensorflow XLA Service Buffer优化详解

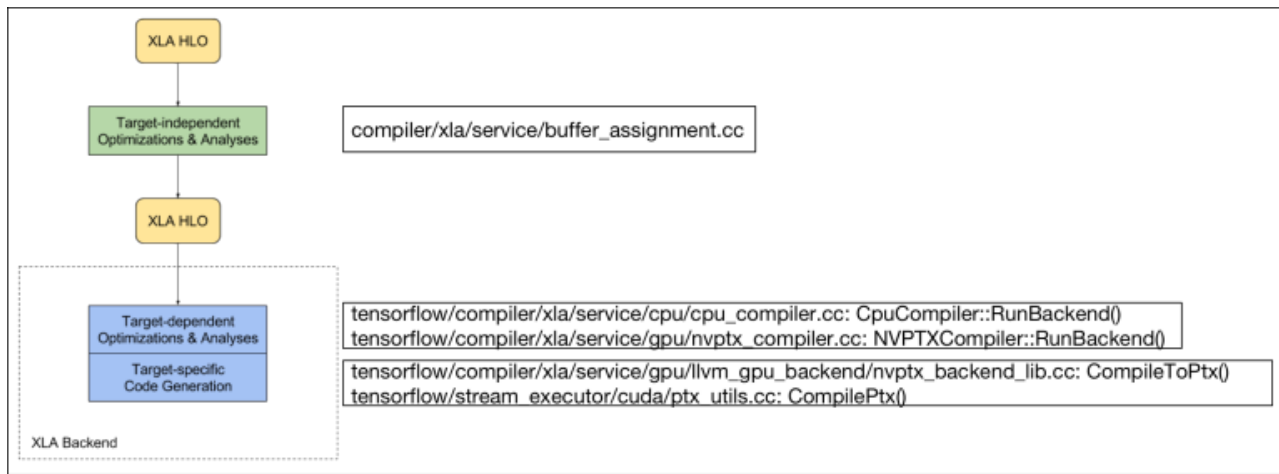
1820 Jiang XIAO

2020年3月8日 at pm6:16 (last edited 2020年6月20日 at am9:52)

下图是Tensorflow 架构图，以及XLA在Tensorflow中的位置



通过层层转换(参考[Tensorflow XLA Service 详解 I](#))，Graph在进入XLA Service前已经被表达为HloModule的形式，而作为图编译器的核心，XLA Service就负责将HloModule表达的计算图编译为可以直接在不同硬件平台(Backend)执行的程序，而编译的核心，就是优化代码，包括设备无关的优化和设备相关的优化：



1. 优化HloModule所表示的计算图，并将其转化为LLVM HLO
2. 基于LLVM，生成硬件相关的二进制

作为通用的编译器框架，LLVM 会对LLVM HLO做大量的优化，生成高效的Target binary, 所以作为XLA的开发者，主要关注**阶段1**的优化: 设备无关的图优化算法。Compiler是XLA适配硬件的接口类，每个适配XLA的硬件都必须实现其中的方法。尤其是RunBackend(), 参考[Tensorflow XLA Service 详解 I](#)一文，该接口是进行图优化和编译的入口，统领整个优化和编译过程。

同时，XLA Service还实现了一组通用的优化方法(各种Schedule策略，各种Memory优化算法)供各个硬件平台的编译器使用，当然，主要是供给RunBackend()调用。以XLA的GPU平台的编译优化流程为例:

```

1.  NVPTXCompiler::RunBackend()
2.    hlo_schedule = GpuHloSchedule::Build(*module, *stream_assignment, pointer_size_)
3.    BufferAssigner::Run(hlo_schedule->ConsumeHloOrdering()...)
4.    entry_computation->Accept(&ir_emitter)
5.    CompileToPtx()
6.    CompilePtxOrGetCachedResult()
  
```

- 1- 从XLA Service通用层中选择适合GPU的Schedule策略
- 3- 基于Schedule策略，进行设备无关的Buffer优化，主要关注尽可能的减少Buffer的大小。注意，这里是设备无关的优化，是无法利用硬件Memory特性的。
- 4- 将HloModule转化为LLVM IR
- 5,6- 利用LLVM框架，将LLVM IR编译为二进制代码。

**本文主要关注-3-，是XLA优化的核心。**

对BufferAssigner::Run()进一步分解。

```

1.  NVPTXCompiler::RunBackend()
2.    hlo_schedule = GpuHloSchedule::Build(*module, *stream_assignment, pointer_size_)
3.    //this analysis figures out which temp buffers are required to run the computation
4.    BufferAssigner::Run(hlo_schedule->ConsumeHloOrdering()...)
5.    assigner.CreateAssignment(HloModule, hlo_ordering, buffer_size)
6.    liveness = BufferLiveness::Run()
7.    assignment = new BufferAssignment(module, liveness, ...)
8.    set<LogicalBuffer*> colocated_buffers
9.    set<BufferAllocation::Index> colocated_allocations
10.   vector<ColocatedBufferSet> colocated_buffer_sets
11.   BuildColocatedBufferSets(&colocated_buffer_sets)
  
```

```

12.     colorer_ (assignment->liveness())
13.     AssignColocatedBufferSets(colocated_buffer_sets, assignment, &colocated_buffers
14.     GatherComputationsByAllocationType(module, &thread_local_computations, &global_
15.     for computation : global_computations:
16.         AssignBufferForComputation(computation, false, buffers_to_assign_sequentially
17.     AssignBuffersWithSequentialOrdering(buffers_to_assign_sequentially, assignment
18.     for computation : thread_local_computations:
19.         AssignBuffersForComputation()
20.     for buffer : assignment->liveness().maybe_live_out_buffers():
21.         if assignment->HasAllocation(buffer):
22.             assignment->GetMutableAssignedAllocation(buffer).set_mayby_live_out(true)
23.     assignment->CombineTempAllocations()
24.     return std::move(assignment)
25.     entry_computation->Accept(&ir_emitter)
26.     CompileToPtx()
27.     CompilePtxOrGetCachedResult()

```

-6- 进行BufferLiveness分析, 分析整个HloModule的LogicalBuffer的干涉关系, 为后续优化提供依据

-11- BuildColocatedBufferSets, 依据Bufferliveness的分析, 将所有的LogicalBuffer分为几个Bufferset, 并进行初步的Set融合, 每个Bufferset内

参照注释, colocated buffer sets, 每个set都是一组可以共享BufferAllocation的LogicalBuffer, 共享Allocation, 意味着共享同一块物理内存(GPU的显存)

-12- colorer\_ 缺省被赋值为BufferLiveness::DefaultColorer(), 所有的LogicalBuffer实例的color都会被设置为0

-13- AssignColocatedBufferSets, 为Bufferset分配BufferAllocation, 每一个LogicalBufferSet 与其关联, 这里用到了buffer\_size\_ 这个函数来判断一个LogicalBuffer大小, LogicalBuffer的大小要和相应的Allocation一样 具体可以参

考t12xla/write\_op.cc t12xla/ir\_op.cc xla/client/builder.cc KConditional代码, 可以看到明显的要求合body的

Shape要一致。通过TEST用例也能确认

-14- GatherComputationsByAllocationType, 根据内含的LogicalBuffer的属性, 将Allocation分为global和thread local两类, 这部分是理解显存优化的关键, 后文详细

-16- AssignBufferForComputation, 关联Allocation和XlaComputation, 此调用点只针对global,temp buffer被收集到buffers\_to\_assign\_sequentially, 延后处理,

## BufferLiveness::Run()

### Liveness

From Wikipedia, the free encyclopedia

In [concurrent computing](#), **liveness** refers to a set of properties of [concurrent systems](#), that require a system to make progress despite the fact that its concurrently executing components ("processes") may have to "take turns" in [critical sections](#), parts of the program that cannot be simultaneously run by multiple processes.<sup>[1]</sup> Liveness guarantees are important properties in operating systems and [distributed systems](#).<sup>[2]</sup>

所以, BufferLiveness, 就是获取Buffer生命周期关系, 以便决定Buffer复用策略. 源码整体调用栈如下, 该阶段进一步也分3个过程: LogicalBufferAnalysis->TuplePointsToAnalysis->BufferLiveness

```

1.     BufferAssigner::Run()
2.         assigner::CreateAssignment()
3.         liveness = BufferLiveness::Run(module, std::move(hlo_ordering)) //class BufferLiv
4.         liveness = new BufferLiveness()
5.         liveness->Analyze()
6.         points_to_analysis_ = TuplePointsToAnalysis::Run()
7.         logical_buffer_analysis = LogicalBufferAnalysis::Run()
8.         analysis = new LogicalBufferAnalysis()
9.         analysis.Analyze()

```

```

10.         return analysis
11.         analysis = new TuplePointsToAnalysis(logical_buffer_analysis)
12.         analysis.Analyze()
13.         return analysis
14.         maybe_live_out_buffers_ = points_to_analysis->GetPointsToSet(root).CreateFla
15.         return liveness
16.         assignment(new BufferAssignment(module, std::move(liveness))

```

## -1- Memory优化入口

-3-16- 获取BufferLiveness, 对比-16-, 获取到的BufferLiveness会用于支撑Buffer优化决策, 用OOP的方式优雅的实现pipeline

-8- 获取LogicalBufferAnalysis

-9- 根据HloInstruction的Shape, 构造一组相应的LogicalBuffer, Shape本身是一个树结构,表示一个HloInstruction的输出形状. 一个LogicalBuffer对应一个Shape的一个节点(Subshape), 表示一块逻辑buffer, 用pair<HloInstruction\*, ShapeIndex>来表示一个LogicalBuffer。

-10- 作为Liveness分析的原材料-LogicalBuffer已经构造完毕, 存储下来, 并返回analysis, 准备进入下一阶段, 进行TuplePointsToAnalysis.

-11- 用存储有所有LogicalBuffer信息的LogicalBufferAnalysis实例构造TuplePointsToAnalysis实例, **Tuple, 用来描述Buffer树状结构的方式**,  $E = \{\%1, \%2, \{\%3, \%4\}\}$  表示这样一个树:深为3, 深1的节点有一个, 树的根, 深为2的节点有3个, %1, %2, 没体现名字的节点暂且叫“X”, 深为3的节点有2个, %3, %4, 这两个节点是上一层中“X”的子节点. **PointsTo, “指向”**, 在前面的例子中, root的“PointsTo”就是%1, %2和“X”, “X”的“PointsTo”就是%3和%4, 所以 **TuplePointsToAnalysis就是分析整个计算图中的LogicalBuffer依赖关系并存储在PointsToSet中**, 后面会详细分析

-12- 执行分析逻辑.

-13- 对计算图中的LogicalBuffer的依赖关系分析完毕, 存储下来, 并返回analysis, 准备进入下一阶段, 进行BufferLiveness

-14- 用TuplePointsToAnalysis实例获取root的alias\_buffer, 也就是潜在的需要传出的HloModule的Buffer, 存储在maybe\_live\_out\_buffers\_中.

-15- 返回liveness实例, TuplePointsToAnalysis实例会存储LogicalBuffer的依赖关系, 但BufferLiveness并不会存储每一个LogicalBuffer的“liveness”, 而是基于TuplePointsToAnalysis封装了一组判断特定LogicalBuffer的函数.

-16- 将BufferLiveness实例传入构造LogicalBuffer与BufferAllocation映射关系的BufferAssignment实例.

上面是整个BufferLiveness分析的主干流程, 下面介绍一些关键细节

## LogicalBufferAnalysis

```

1.     class LogicalBufferAnalysis : public DfsHloVisitorWithDefault

```

LogicalBufferAnalysis “is a” DfsHloVistorWithDefault, 后者是HLO层提供的以**visitor模式**(没错, 就是设计模式中的Visitor模式)遍历HloModule/HloComputation/HloInstruction的接口. 是XLA处理HLO优化的基础工具, 同时对于二次开发而言, 也是理解代码以及Debug的趁手工具. 对于一个“DfsHloVisitorWithDefault”, 我们主要关心就是其各种“Handle”的实现, 同时, 也关心LogicalBufferAnalysis作为子类更加specific的部分.

```

1.     //tensorflow/compiler/xla/service/logical_buffer_analysis.h
2.     Status DefaultAction(HloInstruction* hlo_instruction) override;
3.     Status HandleTuple(HloInstruction* tuple) override;
4.     Status HandleGetTupleElement(HloInstruction* get_tuple_element) override;
5.     Status HandleBitcast(HloInstruction* bitcast) override;
6.     Status HandleDomain(HloInstruction* domain) override;
7.     Status HandleCopy(HloInstruction* copy) override;

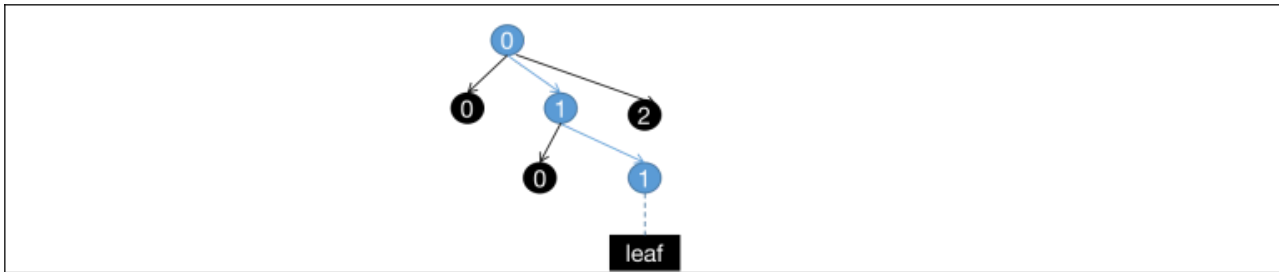
```

```

8.   Status HandleRecvDone(HloInstruction* recv_done) override;
9.   Status HandleSend(HloInstruction* send) override;
10.  Status HandleTupleSelect(HloInstruction* tuple_select) override;

```

-2- 缺省Handle, 遍历HloInstruction的每一个Subshape, 调用NewLogicalBuffer()为其构造LogicalBuffer, 每个LogicalBuffer都有LogicalBufferAnalysis内唯一的id, 由于此时一个Compilation Graph内也只有一个LogicalBufferAnalysis实例, 也就是说这个id是HloModule内唯一的. 构造的LogicalBuffer会统一存储在'logical\_buffers\_'和'output\_buffers\_'中, 二者的主要区别是索引方式不同, logical\_buffers\_以unique\_ptr的方式存储原始的LogicalBuffer实例, 后者以std::pair<HloInstruction\*, ShapeIndex>为Key, 进行Hash查找, 可快速进行LogicalBuffer检索。如果查看了ShapeIndex的实现细节, 可以看到其Index并不是一个单独的"Id"而是一个"absl::InlinedVector<int64, 2> indices\_", 这里indices\_中的每一个int64都代表HloInstruction对应的Shape多叉树的一个节点。ShapeIndex只负责Index的部分, 至于"Leaf"是什么, 由用户决定, 这里, LogicalBufferAnalysis的Leaf就是LogicalBuffer, 使用hash建立ShapeIndex和LogicalBuffer的关系。



-3:11- 显然, 9个"non default"的Handle都是因为缺省DefaultAction的处理方式不适用于自己所负责HloOpcode, 里面的注释很清晰, 嗯, 如果已经理解TuplePointsToAnalysis的话. LogicalBufferAnalysis和TuplePointsToAnalysis分属不同文件主要是由于使用了Visitor的处理方式, 从单个HloInstruction或者单个Buffer的角度, 设计上一脉相承, 通常要一起分析. 这里, 从理解代码的角度, 可以重点关注以下几个Handle:

- 3- HandleTuple, NewLogicalBuffer(tuple, /\*index=\*/{}), 分配以"{}"为索引的LogicalBuffer, "{}", 看起来并不符合"ShapeIndex"的含义, 看起来没有任何"形状"信息, 实际上, 这里只需了解: 以"{}"为索引是因为kTuple只需要一个"top-level buffer", 至于什么是"top-level buffer"呢? 我在TuplePointsToAnalysis中详细分析.
- 5:6- HandleBitcast 和 HandleDomain, 不构造任何LogicalBuffer.
- 7- HandleCopy, 单纯构造一个新LogicalBuffer, 同样, 一个"top-level buffer"
- 8- HandleRecvDone, 两个LogicalBuffer, 位于{}的top-level buffer用于以buffer alias的形式存储接收的数据, 位于{1}的buffer用于存储交互token
- 9- HandleSend, 3个LogicalBuffer, 位于{}的top-level buffer用于以buffer alias的形式存储即将发送的数据, 位于{1}的buffer用于存储交互context, 位于{2}的buffer用于存储交互token

了解了Visitor的部分, 现在来看下Specific的部分.

```

Status LogicalBufferAnalysis::Analyze();
// We filter out fusion computations, and get to them through fusion
// instructions. This is because it's possible to have orphaned (unreachable)
// fusion computations, and we don't want to try to assign buffers to those.
std::vector<HloInstruction*> fusion_instructions;
for (auto* computation : module->MakeNonfusionComputations()) {
  computation->Accept(this);
  for (auto* instruction : computation->instructions()) {
    if (instruction->opcode() != HloOpcode::kFusion) {
      continue;
    }
    GatherFusionInstructions(instruction, &fusion_instructions);
  }
}
for (auto* instruction : fusion_instructions) {
  instruction->fused_expression_root()->Accept(this);
}

```

上面这个函数是执行遍历的入口, 其中对kFusion进行了filter out的处理, 注释写的很清楚, 对计算图进行“剪枝”, 不给 unreachable fusion computations分配LogicalBuffer, 避免浪费内存, 这样的处理在XLA很多代码中都有遇到. 这里有一个经典问题, TF runtime前期进行图处理的时候, 已经将unreachable的节点去掉了, 为什么XLA这里还要处理一遍呢? 理解的关键点在Op->HloInstruction的映射, XLA用少数几个kOpcode就组合成了几百个Op类型, 虽然TF前期已经将unreachable op去掉了, 但XLA将Op组成的计算图转化为HloInstruction组成的计算图之后, 又会出现新的unreachable节点, 这里的节点不再是Op, 而是粒度更小HloInstruction. 但即便粒度不同, TF传统引擎和XLA执行的很多图优化算法是类似的, 而二者并没有在更高的层次上进行图优化算法的抽象, 导致了这种功能重复实现. 有些同学认为从OOP的角度, 这里存在优化空间, 我的理解是恰恰相反, 这里不但不需要进一步抽象引入依赖, XLA最好将自己完全从TF引擎解耦, 以外部插件的形式接入到TF核心层. 大家感兴趣的可以给我留言讨论.

## TuplePointsToAnalysis

理解这个Analysis的主要难点在于理解“**TuplePointsTo**”的含义, 理解了含义, 其内部的设计逻辑就一目了然. 对于一个HloInstruction而言, “**PointsTo**”是表达的内容: 该HloInstruction所依赖的Buffer, “**Tuple**”是表达形式: Buffer和Buffer间的拓扑关系用“{..., {...}}”的形式表达.

源码中, 下面这段注释有很好的解释:

```
// For the subshape at the given index (where index is defined as in
// ShapeUtil::GetSubshape) this method returns the set of HLO instructions
// which may produce the tuple subshape at that index. For example, given:
//
// %tuple1 = tuple(...)
// %tuple2 = tuple(...)
// %select = select(%tuple1, %tuple2)
// %nested_tuple = tuple(%select, %tuple1)
//
// These are the values for tuple_sources() for the PointsToSet of
// %nested_tuple:
//
// tuple_sources({}) = {%nested_tuple}
// tuple_sources({0}) = {%tuple1, %tuple2}
// tuple_sources({1}) = {%tuple1}
//
// tuple_sources() at the index of an array shape (not a tuple) returns the
// empty set. The instructions in the set returned by tuple_sources
// necessarily are either Tuple instructions, constants, or parameters.
```





转化成图示, 左图就是上例图示, 其中蓝色箭头就是"PointsTo"的含义, 表示的buffer索引, 所有的箭头和它对应的buffer们就是kTuple的"PointsToSet", 个人认为, 这里使用kTupleSelect并不适合作为例子, 因为它本身不是大多数情况使用的DefaultAction, 而是和kTuple一样对输入的Buffer进行alias的操作, 所以简单画了右图, 希望能表达清楚。

介绍了"TuplePointsTo"的含义, 接下来从源码角度分析其实现细节。

```
// A class describing the source(s) of the Buffer(s) contained in the output of
// a particular HLO instruction. The structure of PointsToSet mirrors the
// structure of the instruction's shape, which may be an arbitrary tree (eg, a
// nested tuple). Each node in this tree corresponds to a single buffer in the
// instruction's output and contains the set of Buffers which might define
// the corresponding buffer.
class PointsToSet {
private:
    struct Elem {
        BufferList buffers;
        SourceSet tuple_sources;
    };
    ShapeTree<Elem> tree_;
};
// DFS visitor that performs tuple points-to analysis. This analysis determines
// the potential sources of each buffer in each instruction's output.
class TuplePointsToAnalysis : public DfsHloVisitorWithDefault {
    // Information kept per instruction
    struct PerInstruction {
        std::unique_ptr<PointsToSet> points_to_set;
        // Empirically, ~92% of instructions have 1
        // instruction_defined_buffer, and 99% have 0 or 1
        BufferDefinitionVector instruction_defined_buffers;
    };

    // The logical buffers for this module.
    const std::unique_ptr<LogicalBufferAnalysis> logical_buffer_analysis_;

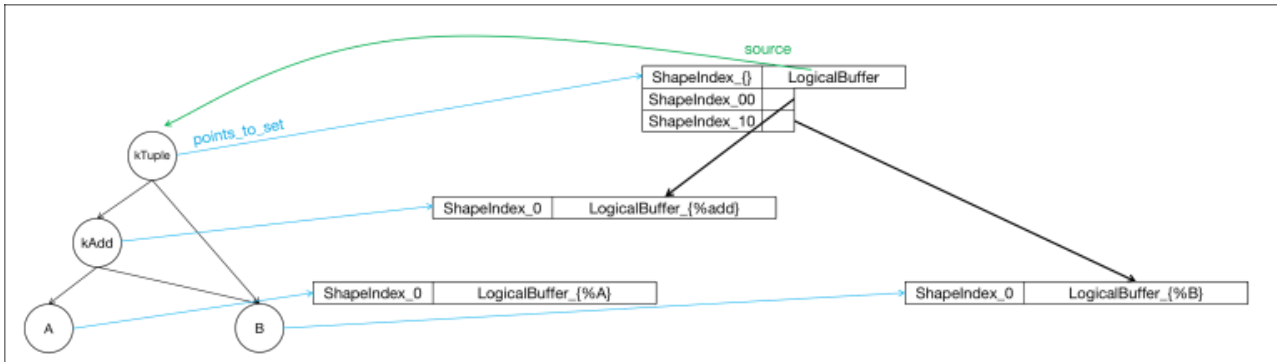
    // A map from instruction->unique_id() to
    absl::flat_hash_map<int, std::unique_ptr<PerInstruction>> per_instruction_;

    // A map from LogicalBuffer->id() to alias information about that logical
    // buffer
    std::vector<BufferAliasVector> logical_buffer_aliases_;
```

同样, TuplePointsToAnalysis也是DfsHloVisitorWithDefault, 而对于所有的Visitor, 我们都主要关注其Handle实现, 细心的同学可能发现了, 这里的Handle名与LogicalBufferAnalysis实现的一模一样, 事实上, TuplePointsToAnalysis的几个Handle实现其功能依赖于LogicalBufferAnalysis中对应的Handle的实现, 所以在看这部分代码时, 如果TuplePointsToAnalysis的某个Handle不理解, 可以去LogicalBufferAnalysis中的同名Handle看一下。

1.	Status DefaultAction(HloInstruction* hlo_instruction) override;
2.	Status HandleTuple(HloInstruction* tuple) override;
3.	Status HandleGetTupleElement(HloInstruction* get_tuple_element) override;
4.	Status HandleBitcast(HloInstruction* bitcast) override;
5.	Status HandleDomain(HloInstruction* domain) override;
6.	Status HandleCopy(HloInstruction* copy) override;
7.	Status HandleRecvDone(HloInstruction* recv_done) override;
8.	Status HandleSend(HloInstruction* send) override;
9.	Status HandleTupleSelect(HloInstruction* tuple_select) override;
10.	Status HandleAddDependency(HloInstruction* add_dependency) override;

- 1- DefaultAction, LogicalBufferAnalysis的DefaultAction()根据Shape构造LogicalBuffer实例, 这里的DefaultAction()取其结果, 将每一个LogicalBuffer都存储在PointToSet的buffers中, 此外, 对于TupleShape的HloInstruction, 还将该HloInstruction作为其上述buffers的"source". 至于什么样XlaOp会将Shape设置为Tuple, 有这些: Infeed, InfeedWithToken, RecvFromHost, RecvWithToken, SendToHost, SendWithToken, 都是用于大量传输数据的XlaOp. 构造PointToSet实例是作为TuplePointsToAnalysis作为DfsHloVisitorWithDefault的主要工作.
- 2- 在LogicalBufferAnalysis的HandleTuple()基础上, 构造PointsToSet实例, 注意到, 1. Handle内获取LogcialBuffer的索引与LogicalBufferAnalysis中存入LogcialBuffer的索引一致, 2. 对于Tuple来说, 它的PointsToSet就是每一个operand的PointsToSet集合, 包括buffers和buffers对应的source. 这一点后面在BufferLiveness中也会用到.
- 3:10- 在LogicalBufferAnalysis相应Handle的基础上, 构造相应的PointToSet实例



共性部分分析完, 同样, 下面来看下TuplePointsToSet的Specific部分. 如前文所述, Analyze()方法是入口.

```

1. TuplePointsToAnalysis::Analyze()
2.   std::vector<HloInstruction*> fusion_instructions
3.   for computation : module->MakeNonfusionComputations():
4.     computation->Accept(this)
5.     PopulateDefinedBuffersAndAliases(computation->instructions())
6.     for instruction : instructions:
7.       PerInstruction* pi = PerInst(instruction)
8.       GatherBuffersDefinedByInstruction(instruction, &pi->instruction_defined_buffers)
9.       points_to_set = GetPointsToSet(instruction)
10.      points_to_set.ForEachElement(
11.        [this, &instruction](const ShapeIndex& index, const PointsToSet::BufferList& p
12.          for (const LogicalBuffer* buffer : pointed_to_buffers):
13.            logical_buffer_aliases_[buffer->id()].emplace_back(instruction, index)
14.        })
15.      if instruction : computation->instructions():
16.        if instruction->opcode() == HloOpcode::kFusion:
17.          GatherFusionInstructions(instruction, &fusion_instructions)
18.          // Run points-to analysis on fusion instructions in 'computation'.
19.          for instruction : fusion_instructions:
20.            instruction->fused_expression_root()->Accept(this)
21.            PopulateDefinedBuffersAndAliases(instruction->fused_instructions())

```

-3- 这里和LogicalBufferAnalysis类似, 同样进行的"剪枝"处理.

-5- 除了通过Accept()执行Handle, TuplePointsToAnalysis还会对Handle处理的结果通过PopulateDefinedBuffersAndAlias()进一步处理, 形成两个存储结构:

'TuplePointsToSet::PerInstruction::instruction\_defined\_buffers' 和

'TuplePointsToAnalysis::logical\_buffer\_aliases\_', 一个instruction\_defined\_buffers是通过

GatherBuffersDefinedByInstruction()构造的, 里面依然是一个对所有SubShape遍历操作, 经过之前的过程, 不是可以轻松的获得一个HloInstruction对应的LogicalBuffer, 为什么这里还要遍历一次呢? 如果您有这样的疑问, 建议看



下HandleTuple()等TuplePointsToAnalysis的“定制”Handle, 对于 **kTuple** 等几个HloOpcode, 它的PointsTo的source可不是本身, 而是其operand的PointsTo的source, 显然, 这类HloInstruction没有自己的“defined buffers”, 而后续进行allocation assignment的前提条件就是“buffer is defined by instruction”, 这里即是挑出会被assigned的“buffer defined by instruction”, 此外, 代码的实现还顺便做了合法性检查: **kTuple** etc类的HloInstruction的所有buffers的source, 都不会是自己. 至于logical\_buffer\_aliases\_, 可以理解为一个反向检索, Alias, 别名, 也是这个意思, 可以方便用户根据LogicalBuffer->id, 快速的索引到对应的HloInstruction和ShapeIndex.

## BufferLiveness

经过LogicalBufferAnalysis构造原材料(LogicalBuffer), TuplePointsToAnalysis构造原材料之间的关系(PointsToSet), 大部分苦活已经完成, BufferLiveness进过封装就可以对外提供判断一个LogcialBuffer的Liveness的接口, 主要以下几个:

```
// Returns true if the live range of the buffer containing the output of 'a'
// may overlap with the live range of the buffer of 'b'. If instruction 'a'
// interferes with instruction 'b' then they cannot share the same buffer.
bool MayInterfere(const LogicalBuffer& a, const LogicalBuffer& b) const;

// Returns true if the buffer for the given instruction may be live out of the
// module. That is, the instruction's buffer may be included in the output of
// the entry computation.
bool MaybeLiveOut(const LogicalBuffer& buffer) const;

// Returns the complete set of buffers that may be live out of the module.
const PointsToSet::BufferSet& maybe_live_out_buffers() const;
// Returns the underlying points-to analysis used for this liveness analysis.
const TuplePointsToAnalysis& points_to_analysis() const;
```

这里还是关注几个小细节.

```
1. Status BufferLiveness::Analyze() {
2.     points_to_analysis_ = TuplePointsToAnalysis::Run(module_)
3.     for computation : module_>computations():
4.         if computation->IsFusionComputation():
5.             continue;
6.         // Gather all instructions whose buffers might alias other instructions into
7.         // the set aliased_buffers_. This includes those contained as a tuple
8.         // element in other instruction's output.
9.         for instruction : computation->instructions():
10.            for LogicalBuffer* aliased_buffer : points_to_analysis_>GetPointsToSet(instruc
11.                if (aliased_buffer->instruction() != instruction):
12.                    aliased_buffers_.insert(aliased_buffer);
13.            if computation == module_>entry_computation():
14.                HloInstruction* root = computation->root_instruction();
15.                maybe_live_out_buffers_ = points_to_analysis_>GetPointsToSet(root).CreateFlatt
```

-11- 如果一个HloInstruction->PointsToSet中->LogicalBuffer的source不是该HloInstruction, 就表示该HloInstruction是类Tuple的, 没有实质的LogicalBuffer, 也不是任何LogicalBuffer的source, 将其存储在alias\_buffers\_

-15- live\_out\_buffer就是生命周期长于entry\_computation(或者理解为HloModule)的Buffer, 就是承载整个entry\_computation输出数据的Buffer, root作为整个程序的总入口和总出口, 其PointsToSet的LogicalBuffers就是整个程序总输出Buffer不难理解, 但这里又为什么叫“maybe”? 代码的注释是, 后续进行执行流优化的时候, 会对hlo\_ordering调整, 那么就有可能root不再作为最后的HloInstruction, 所以叫“maybe”. 但根据我对XLA的理解, 没有

想到哪些优化会导致root被提前, 也许只是这个文件的程序员给自己留个后路, 毕竟万一有呢, 当然, 更可能是我忽略了一些细节, 理解不深(TODO).

下面两个是BufferLiveness和核心方法, 在后续的优化被频繁调用, 逻辑简单, 就不展开

```
// Returns true if the live range of the buffer containing the output of 'a'
// may overlap with the live range of the buffer of 'b'. If instruction 'a'
// interferes with instruction 'b' then they cannot share the same buffer.
bool MaybeInterfere(const LogicalBuffer& a, const LogicalBuffer& b) const;
// Returns true if the buffer for the given instruction may be live out of the
// module. That is, the instruction's buffer may be included in the output of
// the entry computation.
bool MaybeLiveOut(const LogicalBuffer& buffer) const;
```

## BuildColocatedBufferSets()

这个方法的核心是将可以共享的Memory Chunk的LogicalBuffer挑出来, 存在同一个BufferSet里

```
// Builds sets of buffers in 'colocated_buffer_sets' which should be colocated
// in the same allocation (currently just supports kWhile, kCall, and
// kConditional and input output aliasing).
void BufferAssigner::BuildColocatedBufferSets()
```

这句注释能看出几点信息:

1. 相同的buffer set会共享同一个allocation, 即共享同一块物理内存
2. 当前只针对kWhile, kCall, 和kConditional, 而这三个HloOpcode的共同点是都有called\_computations — 是连接entry\_computation和其他computation的”铰链”
3. 该函数只会进行computation-wide 的合并, 即: 不会有一个ColocatedBufferSets会cross-computation ColocatedBufferSet. 这一点在很多地方都能体现。

```
1. BuildColocatedBufferSets()
2.   module->input_output_alias_config().ForEachAlias({...})
3.   for computation : module->MakeComputationPostOrder():
4.     if computation->IsFusionComputation():
5.       continue
6.     for instruction : computation->MakeInstructionPostOrder():
7.       if kWhile:
8.         ...
9.       else if kCall:
10.        ...
11.       else if kConditional:
12.        ...
13.
14.   MergeColocatedBufferSets()
15.   is_readonly_entry_parameter = [](){}
16.   set_can_be_merged              //这是一个vector<bool>
17.   for i : colocated_buffer_sets.size():
18.     for buffer : colocated_buffer_sets[i]:
19.       if buffer_liveness.MaybeLiveOut(buffer) || is_readonly_entry_parameter(buffer)
20.         set_can_be_merged[i] = false
21.         break
22.
23.   cannot_merge_buffer_sets = [](){};
24.   interference_map
25.   for i : colocated_buffer_sets.size():
```

```

26.         for j : colocated_buffer_sets.size():
27.             if cannot_merge_buffer_sets(i, j):
28.                 interference_map[i].push_back(j)
29.                 interference_map[j].push_back(i)
30.         assigned_colors = ColorInterferenceGraph(interference_map)
31.         c_sort
32.         assigne_colors
33.         for node : nodes:
34.             for
35.                 for
36.                     assigned_colors[node] = color
37.                 return assigned_colors
38.         for
39.             new_colocated_buffer_sets

```

-2- 处理Alias的buffer，用ShapeIndex来索引，确保相同Set内的LogicalBuffer的Shape相同，相应的Buffersize也相同

-7,9,11- 将kWhile, kCall, kConditional的buffer都按照”相同ShapeIndex的Buffer归属到相同的BufferSet”的原则构造colocated\_set

-15- MergeColocatedBufferSets, 前面都是构造BufferSets，这里尝试能否将多个BufferSet进一步合并，之前每个BufferSet的构造都是独立的，

-23 cannot\_merge\_buffer\_sets, 判断两个BufferSet能否合并，是两个bufferSet能合并的”充要条件”

-24- interference\_map, 标记BufferSet之间的”cannot be merged”的关系

-24:29- 使用邻接矩阵法表示两个BufferSet是否有干涉。比如，以下面的inference\_map为例:

```

0:1,2
1:0,3
2:0
3:1

```

该map表示BufferSet0和BufferSet1, BufferSet2 都是cannot be merged, BufferSet3同理

-30:39- ColorInterferenceGraph, 为inference\_map里的set标记color, 相同color的表示可以merge, 用color是为了计算最终需要多少BufferSet。最终就会生成merge好的colocated\_buffersets, 注意，目前为止，这里还只有kWhile, kConditional, kCall的BufferSet。一开始进行按照冲突的set的数量进行排序，优先处理冲突最多的。进

”，按照冲突的多寡排序，会保证冲突多的color值比较小，color值越大，冲突越小，而按照相同的color合并，前面的能合并的少，后面的能合并的多。举个例子:

```

1:{4,5}
2:{4,5}
3:{4}
4:{1,2,3}
5:{1,2}

```

问，4, 5能否合并

答，这个其实看情况，如果3 的color比1, 2小，那么4和5就不会是同一个color，也就不会合并，而如果不是这样，那么4和5就是相同的color，也就可以合并4的color和1, 2怎么确定呢？这里将整体排序，按照冲突程度排序，冲突越大，color越小在本例中，4的color最小，5和4的color一样将冲突大的往前排，先分配好color，这样对于4和5这种情况，会使4和5在一个set中

## AssignColocatedBufferSets()

在Computation内部，为ColotatedBufferSet分配BufferAllocation，此时，ColocatedBufferSet只有Alias，kCall，kWhile，kConditional几个HloOpcode

```
// For each buffer set in 'colocated_buffer_sets', assigns all buffers in the
// same set to the same buffer allocation in 'assignment'.
void AssignColocatedBufferSets(
```

```
1. AssignColocatedBufferSets() //Assign, 赋值, Logical->Allocatio
2. for colocated_buffer_set : colocated_buffer_sets:
3.     BufferAllocation* allocation = nullptr
4.     for buffer : colocated_buffer_set:
5.         if instruction->opcode() == HloOpcode::kParameter && computation == computation
6.             entry_parameter...
7.         else if instruction->opcode() == HloOpcode::kConstant:
8.             is_constant = true
9.         for buffer : colocated_buffer_set:
10.            if allocation == nullptr:
11.                // TODO(b/32491382) Avoid current trivial solution of using new
12.                // allocations for each colocated buffer set. When liveness has
13.                // module-level scope, we can allow buffers to be shared across
14.                // computations (in some cases).
15.                allocation = assignment->NewAllocation(*buffer, buffer_size)
16.                if is_constant:
17.                    allocation->set_constant(true)
18.                    colocated_allocations->insert(allocation->index())
19.                else:
20.                    assignment->AddAssignment(...offset = 0, buffer_size)
21.                    colocated_buffers->insert(buffer)
22.            if entry_parameter_number >= 0:
23.                parameter_has_alias = ...
24.                allocation->set_entry_computation_paramter()
```

-2- 遍历 colocated\_buffer\_sets 中的每一个colocated\_buffer\_set中的每一个Logical buffer

-3- BufferAllocation\* 每个colocated\_buffer\_set共享一个BufferAllocation

-5,7- 如果一个ColocatedBufferSet中有一个用于kParameter和kConstant的LogicalBuffer，那么整个BufferAllocation都要配置相应的域

-9- 遍历每一个 colocated\_buffer\_set, 如果是Set中的第一个LogicalBuffer，就通过NewAllocation创建BufferAllocation实例并保存在assignment中，否则就将LogicalBuffer关联已经构造好的BufferAllocation实例中。这里注意，这里的offset都是0，意味着同一个ColocatedBufferSet中的LogicalBuffer都是每块Memory的0地址开始使用，这也符合ColocatedBufferSet的定义，同一个ColocatedBufferSet内的任意两个Buffer都在时间上不干涉的，配置同一个地址，没有任何问题。注释：由于当前BufferLiveness的分析都是在HloComputation内部，所以基于BufferLiveness的ColocatedBufferSets分析和BufferAllocation的构造，也不会computation-level scope的，这块还有优化空间，可以进一步压缩BufferAllocation的需求。

## GatherComputationsByAllocationType()

HloComputation 在概念上类比“函数”，这个接口按照其返回reference和value的不同，将其分为global和thread-local两类，划分标准十分粗暴，还有很大的优化空间：

1. 如果一个HloComputation实例返回reference, 那么该实例就是global, 其内部的所有BufferAllocation, 无论是否是"局部变量", 均当做"全局变量"处理(set\_thread\_local(false)), 一同分配, 优化, 生命周期也和全局变量一样。体现在BufferAllocation的属性上, 就是is\_thread\_local == false。这种划分方法会导致global的HloComputation内的生命周期较实际需要的状况被延长了, 造成浪费, 但能大幅简化优化逻辑。
2. 如果一个HloComputation实例返回value, 那么该实例就是thread-local, 其内部的所有BufferAllocation, 必须全部是"局部变量"(set\_thread\_local(true)), 简化了Buffer生命周期的管理, 随着HloComputation的执行结束, "局部变量"会被回收。也正是由于栈变量可以反复的申请释放, 也就无需像global一样进行全局的优化。全部是"局部变量"的约束, 显得僵硬, 但能大幅简化优化逻辑。

可以类比线程来理解这两个概念: global\_computations -> 主线程, thread\_local\_computations -> 子线程。但为什么要做这个划分呢? 笔者认为可以从以下2点来理解:

1. 当前的软件层是HLO: High Level Optimization, 是硬件无关的, 而无论什么执行硬件, 都有"执行流"的概念, 也有"主执行流"的概念, 只不过具体实现不同, 比如CPU中有主线程和子线程, GPU中有CUDA Stream。对这些硬件公共的执行流进行抽象, 就是HLO的工作, 所以, 这里既没有说进程, 线程, CUDA Stream, 而只是说global和thread\_local, 是一个高度抽象的执行流
2. HLO要做显存优化, 一方面要知道两个Buffer是否存在必需的依赖关系(BufferLiveness), 如果没有依赖关系, 就要考察两个Buffer是否在同一个执行流, 简单的理解, 在同一个执行流的Buffer, 由于存在先后顺序, 通常是可以复用的。

```
1. // Walk the call graph of the HLO module and place each computation into either
2. // thread_local_computations or global_computations depending upon whether the
3. // computation requires thread-local allocations or global allocations. The
4. // elements in thread_local_computations and global_computations are in post
5. // order (if computation A has an instruction which calls computation B, then A
6. // will appear after B in the vector). --> thread_local 就是per thread 变量, 线程不安全
7. GatherComputationsByAllocationType()
8. while !worklist.empty():
9.     if is_thread_local && in_global_set ||
10.        ...
11.     if is_thread_local:
12.         thread_local_set...
13.     else:
14.         global_set...
15.     for instruction : computation->instructions():
16.         for subcomputation : instruction->called_computation():
17.             switch instruction->opcode():
18.                 case HloOpcode::kCall:
19.                 case HloOpcode::kConditional:
20.                 case HloOpcode::kWhile:
21.                     worklist.push_back(std::make_pair(subcomputation, false))
22.                     break;
23.                 case HloOpcode::kAllReduce:
24.                 case HloOpcode::kMap:
25.                 case HloOpcode::kReduce:
26.                 case HloOpcode::kReduceWindow:
27.                 case HloOpcode::kScatter:
28.                 case HloOpcode::kSelectAndScatter:
29.                 case HloOpcode::kSort:
30.                 case HloOpcode::kFusion:
31.                     worklist.push_back(std::make_pair(subcomputation, true))
32.             default:
33.                 return InternalError()
34.
35. for computation : module->MakeComputationPostOrder():
36.     if thread_local_set.contains(computation):
```

```

37.         thread_local_computation->push_back(computation)
38.     else if global_set.contains(computation):
39.         global_computations->push_back(computation)

```

-26- 只处理有called\_computation的情况

-28:30- 如果instruction是kCall, kConditional, kWhile, 就是global, 与AssignColocatedBufferSets()相呼应

-32:39- 对于kAllreduce, kMap, kReduce, kReduceWindow, kScatter, kSelectAndScatter, kSort, kFusion 都会新开thread\_local

## AssignBuffersForComputation()

为剩下的LogicalBuffer找Allocation, hlo\_opcode.h有97个HloCode, 之前主要是围绕Alias, kCall, kWhile, kConditional准备Buffer, 现在要处理剩下的HloOpcode, 尽可能复用之前已经分配的BufferAllocation, 复用的算法就是已经分配好global\_computation和thread\_local\_computation

1. global\_computation是"main stream", 是顺序的, 尽可能复用已有的BufferAllocation
2. thread\_local\_computation 要分配自己的

```

1. // Assigns buffers to the instructions in the given computation. "assignment"
2. // is modified to reflect the new buffer assignments. If is_thread_local is
3. // true, then all assigned buffers have the is_thread_local flag set to
4. // true.
5. AssignBuffersForComputation
6.     for instruction : computation->instructions():
7.         for buffer : assignment->points_to_analysis().GetBufferDefinedByInstruction()
8.             sorted_buffers.push_back()
9.         for instruction : computation->MakeInstructionPostOrder():
10.            post_order_position.emplace(instruction, position)
11.            position++
12.        absl::c_sort(sorted_buffers,...)
13.        for buffer : sorted_buffer:
14.            if colocated_buffers.contains(buffer):
15.                continue
16.            if instruction->opcode() == HloOpcode::kConstant:
17.                BufferAllocation* allocation = assignment->NewAllocation(*buffer, buffer_size)
18.                allocation->set_constant(true)
19.                continue
20.            if is_entry_parameter:
21.                ...
22.                continue
23.            if is_thread_local:
24.                allocation->set_is_thread_local(true)
25.                continue
26.            //global
27.            if (buffer->IsTopLevel() && !buffer->IsTuple()):
28.                ...
29.                ...
30.            //temp
31.            if (!assignment->HasAllocation(*buffer) && has_sequential_order &&
32.                !liveness.MaybeLiveOut(*buffer)):
33.                (*buffers_to_assign_sequentially)[computation].insert(buffer)
34.
35.            if (!assignment->HasAllocation(*buffer)):
36.                BufferAllocation* allocation = assignment->NewAllocation(*buffer, buffer_size)

```



- 14- 如果是ColocatedBuffer, 那么continue, 因为在GatherComputationsByAllocationType()已经为这类(Alias, kCall, kWhile)分配了BufferAllocation(而且还是offset为0的)
- 16- 处理kContant, 分配新BufferAllocation
- 20- 处理kParameter, 分配新BufferAllocation
- 23- 对于thread local, 分配新的BufferAllocation
- 24:27- 对于global, 尽可能复用
- 28- 对于一直没有assign的temp buffer, 先登记, 延后处理
- 33- 剩下依然没有办法复用的, 分配新的BufferAllocation

## AssignBuffersWithSequentialOrdering()

对于前文规划好的一组顺序分配的BufferAllocation(global), 利用HeapSimulation算法, 将它们压缩到一块完整的BufferAllocation中, 主要针对 global 和 temp

```
// Assigns 'buffers_to_assign_sequentially' using heap simulation, assuming
// the HLO instructions will be executed in the sequential order given by
// assignment->liveness().hlo_ordering().SequentialOrder. If
// 'run_whole_module_heap_simulation' is true, the heap simulation will be run
// assuming all global computations are sequentially ordered.
AssignBuffersWithSequentialOrdering()
  AssignBuffersFromHeapSimulator()
  BufferAllocation* allocation = assignment->NewEmptyAllocation(result.heap_size, color);
```

## CombineTempAllocations()

优化的最后一步, 优化Temp的BufferAllocation

```
// Combines allocations of temporary buffers of the same color into one big
// BufferAllocation.
CombineTempAllocations():
```

### Related:

[Tensorflow XLA Service 详解 II](#)

[Tensorflow XLA Service 详解 I](#)

[Tensorflow XLA Client | HloModuleProto 详解](#)

[Tensorflow XlaOpKernel | tf2xla 机制详解](#)

[Tensorflow JIT 技术详解](#)

[Tensorflow JIT/XLA UML](#)

[Tensorflow Forums](#)

📁 技术 🔗 Tensorflow, XLA, 技术

[◀ BERT Training By Example](#)