

GPU programming: Warp-synchronous programming


Sylvain Collange
Inria Rennes – Bretagne Atlantique
sylvain.collange@inria.fr

Bypassing SIMT limitations

Multi-thread + SIMD



SIMT model



- Across SIMD lanes of same thread
 - ◆ Direct communication without synchronization
- Between threads
 - ◆ All synchronization primitives, wait-based and lock-based algorithms
- If we know the warp/thread organization, we can pretend we program a multi-thread + SIMD machine
 - ◆ Drawback: give up automatic divergence management
- Common denominator of operations allowed in SIMD and multi-thread
 - ◆ No direct communication
 - ◆ No blocking algorithm

Outline

- Inter-thread communication
- Warp-based execution and divergence
- Intra-warp communication:
warp-synchronous programming

Inter-thread/inter-warp communication

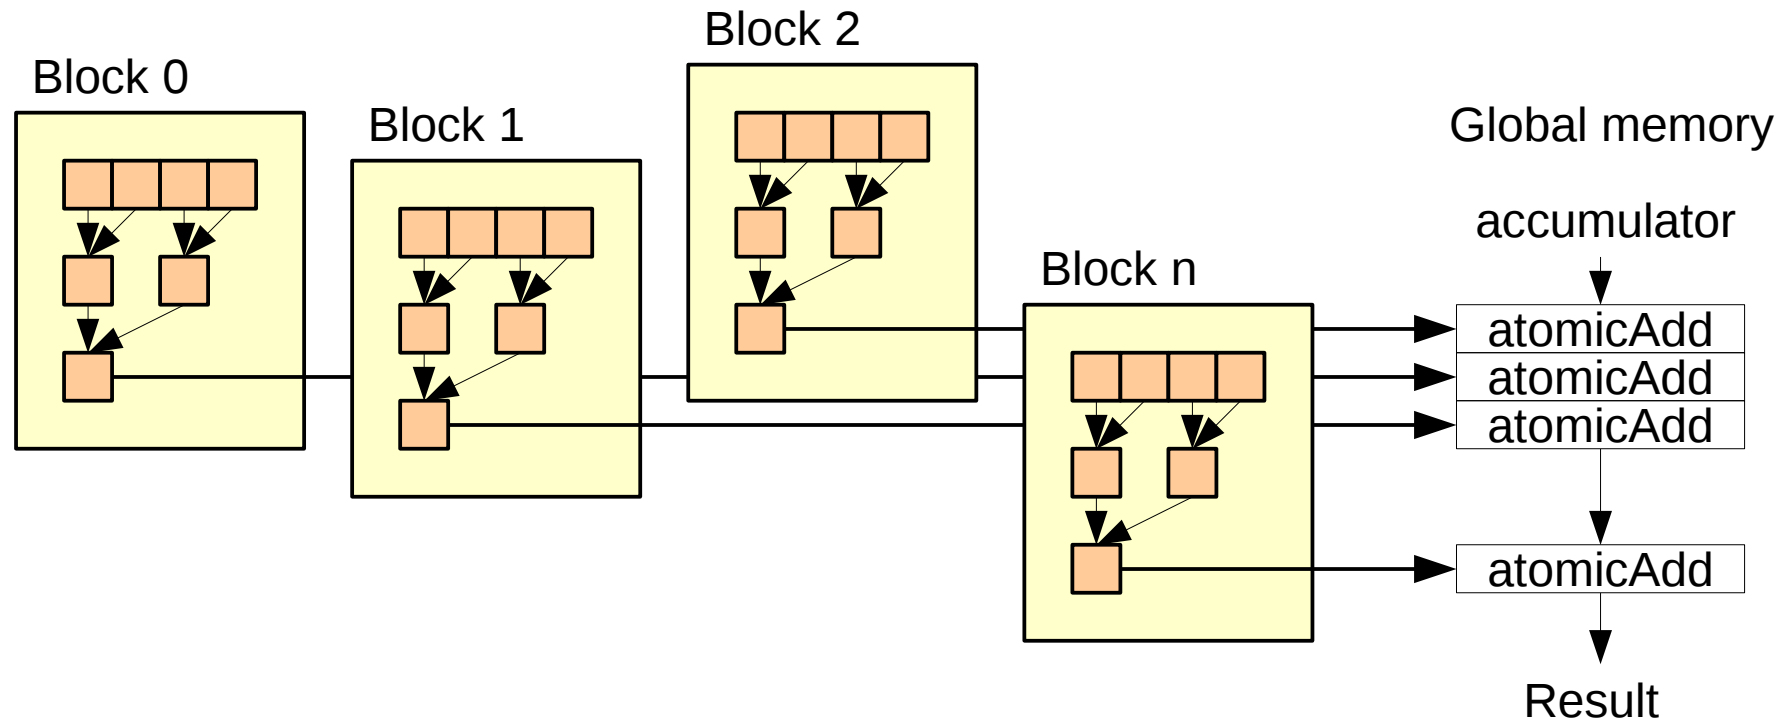
- Barrier: simplest form of synchronization
 - ◆ Easy to use
 - ◆ Coarse-grained
- Atomics
 - ◆ Fine-grained
 - ◆ Can implement *wait-free* algorithms
 - ◆ May be used for blocking algorithms (locks) among **warps** of the **same block**, but not recommended
- Communication through memory
 - ◆ Beware of consistency

Atomics

- Read, modify, write in one operation
 - ◆ Cannot be mixed with accesses from other thread
- Available operators
 - ◆ Arithmetic: `atomic{Add,Sub,Inc,Dec}`
 - ◆ Min-max: `atomic{Min,Max}`
 - ◆ Synchronization primitives: `atomic{Exch,CAS}`
 - ◆ Bitwise: `atomic{And,Or,Xor}`
- On global memory, and shared memory
- Performance impact in case of contention
 - ◆ Atomic operations to the same address are serialized

Example: reduction

- After local reduction inside each block, use atomics to accumulate the result in global memory



- Complexity?
- Time including kernel launch overhead?


Example: compare-and-swap

- Use case: perform an arbitrary associative and commutative operation atomically on a single variable
- `atomicCAS(p, old, new)` does atomically
 - ◆ if `*p == old` then assign `*p ← new`, return old
 - ◆ else return `*p`

```
__shared__ unsigned int data;
unsigned int old = data;           // Read once
unsigned int assumed;
do {
    assumed = old;
    newval = f(old, x);             // Compute
    old = atomicCAS(&data, old, newval); // Try to replace
} while(assumed != old);           // If failed, retry
```

Memory consistency model

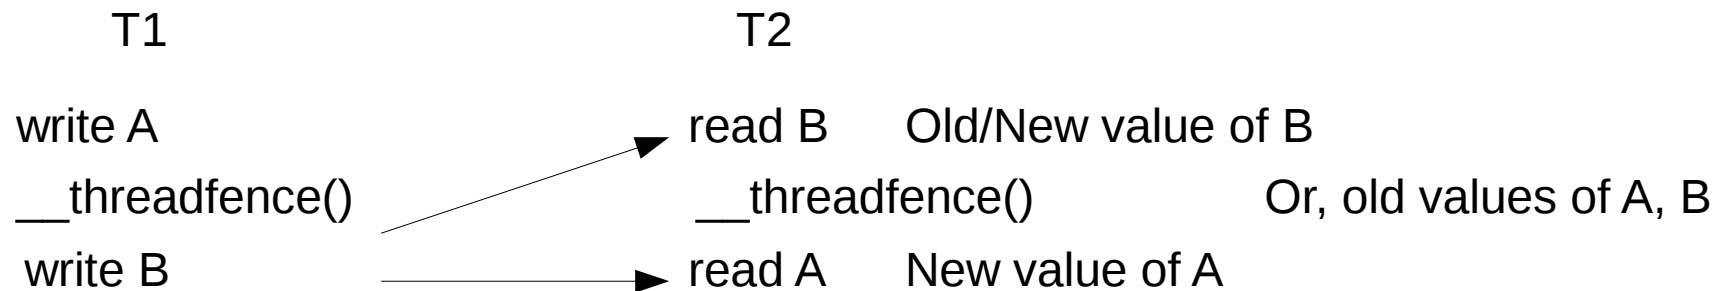
- Nvidia GPUs (and compiler) implement a relaxed consistency model
 - ◆ No global ordering between memory accesses
 - ◆ Threads may not see the writes/atomics in the same order



- Need to enforce explicit ordering

Enforcing memory ordering: fences

- `__threadfence_block`
`__threadfence`
`__threadfence_system`
 - ◆ Make writes preceding the fence appear before writes following the fence for the other threads at the block / device / system level
 - ◆ Make reads preceding the fence happen after reads following the fence



- Declare shared variables as **volatile** to make writes visible to other threads (prevents compiler from removing “redundant” read/writes)
- `__syncthreads` implies `__threadfence_block`

Outline

- Inter-thread communication
- Warp-based execution and divergence
- Intra-warp communication:
warp-synchronous programming

Warp-based execution

Reminder



- Threads in a warp run in lockstep
- On current NVIDIA architectures, warp is 32 threads
- A block is made of warps
 - ◆ Warps do not cross block boundaries
 - ◆ Block size multiple of 32 for best performance

Branch divergence

- Conditional block

```
if(c) {  
    // A  
}  
else {  
    // B  
}
```

- When all threads of a warp take the same path:




With imaginary 4-thread warps

Branch divergence

- Conditional block

```
if(c) {  
    // A  
}  
else {  
    // B  
}
```

- When threads in a warp take different paths:



- Warps have to go through both **A** and **B**: lower performance

Avoiding branch divergence

- Hoist identical computations and memory accesses outside conditional blocks

```
if(tid % 2) {  
    s += 1.0f/tid;  
}  
else {  
    s -= 1.0f/tid;  
}
```

```
float t = 1.0f/tid;  
if(tid % 2) {  
    s += t;  
}  
else {  
    s -= t;  
}
```


- When possible, re-schedule work to make non-divergent warps

```
// Compute 2 values per thread  
int i = 2 * tid;  
s += 1.0f/i - 1.0f/(i+1);
```

- What if I use C's ternary operator (?:) instead of if?
(or tricks like ANDing with a mask, multiplying by a boolean...)

Ternary operator ? good : bad

- Run both branches and select: $R = c ? A : B;$
 - ◆ No more divergence?
- All threads have to take both paths
No matter whether the condition is divergent or not




- Does **not** solve divergence: we lose in all cases!
- Only benefit: fewer instructions
 - ◆ May be faster for short, often-divergent branches
- Compiler will choose automatically when to use predication
 - ◆ Advice: keep the code readable, let the compiler optimize

Barriers and divergence

- Remember barriers cannot be called in divergent blocks
 - ◆ All threads or none need to call `__syncthreads()`

- If: need to split


```
if(p) {  
    ...  
    // Sync here?  
    ...  
}
```



```
if(p) {  
    ...  
}  
__syncthreads();  
if(p) {  
    ...  
}
```

- Loops: what if trip count depends on data?

```
while(p) {  
    ...  
    // Sync here?  
    ...  
}
```




?

Barriers instructions

- Barriers with boolean reduction
 - ◆ `__syncthreads_or(p) / __syncthreads_and(p)`
Synchronize, then returns the boolean OR / AND of all thread predicates p
 - ◆ `__syncthreads_count(p)`
Synchronize, then returns the number of non-zero predicates
- Loops: what if trip count depends on data?
 - ◆ Loop while at least one thread is still in the loop

```
while(p) {  
    ...  
    // Sync here?  
    ...  
}
```



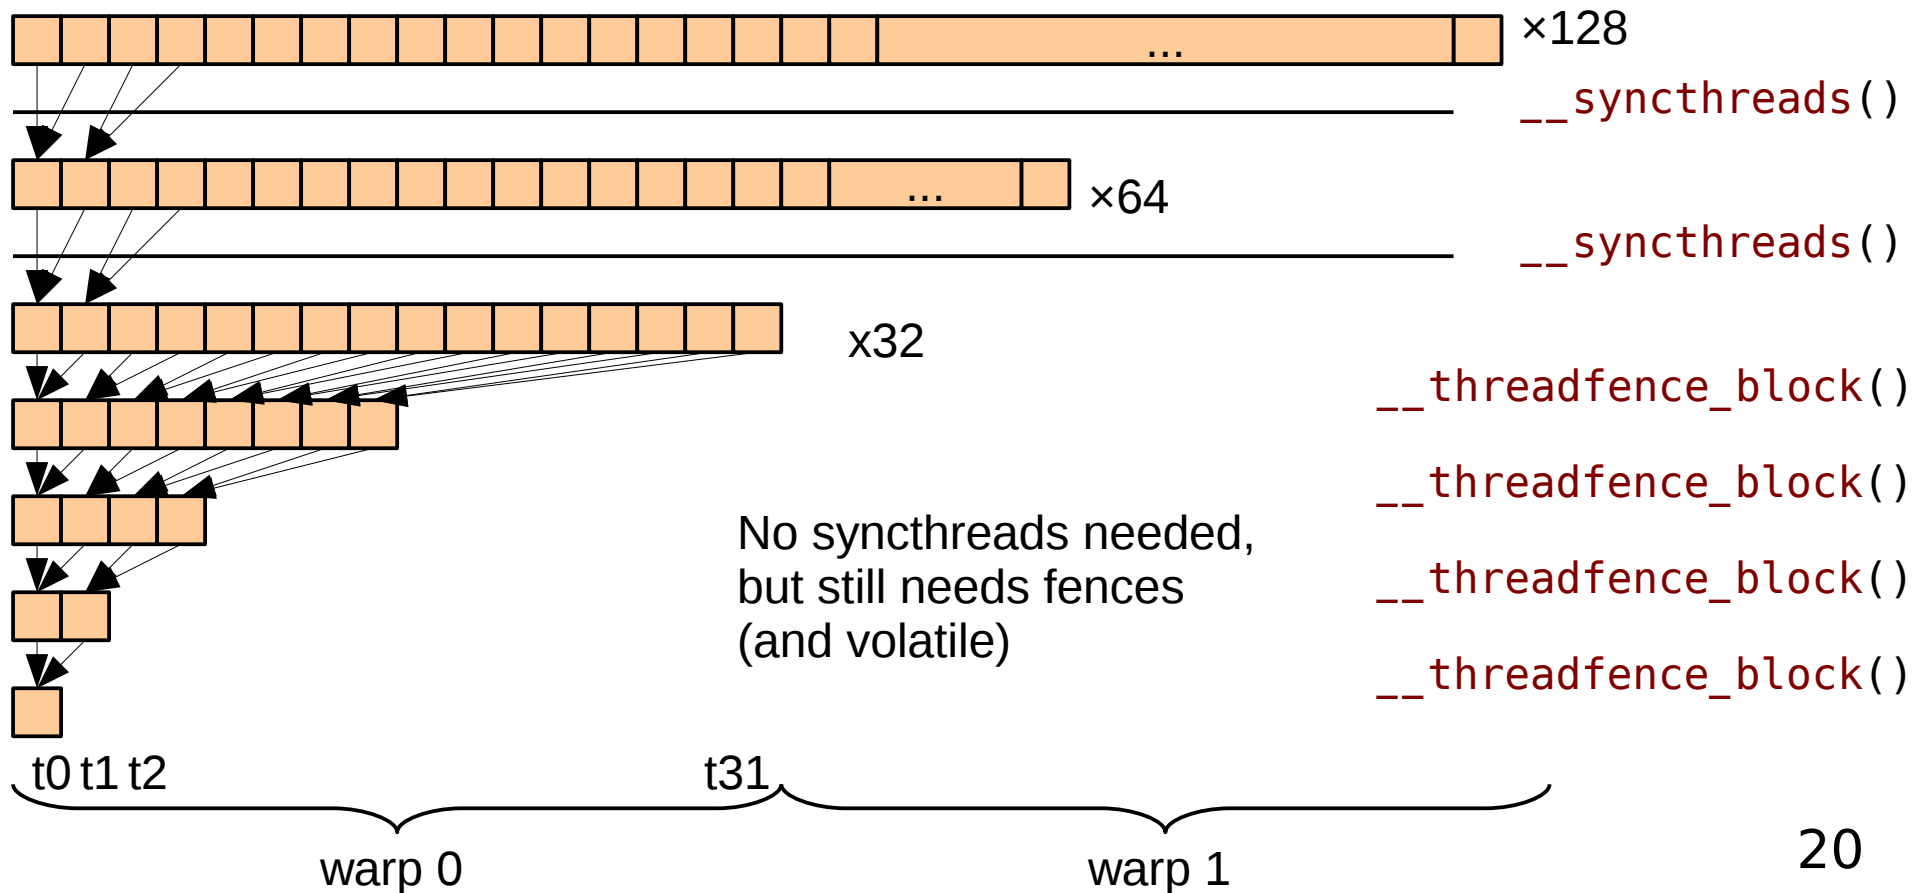
```
while(__syncthreads_or(p)) {  
    if(p) {  
        ...  
    }  
    __syncthreads();  
    if(p) {  
        ...  
    }  
}
```

Outline

- Inter-thread communication
- Warp-based execution and divergence
- Intra-warp communication:
warp-synchronous programming

Warp-synchronous programming

- We know threads in a warp run synchronously
 - ◆ No need to synchronize them explicitly
- Can use SIMD (PRAM-style) algorithms inside warps
- Example: last steps of a reduction



Warp-synchronous programming: tools

We need warp size, and thread ID inside a warp: lane ID

- Predefined variable: `warpSize`
- Lane ID exists in PTX, not in C for CUDA
Can be recovered using inline PTX


```
__device__ unsigned int laneID() {  
    unsigned int lane;  
    asm volatile("mov.u32 %0, %%laneid;" : "=r"(lane));  
    return lane;  
}
```

We can emit PTX using the `asm` keyword,
but remember PTX is not assembly.

Warp vote instructions

- `p2 = __all(p1)`
horizontal AND between predicates p1
of all threads in the warp
- `p2 = __any(p1)`
OR between all p1
- `n = __ballot(p)`
Set bit i of integer n
to value of p for thread i
i.e. get bit mask as an integer

Like `__syncthreads_{and,or}` for a warp
Use: take control decisions for the whole warp



Manual divergence management

- How to write an if-then-else in warp-synchronous style?
i.e. without breaking synchronization

- ◆ Using predication:
execute both sides always

```
if(p) { A(); B(); }  
else { C(); }
```

```
if(p) { A(); }  
// Threads synchronized  
if(p) { B(); }  
// Threads synchronized  
if(!p) { C(); }
```

- ◆ Using vote instructions:
only execute taken paths
Skip block if no thread takes it

```
if(__any(p)) {  
    if(p) { A(); }  
    // Threads synchronized  
    if(p) { B(); }  
}  
if(__any(!p)) {  
    // Threads synchronized  
    if(!p) { C(); }  
}
```

- How to write a while loop?

Ensuring reconvergence

When and where do threads reconverge?

- No hard guarantees today
 - ◆ e.g. Can the compiler alter the control flow?
 - ◆ Will hopefully get standardized [1]
- In practice, we can assume reconvergence:
 - ◆ After structured if-then-else
 - ◆ After structured loops
 - ◆ At function return
- Example: can we assume reconvergence here?

```
if(threadIdx.x != 0) {           ← Divergent condition
    if(blockIdx.x == 0) return;  ← Uniform condition
}
// Threads synchronized at this point ?
```

Ensuring reconvergence

When and where do threads reconverge?

- No hard guarantees today
 - ◆ e.g. Can the compiler alter the control flow?
 - ◆ Will hopefully get standardized [1]
- In practice, we can assume reconvergence:
 - ◆ After structured if-then-else
 - ◆ After structured loops
 - ◆ At function return
- Example: can we assume reconvergence here?


```
if(threadIdx.x != 0) {           ← Divergent condition
    if(blockIdx.x == 0) return;  ← Uniform condition
}
// Threads synchronized at this point ?
```

- ◆ **No.** Return creates unstructured control flow.
Will or won't reconverge depending on the architecture.

Shuffle


Exchange data between lanes

- `__shfl(v, i)`
Get value of thread i in the warp
 - ◆ Use: 32 concurrent lookups in a 32-entry table
 - ◆ Use: arbitrary permutation...
- `__shfl_up(v, i) \approx __shfl(v, tid-i),`
`__shfl_down(v, i) \approx __shfl(v, tid+i)`
Same, indexing relative to current lane
 - ◆ Use: neighbor communication, shift
- `__shfl_xor(v, i) \approx __shfl(v, tid ^ i)`
 - ◆ Use: exchange data pairwise: “butterfly”



Example: reduction + broadcast

- Naive algorithm



$$a[2*i] \leftarrow a[2*i] + a[2*i+1]$$

$$a[4*i] \leftarrow a[4*i] + a[4*i+2]$$


$$a[8*i] \leftarrow a[8*i] + a[8*i+4]$$

$$a[i] \leftarrow a[0]$$

- Let's rewrite it using shuffle

Example: reduction + broadcast

- With shuffle



```
ai += __shfl_xor(ai, 1);
```


```
ai += __shfl_xor(ai, 2);
```

```
ai += __shfl_xor(ai, 4);
```

- Exercise: implement complete reduction without `__syncthreads`
 - Hint: use shared memory atomics

Other example: parallel prefix

- Remember our PRAM algorithm



Σ_{i-j} is the sum $\sum_{k=i}^j a_k$

```
s[i] ← a[i]
if i ≥ 1 then
    s[i] ← s[i-1] + s[i]
```


```
if i ≥ 2 then
    s[i] ← s[i-2] + s[i]
```

```
if i ≥ 4 then
    s[i] ← s[i-4] + s[i]
```

```
Step d: if i ≥ 2d then
    s[i] ← s[i-2d] + s[i]
```

Other example: parallel prefix

- Using warp-synchronous programming



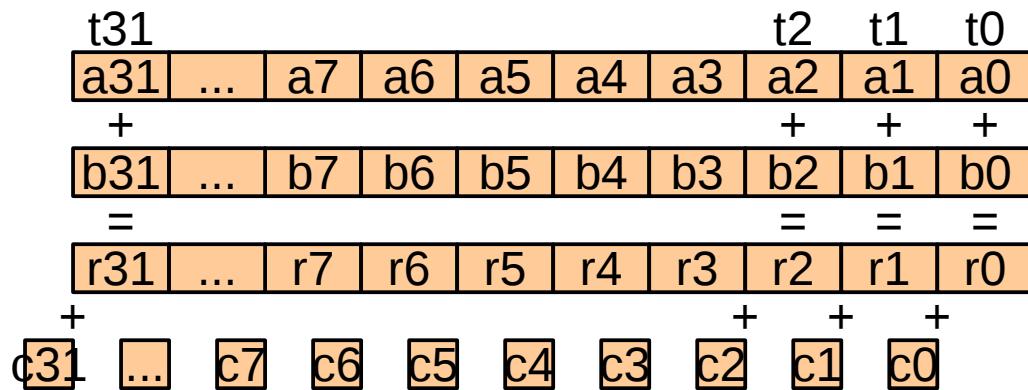
Σ_{i-j} is the sum $\sum_{k=i}^j a_k$

```
s = a;
n = __shfl_up(s, 1);
if(laneid >= 1)
    s += n;
n = __shfl_up(s, 2);
if(laneid >= 2)
    s += n;
n = __shfl_up(s, 4);
if(laneid >= 4)
    s += n;
```

```
for(d = 1; d <= 5; d *= 2) {
    n = __shfl_up(s, d);
    if(laneid >= d)
        s += n;
}
```

Example: multi-precision addition


- Do an addition on 1024-bit numbers
- Represent numbers as vectors of 32×32-bit
 - ◆ A warp works on a vector
- First step: add elements of the vectors in parallel and recover carries



```
uint32_t a = A[tid],  
          b = B[tid], r, c;  
  
r = a + b;    // Sum  
  
c = r < a;    // Get carry
```

Second step: propagate carries

- This is a parallel prefix operation
 - ◆ We can do it in $\log(n)$ steps
- But in most cases, one step will be enough
 - ◆ Loop until all carries are propagated




```
uint32_t a = A[tid],  
         b = B[tid], r, c;
```

```
r = a + b;    // Sum  
c = r < a;    // Get carry  
while(__any(c)) { // Carry left?  
    c = __shfl_up(c, 1); // Move left  
    if(laneid == 0) c = 0;  
    r = r + c;    // Sum carry  
    c = r < c;    // New carry?  
}  
R[tid] = r;
```

Mixing SPMD and SIMD

- Most language features work in either SPMD mode or SIMD mode



- Avoid divergent branches when writing SIMD code
- Consequence: SIMD code may call SPMD functions, not the other way around

Takeaway

- Two ways to program an SIMT GPU
 - ◆ With independent threads, grouped in warps
 - ◆ With warps operating on vectors
- 3 levels
 - ◆ Blocks in grid: independent tasks, no synchronization
 - ◆ Warps in block: concurrent “threads”, explicitly synchronizable
 - ◆ Threads in warp: implicitly synchronized
- Warp-synchronous still somehow uncharted territory in CUDA, but widely used in libraries
- If you know warp-synchronous programming in CUDA, you know SIMD programming with masking (Xeon Phi, AVX-512...)

Device functions

- Kernel can call functions
- Need to be marked for GPU compilation

```
__device__ int foo(int i) {  
}
```

- A function can be compiled for both host and device

```
__host__ __device__ int bar(int i) {  
}
```

- Device functions can call device functions
 - ◆ Newer GPUs support recursion

Local memory

- Registers are fast but
 - ◆ Limited in size
 - ◆ Not addressable
- Local memory used for
 - ◆ Local variables that do not fit in registers (*register spilling*)
 - ◆ Local arrays accessed with indirection

```
int a[17];  
b = a[i];
```

- **Warning:** local is a misnomer!
 - ◆ Physically, local memory usually goes off-chip
 - ◆ About same performance as coalesced access to global memory

Loop unrolling

- Can improve performance
 - ◆ Amortizes loop overhead over several iterations
 - ◆ May allow constant propagation, common sub-expression elimination...
- Unrolling is **necessary** to keep arrays in registers

Not unrolled

```
int a[4];  
for(int i = 0; i < 4; i++) {  
    a[i] = 3 * i;  
}
```

Indirect addressing:
a in local memory

Unrolled

```
int a[4];  
a[0] = 3 * 0;  
a[1] = 3 * 1;  
a[2] = 3 * 2;  
a[3] = 3 * 3;
```

Static addressing:
a in registers

Trivial
computations:
optimized away

- The compiler can unroll for you

```
#pragma unroll  
for(int i = 0; i < 4; i++) {  
    a[i] = 3 * i;  
}
```

Device-side functions: C library support

- Extensive math function support
 - ◆ Standard C99 math functions in single and double precision:
e.g. `sqrtf`, `sqrt`, `expf`, `exp`, `sinf`, `sin`, `erf`, `j0`...
 - ◆ More exotic math functions:
`cospi`, `erfcinv`, `normcdf`...
 - ◆ Complete list with error bounds in the CUDA C programming guide
- Starting from CC 2.0
 - ◆ `printf`
 - ◆ `memcpy`, `memset`
 - ◆ `malloc`, `free`: allocate global memory on demand

Device-side intrinsics

- C functions that translate to one/a few machine instructions
 - ◆ Access features not easily expressed in C
- Hardware math functions in single-precision
 - ◆ GPUs have dedicated hardware for reverse square root, inverse, log2, exp2, sin, cos in single precision
 - ◆ Intrinsics: `__rsqrt`, `__rcp`, `exp2f`, `__expf`, `__exp10f`, `__log2f`, `__logf`, `__log10f`, `__sinf`, `__cosf`, `__sincosf`, `__tanf`, `__powf`
 - ◆ Less accurate than software implementation (except `exp2f`), but much faster
- Optimized arithmetic functions
 - ◆ `__brev`, `__popc`, `__clz`, `__ffs`...
Bit reversal, population count, count leading zeroes, find first bit set...
 - ◆ Check the CUDA Toolkit Reference Manual
- We will see other intrinsics in the next part