



ECE 508

Manycore Parallel Algorithms

Lecture 7: Privatization for Graph Search

Objective

- to learn the role of privatization in algorithm scalability
- to learn how to implement queue structures that support efficient dynamic data extraction
- to learn efficient kernel structures to accommodate dynamic variation of working set size

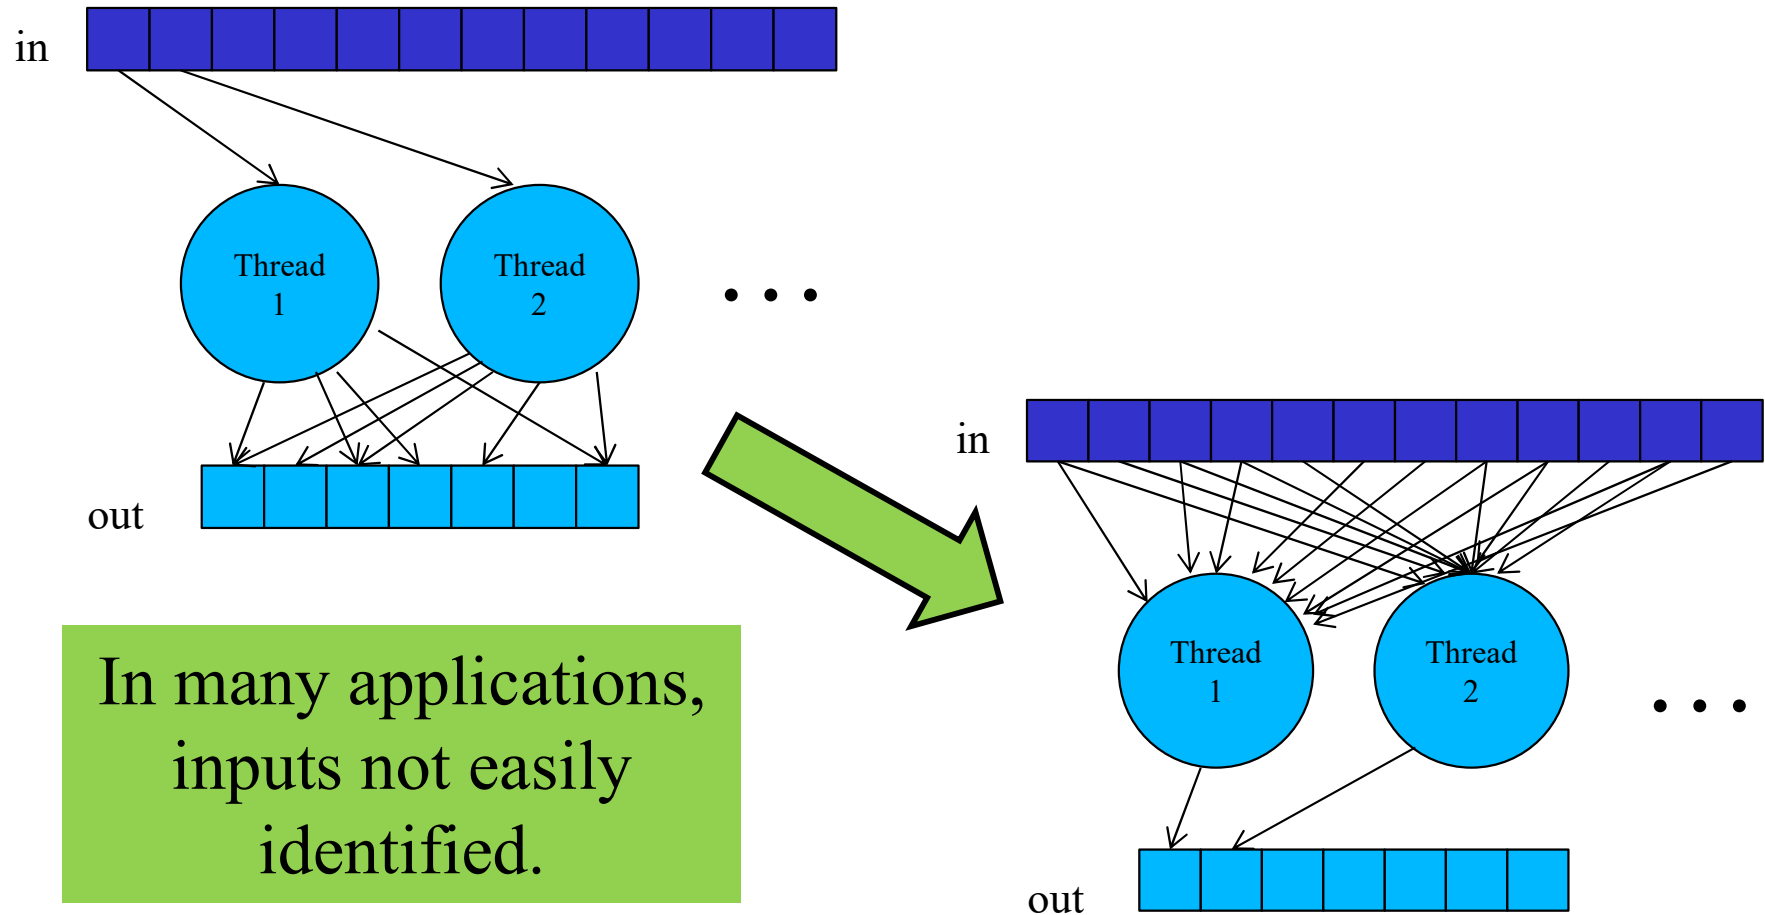
Started with Data-Independent Problems

We **started with** applications

- with **data-independent access patterns**.
- In DCS and GEMM,
 - all accesses are known
 - once you know the problem parameters.

So we **focused on optimizing** the **access** patterns.

Scatter to Gather Transformation



Then Examined Data Reorganization
Using a **cutoff** algorithm was **more challenging**.

But binning sufficed to **reorganize the data**
– so that accesses are again defined
– **before** we start **computing** grid potentials.

There's still **some variation** with variable-sized bins,
but the access patterns are **mostly well-defined**.

Today, Look at Highly Data-Dependent Accesses

Now we look at

- **highly data-dependent access patterns**
- that **cannot be known in advance**.

In other words,

- the only way **to know the** access **pattern**
- is to **perform the computation**.
- No competitive algorithm is known to find the pattern in a different way (to enable, for example, data reorganization).

Example: Finding a Shortest Path in a Graph

Consider finding a shortest path in a graph.

- Finding a path depends on the graph itself.
- Which nodes to examine depends on
 - the edges from the starting point, and
 - the edges from the neighbors, and
 - so on.

Alternatives are Not Competitive and Not Scalable

Can we precompute?

Yes! **Can precompute shortest paths for all pairs**, then optimize accesses for a given graph.

- That's **not** a **competitive** algorithm if all we want is one shortest path.
- And it's **not scalable** regardless.
 - Google maps does not store shortest paths for all pairs of points in the world,
 - Although they probably do maintain derived hierarchical representations of the world map.

Dynamic Extraction: Step Identifies Accesses for Next Step

- We need to solve
 - a problem of **dynamic extraction**:
 - **at every step**, the algorithm **identifies which part** of the data must be **considered in** the **next** step.
- This kind of problem
 - is **extremely difficult for practical architectures**, as it
 - **stresses** the **memory** system **and**
 - **offers** widely **varying** levels of **parallelism** in each step.

(DARPA still has big graph challenges to address this mismatch between algorithms and architectures.)

Start with Breadth-First Search

- We **start** simple, **with BFS**.
 - BFS is the **entry-level computation** for graphs:
 - one of the **simplest** problems,
 - **but** gives **good insights**
 - on how **to address irregular, dynamic behavior**.

Queues and Privatization will Improve Performance


- As part of our discussion,
 - we **will introduce queues**,
 - not so different than queues you've seen,
 - but **with extremely high throughput**.
- The only way to achieve the necessary throughput
 - is **through privatization**:
 - replicating data to expose parallelism.

Privatization is Not a Panacea for Performance

- **Privatization** is probably not new to you:
 - it's **used with histograms** (in 408)
 - **to transform global atomic** operations
 - **into atomic operations on shared memory.**
- In histogram,
 - privatization is relatively easy
 - as the operation is associative and commutative.
- **Privatization doesn't solve all such problems.**

Dynamic Data Extraction Necessary


- For good GPU performance, we **need massive parallelism**.
- With multidimensional data, array indices suffice to enable each thread to find values in parallel.
- With binning, we reorganized data to enable parallel access.
- Now that's not possible: **data must be extracted dynamically**.



Graph Algorithms Often Require Dynamic Extraction

- **Data** for each phase are **dynamically**
 - **determined and extracted** from a bulk data structure.
 - **May** also **need to reorganize** the data for access!
- Graph algorithms are popular examples, and are
 - widely used in EDA (electronic design automation / computer-aided design) and large-scale optimization.
 - We use Breadth-First Search (BFS) as an example.

Graphs Represented as Adjacency Matrices



easier to modify, but
slower to access

	0	1	2	3	4	5	6	7	8
0		1	1						
1				1	1				
2						1	1	1	
3					1				1
4						1			1
5							1		
6									1
7	1						1		
8									

faster to access, but
harder to modify

Main Challenges of Dynamic Data

1. **organize** data **for locality, coalescing, and avoiding contention** during execution
2. available **parallelism may grow and shrink dynamically** during execution
 - different kernel strategies fit different data sizes
 - difficult to encode into a single CUDA kernel launch, in which neither code nor parallelism can change*

*Unless using CUDA's dynamic parallelism, which we discuss later.

Graphs Represented as Adjacency Matrices

Graphs typically **sparse**.

Adjacency matrix can be **compacted into CSR format**—we know how to use that!

	0	1	2	3	4	5	6	7	8
0		1	1						
1				1	1				
2						1	1	1	
3					1				1
4						1			1
5							1		
6									1
7	1						1		
8									

weights[15]

1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

destination[15]

1	2	3	4	5	6	7	4	8	5	8	6	8	0	6
---	---	---	---	---	---	---	---	---	---	---	---	---	---	---

edge_ptrs[10]

0	2	4	7	9	11	12	13	15	15
---	---	---	---	---	----	----	----	----	----

Sparse Matrix Format and Use is Important

- That's why we teach **sparse matrix format and use** in 408: they are **important for all types of graph algorithms** (and many other applications).
- Techniques usually transfer between sparse matrix uses for algebra and for graphs.
- We'll discuss some matrix-based approaches to BFS later.


CSR is the Basic Form for Graphs

- CSR format (as shown) is convenient if you only need to follow the edge directions.
- Can just as easily build CSR for reversed edges (sometimes useful, as we'll see).
- **Not trivial to** transpose / **use in reverse**.
- More complex representations are possible, and trade additional data for access efficiency.

Definition of Breadth-First Search


BFS:

- Given a source node,
- calculate the minimum number of edges that must be traversed to reach a destination node (or to any node).



BFS Used Heavily in VLSI Routing

(Maze Routing example)



How does one find the path?

Backtrack through distance, or use a heuristic to pick.

Example: Prefer straight lines.

More BFS Applications in VLSI CAD


- reachability analysis
- finding connected components
- logic simulation/timing analysis
- logic synthesis
- and so forth

Sequential BFS is Efficient!

ECE students see this algorithm in their first class!

- **Use a queue** to explore the graph (building a logical tree).
 1. Add source node to queue.
 2. Explore queue nodes in order.
 3. For each queue node, add all unvisited neighbors to queue (recording distance and/or predecessor).


Initialize BFS with the Source



distance 0





Exploring the Node 0 to Find 1 and 2






distance 0




distance 1






Exploring the Node 1 to Find 3 and 4



Explore Node 2 to Find 5, 6, and 7




Explore Node 5 to Find 8



BFS Outcome for Source Node 0


Complexity $O(V+E)$ —linear in size of graph!

We also need per-node distance labels (initially -1).



BFS Outcome for Source Node 2

- Execution from source node 2 is entirely different!
(Explored differently, producing different results!)




And Now for a Volunteer from Our Audience...

So ...

Where is the parallelism?

Within one level,
can explore
nodes in parallel.


Subset of nodes
at a level is
called a frontier.



Parallelism Varies Over Algorithm Execution

Even in a tiny graph, **parallelism**

- **starts out small**,
- **grows**, then
- **shrinks again**
- (highly variable).




One Kernel per Frontier


What happens between frontiers?

Simplest answer:

- Wait until frontier complete (**synchronize**),
- but >1 thread block, so
- launch **one kernel per frontier**.



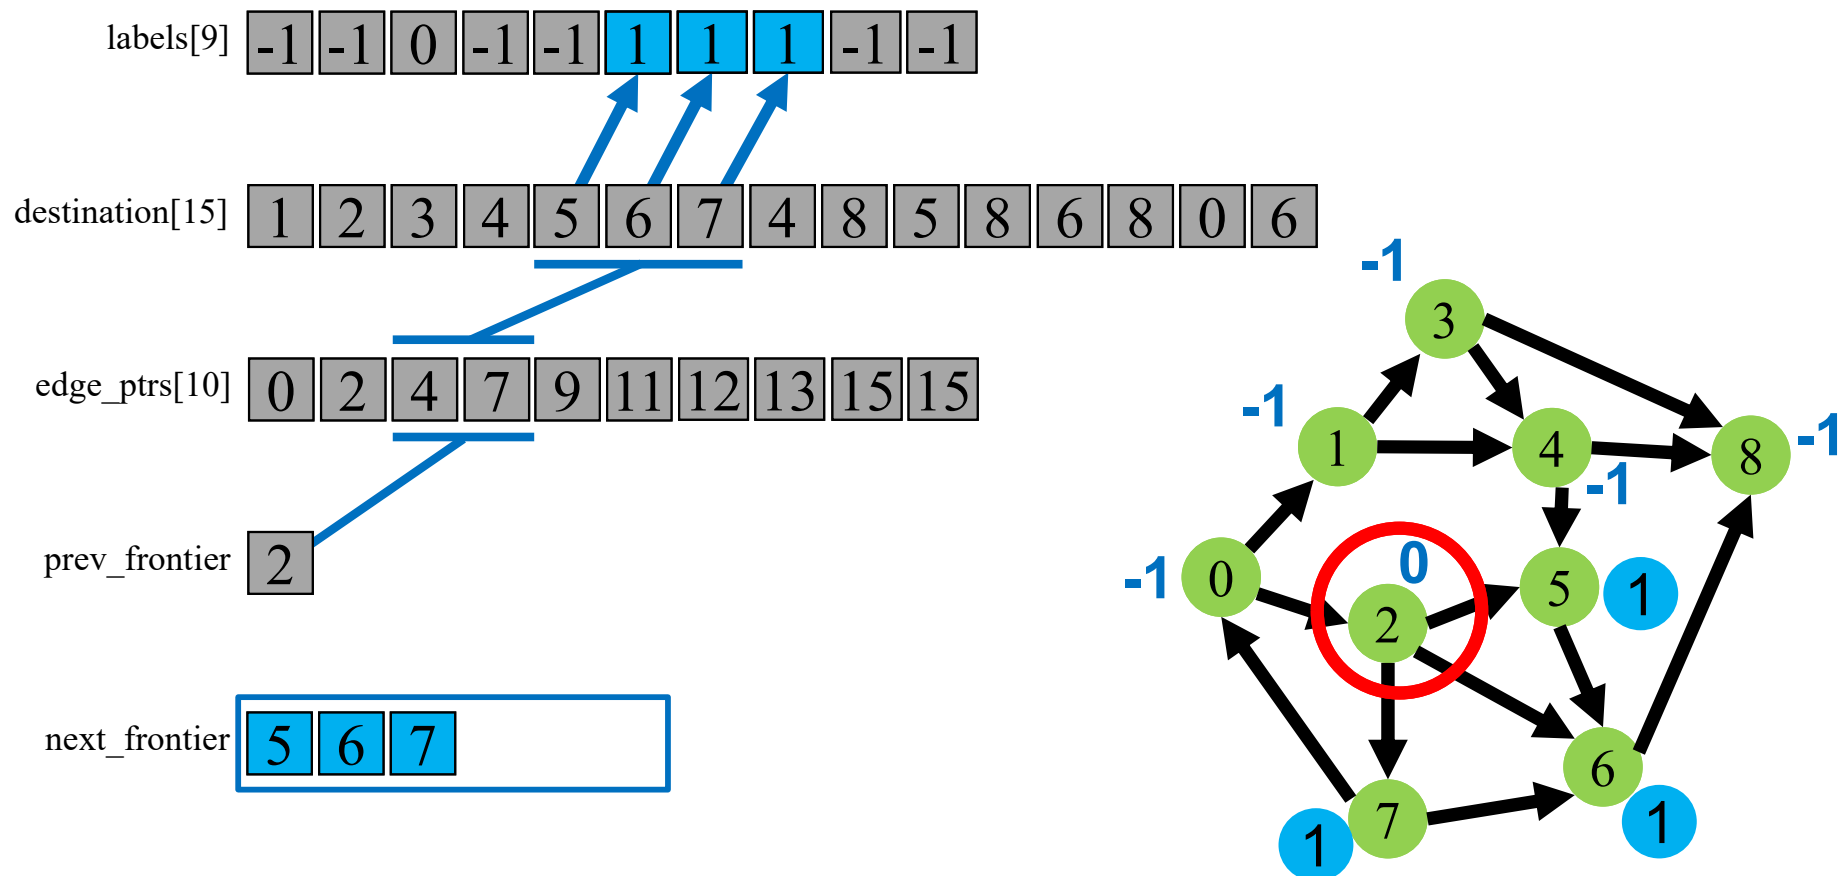
Avoid Backtracking with Distance Values



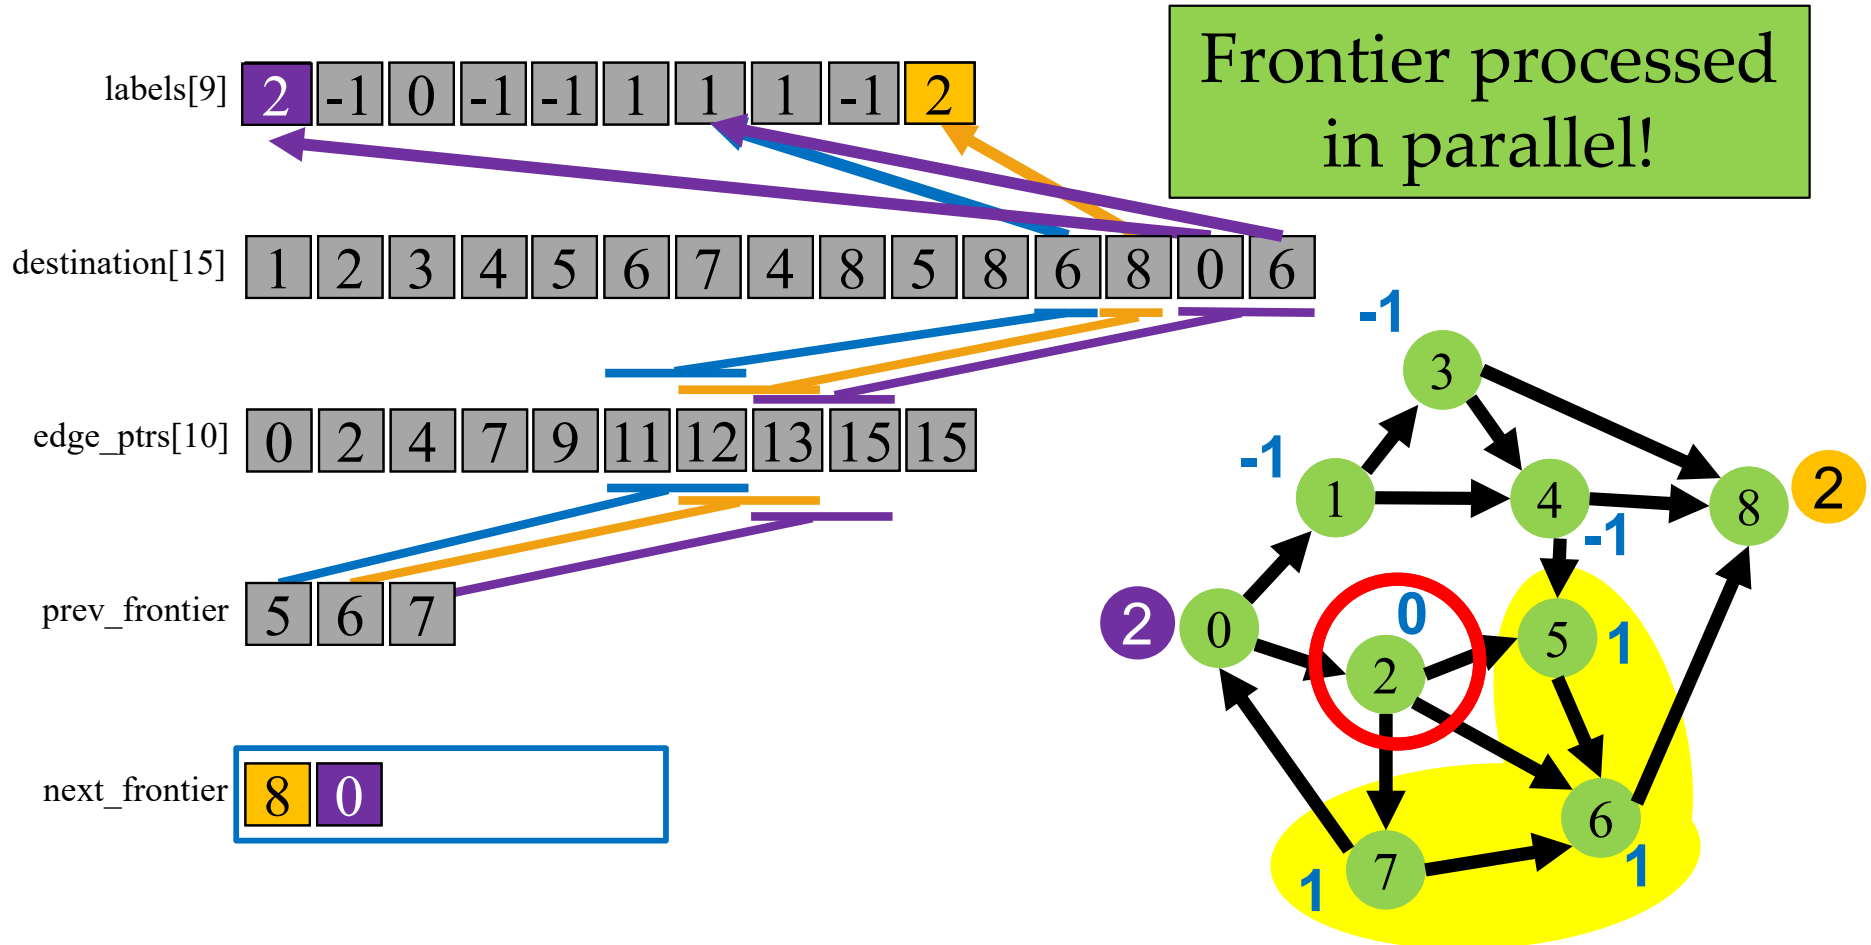
Is there a queue?

- **No**, need to **look at** edge destination **distance**.
- Examples:
 - **7** reaches **0**, but node **0** has distance below **3**
 - **3**, **4**, and **6** reach **8** ... better **add to next frontier atomically**!

Example: Source Node 2, 0-Hop Frontier



Example: Source Node 2, 1-Hop Frontier



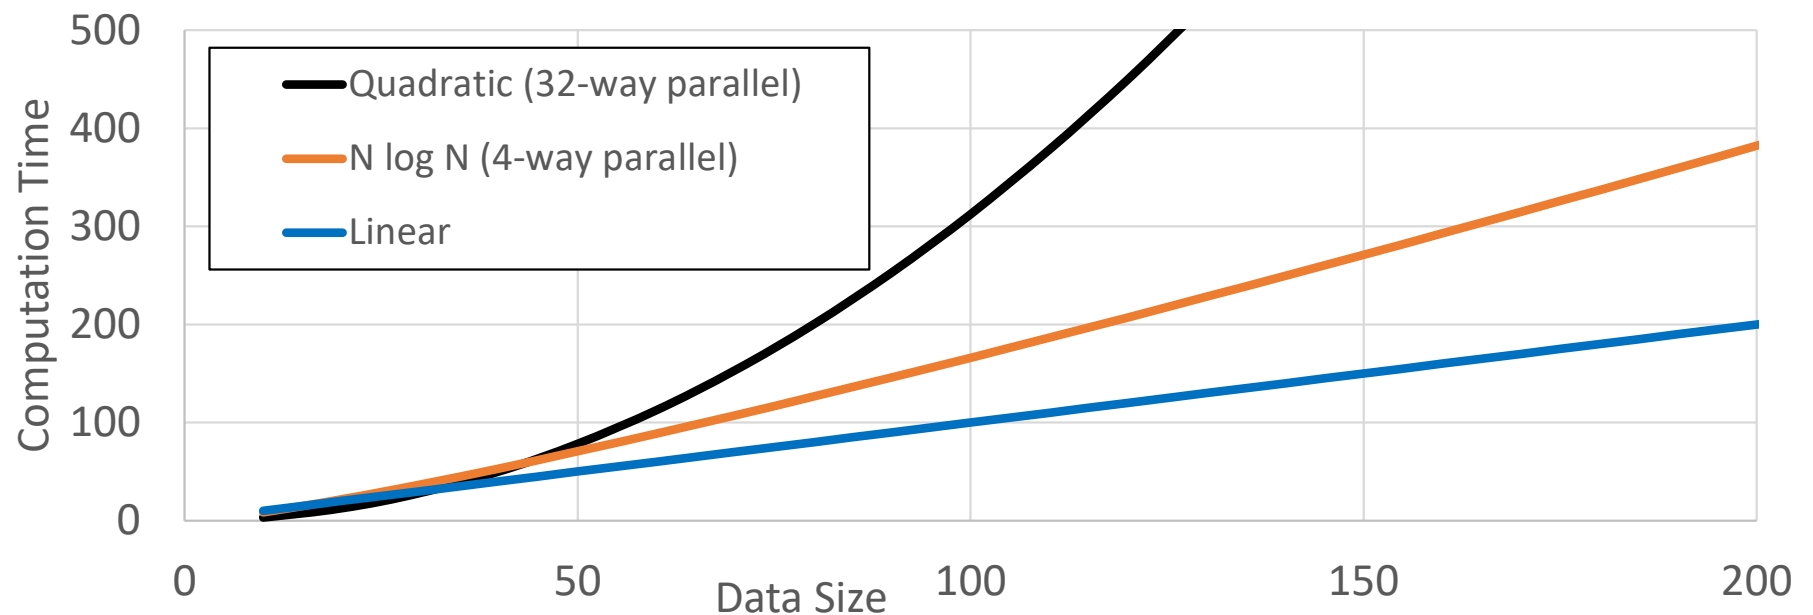
Hard to Map BFS to Our Usual Strategies

The implications of dynamic extraction and variable parallelism?

- Must **find enough parallelism** to keep GPU busy.
- **Not easy to coalesce** accesses to graph data.
- Threads in charge of different nodes from frontier to frontier.
- **Not easy to amortize access cost** through reuse.

Potential Pitfall of Parallel Algorithms

Remember: best sequential code is the baseline!




Massively **parallel $O(N \log N)$** algorithm still **slower than $O(N)$ sequential** approach **on large data** sets.

Narrow Parallelism Fits within a Thread Block

Do we really need one kernel per frontier?

Not quite.

When little parallelism
available, **can work
within one thread block**
and use barriers.



Use CPU to Expand First Few Levels

One thread block makes sense at the start,

- but the **CPU is probably faster, and**
- is typically **used in practice** to start.

The **end is trickier**:


- narrow parallelism doesn't imply completion, so
- may **not** be **easy to shift models**.

Correctness Does Not Require Synchronization

But ... do we really even need to synchronize?

Not exactly.

Updating a node with an overly long distance **creates extra work, but we will eventually fix it.**



Many Challenges if Synchronization Dropped

The challenges are...

- to **avoid** creating **too much useless work**,
- to **prioritize important work** (even if we can prioritize, how can we express priority in CUDA?),
- to **squash useless work** as quickly as possible, and
- to **obey CUDA's** memory consistency **rules**.

We won't try it in lecture.

Always Consider the Big Picture

Always **important to see the big picture.**

- Need to solve many graphs?
Let parallelism interleave naturally to occupy GPU.
- Many sources on the same graph?
Faster to solve all pairs shortest-paths?
- One source on one graph ASAP?
Isn't possibly useless work better than idling?

Early Attempts to Parallelize BFS

- **Node-Oriented Parallelization**
threads assigned statically to nodes
- **Matrix approaches**
 - Edge-Oriented, Adjacency-Matrix-Based Parallelization
 - matrix multiplication for frontier propagation

One Thread per Node on 8800 GTX

Early (**2007!**) work: BFS with **parallelization over nodes**.*

- Arrays track whether nodes are in frontier and visited.
- In an iteration, **each thread tests** whether **node** is **in frontier**.
- If so, thread
 - **Marks node visited** and **removes** node **from frontier**.
 - Goes over **all unvisited neighbors**
 - **Adding to frontier** and **updating distance**.
- Claims to predate atomics, but ... algorithm requires atomics for correctness (and doesn't use them).

*P. Harish, P.J. Narayanan, “Accelerating Large Graph Algorithms on the GPU Using CUDA,” HiPC 2007, LNCS 4873, pp. 197-208, 2007.

Complexity Too High for Many Graphs

- **L iterations** required,
 - where **L** is the maximum shortest path length,
 - $O(\text{diameter of graph})$, 1000+ for large graphs.
- **complexity $O(VL + E)$**
 - too much extra work!
 - especially for sparse / high diameter graphs
 - usually **slower than CPU for large graphs**
- And, if you add atomics...

BFS Using Sparse Matrix-Vector Multiply


CAD community, **2009***

- **BFS via** sparse matrix-vector multiply (SpMV):
- matrix is “transposed” adjacency in **CSR**:
directed edges from columns to rows, and
- vector is **reachable set** (in some # of hops)
with path count per element.
- SpMV gives next reachable set with path count per element.
- Path counts also have value in EDA.

*Y. Deng, B.D. Wang, S. Mu, “Taming Irregular EDA Applications on GPUs,” ICCAD '09, pp. 539-546, Nov. 2009.

Example: Source Node 2 SpMV

0	0	0	1	0	0
1	0	0	0	1	0
2	1	0	0	0	0
3	0	0	0	0	1
4	0	0	0	0	1
5	0	1	0	0	0
6	0	1	2	0	0
7	0	1	0	0	0
8	0	0	1	2	0
	v	Av	A^2v	A^3v	A^4v



Slower than CPU for Large Graphs

Complexity $O((V+E)L)$ (L executions of SpMV)

- **Slower than CPU** for large graphs.
- Can precompute matrix powers to speed up,
 - but matrix-matrix is slow, and
 - matrix powers are larger.

$A^k \mathbf{v}$ is a superset of nodes at distance k .

- If desired, **remove nodes** at distance $< k$
between multiplications.
- Added algorithm-specific work makes
algebraic solutions less attractive.

Simplification Produces a Gather Pattern

Can also **reduce computation** (losing path counts):

1. **Start with vector** of “**infinity**” (except for source).
2. Check vector value for each row and
skip computation of non-infinite row values,
3. Check that edge destinations have value “infinity”; **stop if non-infinite value found**, producing (dest value + 1).

With these changes,

- computation is a **Gather pattern** (no atomics!):
- in each iteration, elements “gather” from neighbors.

Complexity still $O(V+EL)$, but constants smaller.

Need More Flexibility and Information

To address the problem more generally,
need to be more flexible.

- Keep the big picture in mind:
 - What needs to get done?
 - What are the characteristics of the data?
 - Where are the data?
 - What are the capabilities of the hardware?



Make Timely Choices from Multiple Algorithms

In practice, often need several algorithms.

Sometimes, we can **choose** the right one
when writing the application.

Sometimes, we need to
determine the right one **dynamically.**

Fortunately, this **approach is well understood!**


Example ca. 1990: CMSSL (CM-2)

Lennart Johnson put together the Connection Machine Scientific Software Library with these goals in mind.

- **Track** dynamic information about **distribution of data** (in a distributed memory machine).
- **Given** the **next** desired **operation**,
 - **compute cost models** for several algorithms based on data distribution, then
 - **choose** the **fastest**.
- Execute and **update info on** dynamic data **distribution**.
- Could optimize over several operations, too.



Examples Include FFTW and NVIDIA Libraries

- **FFTW** is a more **recent example** focused on FFT (DFT). 
- **NVIDIA's libraries** (ex: Cub) are also **examples**.
 - Use C++ templates to reduce overhead.
 - Hard to read the code until you've seen some of the techniques explained.






Another Example: Matrix Multiplication

Algorithm variations come up in 408 projects, too.

Consider **dense matrix multiplication**:

- **approaches** we've discussed
- **work** well for **square matrices...**






Certain Shapes Require More Parallelism

What about this shape?

Not many outputs to parallelize!

Instead, parallelize over
inner-product terms and
use reductions per output.



Be Sure to Examine Known Techniques

What about the architecture?

Important to realize that

- distributed resources in modern architectures are an outcome of the laws of physics.
- They haven't changed recently.
- **Techniques from distributed memory machines**
 - **can** also **be applied in GPUs**, as
 - “shared” memory is actually private scratchpads distributed amongst SMs, as you know.

Alternate Between Local and Global

Distributed memory demands hierarchy.

Fast data are implicitly private,
so **privatization is key.**

- Need to **execute fast locally.**
- Need to **share information globally.**
- Most effective algorithms **go back and forth (local/global)** as few times as possible.

Example: Sorting on a Cluster

A fun variant of these ideas

SroNtwo

- won the world record for sorting back in 1997.*
- The context—a cluster of workstations—forces privatization and hierarchy.
- Sample globally for splitters, split locally, exchange data globally, sort locally.

*A.C. Arpaci-Dusseau, R.H. Arpaci-Dusseau, D.E. Culler, J.M. Hellerstein, D.A. Patterson, “High-Performance Sorting on Networks of Workstations,” SIGMOD ’97, May, 1997.

Need a Queue to Adapt Sequential BFS

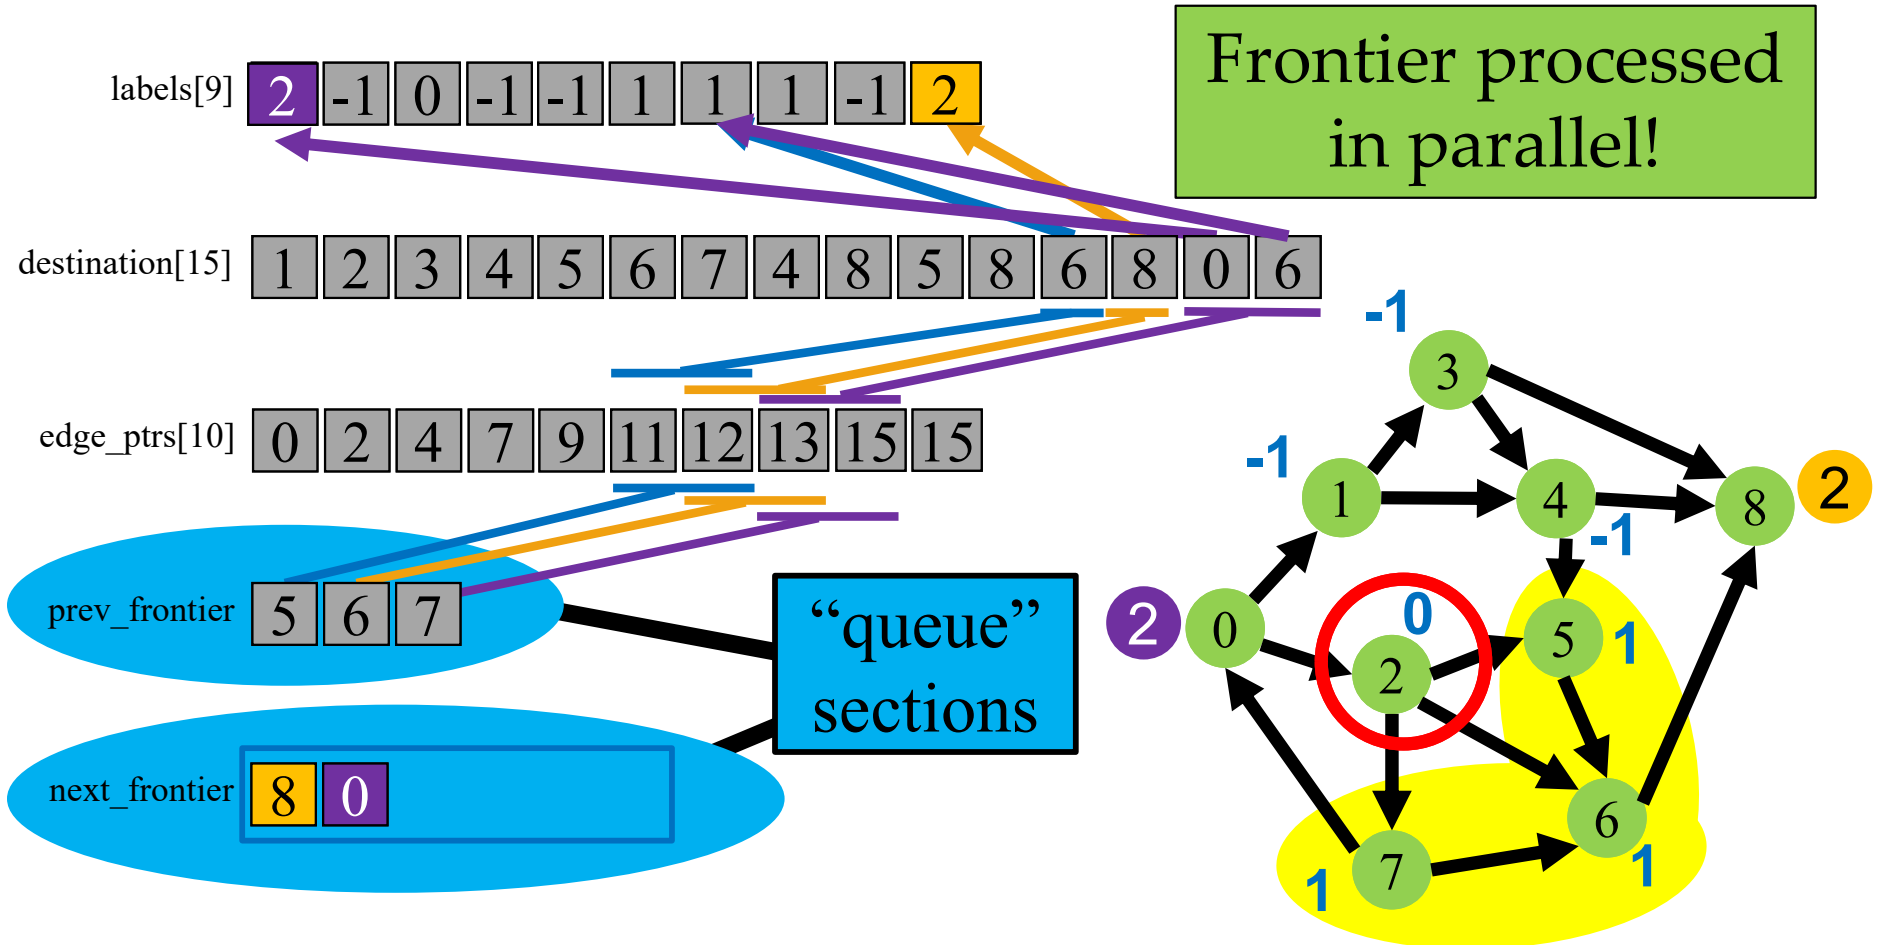
Back to the problem at hand: BFS.

We have...

- **irregular data** (somewhat tamed with CSR),
- **irregular access** patterns (data-dependent), and
- **varying** levels of **parallelism**.

Sequential **solution relies on queue**.

Review Execution on a Simple Graph



Need a More General Technique

- **Why “queue” instead of queue?**
 - **Serializable** parallel queue
 - Parallel operations **“appear”** to have executed **in some serial order** on a sequential queue.
 - Same order **observed by all users**.
- **BFS** has little computation—**requires** efficiency:
 - **hierarchical scalable queue** implementation,
 - and **hierarchical kernels**.


Global Queue is Too Slow

First attempt: serializable queue in global memory.

- Parallelize over nodes in frontier.
- Dequeue in parallel.
- Enqueue using global atomic operations:
 - poor coalescing, and
 - low global atomic throughput.
- Complexity $O(V+E)$.
- No speedup.

prev_frontier

5	6	7
---	---	---



next_frontier



Can Reduce Pressure by Compacting Later

GPU community did try alternatives.*

- **Separate generation and compaction**, using a
- **fixed-size output array** per thread, thus
- **no atomics** and no contention.
- Overhead of **compaction may outweigh gains** in some cases (with little work per node, as in BFS).

*C.Lauterbach, M. Garland, S. Sengupta, D. Luebke, D. Manocha,
“Fast BVH Construction on GPUs,” EUROGRAPHICS 2009, 2009.


Leverage Throughput of Shared Memory

Maxwell generation of GPUs

- substantially **increased throughput**
- **for shared memory atomics,**
- enabling GPUs to leverage older techniques.

Graph algorithms were a big reason, along with image processing, vision, and feature extraction.

Privatization Reduces Contention




Reduce contention by

- **replicating data**,
- accessing with fewer (or 1) threads,
- and **aggregating** more efficiently **later**.

This technique is **called privatization**.

Privatization Enables Use of Private Memories




Privatization **requires**


- **extra storage**, but
- **enables** copies to reside in **private memories**, such as
- shared memory in GPUs.

Privatization Applies to BFS Queues

- For BFS, each **thread processes one** (or more) **node**.
- When **node found** for next frontier,
 - **insert to local** queue in shared memory
 - (shared among threads in block).
- When **all threads** in block **done**,
 - **allocate space** (one atomic) in global queue, then
 - **copy** from **local to global**.



Two-level Hierarchy



Result is **two-level hierarchy**:


- block queue (**b-queue**)
 - in **shared memory**
 - **used by** threads in a **block**,
- global queue (**g-queue**)
inserted only **when block completes**.

B-Queue Must Fit in Shared Memory

Shared memory is limited.

One **option: overflow into global** queue.

Another **option: flush and continue**. May be **tricky to synchronize**—want all threads to cooperate in flush.



Benefits of Two-Level Hierarchy

privatization benefits:

- **contention on global queue**
reduced to one operation per block
- **contention on block queues**
still a problem (but **less** so)


shared memory benefits:

- **higher throughput** atomic queue insertions
- cooperative, **coalesced writes** to global queue.

Threads in Warp ALWAYS Collide!

b-queue contention is a problem.

- The worst part? **SIMD execution** to one resource!
- Threads executed together in a cycle are **guaranteed to collide**.




Collisions within a Warp are Slow


What happens?

Threads normally executed in a cycle are **serialized across many cycles** / issue slots.

All threads inserting to b-queue must do so in separate cycles.



Solution: Add a Warp-level Queue



- **Solution? Add** a level of **hierarchy!**
- Provide **queues for SIMD lanes**:
 - shared across warps in a block;
 - **reduces conflicts** in a cycle;
 - more queues, fewer conflicts (up to warp size).

Warp Queue Reduces SIMD Contention

Warp queues (**w-queues**)

- placed **in shared memory** and
- **split across SIMD lanes** in warp (executed in same cycle), but
- **shared across warps** in thread block (less likely to conflict).
- Example: use **wQueue[threadIdx.x % 8]**.
- **Must still use atomic** operations (different warps *can* execute simultaneously).

BFS Extended to Use Warp Queues

- For BFS, each **thread processes one** (or more) **node**.
- When **node found** for next frontier,
 - **insert to warp queue** in shared memory
 - (shared among threads in block).
- When **all threads** in block **done**,
 - **consolidate warp queues into block queue**,
 - allocate space (one atomic) in global queue, then
 - copy from local to global.

Need to Think About Bank Conflicts

Problem: bank conflicts!

Multi-ported memory is expensive.

- No one builds memory that way.
- Instead, SM's **shared memory has many banks**.

Adjacent memory **locations**

- placed **in different banks**, and
- serviced in parallel.

Strided memory **locations**

- may hit the same bank(s) and
- **create bank conflicts**.

Need to Think About Bank Conflicts

Why didn't we mention earlier (in 408)?

- Dropped after first few years, as
- usually **less important for performance**.

What does this have to do with w-queues?


- Likely to use 2^k w-queues of length 2^m , but
- **power-of-2 strides of lead to bank conflicts**.

Solution: interleave (swap dimensions)!

Hierarchical Queue Memory Management

Shared Memory

- divided amongst w-queues and b-queues.
- interleaved w-queue layout
reduces bank conflicts
(example with $N = 8$)




Global Memory holds g-queue;

writes from b-queue to g-queue coalesced


Texture memory stores graph structure

(random access, no coalescing)


Three-Level Hierarchy Illustration



Interleaved Warp-Level Queue Layout



Conceptual



Physical

Notice the sparse filling due to imbalance between w-queues.

Dividing up the Resources

How many w-queues?

More w-queues

- requires **more space** (or more overflows)
- leads to **more load imbalance**, and
- causes **more bank conflicts** (serialization).

Fewer w-queues causes

more SIMD atomic conflicts (serialization).

Be sure to allow >1 thread block per SM!

Benefits of Three-Level Hierarchy

privatization benefits:

- **contention on g-queue** and **b-queues reduced** to one operation per block
- **less contention on w-queues**

shared memory benefits:

- **higher throughput** atomic queue insertions
- cooperative, **coalesced writes** to global queue.

Queue Strategy Reuse and Limitations

Techniques

- **applicable to any** sequential **data structure**
- **provided that operations can be reordered**
- (harder for a heap, for example, hence many more BFS papers than shortest path papers!).

Local queues

- **limited by** capacity of **shared memory**.
- If node degree is bounded, can estimate well.
- Later, we'll use dynamic parallelism to bound.

Thoughts on the Details

Why keep the b-queue?

- w-queues overflow to b-queue, which
- **Smooths load** before overflow to g-queue.

Why not copy w-queue directly to g-queue?

- That is, scan across w-queues and b-queue, allocate space for all once, then copy all using scan results.
- **Not clear.** Possibly control complexity. **Try it?**

More Thoughts on the Details


Why not use Cub's warp scans to avoid conflicts instead of adding w-queues?

Fairly new ... try it?

Why not use register tiling to accumulate 4-8 nodes first?

Try it? (*Hint: loop must be unrolled...*)

Using CPU and Kernels to Deal with Variable Parallelism




How can we reduce the cost of kernel launches?

- As mentioned before, **use CPU** to explore first few levels.
- Then **use multiple kernels** (and dynamic parallelism) to adjust to available parallelism.

Hierarchical Kernel Arrangement

- **Customize** kernels **based on** the **size of frontiers**.
- Use barrier synchronization (`__syncthreads()`) when frontier is small.




First Kernel: So Long as Frontier Fits a Block

Kernel 1: small-sized frontiers

- **one thread block**
- **`__syncthreads()`**
between frontiers
- uses **only w-queues**
and b-queue

Switch kernels when b-queue overflows to g-queue.



Second Kernel Handles Large Frontiers

Kernel 2: large frontiers

- **multiple thread blocks**
- **kernel re-launch** to synchronize (overhead acceptable compared to large frontier processing time)
- uses full hierarchy of queues

Or use **dynamic parallelism** (Kepler or later GPUs)

- **to launch Kernel 2** from Kernel 1
- **when** frontier size **threshold exceeded**.

Still Need Global Atomics ... ?

Actually, we didn't eliminate global atomics.

- Using **hierarchical queues**
 - reduces contention and
 - improves throughput.
 - However, it also
decouples the threads' actions.

Competing for Insertion Rights is Atomic

**What if two nodes in a frontier
share an unvisited neighbor?**

We **want** only **one copy in** the **next frontier!**

Must coordinate insertion.

Here we expect little contention,

- so continue to **use global atomics**
- **to guarantee** desired behavior.

(Use atomic exchange on visited array in Lab 5.)

Performance Data: Not as Easy as It Sounds...

Lumetta's implementations

We'll do this as times to beat...

10^6 nodes of degree 8	0.1029 msec
10^8 nodes of degree 4	17.2 msec
10^8 nodes of degree 8	30.8 msec

Some hints: (use exclusive queue for final timing!)

- lots of threads
- plenty of queue
- don't think you're going to fit everything at once
- just use bqueues
- try inventing some tricks...



ANY QUESTIONS?