

ECE408/CS483/CSE408 Spring 2020

Applied Parallel Programming

Lecture 4: Memory Model

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

1

1

Objective


- To learn the basic features of the memories accessible by CUDA threads
- To prepare for MP-2 - basic matrix multiplication
- To learn to evaluate the performance implications of global memory accesses

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

2

2

The Von-Neumann Model



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

3

3

Instructions are Stored in Memory

- Every instruction needs to be fetched from memory, decoded, then executed.
- Instruction processing breaks into steps:
Fetch | Decode | Execute | Memory
- Instructions come in three flavors:
Operate, Data Transfer, and Control Flow.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

4

4

Example: Processing an Add Instruction

- Example of an (LC-3) operate instruction:
ADD R1, R2, R3
- meaning:
 - read R2 and R3
 - add them as unsigned/2's complement
 - write sum to R1
- Instruction processing for an operate instruction:
Fetch | Decode | Execute | Memory

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

5

5

Example: Processing a Load Instruction

- Example of an (LC-3) data transfer instruction:
LDR R4, R6, #3 ; a load
- meaning:
 - read R6
 - add the number 3 to it
 - load the contents of memory at the resulting address
 - write the bits to R4
- Instruction processing for a load instruction:
Fetch | Decode | Execute | Memory


© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

6

6

Registers vs Memory

- Registers
 - Fast: 1 cycle; no memory access required
 - Few: hundreds for CPU, O(10k) for GPU SM
- Memory
 - Slow: hundreds of cycles
 - Huge: GB or more



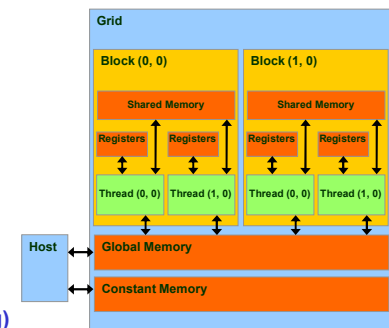
© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

7

7

Programmer View of CUDA Memories

- Each thread can:
 - read/write per-thread **registers (~1 cycle)**
 - read/write per-block **shared memory (~5 cycles)**
 - read/write per-grid **global memory (~500 cycles)**
 - read/only per-grid **constant memory (~5 cycles with caching)**



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

8

8

CUDA Variable Type Qualifiers

Variable declaration	Memory	Scope	Lifetime
<code>int LocalVar;</code>	register	thread	thread
<code>__device__ __shared__ int SharedVar;</code>	shared	block	block
<code>__device__ int GlobalVar;</code>	global	app.	application
<code>__device__ __constant__ int ConstantVar;</code>	constant	app.	application

- `__device__`
 - optional with `__shared__` or `__constant__`
 - not allowed by itself within functions
- Automatic variables with no qualifiers
 - in registers for primitive types and structures
 - in global memory for per-thread arrays

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

9

9

Next Application: Matrix Multiplication

- Given two $Width \times Width$ matrices, M and N,
 - we can multiply M by N
 - to compute a third $Width \times Width$ matrix, P:
 - $P = MN$

In terms of the elements of P, matrix multiplication implies computing...

$$P_{ij} = \sum_{k=1}^{Width} M_{ik} N_{kj}$$

© Steven S. Lumetta, 2020
ECE408/CS483/University of Illinois at Urbana-Champaign


10

10

Matrix Multiplication

$$P_{ij} = \sum_{k=1}^{Width} M_{ik} N_{kj}$$

- Graphically, imagine
 - taking each element in a row of M,
 - multiplying it by the corresponding element in a column of N, and
 - summing up the products.
- Do that for every row and every column to produce P.



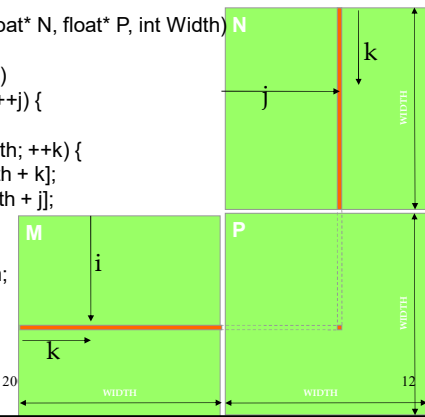
© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
University of Illinois at Urbana-Champaign

11

Matrix Multiplication Example A Simple Host Version in C

// Matrix multiplication on the (CPU) host in single precision

```
void CPUMatMul(float* M, float* N, float* P, int Width)
{
    for (int i = 0; i < Width; ++i)
        for (int j = 0; j < Width; ++j) {
            float sum = 0;
            for (int k = 0; k < Width; ++k) {
                float a = M[i * Width + k];
                float b = N[k * Width + j];
                sum += a * b;
            }
            P[i * Width + j] = sum;
        }
}
```



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
University of Illinois at Urbana-Champaign

12

Parallelize Elements of P

- What can we parallelize?
 - start with the **two outer loops**
 - parallelize **computation of elements of P**
- What about the inner loop?
 - Technically, floating-point is NOT associative.
 - The **parallel sum** is called a **reduction**—we'll come back to it in a few weeks.
 - For now, **use a single thread for each P_{ij}** .

© Steven S. Lumetta, 2020
ECE408/CS483/University of Illinois at Urbana-Champaign

13

13

Tricky Questions about Floating-Point

A question for you:

What is $2^{-30} + (1 - 1)$?

Quite tricky, I know. But yes, it's **2^{-30}** .

Another question for you:

What is $(2^{-30} + 1) - 1$?

That's right. It's **0**.

At least it is with floating-point.

ECE 120: Introduction to Computing

© 2016 Steven S. Lumetta. All rights reserved. 14

14

IEEE Floating-Point is Not Associative

Why?

Take CS357: Numerical Methods

Our first sum was **$(2^{-30} + 1)$** .

To hold the integer 1, the bit pattern's exponent must be **2^0** .

But, the mantissa for single-precision floating point has only **23 bits**.

And thus represents powers down to **2^{-23}** .

The **2^{-30}** term is lost, giving **$(2^{-30} + 1) = 1$** .

So $2^{-30} + (1 - 1) \neq (2^{-30} + 1) - 1$.

ECE 120: Introduction to Computing

© 2016 Steven S. Lumetta. All rights reserved. 15

15

Compute Using 2D Blocks in a 2D Grid

- **P** is 2D, so organize threads in 2D as well:
 - Split the output **P** into square **tiles**
 - of size **$\text{TILE_WIDTH} \times \text{TILE_WIDTH}$**
 - (a preprocessor constant).
 - **Each thread block produces one tile** of TILE_WIDTH^2 elements.
 - Create **$\lceil \text{Width} / \text{TILE_WIDTH} \rceil^2$** thread **blocks** to cover the output matrix.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

16

16

Example: Width 8, TILE_WIDTH 2

P _{0,0}	P _{0,1}	P _{0,2}	P _{0,3}	P _{0,4}	P _{0,5}	P _{0,6}	P _{0,7}
P _{1,0}	P _{1,1}	P _{1,2}	P _{1,3}	P _{1,4}	P _{1,5}	P _{1,6}	P _{1,7}
P _{2,0}	P _{2,1}	P _{2,2}	P _{2,3}	P _{2,4}	P _{2,5}	P _{2,6}	P _{2,7}
P _{3,0}	P _{3,1}	P _{3,2}	P _{3,3}	P _{3,4}	P _{3,5}	P _{3,6}	P _{3,7}
P _{4,0}	P _{4,1}	P _{4,2}	P _{4,3}	P _{4,4}	P _{4,5}	P _{4,6}	P _{4,7}
P _{5,0}	P _{5,1}	P _{5,2}	P _{5,3}	P _{5,4}	P _{5,5}	P _{5,6}	P _{5,7}
P _{6,0}	P _{6,1}	P _{6,2}	P _{6,3}	P _{6,4}	P _{6,5}	P _{6,6}	P _{6,7}
P _{7,0}	P _{7,1}	P _{7,2}	P _{7,3}	P _{7,4}	P _{7,5}	P _{7,6}	P _{7,7}

Each block has
2*2 = 4 threads.

WIDTH/TILE_WIDTH = 4
Use 4*4 = 16 blocks.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/ University of Illinois at Urbana-Champaign

17

17

Example: Same Matrix, Larger Tiles (Width 8, TILE_WIDTH 4)

P _{0,0}	P _{0,1}	P _{0,2}	P _{0,3}	P _{0,4}	P _{0,5}	P _{0,6}	P _{0,7}
P _{1,0}	P _{1,1}	P _{1,2}	P _{1,3}	P _{1,4}	P _{1,5}	P _{1,6}	P _{1,7}
P _{2,0}	P _{2,1}	P _{2,2}	P _{2,3}	P _{2,4}	P _{2,5}	P _{2,6}	P _{2,7}
P _{3,0}	P _{3,1}	P _{3,2}	P _{3,3}	P _{3,4}	P _{3,5}	P _{3,6}	P _{3,7}
P _{4,0}	P _{4,1}	P _{4,2}	P _{4,3}	P _{4,4}	P _{4,5}	P _{4,6}	P _{4,7}
P _{5,0}	P _{5,1}	P _{5,2}	P _{5,3}	P _{5,4}	P _{5,5}	P _{5,6}	P _{5,7}
P _{6,0}	P _{6,1}	P _{6,2}	P _{6,3}	P _{6,4}	P _{6,5}	P _{6,6}	P _{6,7}
P _{7,0}	P _{7,1}	P _{7,2}	P _{7,3}	P _{7,4}	P _{7,5}	P _{7,6}	P _{7,7}

Each block has
4*4 = 16 threads.

WIDTH/TILE_WIDTH = 2
Use 2*2 = 4 blocks.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

18

18

Kernel Invocation (Host-side Code)

```
// TILE_WIDTH is a #define constant
dim3 dimGrid(ceil((1.0*Width)/TILE_WIDTH),
             ceil((1.0*Width)/TILE_WIDTH), 1);
dim3 dimBlock(TILE_WIDTH, TILE_WIDTH, 1);

// Launch the device computation threads!
MatrixMulKernel<<<dimGrid, dimBlock>>>
    (Md, Nd, Pd, Width);
```

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/ University of Illinois at Urbana-Champaign

19

19

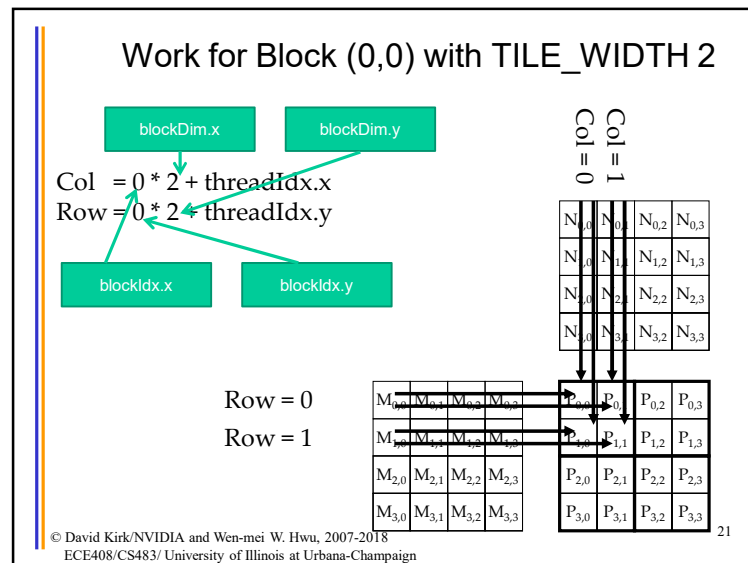
Kernel Function

```
// Matrix multiplication kernel – per thread code
__global__ void MatrixMulKernel(float* d_M, float* d_N, float* d_P, int Width)
{
    // Pvalue is used to store the element of the matrix
    // that is computed by the thread
    float Pvalue = 0;
```

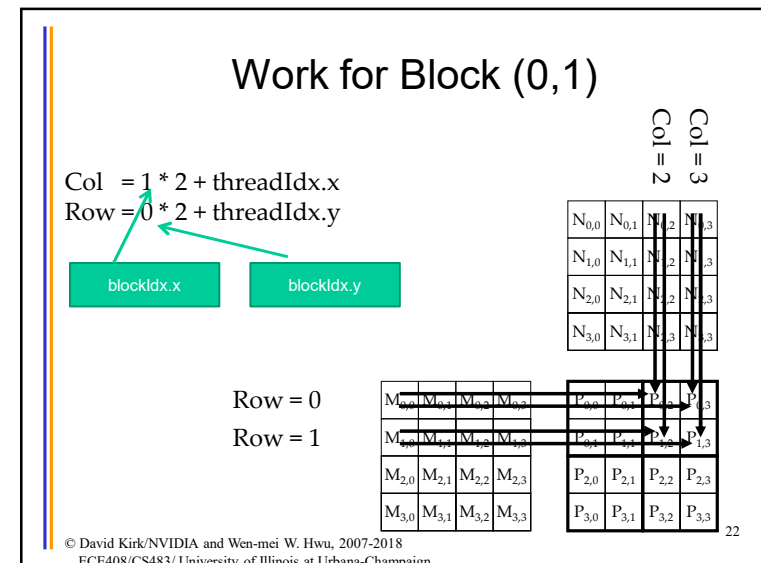
© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/ University of Illinois at Urbana-Champaign

20

20



21



22

A Simple Matrix Multiplication Kernel

```

global__ void MatrixMulKernel(float* d_M, float* d_N, float* d_P, int Width)
{
    // Calculate the row index of the d_P element and d_M
    int Row = blockIdx.y*blockDim.y+threadIdx.y;
    // Calculate the column idenx of d_P and d_N
    int Col = blockIdx.x*blockDim.x+threadIdx.x;

    if ((Row < Width) && (Col < Width)) {
        float Pvalue = 0;
        // each thread computes one element of the block sub-matrix
        for (int k = 0; k < Width; ++k)
            Pvalue += d_M[Row*Width+k] *
                      d_N[k*Width+Col];
        d_P[Row*Width+Col] = Pvalue;
    }
}
  
```

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/ University of Illinois at Urbana-Champaign

23

Memory Bandwidth is Overloaded!

- That's a **simple implementation**:
 - GPU kernel is the **CPU code**
 - with the **outer loops replaced**
 - with per-thread **index calculations!**
- Unfortunately, performance is quite bad.
- Why?
- With the given approach,
 - global **memory bandwidth**
 - can't** supply enough data
 - to **keep the SMs busy!**

© Steven S. Lumetta, 2020
ECE408/CS483/ University of Illinois at Urbana-Champaign

24

Where Do We Access Global Memory?

```
global__ void MatrixMulKernel(float* d_M, float* d_N, float* d_P, int Width)
{
    // Calculate the row index of the d_P element and d_M
    int Row = blockIdx.y*blockDim.y+threadIdx.y;
    // Calculate the column index of d_P and d_N
    int Col = blockIdx.x*blockDim.x+threadIdx.x;

    if ((Row < Width) && (Col < Width)) {
        float Pvalue = 0;
        // each thread computes one element of the block sub-matrix
        for (int k = 0; k < Width; ++k)
            Pvalue += d_M[Row*Width+k] * d_N[k*Width+Col];
        d_P[Row*Width+Col] = Pvalue;
    }
}
```

accesses to global memory

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

25

25

Each Thread Requires 4B of Data per FLOP

- Each threads access global memory
 - for elements of **M** and **N**:
 - 4B each, or 8B per pair.
 - (And once TOTAL to **P** per thread—ignore it.)
- With each pair of elements,
 - a thread does a single multiply-add,
 - 2 FLOP—floating-point operations.
- So for every FLOP,
 - a thread needs 4B from memory:
 - 4B / FLOP.

© Steven S. Lumetta, 2020
ECE408/CS483/University of Illinois at Urbana-Champaign

26

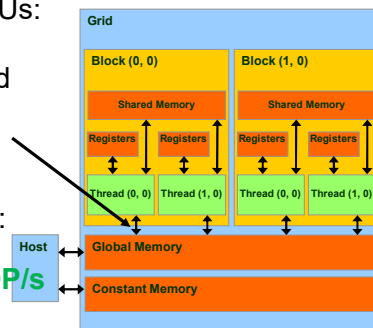
26

150 GB/s Bandwidth Implies 37.5 GFLOPs

- One generation of GPUs:
 - 1,000 GFLOP/s of compute power, and
 - 150 GB/s of memory bandwidth.
- Dividing bandwidth by memory requirements:

$$\frac{150 \text{ GB/s}}{4 \text{ B/FLOP}} = 37.5 \text{ GFLOP/s}$$

which limits computation!



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

27

27

What to Do? Reuse Memory Accesses!

But 37.5 GFLOPs is a limit.

In an actual execution,

- memory is not busy all the time, and
- the code runs at about 25 GFLOPs.

To get closer to 1,000 GFLOPs

- we need to drastically cut down
 - accesses to global memory.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

28

28

Two vertical lines, one blue and one orange, are positioned on the left side of the slide.

**ANY MORE QUESTIONS?
READ CHAPTER 4!**

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/ University of Illinois at Urbana-Champaign

29