

ECE408/CS483/CSE408 Spring 2020

Applied Parallel Programming

Lecture 5: Locality and Tiled Matrix Multiplication

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

1

1

Objective

- To learn to evaluate the performance implications of global memory accesses
- To prepare for MP-3: tiled matrix multiplication
- To learn to assess the benefit of tiling

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

2

2

The Problem: Accesses to Global Memory

```
global__ void MatrixMulKernel(float* d_M, float* d_N, float* d_P, int Width)
{
    // Calculate the row index of the d_P element and d_M
    int Row = blockIdx.y*blockDim.y+threadIdx.y;
    // Calculate the column index of d_P and d_N
    int Col = blockIdx.x*blockDim.x+threadIdx.x;

    if ((Row < Width) && (Col < Width)) {
        float Pvalue = 0;
        // each thread computes one element of the block sub-matrix
        for (int k = 0; k < Width; ++k)
            Pvalue += d_M[Row*Width+k] * d_N[k*Width+Col];
        d_P[Row*Width+Col] = Pvalue;
    }
}
```

accesses to global memory

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

3

3

Review: 4B of Data per FLOP

- Each threads access global memory
 - for elements of **M** and **N**:
–**4B each**, or **8B per pair**.
 - (And once TOTAL to **P** per thread—ignore it.)
- With each pair of elements,
 - a thread does a single multiply-add,
–**2 FLOP**—floating-point operations.
- So for every FLOP,
 - a thread needs 4B from memory:
–**4B / FLOP**.

© Steven S. Lumetta, 2020
ECE408/CS483/University of Illinois at Urbana-Champaign

4

4

Review: Extremely Poor Performance

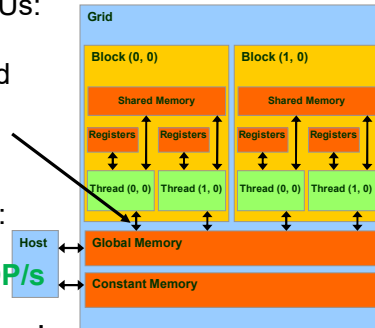
- One generation of GPUs:

- **1,000 GFLOP/s** of compute power, and
- **150 GB/s** of memory bandwidth.

- Dividing bandwidth by memory requirements:

$$\frac{150 \text{ GB/s}}{4 \text{ B/FLOP}} = 37.5 \text{ GFLOP/s}$$

which **limits computation!**



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/ University of Illinois at Urbana-Champaign

5

The Solution? Reuse Memory Accesses!

But **37.5 GFLOPs** is a limit.

In an **actual execution**,

- memory is not busy all the time, and
- the code **runs at** about **25 GFLOPs**.

To get closer to 1,000 GFLOPs

- we **need to** drastically **cut down**
- accesses to global memory.**

But ... how?

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

6

Tiled Matrix-Matrix Multiplication using Shared Memory

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

7

A Common Programming Strategy

- The dilemma:
 - Matrices **M** and **N** are large.
 - They **fit** easily **in global memory**, but that's **slow**.
 - Shared memory** is **fast**, but **M and N don't fit**.
- The solution:
 - Break M and N into tiles**
 - (called blocks in the much older CPU literature).
 - Read a tile** into shared memory.
 - Use the tile** from shared memory.
 - Repeat** until done.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

8

8

A Common Programming Strategy

- In a GPU, **only threads in a block** **can** use **shared** memory.
- Thus, each **block** operates on **separate tiles**:
 - **Read tile(s)** into shared memory **using multiple threads** to exploit memory-level parallelism.
 - **Compute** based on shared memory tiles.
 - **Repeat.**
 - **Write results** back **to global memory.**

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

9

9

Declaring Shared Memory Arrays


```
__global__ void MatrixMulKernel(float* M, float* N, float* P, int Width)
{
    __shared__ float subTileM[TILE_WIDTH][TILE_WIDTH];
    __shared__ float subTileN[TILE_WIDTH][TILE_WIDTH];
}
```

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/University of Illinois at Urbana-Champaign

10

10

Shared Memory Tiling Basic Idea



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

11

11

Outline of Technique

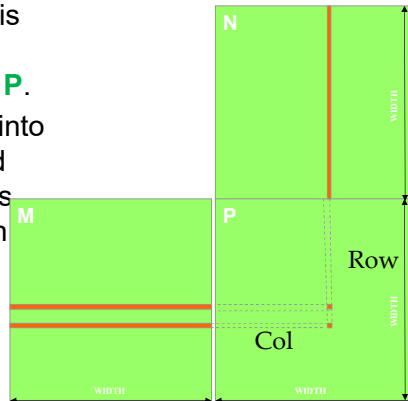
- Identify a tile of global data that are accessed by multiple threads
- Load the tile from global memory into on-chip memory
- Have the multiple threads to access their data from the on-chip memory
- Move on to the next block/tile

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

12

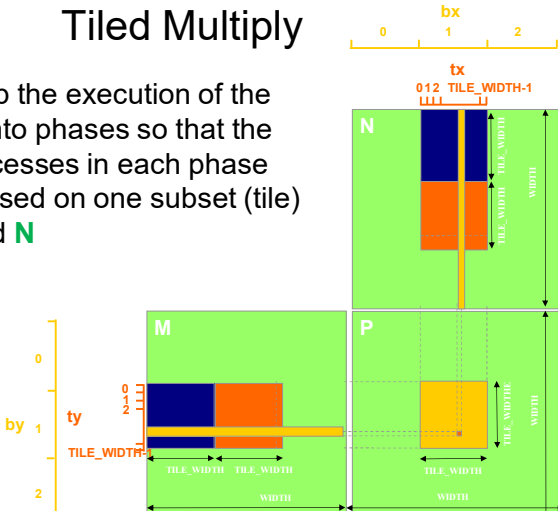
12

- Each input element is used to calculate **WIDTH** elements of **P**.
- Load each element into Shared Memory and have several threads use the local version to reduce memory bandwidth.



13

- Break up the execution of the kernel into phases so that the data accesses in each phase are focused on one subset (tile) of **M** and **N**

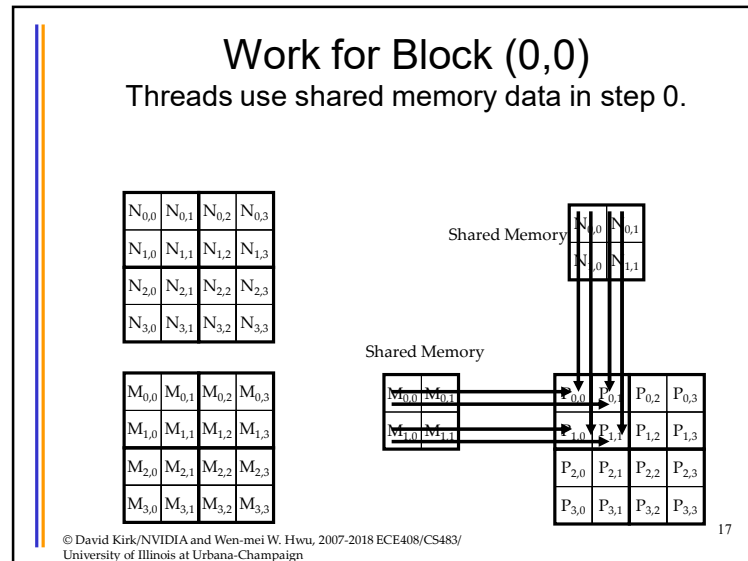


14

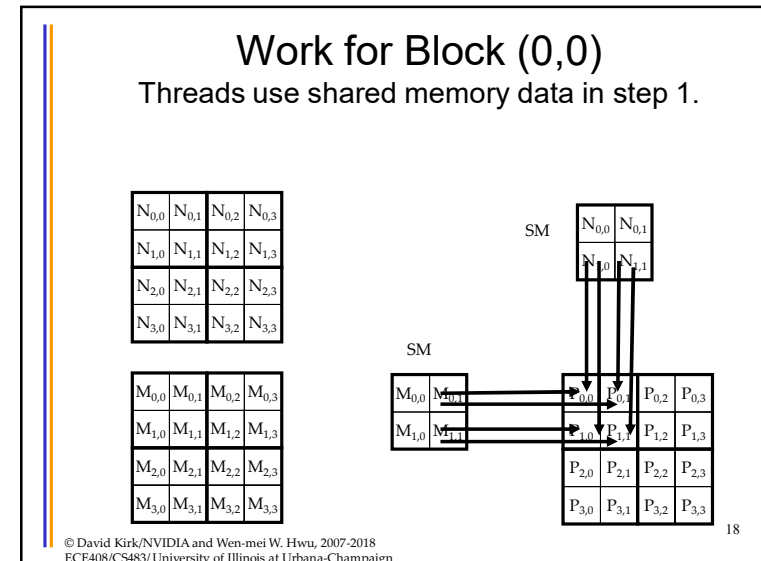
- All threads in a block participate
 - Each thread loads
 - one **M** element and
 - one **N** element
 - in basic tiling code.
- Assign the loaded element to each thread such that the accesses within each warp is coalesced (more later).

15

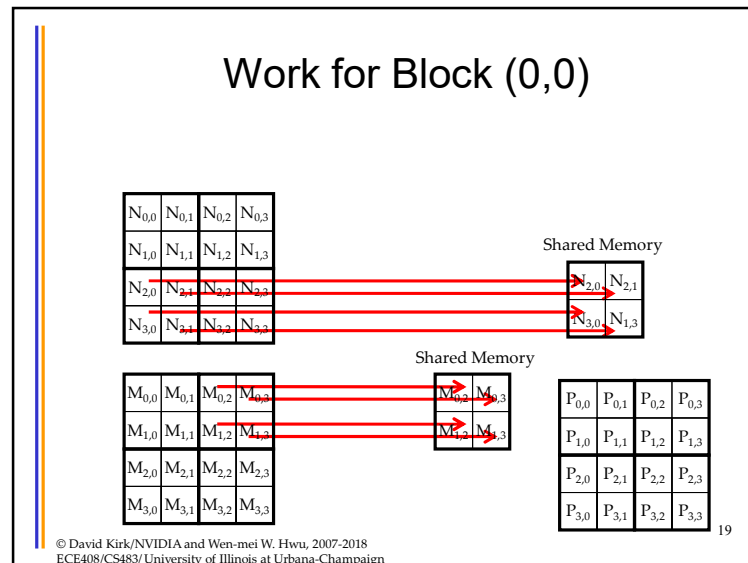
16



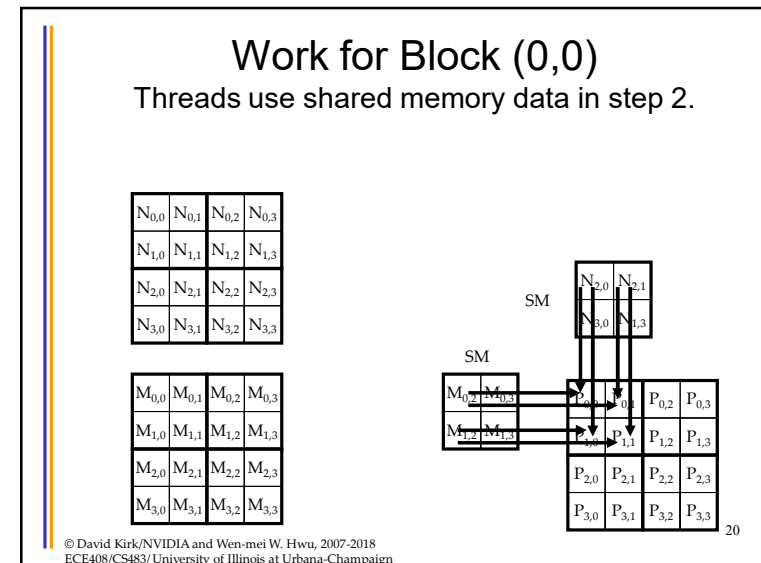
17



18



19

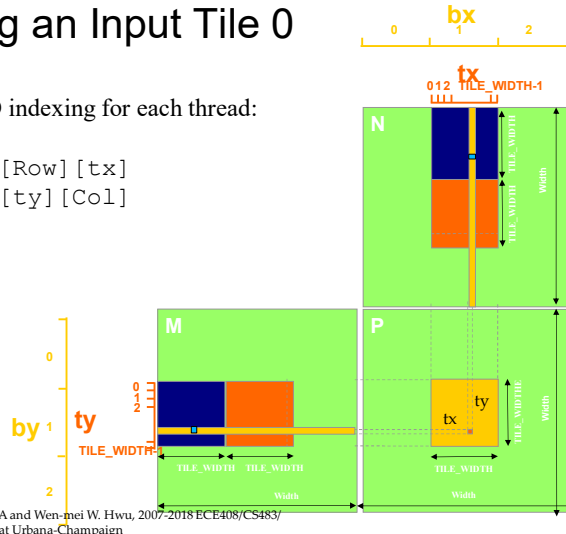


20

Loading an Input Tile 0

Tile 0 2D indexing for each thread:

$M[\text{Row}][\text{tx}]$
 $N[\text{ty}][\text{Col}]$



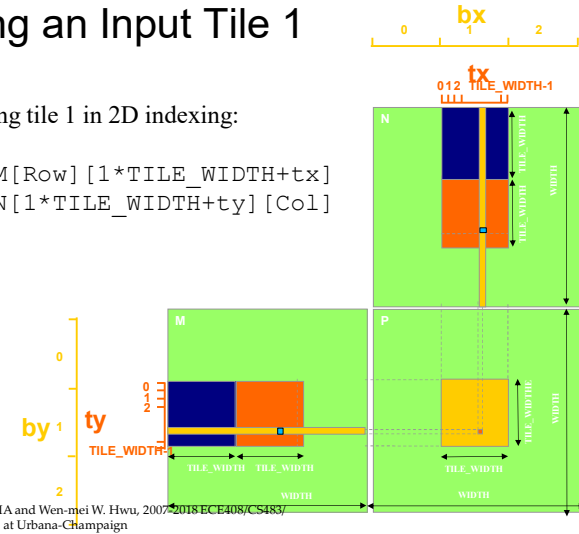
© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
 University of Illinois at Urbana-Champaign

21

Loading an Input Tile 1

Accessing tile 1 in 2D indexing:

$M[\text{Row}][1 * \text{TILE_WIDTH} + \text{tx}]$
 $N[1 * \text{TILE_WIDTH} + \text{ty}][\text{Col}]$



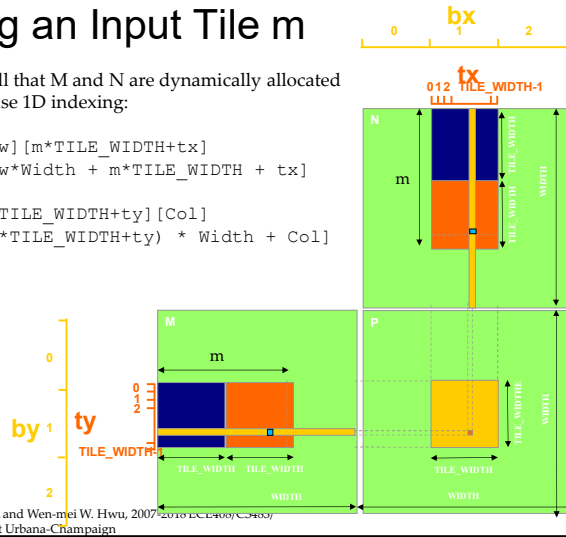
© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
 University of Illinois at Urbana-Champaign

22

Loading an Input Tile m

However, recall that M and N are dynamically allocated
 and can only use 1D indexing:

$M[\text{Row}][m * \text{TILE_WIDTH} + \text{tx}]$
 $M[\text{Row} * \text{Width} + m * \text{TILE_WIDTH} + \text{tx}]$
 $N[m * \text{TILE_WIDTH} + \text{ty}][\text{Col}]$
 $N[(m * \text{TILE_WIDTH} + \text{ty}) * \text{Width} + \text{Col}]$



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
 University of Illinois at Urbana-Champaign


23

Accessing a Tile

To perform the k^{th} step of the product within the tile:

$\text{subTileM}[\text{ty}][k]$

$\text{subTileN}[k][\text{tx}]$



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
 University of Illinois at Urbana-Champaign

24

We're Not There Yet!

- But ...
- **How can a thread know ...**
 - **That another thread has finished** its part of the tile?
 - Or that another thread has finished using the previous tile?

We need to synchronize!

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

25

25

Leveraging Parallel Strategies

- **Bulk synchronous execution:**
threads execute roughly in unison
 1. Do some work
 2. Wait for others to catch up
 3. Repeat
- **Much easier programming model**
 - Threads only parallel within a section
 - Debug lots of little programs
 - Instead of one large one.
- **Dominates high-performance applications**

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

26

26

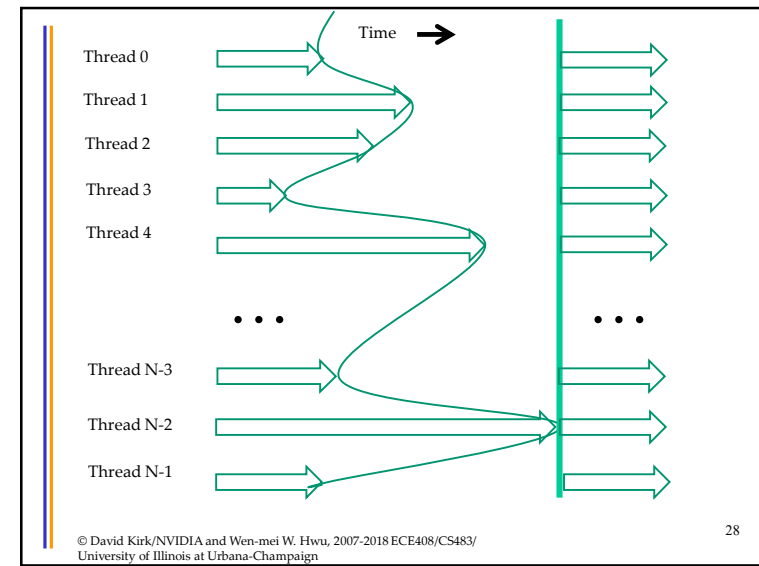
Bulk Synchronous Steps Based on Barriers

- **How does it work?**
Use a barrier to wait for thread to 'catch up.'
- A barrier is a synchronization point:
 - **each thread calls a function** to enter barrier;
 - **threads block** (sleep) in barrier function **until all threads have called**;
 - **after last thread calls** function, **all threads continue** past the barrier.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

27

27



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

28

28

Use __syncthreads for CUDA Blocks

- **How does it work in CUDA?**
Only **within thread blocks**!
- The function: **void __syncthreads(void) ;**
- N.B.
 - **All threads** in block **must enter** (no subsets).
 - All threads must enter the **SAME static call** (not the same as all threads calling function!).

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

29

29

Barrier Trauma: What's Actually Done?

- **What exactly is guaranteed** to have finished?
 - Are **shared memory** operations before a barrier (e.g., stores) guaranteed to have completed?
 - What about **global memory** ops?
 - What about **atomic ops** with no return values?
 - What about I/O operations?
- CUDA manual: all global and shared memory ops (which presumably includes atomic variants) have completed.
- **Avoid assumptions about I/O** (such as printf).

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

30

30

Tiled Matrix Multiplication Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, int Width)
{
    1. __shared__ float subTileM[TILE_WIDTH][TILE_WIDTH];
    2. __shared__ float subTileN[TILE_WIDTH][TILE_WIDTH];

    3. int bx = blockIdx.x; int by = blockIdx.y;
    4. int tx = threadIdx.x; int ty = threadIdx.y;

    // Identify the row and column of the P element to work on
    5. int Row = by * TILE_WIDTH + ty; // note: blockDim.x == TILE_WIDTH
    6. int Col = bx * TILE_WIDTH + tx; //      blockDim.y == TILE_WIDTH
    7. float Pvalue = 0;

    // Loop over the M and N tiles required to compute the P element
    // The code assumes that the Width is a multiple of TILE_WIDTH!
    8. for (int m = 0; m < Width/TILE_WIDTH; ++m) {
        // Collaborative loading of M and N tiles into shared memory
        9. subTileM[ty][tx] = M[Row*Width + m*TILE_WIDTH+tx];
        10. subTileN[ty][tx] = N[(m*TILE_WIDTH+ty)*Width+Col];
        11. __syncthreads();
        12. for (int k = 0; k < TILE_WIDTH; ++k)
        13.     Pvalue += subTileM[ty][k] * subTileN[k][tx];
        14. __syncthreads();
    15. }
    16. P[Row*Width+Col] = Pvalue;
}
```

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

31

31

Compare with Basic MM Kernel

```
__global__ void MatrixMulKernel(float* M, float* N, float* P, int Width)
{
    // Calculate the row index of the P element and M
    int Row = blockIdx.y * blockDim.y + threadIdx.y;
    // Calculate the column index of P and N
    int Col = blockIdx.x * blockDim.x + threadIdx.x;

    if ((Row < Width) && (Col < Width)) {
        float Pvalue = 0;

        // each thread computes one element of the block sub-matrix
        for (int k = 0; k < Width; ++k)
            Pvalue += M[Row*Width+k] * N[k*Width+Col];

        P[Row*Width+Col] = Pvalue;
    }
}
```

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

32

32

Use of Large Tiles Shifts Bottleneck

- Recall our example GPU: **1,000 GFLOP/s, 150 GB/s**
- 16x16 tiles** use each operand for 16 operations
 - **reduce global** memory **accesses by** a factor of **16**
 - **150GB/s** bandwidth supports $(150/4)*16 = \mathbf{600\ GFLOPS!}$
- 32x32 tiles** use each operand for 32 operations
 - **reduce global** memory **accesses by** a factor of **32**
 - **150 GB/s** bandwidth supports $(150/4)*32 = \mathbf{1,200\ GFLOPS!}$
 - **Memory bandwidth is no longer the bottleneck!**

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

33

33

Also Need Parallel Accesses to Memory

- Shared memory size
 - implementation dependent
 - **64kB** per SM in Maxwell (48kB max per block)
- Given **TILE_WIDTH of 16** (256 threads / block),
 - each thread block uses $2*256*4B = 2kB$ of shared memory,
 - which limits active blocks to 32;
 - max. of 2048 threads per SM,
 - which limits blocks to 8.
 - Thus up to $8*512 = \mathbf{4,096\ pending\ loads}$ (2 per thread, 256 threads per block)

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

34

34

Another Good Choice: 32x32 Tiles

- Given **TILE_WIDTH of 32** (1,024 threads / block),
 - each thread block uses $2*1024*4B = 8kB$ of shared memory,
 - which limits active blocks to 8;
 - max. of 2,048 threads per SM,
 - which limits blocks to 2.
 - Thus up to $2*2,048 = \mathbf{4,096\ pending\ loads}$ (2 per thread, 1,024 threads per block)

(same memory parallelism exposed)

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

35

35

Current GPU? Use Device Query

- Number of devices in the system

```
int dev_count;
cudaGetDeviceCount( &dev_count);
```
- Capability of devices

```
cudaDeviceProp dev_prop;
for (i = 0; i < dev_count; i++) {
    cudaGetDeviceProperties( &dev_prop, i);

    // decide if device has sufficient resources and capabilities
}
```
- `cudaDeviceProp` is a built-in C structure type
 - `dev_prop.dev_prop.maxThreadsPerBlock`
 - `Dev_prop.sharedMemoryPerBlock`
 - ...

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

36

36

A decorative element consisting of two vertical lines, one blue and one orange, positioned to the left of the main text.

**ANY MORE QUESTIONS?
READ CHAPTER 4!**

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018 ECE408/CS483/
University of Illinois at Urbana-Champaign

37

37