

ECE408/CS483/CSE408 Fall 2020

Applied Parallel Programming

Lecture 9: Tiled Convolution

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498al University of Illinois, 2007-2018

1

1

Objective

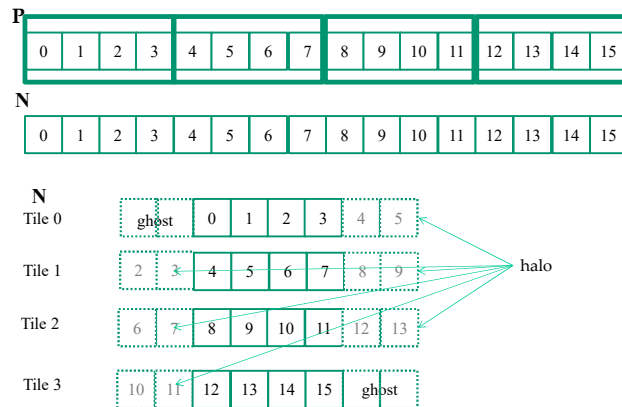
- To learn about tiled convolution algorithms
 - Some intricate aspects of tiling algorithms
 - Output tiles versus input tiles
 - Three different styles of input tile loading
 - To prepare for MP-4

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498al University of Illinois, 2007-2018

2

2

Tiled 1D Convolution Basic Idea



© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498al University of Illinois, 2007-2018

3

3

What Shall We Parallelize?

In other words,

What should one thread do?

One answer:

- (same as with vector sum and matrix multiply)
- **compute an output element!**

© Steven S. Lumetta
ECE408/CS483/ECE498al University of Illinois, 2020

4

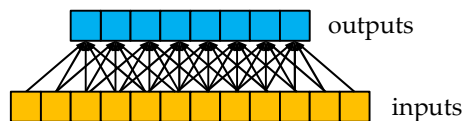
4

Should We Use Shared Memory?

In other words,

Can we reuse data read from global memory?

Let's look at the computation again...



Reuse reduces global memory bandwidth,
so **let's use shared memory**.

© Steven S. Lumetta
ECE408/CS483/ECE498al University of Illinois, 2020

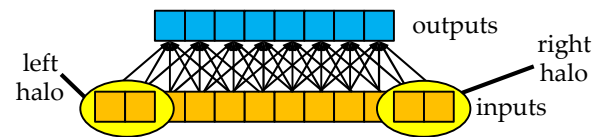
5

5

What About the Halos?

In other words,

Do we also copy halos into shared memory?



Let's **consider both** possible answers.

© Steven S. Lumetta
ECE408/CS483/ECE498al University of Illinois, 2020

6

6

How Much Reuse is Possible?



- Element 2 is used by thread 4 (1×)
- Element 3 is used by threads 4, 5 (2×)
- Element 4 is used by threads 4, 5, 6 (3×)
- Element 5 is used by threads 4, 5, 6, 7 (4×)
- Element 6 is used by threads 4, 5, 6, 7 (4×)
- Element 7 is used by threads 5, 6, 7 (3×)
- Element 8 is used by threads 6, 7 (2×)
- Element 9 is used by thread 7 (1×)

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498al University of Illinois, 2007-2018

7

7

Can Access Halo from Global Mem.

One answer: no,

- threads **read halo values**
- directly **from global memory**.

bad idea

Advantage:

- optimize reuse of shared memory
- (halo reuse is smaller).

Disadvantages:

- **Branch divergence!** (shared vs. global reads)
- Halo **too narrow to fill** a memory **burst**

© Steven S. Lumetta
ECE408/CS483/ECE498al University of Illinois, 2020

8

8

Can Load Halo to Shared Mem.

Better answer: yes,

load halos to shared memory.

Let's try it!

Advantages:

- **Coalesce global memory accesses.**
- **No branch divergence during computation.**

Disadvantages:

- Some threads must do >1 load, so **some branch divergence** in reading data.
- Slightly more shared memory needed.

© Steven S. Lumetta
ECE408/CS483/ECE498a University of Illinois, 2020

9

9

Allocate and Initialize Variables

```
__global__ void convolution_1D_tiled_kernel
(float *N, float *P, int Width)
{
    // shared tile with space for both halos
    __shared__ float tile[TILE_SIZE + MASK_WIDTH - 1];

    int radius = MASK_WIDTH / 2; // a useful constant

    // this thread's index into output P
    int i = blockIdx.x * blockDim.x + threadIdx.x;

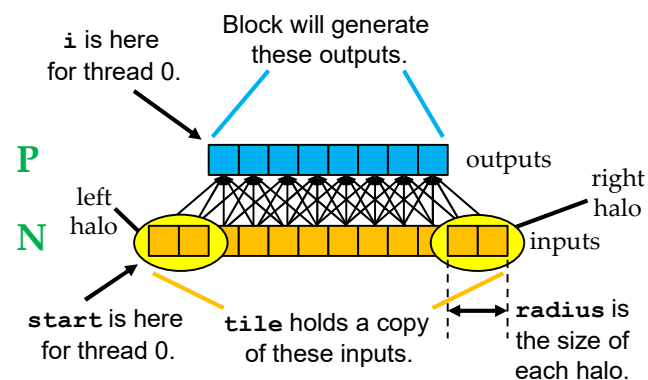
    // this thread's starting element of N
    int start = i - radius;
```

© Steven S. Lumetta
ECE408/CS483/ECE498a University of Illinois, 2020

10

10

Variable Meanings for a Block



© Steven S. Lumetta
ECE408/CS483/ECE498a University of Illinois, 2020

11

11

Load the Input Data

```
if (0 <= start && Width > start) { // all threads
    tile[threadIdx.x] = N[start];
} else {
    tile[threadIdx.x] = 0.0f;
}
if (MASK_WIDTH - 1 > threadIdx.x) { // some threads
    start += TILE_SIZE;
    if (Width > start) {
        tile[threadIdx.x + TILE_SIZE] = N[start];
    } else {
        tile[threadIdx.x + TILE_SIZE] = 0.0f;
    }
}
__syncthreads(); // OUTSIDE of if's
```

© Steven S. Lumetta
ECE408/CS483/ECE498a University of Illinois, 2020

12

12

And Compute an Output Element

```
if (i < Width) { // only threads computing outputs

    float Pvalue = 0;    // running sum

    // compute output element
    for (int j = 0; MASK_WIDTH > j; j++) {
        Pvalue += tile[threadIdx.x + j] * Mc[j];
    }

    // write to P
    P[i] = Pvalue;
}
```

© Steven S. Lumetta
ECE408/CS483/ECE498al University of Illinois, 2020

13

13

Review: What Shall We Parallelize?

In other words,

What should one thread do?

One answer:

- (same as with vector sum and matrix multiply)
- **compute an output element!**

Is that our only choice?

© Steven S. Lumetta
ECE408/CS483/ECE498al University of Illinois, 2020

14

14

Parallelize Loading of a Tile

Alternately,

- **each thread loads** one input element, and
- **some threads compute** an output.

(compared with previous approach)

Advantage:

- **No branch divergence for load** (high latency).
- **Avoid narrow global access** ($2 \times$ halo width).

Disadvantage:

- Branch **divergence for compute** (low latency).

© Steven S. Lumetta
ECE408/CS483/ECE498al University of Illinois, 2020

15

15

2D Example of Loading Parallelization

Let's do an example for 2D convolution.

- Thread block matches input tile size.
- Each thread loads one element of input tile.
- Some threads do not participate in calculating output,

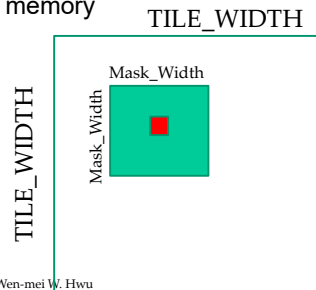
© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498al University of Illinois, 2007-2018

16

16

Parallelizing Tile Loading

- Load a tile of N into shared memory
 - All threads participate in loading
 - A subset of threads then use each N element in shared memory

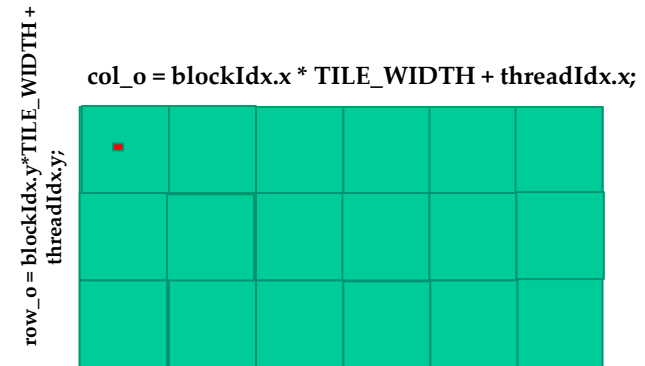


© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498a University of Illinois, 2007-2018

17

17

Output Tiles Still Cover the Output!

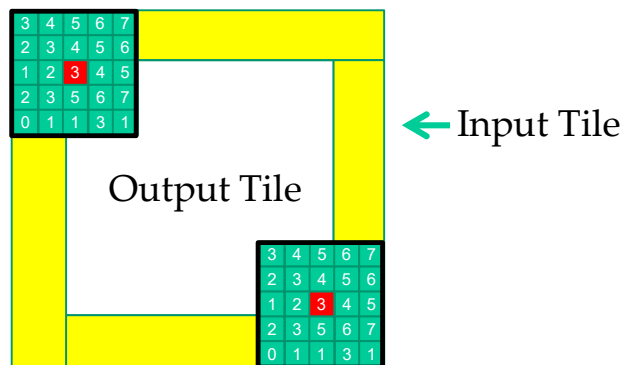


© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498a University of Illinois, 2007-2018

18

18

Input tiles need to be larger than output tiles.



© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498a University of Illinois, 2007-2018

19

19

Setting Block Width

```
dim3 dimBlock(TILE_WIDTH+4, TILE_WIDTH+4, 1);
```

In general, block width should be
 $\text{TILE_WIDTH} + (\text{MASK_WIDTH} - 1)$

```
Dim3 dimGrid(ceil(Width/(1.0*TILE_WIDTH)),  
             ceil(Width/(1.0*TILE_WIDTH)), 1)
```

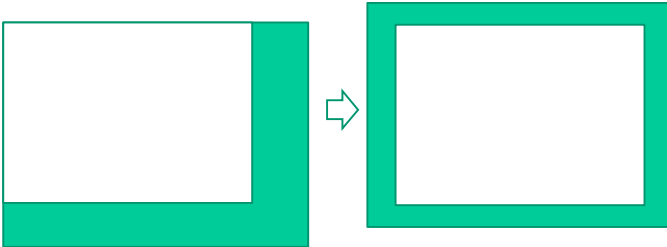
There need to be enough thread blocks
to generate all P elements.

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498a University of Illinois, 2007-2018

20

20

Shifting from output coordinates to input coordinates



© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498al University of Illinois, 2007-2018

21

21

Shifting from output coordinates to input coordinates

```
int tx = threadIdx.x;
int ty = threadIdx.y;
int row_o =
    blockIdx.y * TILE_WIDTH + ty;
int col_o =
    blockIdx.x * TILE_WIDTH + tx;

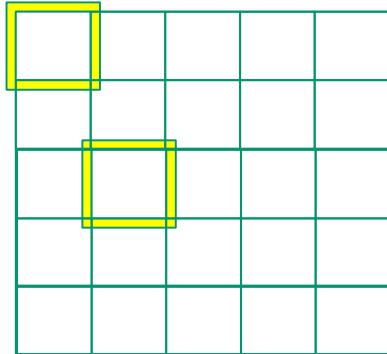
int row_i = row_o-2; // MASK_WIDTH / 2
int col_i = col_o-2; // (radius in
                    // prev. code)
```

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498al University of Illinois, 2007-2018

22

22

Threads that loads halos outside N should return 0.0



© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498al University of Illinois, 2007-2018

23

23

Taking Care of Boundaries

```
float Pvalue = 0.0f;
if((row_i >= 0) && (row_i < Width) &&
    (col_i >= 0) && (col_i < Width)) {
    tile[ty][tx] =
        N[row_i*Width + col_i];
} else {
    tile[ty][tx] = 0.0f;
}
__sync_threads (); // wait for tile
```

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498al University of Illinois, 2007-2018

24

24

Not All Threads Calculate Output

```
if (ty < TILE_WIDTH && tx < TILE_WIDTH) {  
    for (i = 0; i < 5; i++) {  
        for (j = 0; j < 5; j++) {  
            Pvalue += Mc[i][j] *  
                    tile[i+ty][j+tx];  
        }  
    }  
    // if continues on next page
```

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498a University of Illinois, 2007-2018

25

25

Not All Threads Write Output

```
if (row_o < Width && col_o < Width)  
    P[row_o * Width + col_o] = Pvalue;  
}  
} // end of if selecting output  
// tile threads
```

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498a University of Illinois, 2007-2018

26

26

**ANY MORE QUESTIONS?
READ CHAPTER 7**

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498a University of Illinois, 2007-2018

27

27