

ECE408 Lecture 11

Machine Learning and Deep Learning

ECE408 / CS483 / CSE 408
Spring 2020
(by Carl Pearson)

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

1

Objective

- To understand the application areas for machine learning.
- To learn the basic strategy for machine learning applications.
- To understand the extension to deep learning (mostly a research pitch).

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

2

2

Perspective is Important

Chips are cheaper than ever.

Unlike humans, digital systems offer

- **high-speed computation**,
- **low capital investment**
(purchase vs. training a human), and
- **negligible operations cost** (no salary!).

If computer outperforms (or even matches)
a **human, use a computer.**

Industry has done so for 40-50 years now.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

3

What is Machine Learning?

machine learning: important method of building applications whose logic is not fully understood

Typically **by example**:

- **use labeled data** (matched input-output pairs)
- **to represent** desired **relationship**.

Iteratively adjust program logic to produce desired/approximate answers (called **training**).

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

4

Types of Learning Tasks

- classification
 - Map each input to a category
 - Ex: object recognition, chip defect detection
- regression
 - Numerical prediction from a sequence
 - Ex: predict tomorrow's temperature
- transcription
 - Unstructured data into textual form
 - Ex: optical character recognition

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

5

More Advanced Learning Tasks

- translation
 - Convert a sequence of symbols in one language to a sequence of symbols in another
- structured output
 - Convert an input to a vector with important relationships between elements
 - Ex: natural language sentence into grammatical structure
- others
 - Anomaly detection, synthesis, sampling, imputation, denoising, density estimation, genetic variant calling

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

6

Test Cycle Time is Important

You've all written code...

- code, test, code, test, code, test
- integrate, test, test, test
- and test again!

But how long is the code, test cycle?

Depends what you're building.

What's your longest?

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

7

Your Cycle Times are Probably Small

In college, **10k lines** took **½ hour** to compile on my PC.

In grad. school, **100k lines** took

- **½ hour** to compile on my workstation, or
- **2 minutes** on our cluster (research platform).

In ECE435 (networking lab), students needed

- **½ hour** to reinstall Linux after a bad bug.
- (Ever had a good bug?)

Gene sequencing / applications can take **two weeks**.

We're all a little spoiled...

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

8

Why Machine Learning Again?

In 2007, **programmable GPUs accelerated the training cycle.**

Today, **new chip designs** for learning applications **have further accelerated.**

Led to a resurgence of interest

- in Computer Vision, Speech Recognition, Document Translation, Self Driving Cars, Data Science...
- all tasks that **human brains solve regularly, but** for which **we** have **struggled to express solutions** systematically.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

9

Many Problems are Still Hard

Speed is not a panacea.

- Many **tasks still require human insight**
 - for network structure and feature selection
 - for effective input and output formats, and
 - for production of high-quality labeled data.
- Other trends sometimes help: ubiquitous computing enables crowdsourcing, for example.


© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

10

Many Problems Have Systematic Solutions

Example: building a Boolean function from a truth table.

| Input | | | output |
|-------|---|---|--------|
| a | b | c | |
| 0 | 0 | 0 | 0 |
| 0 | 0 | 1 | 1 |
| 0 | 1 | 0 | 1 |
| 0 | 1 | 1 | 0 |
| 1 | 0 | 0 | 1 |
| 1 | 0 | 1 | 0 |
| 1 | 1 | 0 | 0 |
| 1 | 1 | 1 | 1 |



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

11

What if We Lack a Truth Table?

- Make enough observations to construct a rule
 - $000 \rightarrow 0$
 - $011 \rightarrow 0$
 - $100 \rightarrow 1$
 - $110 \rightarrow 0$
- If we cover all input patterns, we can construct a truth table!

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

12

Many Problems are Too Large

- The logic formulation of a 32x32-pixel (small) image recognition problem involves
 - 1024*8 bit input,
 - which will have a truth table of 2^{8196} entries
- If we managed to collect and label 1 billion ($\sim 2^{32}$) images as training data
 - We cover only $2^{32} / 2^{8196} = 1 / 2^{8164}$ of the truth table
 - Solution - learning processes that exploits features

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

13

Features in our logic example


| Input | | | output |
|-------|---|---|--------|
| a | b | c | |
| 0 | 0 | 0 | 0 |
| 0 | 0 | 1 | 1 |
| 0 | 1 | 0 | 1 |
| 0 | 1 | 1 | 0 |
| 1 | 0 | 0 | 1 |
| 1 | 0 | 1 | 0 |
| 1 | 1 | 0 | 0 |
| 1 | 1 | 1 | 1 |

Feature 1: bit patterns with odd number of 1's result in output 1
Feature 2: bit patterns with even number of 1's result in output 0

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

14

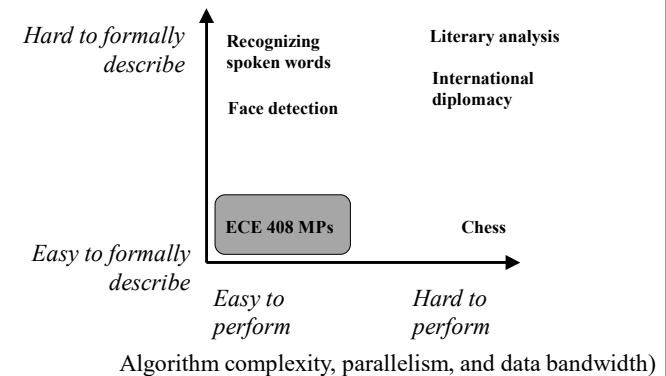
Types of Problems



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

15

Types of Problems




© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

16

Chess as an AI Success (1)

- Easy to formalize
 - 64 locations, 32 pieces
 - Well-defined, allowable moves
- Score each leaf in a tree of possible board positions
- Proceed down path that results in best position



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

17

Chess as an AI Success (2)



Deep Blue defeated Gary Kasparov in 1997


- Hard to perform
 - ~30 legal moves per position
 - 1,015 moves for 10-ply lookahead
 - 30 years of compute at 1M positions/sec
- Heuristics, pruning, parallel search, fast computers

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

18

Cyc: Extending Rule-based Systems to the Real World

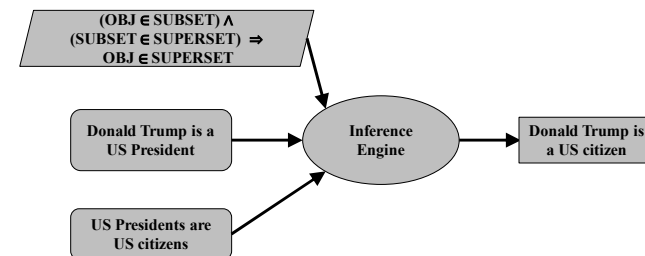
- Comprehensive ontology and knowledge base of common sense
- Cyc reasons about formal statements about the world



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

19


Cyc: A Simple Example



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

20

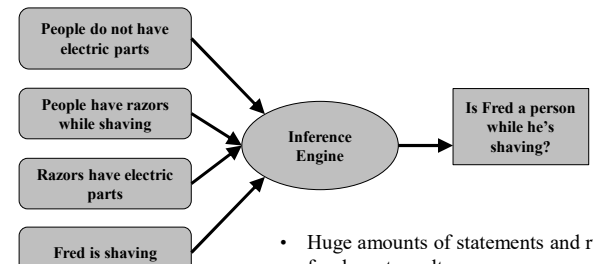
Cyc: FredWhileShaving



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

21

Cyc: FredWhileShaving

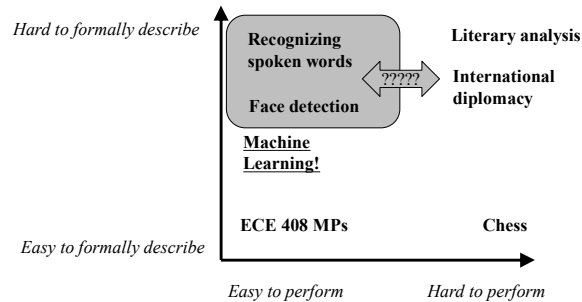


- Huge amounts of statements and rules for decent results
- Cannot learn new rules or statements on its own

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

22

Types of Problems



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

23

The “Machine Learning” Approach

Challenge

Hard to formalize the problem.

Solution

Don't formalize the problem.

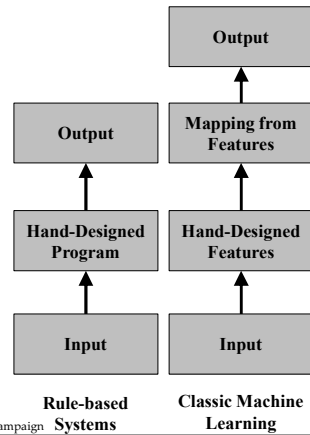
Let the machine learn from experience.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

24

Classic Machine Learning

- Humans choose features
- Learn how features are associated with outputs



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

25

You may have heard of...


- Naïve Bayes:
features as independent contributors to output
- Logistic Regression:
– learn how to weight each feature's contribution to output,
– usually through gradient descent*

*more on this topic later in these slides

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign


26

Data Representation is important!



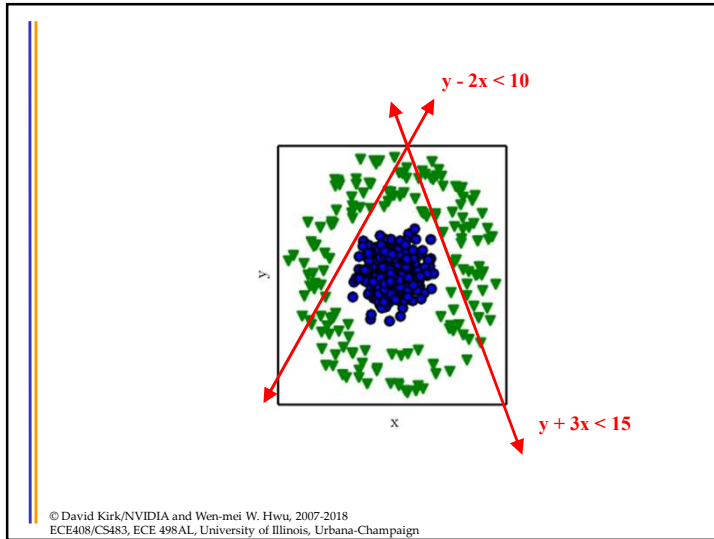
© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

27

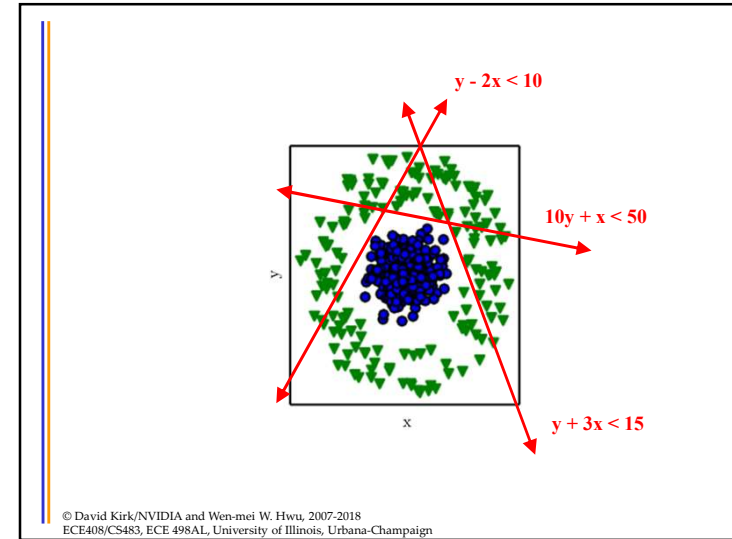


© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

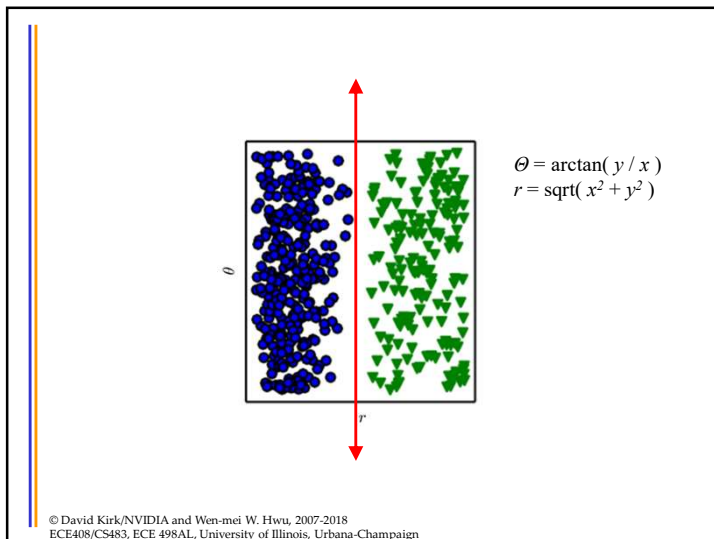
28



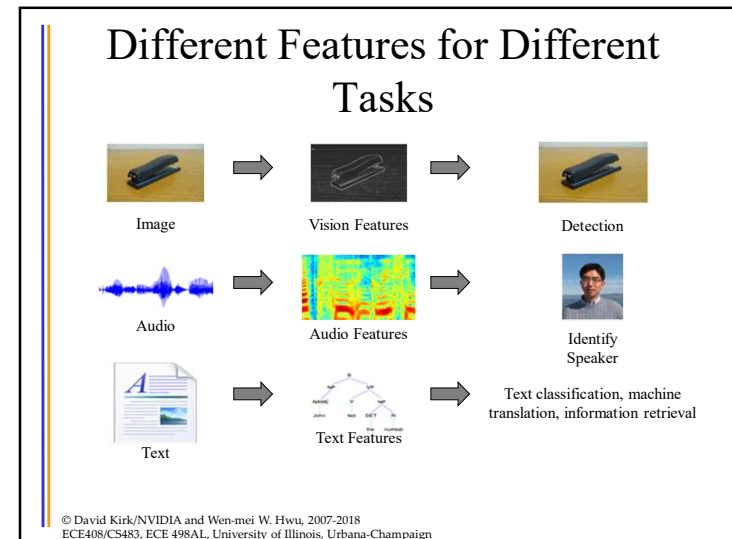
29



30



31



32

Which Data Features are Relevant

- Detecting a car in an image
- Cars have wheels ➡ presence of a wheel?
- Can we describe pixel values that make up a wheel?
 - Circle-shaped?
 - Dark around perimeter?

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

33

Which Data Features are Relevant

- Detecting a car in an image
- Cars have wheels ➡ presence of a wheel?
- Can we describe pixel values that make up a wheel?
 - Circle-shaped?
 - Dark around perimeter?
- But what about?
 - Occlusion, perspective, shadows, white-walled tires, ...

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

34

Identify Factors of Variation that Explain Data

- Unobserved objects or forces that affect observed quantities
- Mental constructs that provide simplifying explanations or inferred causes
- Ex: speech
 - Age, sex, accent, words being spoken
- Ex: car
 - Position, color, angle of sun
- Many factors influence each piece of observed data

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

35

Representation Learning Approach

Challenge

Which data features are relevant?

Solution

Learn the features too!

(Looking ahead)

Deep Learning: a deep hierarchy of features

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

36

Machine Learning


- Ability to acquire knowledge by extracting patterns from data

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

37

Deep Learning

- A type of representation learning
- Representations expressed in terms of other representations



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

38

Deep Learning Approach

Challenge

Hard to formalize the problem?

Which data features are relevant?

Solution

Don't formalize the problem

Let the machine learn from experience

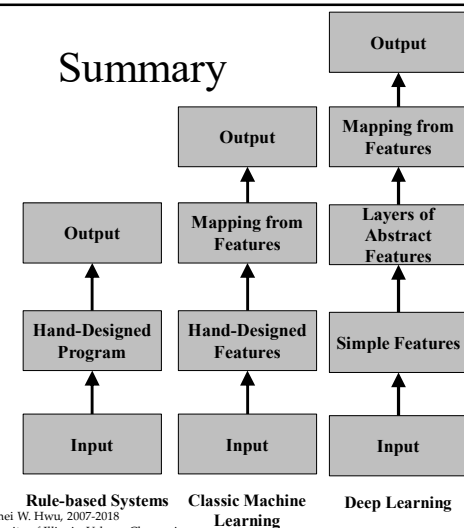
Hierarchy of concepts to capture simple and complicated features

Learn the hierarchy too!

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

39

Summary



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

40