

# ECE408 Lecture 12

## Feed-Forward Networks and Gradient-Based Training

ECE408 / CS483 / CSE 408  
Spring 2020  
(by Carl Pearson)

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

1

## Objective

- To learn the basic approach to feedforward neural networks:
  - neural model
  - common functions
  - training through gradient descent

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

2

## Let's Look at Classification

In a **classification problem**, we model

- a function mapping an input vector to a set of  $C$  categories:  $F: \mathbb{R}^N \rightarrow \{1, \dots, C\}$ ,
- where the function  $F$  is **unknown**.

We **approximate**  $F$  using a set of functions  $f$

- parametrized by a (large) set of weights,  $\theta$
- that map from a vector of  $N$  real values\* to an integer value representing a category:
- for category  $i$ ,  $\text{prob}(i) = f(x, \theta)$

\*floating-point values

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

3



## Perceptron is a Simple Example

- Example: a **perceptron**

$$y = \text{sign}(W \cdot x + b) \quad \theta = \{W, b\}$$

The perceptron

The neuron



- Dot product:  $y = W \cdot x$  (output)
- Scalar addition:  $+ b$  (input, bias, weight)

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign


4

## One Perceptron is not Enough

Some functions are non-linear.

What can we do?

● FALSE  
● TRUE



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign


5

## Multiple Layers Solve More Problems

What if input dimensions are AND and OR?

Now we can divide with one line.

● FALSE  
● TRUE



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign


6

| A | B | OR | AND | XOR |
|---|---|----|-----|-----|
| 0 | 0 | -1 | -1  | -1  |
| 0 | 1 | 1  | -1  | 1   |
| 1 | 0 | 1  | -1  | 1   |
| 1 | 1 | 1  | 1   | -1  |

$$\text{AND} = \text{sign}(x[0] + x[1] - 1.5)$$

$$\text{OR} = \text{sign}(x[0] + x[1] - 0.5)$$


$$\text{XOR} = \text{sign}(2 * \text{OR} - \text{AND} - 2)$$



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

7

## Generalize to Fully-Connected Layer




Linear Classifier:  
Input vector  $x$   $\times$  weight  
vector  $w$  to produce  
scalar output  $y$

Fully-connected:  
Input vector  $x$   $\times$  weight  
matrix  $w$  to produce  
vector output  $y$

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

8

## Multilayer Terminology



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
 ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

9

## Example: Digit Recognition

Let's consider an example.

- **handwritten digit recognition:**
  - given a  **$28 \times 28$  grayscale image**,
  - produce a **number from 0 to 9**.
- Input dataset
- **60,000** images
  - Each labeled by a human with correct answer.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
 ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

10

## How Do We Determine the Weights?

**First layer** of perceptrons

- **784** ( $28^2$ ) inputs, **1024** outputs, **fully connected**
- **$[1024 \times 784]$**  weight matrix  **$W$**
- **$[1024 \times 1]$**  bias vector  **$b$**

**Use labeled training data to pick weights.**

Idea:

- given enough labeled input data,
- we can **approximate the input-output function**.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
 ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

11

## Forward and Backward Propagation

Forward (**inference**):

- given input  **$x$**  (for example, an image),
- **use parameters  $\Theta$**  ( **$W$**  and  **$b$**  for each layer)
- **to compute probabilities  $k[i]$**  (ex: for each digit  $i$ ).

Backward (**training**):

- given input  **$x$** , parameters  **$\Theta$** , and outputs  **$k[i]$** ,
- **compute error  $E$**  based on target label  **$t$** ,
- then **adjust  $\Theta$**  proportional to  **$E$**  to reduce error.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
 ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

12

## Neural Functions Impact Training

Recall perceptron function:  $\mathbf{y} = \text{sign}(\mathbf{W} \cdot \mathbf{x} + \mathbf{b})$

To propagate error backwards,

- use chain rule from calculus.
- Smooth functions are useful.

Sign is not a smooth function.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

13

## One Choice: Sigmoid/Logistic Function

Until about 2017,

- **sigmoid / logistic function** most popular

$$f(x) = \frac{1}{1+e^{-x}} \quad (f: \mathbb{R} \rightarrow (0,1))$$

for replacing sign.

- Once we have  $f(x)$ , finding  $df/dx$  is easy:

$$\frac{df(x)}{dx} = \frac{e^{-x}}{(1+e^{-x})^2} = f(x) \frac{e^{-x}}{(1+e^{-x})} = f(x)(1-f(x))$$

(Our example used this function.)

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

14

## Today's Choice: ReLU

In 2017, most common choice became

- **rectified linear unit / ReLU / ramp function**  
 $f(x) = \max(0, x)$  ( $f: \mathbb{R} \rightarrow \mathbb{R}^+$ )  
which is much faster (no exponent required).
- A smooth approximation is **softplus/SmoothReLU**  
 $f(x) = \ln(1 + e^x)$  ( $f: \mathbb{R} \rightarrow \mathbb{R}^+$ )  
which is the integral of the logistic function.
- Lots of variations exist. See Wikipedia for an overview and discussion of tradeoffs.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

15

## Use Softmax to Produce Probabilities

**How can sigmoid / ReLU produce probabilities?**

They can't.

- Instead, given output vector  $\mathbf{Z} = (z[0], \dots, z[C-1])^*$ ,
- we produce a second vector  $\mathbf{K} = (k[0], \dots, k[C-1])$
- using the **softmax function**

$$k[i] = \frac{e^{z[i]}}{\sum_{j=0}^{C-1} e^{z[j]}}$$

Notice that **the  $k[i]$  sum to 1.**

\*Remember that we classify into one of  $C$  categories.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

16

## Softmax Derivatives Needed to Train

We also need the **derivatives of softmax**,


$$\frac{dk[i]}{dz[m]} = k[i](\delta_{i,m} - k[m]),$$

where  $\delta_{i,m}$  is the Kronecker delta (1 if  $i = m$ , and 0 otherwise).

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

17

## Forward and Backward Propagation



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

18

## Choosing an Error Function

Many error functions are possible.

For example, **given label  $T$**  (digit  $T$ ),

- $E = 1 - k[T]$ ,
- the **probability of not classifying as  $t$** .

**Alternatively**, since our categories are numeric, we can **penalize quadratically**:

$$E = \sum_{j=0}^{C-1} k[j](j - T)^2$$

Let's **go with the latter**.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

19

## Stochastic Gradient Descent

**How do we calculate the weights?**

One common answer: **stochastic gradient descent**.

1. **Calculate**
  - **derivative** of sum of error  $E$
  - **over all** training **inputs**
  - **for** all network parameters  $\theta$ .
2. **Change  $\theta$  slightly** in the opposite direction (to decrease error).
3. **Repeat**.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

20

## Stochastic Gradient Descent


More precisely,

1. **For every input  $X$ ,**
2. evaluate network to **compute  $k[i]$**  (forward),
3. then **use  $k[i]$  and label  $T$**  (target digit) **to compute error  $E$ .**
4. Backpropagate error derivative to **find derivatives for each parameter.**
5. **Adjust  $\theta$  to reduce total  $E$ :  $\theta_{i+1} = \theta_i - \epsilon \Delta \theta$**   
(Update  $\epsilon$  uses most accurate minima estimation.)

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

23

## Parameter Updates and Propagation



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

22

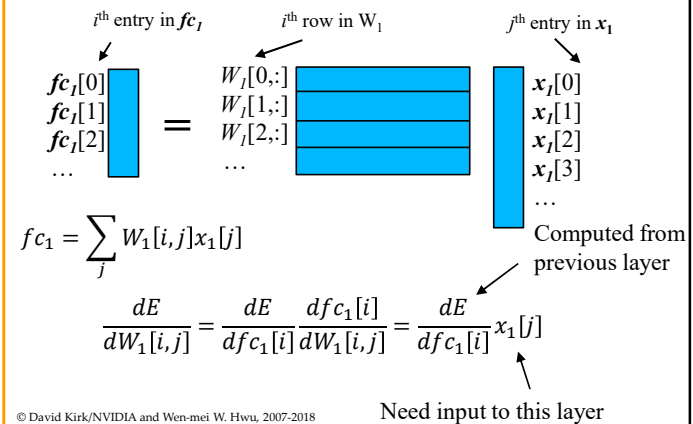
## Example: Gradient Update with One Layer

|  |                                     |                               |
|--|-------------------------------------|-------------------------------|
| $\theta_{i+1} = \theta_i - \epsilon \Delta \theta$       | $W_{i+1} = W_i - \epsilon \Delta W$ | Parameter Update              |
| $y = W \cdot x + b$                                      |                                     | Network function              |
| $\frac{dy}{dW} = x$                                      |                                     | Network weight gradient       |
| $E = \frac{1}{2} (y - t)^2$                              |                                     | Error function                |
| $\frac{dE}{dy} = y - t = Wx + b - t$                     |                                     | Error function gradient       |
| $\Delta W = \frac{dE}{dW} = \frac{dE}{dy} \frac{dy}{dW}$ |                                     | Full weight update expression |
| $W_{i+1} = W_i - \epsilon (Wx + b - t)x$                 |                                     | Full weight update term       |

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

23

## Fully-Connected Gradient Detail



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

24

## Batched Stochastic Gradient Descent

- A training *epoch* (a pass through whole training set)
  - Set  $\Delta\theta = 0$
  - For each labeled image:
    - Read data to initialize input layer
    - Evaluate network to get  $y$  (forward)
    - Compare with target label  $t$  to get error  $E$
    - Backpropagate error derivative to get parameter updates
    - Accumulate parameter updates into  $\Delta\theta$
  - $\theta_{i+1} = \theta_i - \epsilon \Delta\theta$

Aggregate gradient update most accurately reflects true gradient

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

25

## Mini-batch Stochastic Gradient


- For each batch in training set
  - For each labeled image in batch:
    - Read data to initialize input layer
    - Evaluate network to get  $y$  (forward)
    - Compare with target label  $t$  to get error  $E$
    - Backpropagate error derivative to get parameter updates
    - Accumulate parameter updates into  $\Delta\theta$
  - $\theta_{i+1} = \theta_i - \epsilon \Delta\theta$

Balance between accuracy of gradient estimation and parallelism

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

26

## When is Training Done?




Split labeled data into *training* and *test* sets.

- Training data to compute parameter updates.
- Test data to check how model generalizes to new inputs (the ultimate goal!)
- The network can become *too good* at classifying training inputs!

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

27

## How Complicated Should a Network Be?




Intuition: like a polynomial fit. High-order terms improve fit, but add unpredictable swings for inputs outside the training set.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

28

## Overtraining Decreases Accuracy



If network works too well for training data,  
new inputs cause big unpredictable output changes.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

29

## No Free Lunch Theorem

- Every classification algorithm has the same error rate when classifying previously unobserved inputs when averaged over all possible input-generating distributions.
- Neural networks must be tuned for specific tasks

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

30

## Summary (1)

- Classification:
  - $f: \mathbb{R}^N \rightarrow \{1, \dots, C\}$
  - $k[i] = f(x, \theta)$
- Current ML work driven by cheap compute, lots of available data
- Perceptron as a trivial deep network
  - $y = \text{sign}(W \cdot x + b)$
- Forward for inference, backward for training

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

31

## Summary (2)

- Chain rule to compute parameter updates
- Stochastic gradient descent for training

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018  
ECE408/CS483, ECE 498AL, University of Illinois, Urbana-Champaign

32