

ECE408 Spring 2020

Applied Parallel Programming

Lecture 15

Parallel Computation Patterns –
Reduction Trees

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

2

2

Objective

- to learn the basic concept of reductions, one of the most widely used parallel computation patterns
- to learn simple strategies for parallelization of reductions
- to understand the performance issues involved with performing reductions on GPUs

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

3

3

Important Enough to Use in Theory

“... scan operations, also known as prefix computations, can execute in no more time than ... parallel memory references ... greatly simplify the description of many [parallel] algorithms, and are significantly easier to implement than memory references.” —Guy Blelloch, 1989*

*G. Blelloch, “Scans as Primitive Parallel Operations,” IEEE Transactions on Computers, 38(11):1526-1538, 1989.
The idea behind scans for computation goes back another 30+ years.

© Steven S. Lumetta
ECE408/CS483/ECE498aI, University of Illinois, 2020

4

4

Trying to Bridge Theory and Practice

A generic parallel algorithm,

- in which parallel threads access memory arbitrarily,
- is likely to produce an extremely slow access pattern.

Scans

- can be implemented quickly in hardware, and
- form a useful alternative to arbitrary memory accesses.

(His hope was to enable theory without knowledge of microarchitecture.)

© Steven S. Lumetta
ECE408/CS483/ECE498aI, University of Illinois, 2020

5

5

Example Use: Summarizing Results

1. Start with a **large set of things** (examples: integers, social networking user information)
2. **Process** each thing **independently to produce** some **value** (examples: number of friends, timeline posts in last two weeks)
3. **Summarize!**
 - Typically with an associative and commutative operation (+, *, min, max, ...)
 - since things in the set are unordered and independent.

© Steven S. Lumetta
ECE408/CS483/ECE498aI, University of Illinois, 2020

6

6

Focus on Reduction Using a Tree

Pattern is so common that

- people have built frameworks around it!
- examples: Google and Hadoop MapReduce

Let's focus on the summarization, called a **reduction**:

- **no required order** for processing the values (operator is associative and commutative), so
- **partition the data set** into smaller chunks,
- have each thread to process a chunk, and
- **use a tree to compute the final answer.**

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

7

7

Reduction Enables Parallelization

Reduction enables common parallel transformations.

example: **privatization of output**

- **Loop iterations sum into a single output** (examples: inner loops in matrix multiply and convolution).
- To parallelize iterations, must **make private copies** of the output!
- **Use reduction** to sum private copies into the original output.

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

8

8

What Exactly is a Reduction?

Reduce a set of inputs to a single value

- using a binary operator, such as
- sum, product, minimum, maximum,
- or a user-defined reduction operation
 - must be associative and commutative
 - and have an identity value (example: 0 for sum)

Available in most parallel libraries as **collective operations** (like barriers).

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

9

9

Sequential Reduction is $O(N)$

Given binary operator \square and an identity value I^\square

- $I^\square = 0$ for sum
- $I^\square = 1$ for product
- $I^\square = \text{largest possible value}$ for min
- $I^\square = \text{smallest possible value}$ for max

```

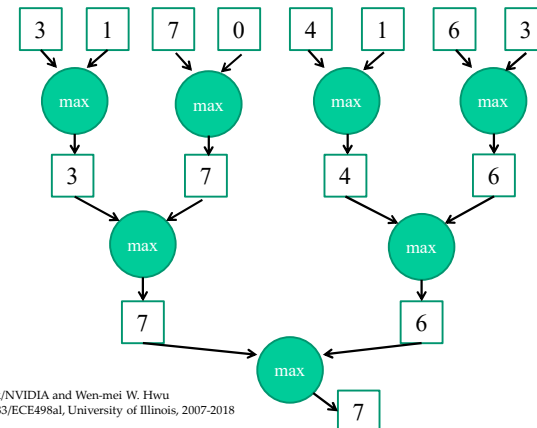
result  $\leftarrow I^\square$ 
for each value  $N$  in input
    result  $\leftarrow \text{result} \square N$ 
    
```

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

10

10

Example: Parallel Max Reduction in $\log(N)$ Steps




© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

11

11

Tournaments Use Reduction with “max”



(A more artful rendition of the reduction tree.)

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

12

12

Algorithm is Work Efficient

For N input values, the number of operations is

$$\frac{1}{2}N + \frac{1}{4}N + \frac{1}{8}N + \dots + \frac{1}{N}N = \left(1 - \frac{1}{N}\right)N = N - 1.$$

The parallel algorithm shown is **work-efficient**:

- requires the **same amount of work** as a sequential algorithm
- (constant overheads, but nothing dependent on N).

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

13

13

Fast if Enough Resources are Available

For N input values, the number of steps is $\log(N)$.

With enough execution resources,

- $N=1,000,000$ takes 20 steps!
- Sounds great!

How much parallelism do we need?

- On average, $(N-1)/\log(N)$.
50,000 in our example.
- But peak is $N/2$!
500,000 in our example.

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

14

14

Diminishing Parallelism is Common

In our parallel reduction,

- the number of operations
- halves in every step.

This kind of narrowing parallelism is common

- from combinational logic circuits
- to basic blocks
- to high-performance applications.

CUDA kernels allow only a fixed number of threads.

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

15

15

Parallel Strategy for CUDA

Let's start simple: N values in device global memory.

Each thread block of M threads

- uses shared memory,
- to reduce chunk of $2M$ values to one value
- ($2M \ll N$ to produce enough thread blocks).

Blocks operate within shared memory

- to reduce global memory traffic, and
- write one value back to global memory.

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

16

16

CUDA Reduction Algorithm

1. Read block of $2M$ values into shared memory.
2. For each of $\log(2M)$ steps,
 - combine two values per thread in each step,
 - write result to shared memory, and
 - halve the number of active threads.
3. Write final result back to global memory.

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

17

17

A Simple Mapping of Data to Threads

Each **thread**

- begins with two **adjacent locations** (**stride of 1**),
- **even index (first)** and an odd index (second).
- Thread 0 gets 0 and 1, Thread 1 gets 2 and 3, ...
- Write **result** back to the **even index**.

After each step,

- **half of active threads** are **done**.
- **Double the stride**.


At the end, **result is at index 0**.

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

18

18

Naïve Data Mapping for a Reduction




© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

19

19

A Sum Example (Values Instead of Indices)



© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

20

20

The Reduction Steps

```
// Stride is distance to the next value being
// accumulated into the threads mapped position
// in the partialSum[] array
for (unsigned int stride = 1;
     stride <= blockDim.x; stride *= 2)
{
    __syncthreads();
    if (t % stride == 0)
        partialSum[2*t] += partialSum[2*t+stride];
}
```

Why do we need `__syncthreads()`?

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498a, University of Illinois, 2007-2018

21

21

Barrier Synchronization

`__syncthreads()` ensures

- all elements of partial sum generated
- before the next step uses them.

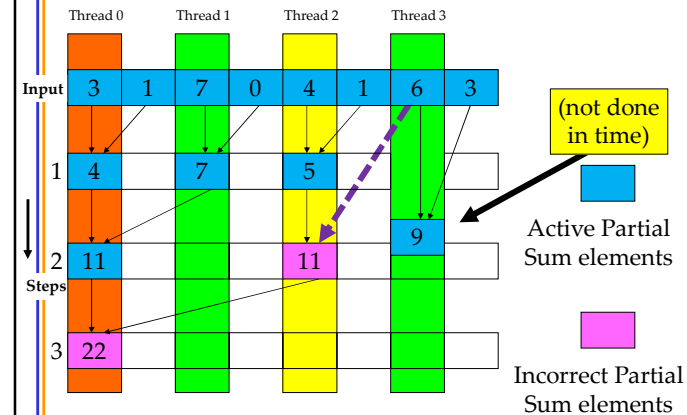
Why do we not need `__syncthreads()` at the end of the reduction loop?

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

22

22

Example Without `__syncthreads`



23

23

Several Options after Blocks are Done

After all reduction steps, **thread 0**

- writes block's sum from `partialSum[0]`
- into global vector indexed by `blockIdx.x`.

Vector has length $N / 2M$.

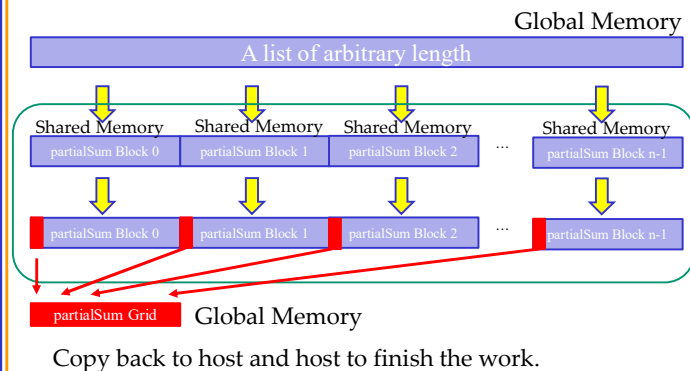
- If small, transfer vector to host and sum it up on CPU.
- If large, launch kernel again (and again).
(Kernel can also accumulate to a global sum using atomic operations, to be covered soon.)

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

24

24

“Segmented Reduction”



25

25

Analysis of Execution Resources

All threads active in the **first step**.

In all **subsequent steps**, two control flow paths:

- **perform addition, or do nothing.**
- Doing nothing still consumes execution resources.

At most half of threads perform addition after first step

- (all threads with odd indices disabled after first step).
- **After fifth step**, entire warps do nothing: **poor resource utilization**, but no divergence.
- **Active warps have only one active thread.**

Up to five more steps (if limited to 1024 threads).

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

26

26

Improve Performance by Reassigning Data

Can we do better?

Absolutely!

How we assign data to threads makes a difference in some algorithms, including reduction.

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

27

27

A Better Strategy

Let's try this approach:

- **in each step,**
- **compact** the partial sums
- **into the first locations**
- in the **partialSum** array

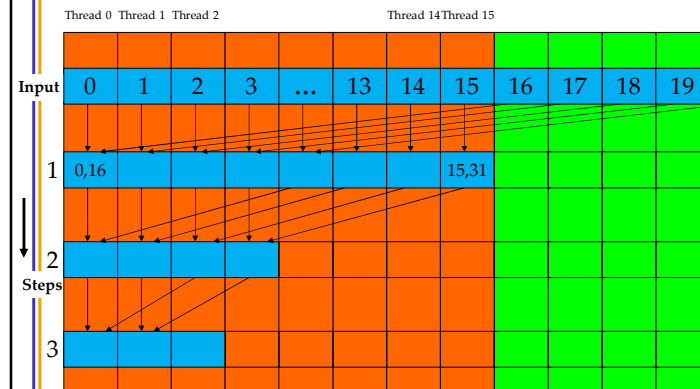
Doing so **keeps the active threads consecutive.**

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

28

28

Illustration with 16 Threads



© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

29

29

A Better Reduction Kernel

```
for (unsigned int stride = blockDim.x;
     stride >= 1;  stride /= 2)
{
    __syncthreads();
    if (t < stride)
        partialSum[t] += partialSum[t+stride];
}
```

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

30

30

Again: Analysis of Execution Resources

Given 1024 threads,

- Block loads 2048 elements to shared memory.
- **No branch divergence** in the **first six steps**:
 - 1024, 512, 256, 128, 64, and 32 consecutive threads active;
 - threads in each warp either all active or all inactive
- **Last six steps** have **one active warp** (branch divergence for last five steps).

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

31

31

Parallel Algorithm Overhead

```
__shared__ float partialSum[2*BLOCK_SIZE];

unsigned int t = threadIdx.x;
unsigned int start = 2*blockIdx.x*blockDim.x;
partialSum[t] = input[start + t];
partialSum[blockDim+t] = input[start+ blockDim.x+t];
for (unsigned int stride = blockDim.x;
     stride >= 1;  stride >>= 1)
{
    __syncthreads();
    if (t < stride)
        partialSum[t] += partialSum[t+stride];
}
```

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

32

32

Parallel Algorithm Overhead

```
__shared__ float partialSum[2*BLOCK_SIZE];


unsigned int t = threadIdx.x;
unsigned int start = 2*blockIdx.x*blockDim.x;
partialSum[t] = input[start + t];
partialSum[blockDim+t] = input[start+ blockDim.x+t];
for (unsigned int stride = blockDim.x;
     stride >= 1;  stride >>= 1)
{
    __syncthreads();
    if (t < stride)
        partialSum[t] += partialSum[t+stride];
}
```

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498aI, University of Illinois, 2007-2018

33


33

Parallel Execution Overhead



Although the number of “operations” is N , each “operation” involves much more complex address calculation and intermediate result manipulation.

If the parallel code is executed on a single-thread hardware, it would be significantly slower than the code based on the original sequential algorithm.



© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498al, University of Illinois, 2007-2018

34

34

Further Improvements

(The slides were meant as extra material for my Wednesday class in Chicago, but that was cancelled.)

Can we further improve reduction?

The **problem is memory-bound**:

- **one operation** for every **4B** value **read**;
- so **focus on memory coalescing** and avoiding poor **computational resource use**.

© Steven S. Lumetta
ECE408/CS483/ECE498al, University of Illinois, 2020

35

35

Make Use of Shared Memory

How much shared memory are we using?

Each block of **1,024** threads reads **2,048** values.

- Let's say **two** blocks per SM,
- so **16 kB** ($= 2,048 \times 2 \times 4B$).

Could **read 4,096 or 8,192** values

- (with **64 kB** per SM)
- to slightly **increase parallelism**.

(For **48 kB** per SM, use **6,144** values and have all threads do a 3-to-1 reduction before the current loop.)

© Steven S. Lumetta
ECE408/CS483/ECE498al, University of Illinois, 2020

36

36

Eliminate the Narrow Parallelism

What about parallelism?

Smaller blocks might seem attractive:

- when one warp is active,
- each SM has one warp per block.

But there are probably better ways. For example,

- **stop reducing at 32 elements** (or at 64, or 128), and
- hand off to the **next kernel**.

© Steven S. Lumetta
ECE408/CS483/ECE498al, University of Illinois, 2020

37

37

Get Rid of the Overhead

Launching kernels is expensive.

- Why bother tearing down and setting up the same blocks on the same SMs?
- Makes no sense.
- Remember that reduction operators are associative and commutative.

Let's **be compute-centric**:

- put **2048 threads** (as two blocks) **on each SM**, and
- just keep them there **until we're done!**

© Steven S. Lumetta
ECE408/CS483/ECE498al, University of Illinois, 2020

38

38

Work Until the Data is Exhausted!

Say there are 8 SMs, so 16 blocks.

1. **Divide** the whole **dataset into 16 chunks**.
2. **Read** enough **to fill shared memory**.
3. **Compute** ... only **until** some **threads not needed**.
4. Then **load more data!**
5. **Repeat** until the data are exhausted,
6. THEN let parallelism drop.
(Gather 16 values on host and reduce them.)

© Steven S. Lumetta
ECE408/CS483/ECE498al, University of Illinois, 2020

39

39

Caveat

I didn't try these ideas.

I'll leave them

- for those of you who feel motivated
- to **try in MP5.1**.

Do

- **save a copy of** your **simpler solution**, though, as
- you will need the partial sums for scan (MP5.2).

© Steven S. Lumetta
ECE408/CS483/ECE498al, University of Illinois, 2020

40

40

Read Chapter 5

ANY MORE QUESTIONS?

© David Kirk/NVIDIA and Wen-mei W. Hwu
ECE408/CS483/ECE498al, University of Illinois, 2007-2018

41

41