

ECE408 / CS483 Spring 2020

Applied Parallel Programming

Lecture 20: GPU as part of the PC Architecture

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

1


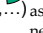
Objectives

- to understand the impact of data transfers on performance when using a GPU as a co-processor
 - speeds and feeds of traditional CPU
 - speeds and feeds when employing a GPU
- to develop a knowledge base for performance tuning for modern GPU's

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

2

Review: Canonical CUDA Program Structure

- Global variables declaration
- Kernel functions
 - `__global__ void kernelOne(...)`
- Main () // host code
 - allocate memory space on the device – `cudaMalloc(&d_GlbVarPtr, bytes)`
 - transfer data from host to device – `cudaMemcpy(d_GlbVarPtr, h_Gl..., ...)`
 - execution configuration setup
 - kernel call – `kernelOne<<<execution configuration>>>(<args...>);` 
 - transfer results from device to host – `cudaMemcpy(h_GlbVarPtr, ..., ...)` 
needed
 - optional: compare against golden (host computed) solution

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

3

Bandwidth: The Gravity of Modern Computer Systems

Bandwidth between key components ultimately **dictates system performance**,


- **especially for GPUs** processing large amounts of data.
- Tricks like buffering, reordering, caching can temporarily defy the rules in some cases.
- Ultimately, performance falls back to what the “speeds and feeds” dictate.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

4

Classic (Historical) PC Architecture

- Northbridge connects 3 components that must communicate at high speed
 - CPU, DRAM, video
 - Video needs first-class access to DRAM
 - Previous NVIDIA cards are connected to AGP, up to 2 GB/s transfers
- Southbridge serves as a concentrator for slower I/O devices




Core Logic Chipset

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

5

(Original) PCI Bus Specification A Humble Beginning

- Connected to the southBridge
 - Originally 33 MHz, 32-bit wide, 132 MB/second peak transfer rate
 - Later, 66 MHz, 64-bit, 528 MB/second peak
 - Upstream bandwidth remain slow for device (~256MB/s peak)
 - Shared bus with arbitration
 - Winner of arbitration becomes bus master and can connect to CPU or DRAM through the southbridge and northbridge




© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

6

PCI as Memory Mapped I/O

- PCI device registers are mapped into the CPU's physical address space
 - Accessed through loads/ stores (kernel mode)
- Addresses are assigned to the PCI devices at boot time
 - All devices listen for their addresses




© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

7

PCI Express (PCIe)

switched, point-to-point connection

- each card has dedicated “link” to the central switch, with no arbitration
- packet switches: messages form virtual channel
- prioritized packets for QoS (such as for real-time video streaming)



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

8

PCIe Generations


- Within a generation, number of lanes in a link can be scaled
 - using distinct physical channels (more bits / wider transfers)
 - $\times 1, \times 2, \times 4, \times 8, \times 16, \times 32, \dots$
- Each new generation aims to double the speed.

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

9

PCIe Gen 3 Links and Lanes

- Each link consists of one or more lanes
 - Each lane is 1-bit wide (4 wires, each 2-wire pair can transmit 8Gb/s in one direction)
 - Upstream and downstream simultaneous and symmetric
 - Each Link can combine 1, 2, 4, 8, 12, 16 lanes- x1, x2, etc.
 - Each byte data is **128b/130b** encoded into 130 bits with equal number of 1's and 0's; net data rate 7.8768 Gb/s per lane each way.
 - Thus, the net data rates are 985 MB/s (x1) 1.97 GB/s (x2), 3.94 GB/s (x4), 7.9 GB/s (x8), 15.8 GB/s (x16), each way



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

10

Foundation: 8/10 bit encoding

- Goal is to maintain DC balance while have sufficient state transition for clock recovery
 - 00000000, 00000111, 11000001 bad
 - 01010101, 11001100 good
- The difference of 1s and 0s in a 20-bit stream should be ≤ 2
- There should be no more than 5 consecutive 1s or 0s in any stream
 - Find 256 good patterns among 1024 total patterns of 10 bits to encode an 8-bit datum
 - a 20% overhead

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

11

Current: 128/130 bit encoding

- Same goal: maintain DC balance while have sufficient state transition for clock recovery
 - Scrambler function: long runs of 0s, 1s vanishingly small
 - Instead of guaranteed run length of 8/10b
 - At least one bit shift every 66 bits
- 1.5% overhead instead of 20%

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

12

Patterns Contain Many 0s and 1s

A question for fun:

- if we need 2^{128} code words
- chosen from all 2^{130} 130-bit patterns
- how many 0s/1s must we consider including?

Answer: 63-67 (of either type)

Thus 128b/130b code words are pretty well-balanced, and have lots of 0-1 transitions (for clock recovery).

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign


13

Recent PCIe PC Architecture

Today, PCIe forms the interconnect backbone within PC.

Northbridge and Southbridge are PCIe switches.

Source: Jon Stokes, PCI Express: An Overview
(<http://arstechnica.com/articles/paedia/hardware/pcie.ars>)



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

14

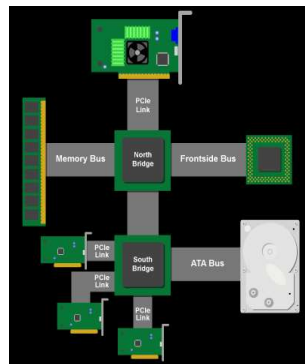
Recent PCIe PC Architecture

How is PCI supported?

- Need a PCI-PCIe bridge, which is
- sometimes included as part of Southbridge, or
- can add as a separate PCIe I/O card.

Current systems integrate PCIe controllers directly on chip with CPU.


Source: Jon Stokes, PCI Express: An Overview
(<http://arstechnica.com/articles/paedia/hardware/pcie.ars>)



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

15

GeForce GTX 1080 (Pascal) GPU Consumer Card Details




© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

16

PCIe Data Transfer using DMA

DMA (Direct Memory Access) is used to fully utilize the bandwidth of an I/O bus

- DMA uses physical address for source and destination
- Transfers a number of bytes requested by OS
- Needs pinned memory



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

17

Pinned Memory

- DMA uses physical addresses
- The OS could accidentally page out the data that is being read or written by a DMA and page in another virtual page into the same location
- Pinned memory cannot not be paged out
- If a source or destination of a `cudaMemcpy()` in the host memory is not pinned, it needs to be first copied to a pinned memory – extra overhead
- `cudaMemcpy` is much faster with pinned host memory source or destination

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

18

Allocate/Free Pinned Memory (a.k.a. Page Locked Memory)

- `cudaHostAlloc()`
 - Three parameters
 - Address of pointer to the allocated memory
 - Size of the allocated memory in bytes
 - Option – use `cudaHostAllocDefault` for now
- `cudaFreeHost()`
 - One parameter
 - Pointer to the memory to be freed

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

19

Using Pinned Memory

- Use the allocated memory and its pointer the same way those returned by `malloc()`;
- The only difference is that the allocated memory cannot be paged by the OS
- The `cudaMemcpy` function should be about 2X faster with pinned memory
- Pinned memory is a limited resource whose over-subscription can have serious consequences

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

20

Important Trends

- Knowing yesterday, today, and tomorrow
 - The PC world is becoming flatter
 - CPU and GPU are being fused together
 - Outsourcing of computation is becoming easier...

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

21

ANY MORE QUESTIONS?

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483/CSE408, ECE 498AL, University of Illinois, Urbana-Champaign

22