

ECE408 / CS483/CSE408 Spring 2020

Applied Parallel Programming

Lecture 23: Joint CUDA-MPI Programming

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

2

2

Objective


- To be familiar with simple MPI-CUDA heterogeneous applications
 - understand the key sections of an MPI application
 - understand explicit communication in parallel computing applications

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

3

3

Blue Waters Computing System



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

4

4

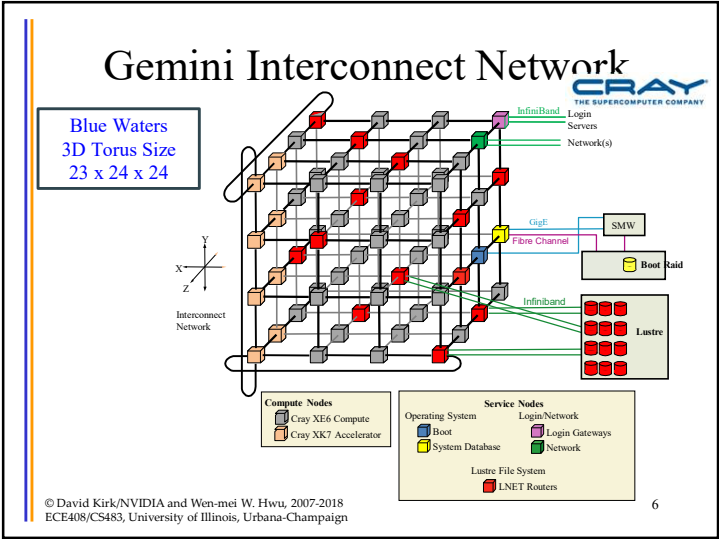
Blue Waters and Titan Computing Systems

System Attribute	NCSA	ORNL
	Blue Waters	Titan
Vendors	Cray/AMD/NVIDIA	Cray/AMD/NVIDIA
Processors	Interlagos/Kepler	Interlagos/Kepler
Total Peak Performance (PF)	11.1	27.1
Total Peak Performance (CPU/GPU)	7.1/4	2.6/24.5
Number of CPU Chips	48,352	18,688
Number of GPU Chips	3,072	18,688
Amount of CPU Memory (TB)	1511	584
Interconnect	3D Torus	3D Torus
Amount of On-line Disk Storage (PB)	26	13.6
Sustained Disk Transfer (TB/sec)	>1	0.4-0.7
Amount of Archival Storage	300	15-30
Sustained Tape Transfer (GB/sec)	100	7

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

5

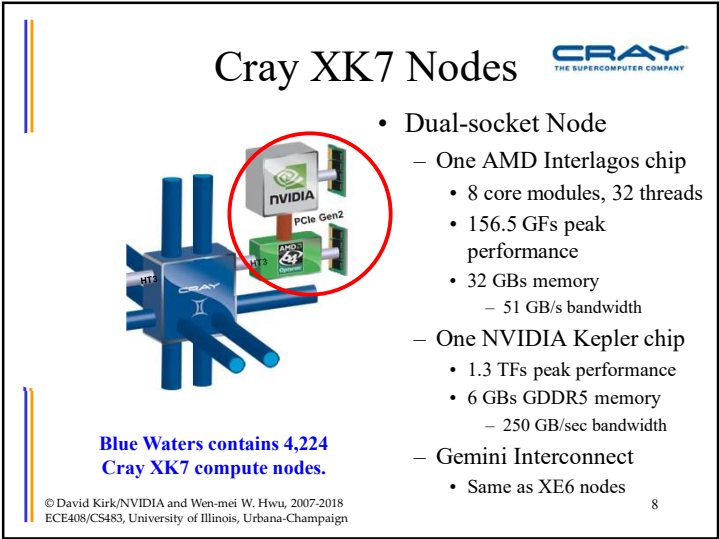
5



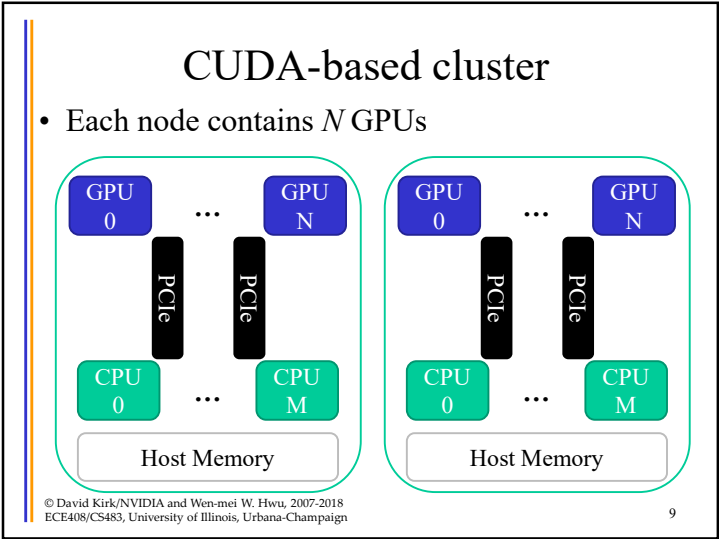
6

Science Area	Number of Teams	Codes	Struct Grids	Unstruct Grids	Dense Matrix	Sparse Matrix	N-Body	Monte Carlo	FFT	PIC	Signal and I/O
Climate and Weather	3	CESM, GCRM, CMU-WRF, HOMME	X	X		X		X			X
Plasmas/Magnetosphere	2	H3DM/VPIC, OSIRIS, Magtail/UPIC	X				X		X		X
Stellar Atmospheres and Supernovae	5	PPM, MAESTRO, CASTRO, SEDONA, ChaNGa, MS-FLUKSS	X			X	X	X		X	X
Cosmology	2	Enzo, pGADGET	X			X	X				
Combustion/Turbulence	2	PSDNS, DISTUF	X						X		
General Relativity	2	Cactus, Ham3D, Luth	X			X					
Molecular Dynamics	4	AMBER, Gromacs, NAMD, LAMMPS				X	X		X		
Quantum Chemistry	2	SIAL, GAMESS, NWChem			X	X	X	X			X
Material Science	3	NEMOS, OMEN, GW, QMCPACK			X	X	X	X			
Earthquakes/Seismology	2	AWP-ODC, HERCULES, PLSQR, SPEC3D	X	X			X				X
Quantum Chromo Dynamics	1	Chroma, MILC, USQCD	X		X	X					
Social Networks	1	EPISIMDEMICS									
Evolution	1	Eve									
Engineering/System of Systems	1	GRIPS, Reviat						X			
Computer Science	1			X	X	X			X	7	X

7




8



9

MPI Model

- Many processes distributed in a cluster



- Each process computes part of the output
- Processes communicate with each other through message passing (not global memory)
- Processes can synchronize through messages

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

10

10

MPI Initialization, Info

- User launches an MPI job with X processes by executing in the command shell
 - `MPIrun -np X`
- `int MPI_Init(int *argc, char ***argv)`
 - Initialize MPI
- `MPI_COMM_WORLD`
 - MPI group formed with all allocated nodes
- `int MPI_Comm_rank(MPI_Comm comm, int *rank)`
 - Rank of the calling process in group of comm
- `int MPI_Comm_size(MPI_Comm comm, int *size)`
 - Number of processes in the group of comm

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

11

11

Vector Addition: Main Process

```
int main(int argc, char *argv[]) {
    int vector_size = 1024 * 1024 * 1024;
    int pid=-1, np=-1;

    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &pid);
    MPI_Comm_size(MPI_COMM_WORLD, &np);

    if(np < 3) {
        if(0 == pid) printf("Needed 3 or more processes.\n");
        MPI_Abort( MPI_COMM_WORLD, 1 ); return 1;
    }
    if(pid < np - 1)
        compute_node(vector_size);
    else
        data_server(vector_size);

    MPI_Finalize();
    return 0;
}
```


© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

12

12

MPI Sending Data

- `int MPI_Send(void *buf, int count, MPI_Datatype datatype, int dest, int tag, MPI_Comm comm)`
 - Buf:** Initial address of send buffer
 - Count:** Number of elements in send buffer (nonnegative integer)
 - Datatype:** Datatype of each send buffer element
 - Dest:** Rank of destination (integer)
 - Tag:** Message tag (integer)
 - Comm:** Communicator (handle)



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

14

14

MPI Receiving Data

- `int MPI_Recv(void *buf, int count, MPI_Datatype datatype, int source, int tag, MPI_Comm comm, MPI_Status *status)`
 - **Buf**: Starting address of receive buffer
 - **Count**: Maximum number of elements in receive buffer (non-negative integer)
 - **Datatype**: Datatype of each receive buffer element
 - **Source**: Rank of source (integer)
 - **Tag**: Message tag (integer)
 - **Comm**: Communicator (handle)
 - **Status**: Status object

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

15

15

Vector Addition: Server Process (I)

```
void data_server(unsigned int vector_size) { // runs on rank 0 only
    int np;
    unsigned int num_bytes = vector_size * sizeof(float);

    /* Set MPI Communication Size */
    MPI_Comm_size(MPI_COMM_WORLD, &np);
    int num_nodes = np - 1;
    unsigned int vector_part = vector_size / num_nodes;

    /* Allocate input data */
    float* input_a = (float *)malloc(num_bytes);
    float* input_b = (float *)malloc(num_bytes);
    float* output = (float *)malloc(num_bytes);
    if(input_a == NULL || input_b == NULL || output == NULL) {
        printf("Server couldn't allocate memory\n");
        MPI_Abort( MPI_COMM_WORLD, 1 );
    }
    /* Initialize input data */
    random_data(input_a, vector_size, 1, 10);
    random_data(input_b, vector_size, 1, 10);
}
```

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

17

17

Vector Addition: Server Process (II)

```
/* Send data to compute nodes */
for(int process = 0; process < num_nodes; process++) {

    MPI_Send(input_a + process * vector_part, vector_part,
             MPI_FLOAT, process, DATA_DISTRIBUTE, MPI_COMM_WORLD);

    MPI_Send(input_b + process * vector_part, vector_part,
             MPI_FLOAT, process, DATA_DISTRIBUTE, MPI_COMM_WORLD);
}

// Example of a barrier in MPI. Here we want to avoid doing a
// barrier, since sends and receives are blocking anyway, and we
// want to avoid synchronizing the workers' attempts to use the
// incoming network bandwidth to the "server," which is the
// communication bottleneck here.

// MPI_Barrier(MPI_COMM_WORLD);
```

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

18

18

Vector Addition: Server Process (III)

```
/* Collect output data */
MPI_Status status;
for(int process = 0; process < num_nodes; process++) {
    MPI_Recv(output + process * vector_part, vector_part,
             MPI_FLOAT, process, DATA_COLLECT, MPI_COMM_WORLD,
             &status);
}

/* Store output data */
store_output(output, vector_size);

/* Release resources */
free(input_a);
free(input_b);
free(output);
}
```

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

19

19

Vector Addition: Compute Process (I)

```
void compute_node(unsigned int vector_size) {
    int np;

    MPI_Comm_size(MPI_COMM_WORLD, &np);
    int server_process = np - 1; // also # of compute processes
    unsigned int vector_part = vector_size / server_process;
    unsigned int num_bytes = vector_part * sizeof(float);

    /* Alloc host memory */
    float* input_a = (float *)malloc(num_bytes);
    float* input_b = (float *)malloc(num_bytes);
    float* output = (float *)malloc(num_bytes);

    /* Get the input data from server process */
    MPI_Status status;
    MPI_Recv(input_a, vector_part, MPI_FLOAT, server_process,
             DATA_DISTRIBUTE, MPI_COMM_WORLD, &status);
    MPI_Recv(input_b, vector_part, MPI_FLOAT, server_process,
             DATA_DISTRIBUTE, MPI_COMM_WORLD, &status);
}
```

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

20

20

Vector Addition: Compute Process (II)

```
/* Compute the partial vector addition */
for(int i = 0; i < vector_part; ++i) {
    output[i] = input_a[i] + input_b[i];
}

/* Or, can offload to GPU here */
/* cudaMalloc(), cudaMemcpy(), kernel launch, SYNCHRONIZE */

// Example of a barrier in MPI (want to avoid here)
// MPI_Barrier(MPI_COMM_WORLD);

/* Send the output */
MPI_Send(output, vector_part, MPI_FLOAT,
         server_process, DATA_COLLECT, MPI_COMM_WORLD);


/* Release memory */
free(input_a);
free(input_b);
free(output);
}
```

21

21

Current Trends

- GPU Direct
 - MPI_Send and MPI_Receive deals directly with GPU memory
 - Requires support from NIC vendors



© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

22

22

QUESTIONS?

© David Kirk/NVIDIA and Wen-mei W. Hwu, 2007-2018
ECE408/CS483, University of Illinois, Urbana-Champaign

23

23