

University of Illinois at Urbana-Champaign
Dept. of Electrical and Computer Engineering

ECE 408 / CS 483 / CSE 408: Applied Parallel Programming

Generalizing Parallelism

ECE 408 / CS 483 / CSE 408:
Applied Parallel Programming

© 2009-2020 Steven S. Lumetta. All rights reserved.

slide 2

2

Objective

- to learn terminology and concepts from the broader high-performance computing community
- to generalize some of the techniques illustrated in class for use with future codes

ECE 408 / CS 483 / CSE 408:
Applied Parallel Programming

© 2009-2020 Steven S. Lumetta. All rights reserved.

slide 3

3

Speedup Measures the Success of Parallelization

Let's start by defining **parallel speedup** (usually just called speedup).

Let's say that

- when I run my program in parallel
- it finishes **X** times faster
- than when I run it sequentially.

Specifically,

- $X = T(\text{sequential}) / T(\text{parallel})$, and
- **X is the speedup** of my parallel code.

Note that speedup assumes a fixed problem size.

ECE 408 / CS 483 / CSE 408:
Applied Parallel Programming

© 2009-2020 Steven S. Lumetta. All rights reserved.

slide 4

4

Speedup Depends on the Best Sequential Code

We have $T(\text{sequential}) / T(\text{parallel})$.

But how do we find $T(\text{sequential})$?

$T(\text{sequential})$ **should measure** the

- **best algorithm** for a sequential machine (may/may not be the algorithm parallelized),
- **optimized** for a sequential machine, with
- **no parallelism support** remnants (no parallel overhead).

ECE 408 / CS 483 / CSE 408:
Applied Parallel Programming

© 2009-2020 Steven S. Lumetta. All rights reserved.

slide 5

5

Find (Don't Write) a Competitive Baseline

Sequential code is what we in Engineering call

- the **baseline design**,
- the alternative against which
- we demonstrate improvements.

As Prof. Hwu once pointed out to me,

- **no one will believe** that **you** worked hard
- to **optimize your baseline...**
- even if you did!

If possible, **compare someone else's best work**.

6

Efficiency Measures Effective Use of Resources

Next is **parallel efficiency**
(or just efficiency).

Efficiency measures how well
a code uses parallel resources.

When executing **on P processors**,

$$\text{efficiency} = \text{speedup on P processors} / P.$$

7

Efficiency is Often Below 1, But Should Not be Tiny

What value should efficiency have?

According to those paying for the machines, 1.

According to most real applications,

- **something non-negligible, near 1**
- but not 1,
- as other bottlenecks come into play.

8

Efficiency is Rarely Above 1

Can efficiency be >1?

Rarely—called **superlinear speedup**.

possible causes:

- certain types of extra resources
(such as caches)
- luck (parallel search happens to
find an answer more quickly).


9

Scalability Measures Effect of Parallel Overheads

Next, **scalability**:

- **for how many processors is speedup linear**, or is efficiency flat?

At some **P**, with fixed problem size, speedup will flatten out.




10

Good Scalability Requires Minimal Parallel Overhead

For larger values of **P**, speedup starts to drop (unless one leaves processors idle).

Good scalability means

- **no falloff** on your machine
- **for maximum measurable value of P.**



11

Efficiency Not So Meaningful When Cores Vary Widely

But what is P for a single GPU?

1?

Number of SMs?

Number of PEs (total)?

We can still measure speedup,

- but for a single GPU,
- we **estimate efficiency**
- **by comparing** resource use
- **with** the GPU's **peak values**.

(As we've done in our class already.)

12

Speedup Measures Improvement for an Input Set

Again, **speedup assumes** a **fixed problem size**.

- For many applications, that's reasonable.
- Users care about their input sets, not about hypothetical inputs.

But that's **not always the best assumption**.

13

For Other Situations, We Need Different Metrics

Sometimes we care about throughput:

- frames per second for video / game quality,
- transactions per second for databases, or
- user operations per second for datacenters.

And sometimes input size

- **is limited** by memory
- or by feasible runtime,
- as in many supercomputing applications.

14

Scaling Problem Size with P Good for Science Apps

Other variants of speedup on P processors:*

scaled speedup:

- problem size is linear in **P**
- (good scaled speedup is **1**)

memory-constrained speedup:

- biggest problem that fits in memory (which scales with **P**)
- only works for **O(N)** algorithms

*J. P. Singh, J. L. Hennessy, A. Gupta, "Scaling Parallel Programs for Multiprocessors: Methodology and Examples," *IEEE Computer*, 26(7):42-50, July 1993.

15

Problem Size Sometimes Chosen Through Practical Means

Other variants of speedup on P processors:*

time-constrained speedup:

- biggest problem that finishes by the time I return from lunch
- sometimes reasonable...
- ...but we could wait overnight for a grand challenge application?

*J. P. Singh, J. L. Hennessy, A. Gupta, "Scaling Parallel Programs for Multiprocessors: Methodology and Examples," *IEEE Computer*, 26(7):42-50, July 1993.

16

Parallel Grain Size is the Work Done per Thread

Parallel grain size is work per thread (task).

- Remember discussing what to parallelize?
- Output elements, input elements, ...

Each source of parallelism has a natural grain size:

- loop body,
- objects in a container,
- rows/columns/blocks/elements in a matrix,
- graph nodes/connected components.

17

Consider Different Sources of Parallelism

Some sources exhibit higher work variance (and branch divergence) than others

- conditionals/inner loops in loop body
- complex per-object methods
- rows in upper/lower diagonal matrix
- matrix elements usually roughly constant
- degree of nodes, size of connected components.

Be sure to consider the alternatives!


18

Amdahl's Law Helps Set Expectations

Amdahl's Law says

- **speedup** is **bounded** above
- by $1 / (\text{sequential fraction})$.

For example, if you parallelize code that takes 75% of the time, you can't get more than 4× speedup.



19

Evaluate Your Work Intelligently and Meaningfully

But, again, for fixed input.

There are other 'laws' as well that view the problem differently.

So what matters most?

- Some apps today are missing/simplified due to resource limits.
- Some apps become possible/more useful with bigger problem sizes.

**Fit evaluation of utility to your app,
not your app to an evaluation metric.**

20

A Few Useful Concepts

Now, I'd like to go over a few useful ideas from high-performance computing.

Most you've seen before, so I'll tie them in to what you've seen and done in our class.

21

Bulk Synchronous Execution Dominates Fast Computing

The **bulk synchronous** style

- dominates HPC and CUDA applications.
- **Barriers separate** temporal **regions of code**
 - usually O(100) lines long
 - interleaving / data **sharing occurs only within regions** (called phases).

Why?

- Simpler to debug regions than whole programs.
- (similar to Stroustrup's view of classes' value).

Bulk synchronous execution **does tend to correlate resource usage**, which is bad.

22

Necessary/Good Sources of Parallel Overhead

Good ways to waste time in parallel;

- push bits around (**communicate**)—a necessary overhead in most parallel codes
- **do some extra work** (to avoid communicating)
 - for example, do pooling after convolution in a CNN kernel to reduce shared-to-global memory traffic
 - another: do extra adds to reduce the number of barriers, as in a Kogge-Stone scan
- bicker about priority (**contend for shared resources**)

23

Bad Sources of Parallel Overhead

Bad ways to waste time in parallel;

- twiddle your thumbs (**wait for long-latency events**)
- **watch others work**
 - example: branch divergence in a GPU
 - example: poor scheduling decisions
- line up single file (**unnecessary serialization**)
 - example: coarse synchronization, lack of privatization
 - example: temporally correlated accesses to shared hardware resource
 - example: use one CUDA stream

24

Dynamic Load Balancing Sometimes Needed

In our class, we have generally

- assigned fixed work per thread.
- Usually, this is the simplest approach
- but may lead to load imbalance.

One common solution—**load balancing**:

- dynamic mapping of work to threads using
- one or more queues of work
 - pull chunk of work from a queue, do it, repeat
 - start with bigger chunks, later grab smaller
 - if queue is empty, **steal work** from another.

25

CUDA Scheduling May Need to Become More Expressive

One last question: kernel/block scheduling.

Most OS schedulers use **time-sharing**:
try to be fair to all of the running programs.

But if you have many processors,
why pay parallel overhead?

Use **space-sharing** instead!

Lots of supercomputers and datacenters do.

How are thread blocks within
CUDA kernels scheduled?