

# The LLVM Compiler Framework and Infrastructure

Chris Lattner  
**[lattner@cs.uiuc.edu](mailto:lattner@cs.uiuc.edu)**

Vikram Adve  
**[vadve@cs.uiuc.edu](mailto:vadve@cs.uiuc.edu)**

**<http://llvm.cs.uiuc.edu/>**

L CPC Tutorial: September 22, 2004

# Acknowledgements

## UIUC Contributors:

- ❖ Tanya Brethour
- ❖ Misha Brukman
- ❖ Cameron Buschardt
- ❖ John Criswell
- ❖ Alkis Evlogimenos
- ❖ Brian Gaeke
- ❖ Ruchira Sasanka
- ❖ Anand Shukla
- ❖ Bill Wendling

## External Contributors:

- ❖ Henrik Bach
- ❖ Nate Begeman
- ❖ Jeff Cohen
- ❖ Paolo Invernizzi
- ❖ Brad Jones
- ❖ Vladimir Merzliakov
- ❖ Vladimir Prus
- ❖ Reid Spencer

## Funding:

This work is sponsored by the NSF Next Generation Software program through grants EIA-0093426 (an NSF CAREER award) and EIA-0103756. It is also supported in part by the NSF Operating Systems and Compilers program (grant #CCR-9988482), the NSF Embedded Systems program (grant #CCR-0209202), the MARCO/DARPA Gigascale Systems Research Center (GSRC), IBM through the DARPA-funded PERCS project, and the Motorola University Partnerships in Research program.

# LLVM Compiler System

## ■ The LLVM Compiler Infrastructure

- ❖ Provides reusable components for building compilers
- ❖ Reduce the time/cost to build a new compiler
- ❖ Build static compilers, JITs, trace-based optimizers, ...

## ■ The LLVM Compiler Framework

- ❖ End-to-end compilers using the LLVM infrastructure
- ❖ C and C++ are robust and aggressive:
  - Java, Scheme and others are in development
- ❖ Emit C code or native code for X86, Sparc, PowerPC

# Three primary LLVM components

- **The LLVM *Virtual Instruction Set***
  - ❖ The common language- and target-independent IR
  - ❖ Internal (IR) and external (persistent) representation
- **A collection of well-integrated libraries**
  - ❖ Analyses, optimizations, code generators, JIT compiler, garbage collection support, profiling, ...
- **A collection of tools built from the libraries**
  - ❖ Assemblers, automatic debugger, linker, code generator, compiler driver, modular optimizer, ...

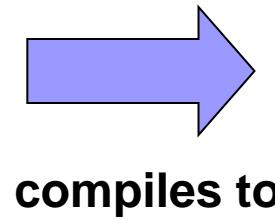
# Tutorial Overview

- **Introduction to the running example**
- **LLVM C/C++ Compiler Overview**
  - ❖ High-level view of an example LLVM compiler
- **The LLVM Virtual Instruction Set**
  - ❖ IR overview and type-system
- **LLVM C++ IR and important API's**
  - ❖ Basics, PassManager, dataflow, ArgPromotion
- **Important LLVM Tools**
  - ❖ opt, code generator, JIT, test suite, bugpoint
- **Example applications of LLVM**

# Running example: arg promotion

Consider use of by-reference parameters:

```
int callee(const int &X) {  
    return X+1;  
}  
  
int caller() {  
    return callee(4);  
}
```



```
int callee(const int *X) {  
    return *X+1; // memory load  
}  
  
int caller() {  
    int tmp; // stack object  
    tmp = 4; // memory store  
    return callee(&tmp);  
}
```

We want:

```
int callee(int X) {  
    return X+1;  
}  
  
int caller() {  
    return callee(4);  
}
```

- ✓ Eliminated load in callee
- ✓ Eliminated store in caller
- ✓ Eliminated stack slot for ‘tmp’

# Why is this hard?

- **Requires interprocedural analysis:**
  - ❖ Must change the prototype of the callee
  - ❖ Must update all call sites → we must **know** all callers
  - ❖ What about callers outside the translation unit?
- **Requires alias analysis:**
  - ❖ Reference could alias other pointers in callee
  - ❖ Must know that loaded value doesn't change from function entry to the load
  - ❖ Must know the pointer is not being stored through
- **Reference might not be to a stack object!**

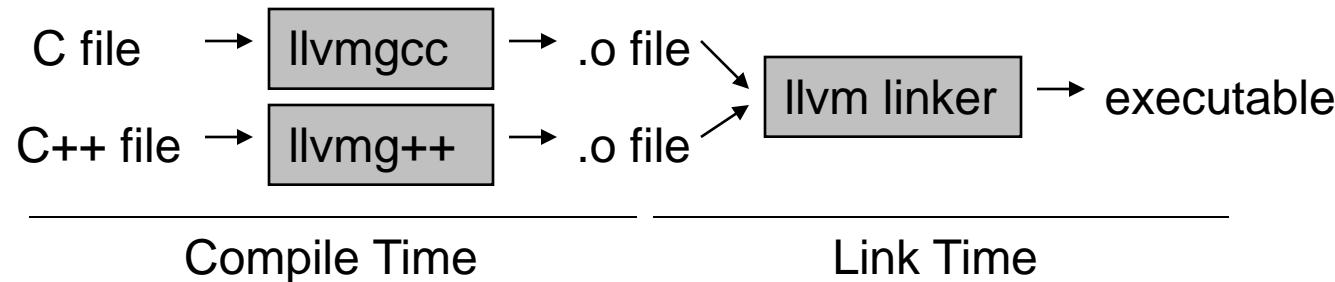
# Tutorial Overview

- **Introduction to the running example**
- **LLVM C/C++ Compiler Overview**
  - ❖ High-level view of an example LLVM compiler
- **The LLVM Virtual Instruction Set**
  - ❖ IR overview and type-system
- **LLVM C++ IR and important API's**
  - ❖ Basics, PassManager, dataflow, ArgPromotion
- **Important LLVM Tools**
  - ❖ opt, code generator, JIT, test suite, bugpoint
- **Example applications of LLVM**

# The LLVM C/C++ Compiler

## ■ From the high level, it is a standard compiler:

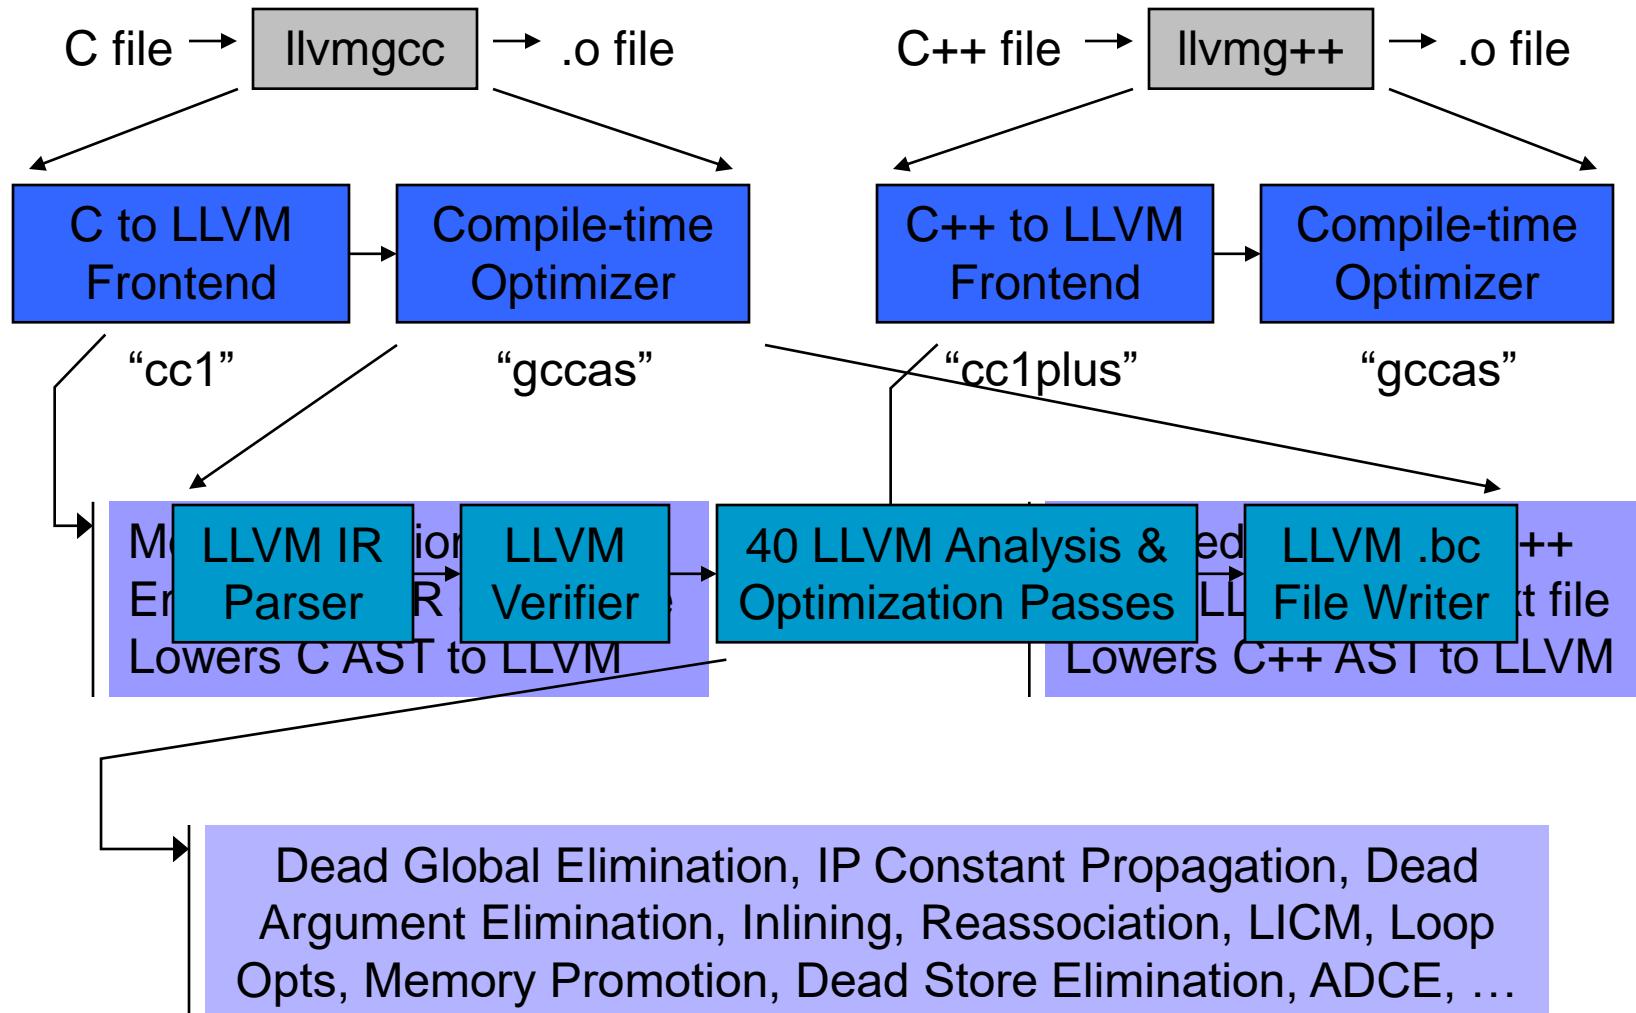
- ❖ Compatible with standard makefiles
- ❖ Uses GCC 3.4 C and C++ parser



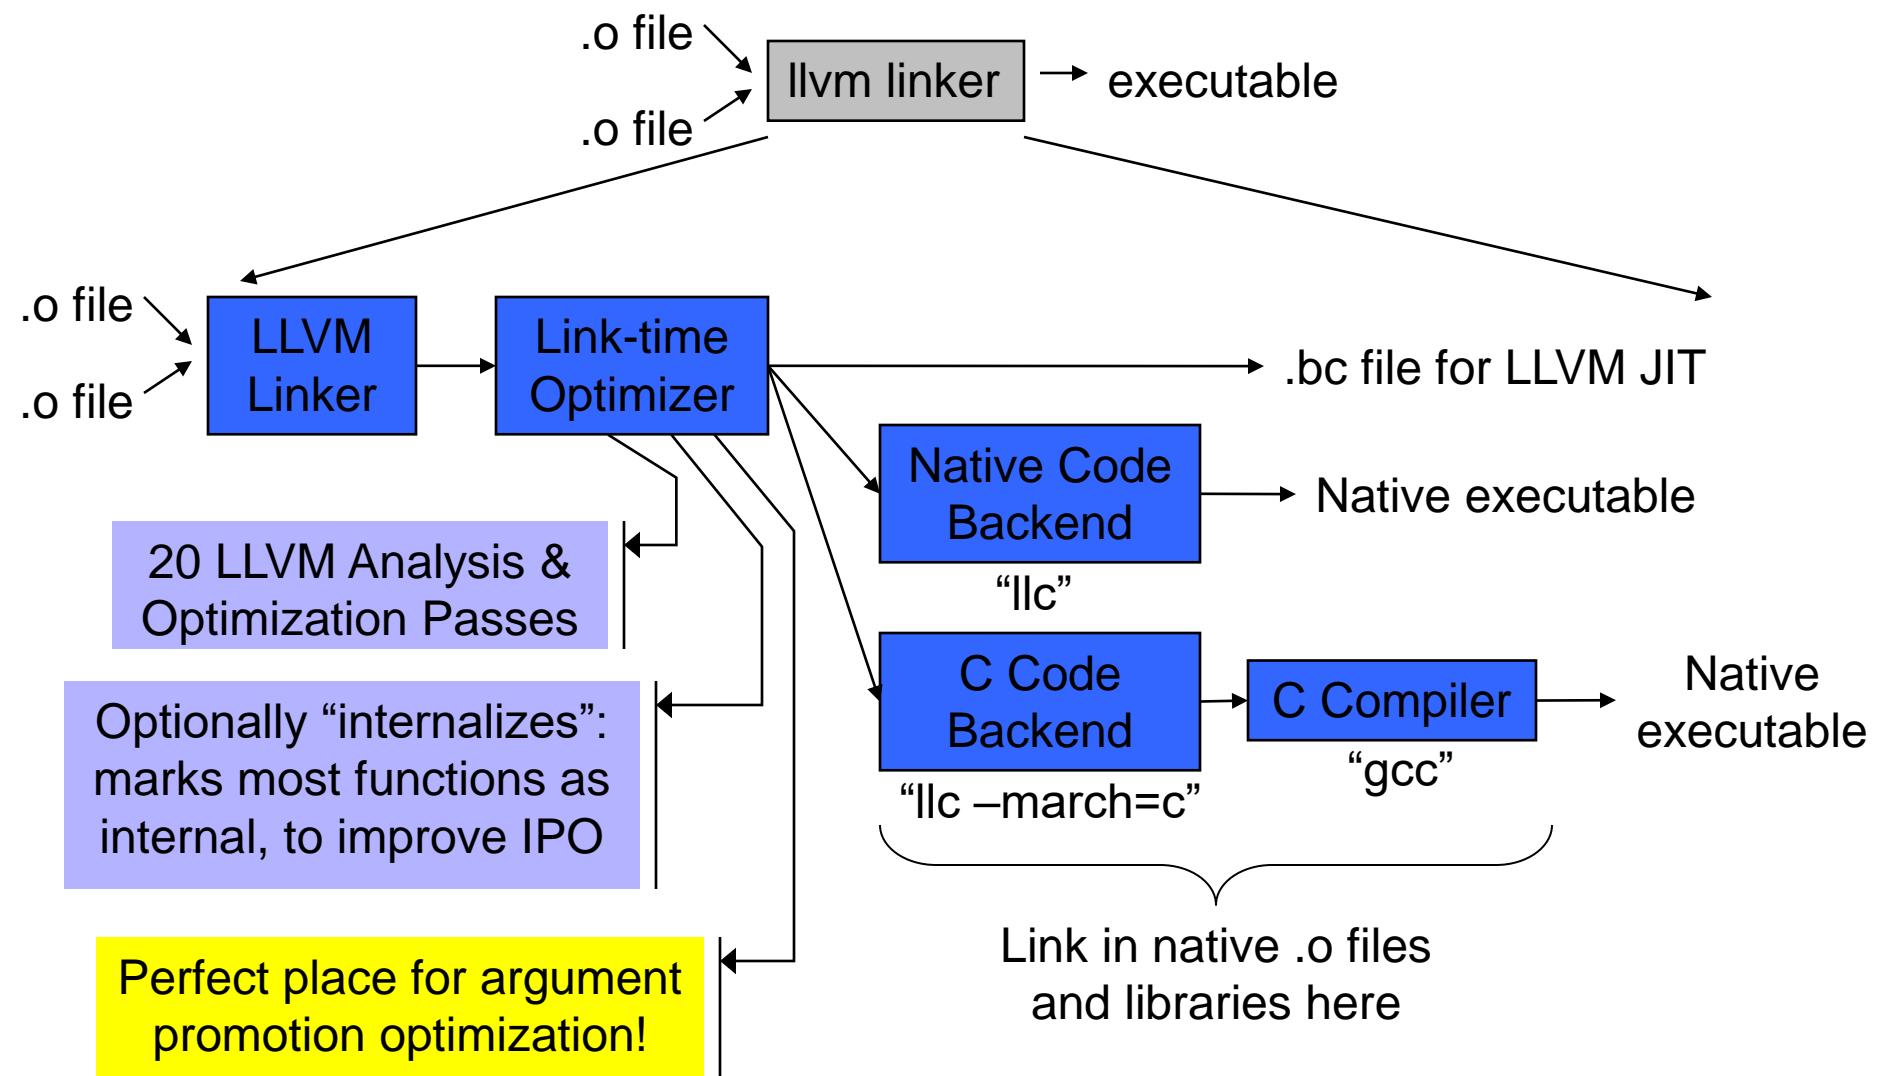
## ■ Distinguishing features:

- ❖ Uses LLVM optimizers, not GCC optimizers
- ❖ .o files contain LLVM IR/bytocode, not machine code
- ❖ Executable can be bytecode (JIT'd) or machine code

# Looking into events at compile-time



# Looking into events at link-time



# Goals of the compiler design

- **Analyze and optimize as early as possible:**
  - ❖ Compile-time opts reduce modify-rebuild-execute cycle
  - ❖ Compile-time optimizations reduce work at link-time (by shrinking the program)
- **All IPA/IPO make an open-world assumption**
  - ❖ Thus, they all work on libraries and at compile-time
  - ❖ “Internalize” pass enables “whole program” optzn
- **One IR (without lowering) for analysis & optzn**
  - ❖ Compile-time optzns can be run at link-time too!
  - ❖ The same IR is used as input to the JIT

***IR design is the key to these goals!***

# Tutorial Overview

- **Introduction to the running example**
- **LLVM C/C++ Compiler Overview**
  - ❖ High-level view of an example LLVM compiler
- **The LLVM Virtual Instruction Set**
  - ❖ IR overview and type-system
- **LLVM C++ IR and important API's**
  - ❖ Basics, PassManager, dataflow, ArgPromotion
- **Important LLVM Tools**
  - ❖ opt, code generator, JIT, test suite, bugpoint
- **Example applications of LLVM**

# Goals of LLVM IR

- **Easy to produce, understand, and define!**
- **Language- and Target-Independent**
  - ❖ AST-level IR (e.g. ANDF, UNCOL) is not very feasible
    - Every analysis/xform must know about ‘all’ languages
- **One IR for analysis and optimization**
  - ❖ IR must be able to support aggressive IPO, loop opts, scalar opts, ... high- *and* low-level optimization!
- **Optimize as much as early as possible**
  - ❖ Can’t postpone everything until link or runtime
  - ❖ No lowering in the IR!

# LLVM Instruction Set Overview #1

## ■ Low-level and target-independent semantics

- ❖ RISC-like three address code
- ❖ Infinite virtual register set in SSA form
- ❖ Simple, low-level control flow constructs
- ❖ Load/store instructions with typed-pointers

## ■ IR has text, binary, and in-memory forms

```
loop:  
    %i.1 = phi int [ 0, %bb0 ], [ %i.2, %loop ]  
    %AiAddr = getelementptr float* %A, int %i.1  
    call void %Sum(float %AiAddr, %pair* %P)  
    %i.2 = add int %i.1, 1  
    %tmp.4 = setlt int %i.1, %N  
    br bool %tmp.4, label %loop, label %outloop  
  
for (i = 0; i < N;  
     ++i)  
    Sum(&A[i], &P);
```

# LLVM Instruction Set Overview #2

## ■ High-level information exposed in the code

- ❖ Explicit dataflow through SSA form
- ❖ Explicit control-flow graph (even for exceptions)
- ❖ Explicit language-independent type-information
- ❖ Explicit typed pointer arithmetic
  - Preserve array subscript and structure indexing

```
for (i = 0; i < N;  
     ++i)  
    Sum(&A[i], &P);  
  
loop:  
    %i.1 = phi int [ 0, %bb0 ], [ %i.2, %loop ]  
    %AiAddr = getelementptr float* %A, int %i.1  
    call void %Sum(float %AiAddr, %pair* %P)  
    %i.2 = add int %i.1, 1  
    %tmp.4 = setlt int %i.1, %N  
    br bool %tmp.4, label %loop, label %outloop
```

# LLVM Type System Details

- **The entire type system consists of:**
  - ❖ Primitives: void, bool, float, ushort, opaque, ...
  - ❖ Derived: pointer, array, structure, function
  - ❖ No high-level types: type-system is language neutral!
- **Type system allows arbitrary casts:**
  - ❖ Allows expressing weakly-typed languages, like C
  - ❖ *Front-ends can implement safe languages*
  - ❖ *Also easy to define a type-safe subset of LLVM*

See also: [docs/LangRef.html](#)

# Lowering source-level types to LLVM

- **Source language types are lowered:**
  - ❖ Rich type systems expanded to simple type system
  - ❖ Implicit & abstract types are made explicit & concrete
- **Examples of lowering:**
  - ❖ References turn into pointers: `T&` → `T*`
  - ❖ Complex numbers: `complex float` → `{ float, float }`
  - ❖ Bitfields: `struct X { int Y:4; int Z:2; } → { int }`
  - ❖ Inheritance: `class T : S { int X; } → { S, int }`
  - ❖ Methods: `class T { void foo(); } → void foo(T*)`
- **Same idea as lowering to machine code**

# LLVM Program Structure

- **Module contains Functions/GlobalVariables**
  - ❖ Module is unit of compilation/analysis/optimization
- **Function contains BasicBlocks/Arguments**
  - ❖ Functions roughly correspond to functions in C
- **BasicBlock contains list of instructions**
  - ❖ Each block ends in a control flow instruction
- **Instruction is opcode + vector of operands**
  - ❖ All operands have types
  - ❖ Instruction result is typed

# Our example, compiled to LLVM

```
int callee(const int *X) {  
    return *X+1; // load  
}  
int caller() {  
    int T;          // on stack  
    T = 4;          // store  
    return callee(&T);  
}
```

Linker “internalizes”  
most functions in most  
cases

```
internal int %callee(int* %X) {  
    %tmp.1 = load int* %X  
    %tmp.2 = add int %tmp.1, 1  
    ret int %tmp.2  
}  
int %caller() {  
    %T = alloca int  
    store int 4, int* %T  
    %tmp.3 = call int %callee(int* %T)  
    ret int %tmp.3  
}
```

# Our example, desired transformation

```
internal int %callee(int* %X) {  
    %tmp.1 = load int* %X  
    %tmp.2 = add int %tmp.1, 1  
    ret int %tmp.2  
}  
  
int %caller() {  
    %T = alloca int  
    store int 4, int* %T  
    %tmp.3 = call int %callee(int* %T)  
    ret int %tmp.3  
}
```

```
internal int %callee(int %X.val) {  
    %tmp.2 = add int %X.val, 1  
    ret int %tmp.2  
}  
  
int %caller() {  
    %T = alloca int  
    store int 4, int* %T  
    %tmp.1 = load int* %T  
    %tmp.3 = call int %callee(%tmp.1)  
    ret int %tmp.3  
}
```

Other transformation  
(-mem2reg) cleans up  
the rest

```
int %caller() {  
    %tmp.3 = call int %callee(int 4)  
    ret int %tmp.3  
}
```

# Tutorial Overview

- **Introduction to the running example**
- **LLVM C/C++ Compiler Overview**
  - ❖ High-level view of an example LLVM compiler
- **The LLVM Virtual Instruction Set**
  - ❖ IR overview and type-system
- **LLVM C++ IR and important API's**
  - ❖ Basics, PassManager, dataflow, ArgPromotion
- **Important LLVM Tools**
  - ❖ opt, code generator, JIT, test suite, bugpoint
- **Example applications of LLVM**

# LLVM Coding Basics

- **Written in modern C++, uses the STL:**
  - ❖ Particularly the vector, set, and map classes
- **LLVM IR is almost all doubly-linked lists:**
  - ❖ Module contains lists of Functions & GlobalVariables
  - ❖ Function contains lists of BasicBlocks & Arguments
  - ❖ BasicBlock contains list of Instructions

- **Linked lists are traversed with iterators:**

```
Function *M = ...  
  
for (Function::iterator I = M->begin(); I != M->end(); ++I) {  
    BasicBlock &BB = *I;  
  
    ...
```

See also: [docs/ProgrammersManual.html](#)

# LLVM Pass Manager

- **Compiler is organized as a series of ‘passes’:**
  - ❖ Each pass is one analysis or transformation
- **Four types of Pass:**
  - ❖ **ModulePass**: general interprocedural pass
  - ❖ **CallGraphSCCPass**: bottom-up on the call graph
  - ❖ **FunctionPass**: process a function at a time
  - ❖ **BasicBlockPass**: process a basic block at a time
- **Constraints imposed (e.g. FunctionPass):**
  - ❖ FunctionPass can only look at “current function”
  - ❖ Cannot maintain state across functions

See also: <docs/WritingAnLLVMPass.html>

# Services provided by PassManager

## ■ Optimization of pass execution:

- ❖ Process a function at a time instead of a pass at a time
- ❖ Example: If F, G, H are three functions in input pgm:  
“FFFFGGGGHHHH” not “FGHFGHFGHFGH”
- ❖ Process functions in parallel on an SMP (future work)

## ■ Declarative dependency management:

- ❖ Automatically fulfill and manage analysis pass lifetimes
- ❖ Share analyses between passes when safe:
  - e.g. “DominatorSet live unless pass modifies CFG”

## ■ Avoid boilerplate for traversal of program

See also: [docs/WritingAnLLVMPass.html](#)

# Pass Manager + Arg Promotion #1/2

## ■ Arg Promotion is a CallGraphSCCPass:

- ❖ Naturally operates bottom-up on the CallGraph
  - Bubble pointers from callees out to callers

```
24: #include "llvm/CallGraphSCCPass.h"
47: struct SimpleArgPromotion : public CallGraphSCCPass {
```

## ■ Arg Promotion requires AliasAnalysis info

- ❖ To prove safety of transformation
  - Works with any alias analysis algorithm though

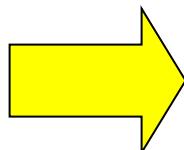
```
48: virtual void getAnalysisUsage(AnalysisUsage &AU) const {
    AU.addRequired<AliasAnalysis>();           // Get aliases
    AU.addRequired<TargetData>();                // Get data layout
    CallGraphSCCPass::getAnalysisUsage(AU); // Get CallGraph
}
```

# Pass Manager + Arg Promotion #2/2

## ■ Finally, implement `runOnSCC` (line 65):

```
bool SimpleArgPromotion::  
runOnSCC(const std::vector<CallGraphNode*> &SCC) {  
    bool Changed = false, LocalChange;  
    do { // Iterate until we stop promoting from this SCC.  
        LocalChange = false;  
        // Attempt to promote arguments from all functions in this SCC.  
        for (unsigned i = 0, e = SCC.size(); i != e; ++i)  
            LocalChange |= PromoteArguments(SCC[i]);  
        Changed |= LocalChange; // Remember that we changed something.  
    } while (LocalChange);  
    return Changed; // Passes return true if something changed.  
}
```

```
static int foo(int ***P) {  
    return ***P;  
}
```



```
static int foo(int P_val_val_val) {  
    return P_val_val_val;  
}
```

# LLVM Dataflow Analysis

- **LLVM IR is in SSA form:**

- ❖ use-def and def-use chains are always available
  - ❖ All objects have user/use info, even functions

- **Control Flow Graph is always available:**

- ❖ Exposed as BasicBlock predecessor/successor lists
  - ❖ Many generic graph algorithms usable with the CFG

- **Higher-level info implemented as passes:**

- ❖ Dominators, CallGraph, induction vars, aliasing, GVN, ...

See also: [docs/ProgrammersManual.html](http://llvm.cs.uiuc.edu/docs/ProgrammersManual.html)

# Arg Promotion: safety check #1/4

## #1: Function must be “internal” (aka “static”)

```
88: if (!F || !F->hasInternalLinkage()) return false;
```

## #2: Make sure address of F is not taken

- ❖ In LLVM, check that there are only direct calls using F

```
99: for (Value::use_iterator UI = F->use_begin();
        UI != F->use_end(); ++UI) {
    CallSite CS = CallSite::get(*UI);
    if (!CS.getInstruction()) // "Taking the address" of F.
        return false;
```

## #3: Check to see if any args are promotable:

```
114: for (unsigned i = 0; i != PointerArgs.size(); ++i)
      if (!isSafeToPromoteArgument(PointerArgs[i]))
          PointerArgs.erase(PointerArgs.begin() + i);
      if (PointerArgs.empty()) return false; // no args promotable
```

# Arg Promotion: safety check #2/4

## #4: Argument pointer can only be loaded from:

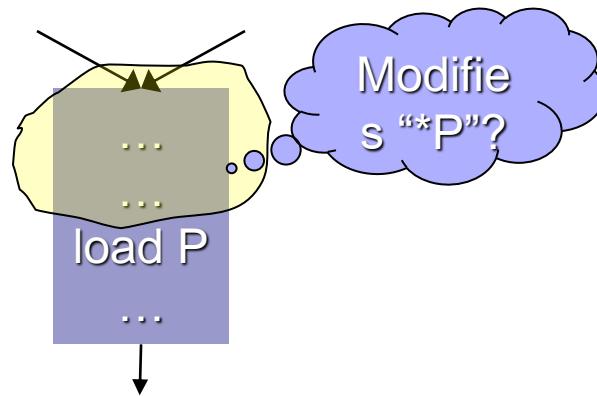
- ❖ No stores through argument pointer allowed!

```
// Loop over all uses of the argument (use-def chains).  
138: for (Value::use_iterator UI = Arg->use_begin();  
        UI != Arg->use_end(); ++UI) {  
    // If the user is a load:  
    if (LoadInst *LI = dyn_cast<LoadInst>(*UI)) {  
        // Don't modify volatile loads.  
        if (LI->isVolatile()) return false;  
        Loads.push_back(LI);  
    } else {  
        return false; // Not a load.  
    }  
}
```

# Arg Promotion: safety check #3/4

## #5: Value of “\*P” must not change in the BB

- ❖ We move load out to the caller, value cannot change!

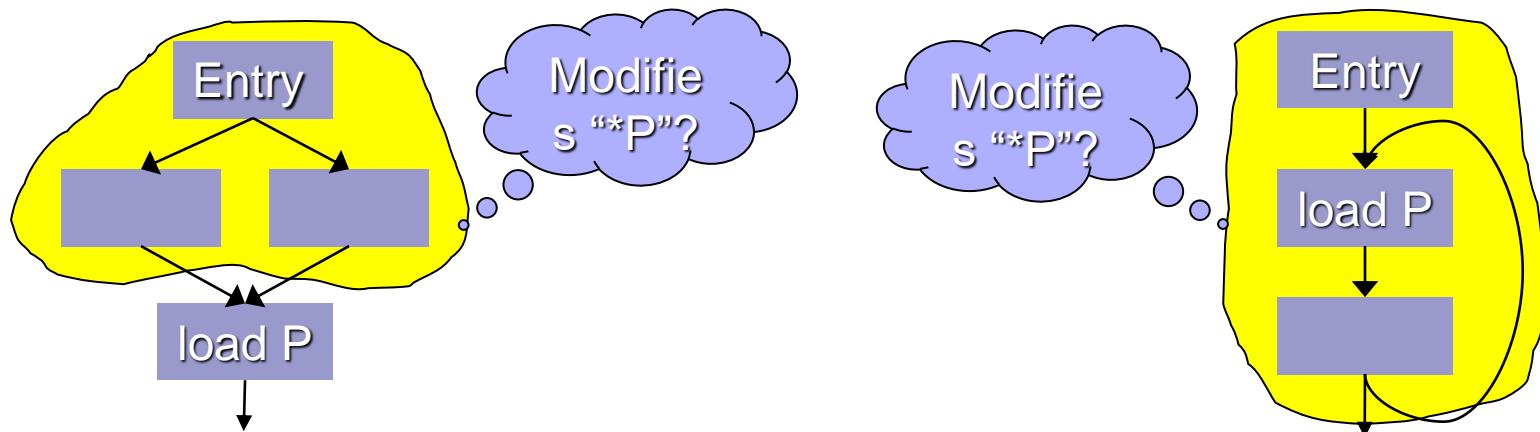


```
// Get AliasAnalysis implementation from the pass manager.  
156: AliasAnalysis &AA = getAnalysis<AliasAnalysis>();  
  
// Ensure *P is not modified from start of block to load  
169: if (AA.canInstructionRangeModify(BB->front(), *Load,  
                                     Arg, LoadSize))  
    return false; // Pointer is invalidated!
```

See also: [docs/AliasAnalysis.html](#)

# Arg Promotion: safety check #4/4

## #6: “\*P” cannot change from Fn entry to BB



```
175: for (pred_iterator PI = pred_begin(BB), E = pred_end(BB);  
        PI != E; ++PI)      // Loop over predecessors of BB.  
        // Check each block from BB to entry (DF search on inverse graph).  
        for (idfa_iterator<BasicBlock*> I = idf_begin(*PI);  
             I != idf_end(*PI); ++I)  
            // Might *P be modified in this basic block?  
            if (AA.canBasicBlockModify(**I, Arg, LoadSize))  
                return false;
```

# Arg Promotion: xform outline #1/4

## #1: Make prototype with new arg types: #197

- ❖ Basically just replaces ‘int\*’ with ‘int’ in prototype

## #2: Create function with new prototype:

```
214: Function *NF = new Function(NFTy, F->getLinkage(),
                                F->getName());
      F->getParent()->getFunctionList().insert(F, NF);
```

## #3: Change all callers of F to call NF:

*// If there are uses of F, then calls to it remain.*

```
221: while (!F->use_empty()) {
      // Get a caller of F.
      CallSite CS = CallSite::get(F->use_back());
```

# Arg Promotion: xform outline #2/4

## #4: For each caller, add loads, determine args

- ❖ Loop over the args, inserting the loads in the caller

```
220: std::vector<Value*> Args;

226: CallSite::arg_iterator AI = CS.arg_begin();
      for (Function::aiterator I = F->abegin(); I != F->aend();
            ++I, ++AI)
      if (!ArgsToPromote.count(I))           // Unmodified argument.
          Args.push_back(*AI);
      else {                                // Insert the load before the call.
          LoadInst *LI = new LoadInst(*AI, (*AI)->getName() + ".val",
                                         Call); // Insertion point
          Args.push_back(LI);
      }
```

# Arg Promotion: xform outline #3/4

## #5: Replace the call site of F with call of NF

```
// Create the call to NF with the adjusted arguments.  
242: Instruction *New = new CallInst(NF, Args, "", Call);  
  
// If the return value of the old call was used, use the retval of the new call.  
if (!Call->use_empty())  
    Call->replaceAllUsesWith(New);  
  
// Finally, remove the old call from the program, reducing the use-count of F.  
Call->getParent()->getInstList().erase(Call);
```

## #6: Move code from old function to new Fn

```
259: NF->getBasicBlockList().splice(NF->begin(),  
                                    F->getBasicBlockList());
```

# Arg Promotion: xform outline #4/4

## #7: Change users of F's arguments to use NF's

```
264: for (Function::aiterator I = F->abegin(), I2 = NF->abegin();  
        I != F->aend(); ++I, ++I2)  
    if (!ArgsToPromote.count(I)) { // Not promoting this arg?  
        I->replaceAllUsesWith(I2); // Use new arg, not old arg.  
    } else {  
        while (!I->use_empty()) { // Only users can be loads.  
            LoadInst *LI = cast<LoadInst>(I->use_back());  
            LI->replaceAllUsesWith(I2);  
            LI->getParent()->getInstList().erase(LI);  
        }  
    }  
}
```

## #8: Delete old function:

```
286: F->getParent()->getFunctionList().erase(F);
```

# Tutorial Overview

- **Introduction to the running example**
- **LLVM C/C++ Compiler Overview**
  - ❖ High-level view of an example LLVM compiler
- **The LLVM Virtual Instruction Set**
  - ❖ IR overview and type-system
- **LLVM C++ IR and important API's**
  - ❖ Basics, PassManager, dataflow, ArgPromotion
- **Important LLVM Tools**
  - ❖ opt, code generator, JIT, test suite, bugpoint
- **Example applications of LLVM**

# LLVM tools: two flavors

## ■ “Primitive” tools: do a single job

- ❖ `llvm-as`: Convert from .ll (text) to .bc (binary)
- ❖ `llvm-dis`: Convert from .bc (binary) to .ll (text)
- ❖ `llvm-link`: Link multiple .bc files together
- ❖ `llvm-prof`: Print profile output to human readers
- ❖ `llvmc`: Configurable compiler driver

## ■ Aggregate tools: pull in multiple features

- ❖ `gccas/gcld`: Compile/link-time optimizers for C/C++ FE
- ❖ `bugpoint`: automatic compiler debugger
- ❖ `llvm-gcc/llvm-g++`: C/C++ compilers

See also: [docs/CommandGuide/](#)

# opt tool: LLVM modular optimizer

## ■ Invoke arbitrary sequence of passes:

- ❖ Completely control PassManager from command line
- ❖ Supports loading passes as plugins from .so files

**opt -load foo.so -pass1 -pass2 -pass3 x.bc -o y.bc**

## ■ Passes “register” themselves:

```
61: RegisterOpt<SimpleArgPromotion> X("simpleargpromotion",  
    "Promote 'by reference' arguments to 'by value'");
```

## ■ From this, they are exposed through opt:

```
> opt -load libsimpleargpromote.so -help  
...  
-sccp           - Sparse Conditional Constant Propagation  
-simpleargpromotion - Promote 'by reference' arguments to 'by value'  
-simplifycfg      - Simplify the CFG  
...
```

# Running Arg Promotion with opt

## ■ Basic execution with ‘opt’:

- ❖ `opt -simpleargpromotion in.bc -o out.bc`
- ❖ Load .bc file, run pass, write out results
- ❖ Use “-load filename.so” if compiled into a library
- ❖ PassManager resolves all dependencies

## ■ Optionally choose an alias analysis to use:

- ❖ `opt -basicaa -simpleargpromotion` (default)
- ❖ Alternatively, `-steens-aa`, `-anders-aa`, `-ds-aa`, ...

## ■ Other useful options available:

- ❖ `-stats`: Print statistics collected from the passes
- ❖ `-time-passes`: Time each pass being run, print output

# Example -stats output (gccas 176.gcc)

```
=====  
          ... Statistics Collected ...  
=====
```

23426 adce	- Number of instructions removed
1663 adce	- Number of basic blocks removed
5052592 bytecodeewriter	- Number of bytecode bytes written
57489 cfgsimplify	- Number of blocks simplified
4186 constmerge	- Number of global constants merged
211 dse	- Number of stores deleted
15943 gcse	- Number of loads removed
54245 gcse	- Number of instructions removed
253 inline	- Number of functions deleted because all callers found
3952 inline	- Number of functions inlined
9425 instcombine	- Number of constant folds
160469 instcombine	- Number of insts combined
208 licm	- Number of load insts hoisted or sunk
4982 licm	- Number of instructions hoisted out of loop
350 loop-unroll	- Number of loops completely unrolled
30156 mem2reg	- Number of alloca's promoted
2934 reassociate	- Number of insts with operands swapped
650 reassociate	- Number of insts reassociated
67 scalarrepl	- Number of allocas broken up
279 tailcallelim	- Number of tail calls removed
25395 tailduplicate	- Number of unconditional branches eliminated

.....

# Example -time-passes (gccas 176.gcc)

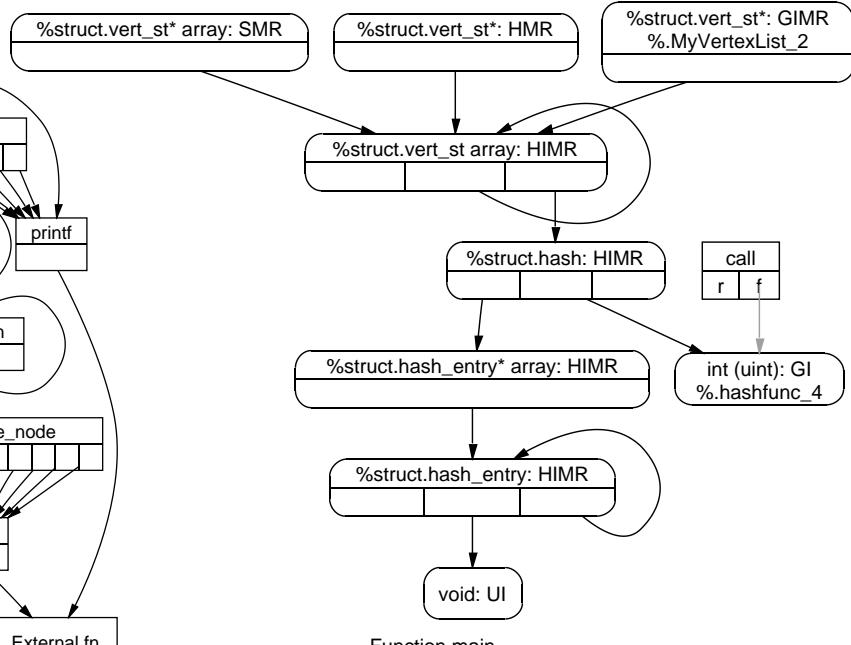
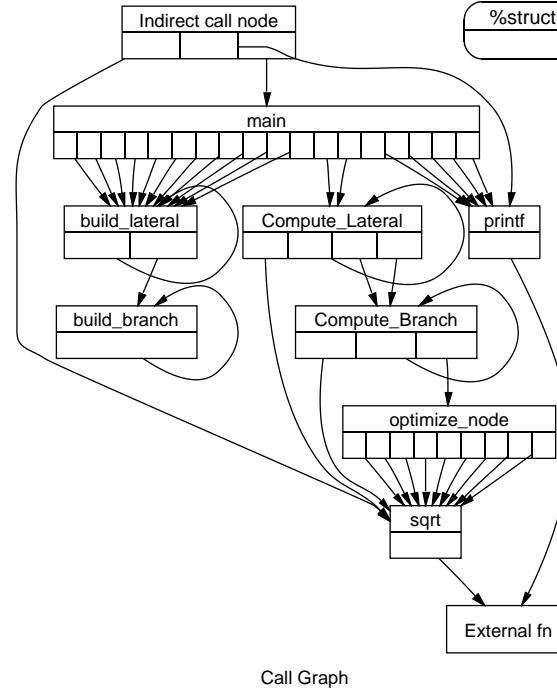
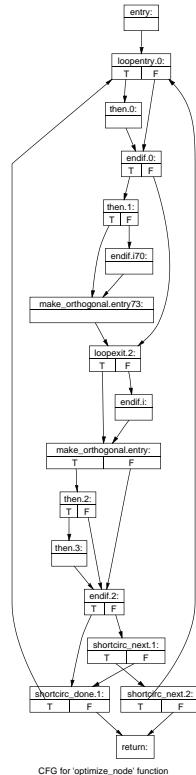
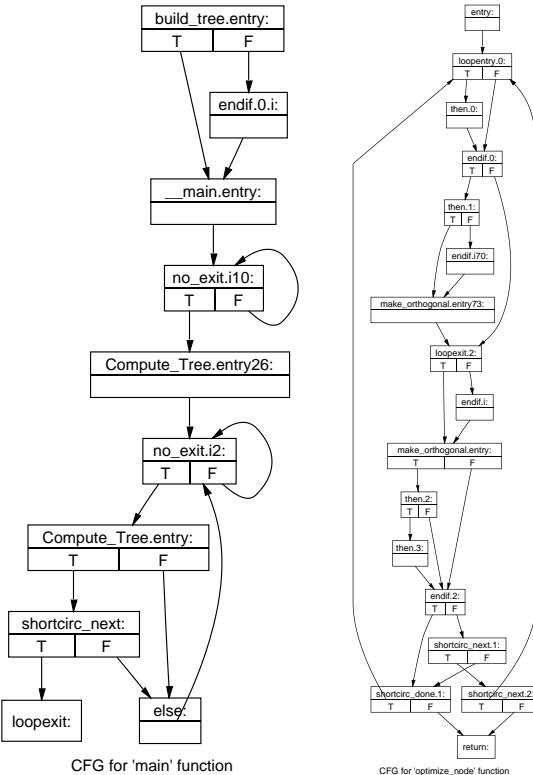
... Pass execution timing report ...

---User Time---	--System Time--	--User+System--	--Wall Time---	--- Name ---
16.2400 ( 23.0%)	0.0000 ( 0.0%)	16.2400 ( 22.9%)	16.2192 ( 22.9%)	Global Common Subexpression Elimination
11.1200 ( 15.8%)	0.0499 ( 13.8%)	11.1700 ( 15.8%)	11.1028 ( 15.7%)	Reassociate expressions
6.5499 ( 9.3%)	0.0300 ( 8.3%)	6.5799 ( 9.3%)	6.5824 ( 9.3%)	Bytecode Writer
3.2499 ( 4.6%)	0.0100 ( 2.7%)	3.2599 ( 4.6%)	3.2140 ( 4.5%)	Scalar Replacement of Aggregates
3.0300 ( 4.3%)	0.0499 ( 13.8%)	3.0800 ( 4.3%)	3.0382 ( 4.2%)	Combine redundant instructions
2.6599 ( 3.7%)	0.0100 ( 2.7%)	2.6699 ( 3.7%)	2.7339 ( 3.8%)	Dead Store Elimination
2.1600 ( 3.0%)	0.0300 ( 8.3%)	2.1900 ( 3.0%)	2.1924 ( 3.1%)	Function Integration/Inlining
2.1600 ( 3.0%)	0.0100 ( 2.7%)	2.1700 ( 3.0%)	2.1125 ( 2.9%)	Sparse Conditional Constant Propagation
1.6600 ( 2.3%)	0.0000 ( 0.0%)	1.6600 ( 2.3%)	1.6389 ( 2.3%)	Aggressive Dead Code Elimination
1.4999 ( 2.1%)	0.0100 ( 2.7%)	1.5099 ( 2.1%)	1.4462 ( 2.0%)	Tail Duplication
1.5000 ( 2.1%)	0.0000 ( 0.0%)	1.5000 ( 2.1%)	1.4410 ( 2.0%)	Post-Dominator Set Construction
1.3200 ( 1.8%)	0.0000 ( 0.0%)	1.3200 ( 1.8%)	1.3722 ( 1.9%)	Canonicalize natural loops
1.2700 ( 1.8%)	0.0000 ( 0.0%)	1.2700 ( 1.7%)	1.2717 ( 1.7%)	Merge Duplicate Global Constants
1.0300 ( 1.4%)	0.0000 ( 0.0%)	1.0300 ( 1.4%)	1.1418 ( 1.6%)	Combine redundant instructions
0.9499 ( 1.3%)	0.0400 ( 11.1%)	0.9899 ( 1.4%)	0.9979 ( 1.4%)	Raise Pointer References
0.9399 ( 1.3%)	0.0100 ( 2.7%)	0.9499 ( 1.3%)	0.9688 ( 1.3%)	Simplify the CFG
0.9199 ( 1.3%)	0.0300 ( 8.3%)	0.9499 ( 1.3%)	0.8993 ( 1.2%)	Promote Memory to Register
0.9600 ( 1.3%)	0.0000 ( 0.0%)	0.9600 ( 1.3%)	0.8742 ( 1.2%)	Loop Invariant Code Motion
0.5600 ( 0.7%)	0.0000 ( 0.0%)	0.5600 ( 0.7%)	0.6022 ( 0.8%)	Module Verifier
...				

# Analyze tool: Visualize analysis results

## ■ Print most LLVM data structures

- ❖ Dominators, loops, alias sets, CFG, call graph, ...
- ❖ Converts most LLVM data structures to ‘dot’ graphs



# LLC Tool: Static code generator

- **Compiles LLVM → native assembly language**
  - ❖ Currently for X86, Sparc, PowerPC (others in alpha)
  - ❖ `llc file.bc -o file.s -march=x86`
  - ❖ `as file.s -o file.o`
- **Compiles LLVM → portable C code**
  - ❖ `llc file.bc -o file.c -march=c`
  - ❖ `gcc -c file.c -o file.o`
- **Targets are modular & dynamically loadable:**
  - ❖ `llc -load libarm.so file.bc -march=arm`

# The LLVM Code Generator

## ■ Target independent:

- ❖ Driven by an algorithm independent target description
  - Data layout, Register, Instruction, Scheduling, ...

## ■ Basic code generator layout:



## ■ All passes are replaceable

(by hand for now)

- ❖ e.g. Trivial to change and add

## ■ Targets do have special passes

Exposes all target-specific details about a function (calling conventions, etc)

special support for FP stack

See also: [docs/CodeGenerator.html](#)

# Porting LLVM to a new target

- **LLVM targets are very easy to write:**
  - ❖ Anecdotal evidence suggests 1 week for a basic port
    - ... for someone familiar with the target machine and compilers in general, but not with LLVM
- **LLVM targets are written with “tablegen” tool**
  - ❖ Simple declarative syntax
  - ❖ Designed to factor out redundancy in target desc
- **Some C++ code is still required**
  - ❖ Primarily in the instruction selector
  - ❖ Continuing work to improve this

See also: [docs/TableGenFundamentals.html](#) and [WritingAnLLVMBackend.html](#)

# LLI Tool: LLVM Execution Engine

- **LLI allows direct execution of .bc files**
  - ❖ E.g.: `lli grep.bc -i foo *.c`
- **LLI uses a Just-In-Time compiler if available:**
  - ❖ Uses same code generator as LLC
    - Optionally uses faster components than LLC
  - ❖ Emits machine code to memory instead of “.s” file
  - ❖ JIT is a library that can be embedded in other tools
- **Otherwise, it uses the LLVM interpreter:**
  - ❖ Interpreter is extremely simple and very slow
  - ❖ Interpreter is portable though!

# C and C++ Program Test Suite

- **Large collection of programs and benchmarks:**
  - ❖ Standard suites (e.g. SPEC 95/2000, Olden, Ptrdist, McCat, Stanford, Freebench, Shootout...)
  - ❖ Individual programs: sgefa, siod, sim, pi, povray, ...
  - ❖ Proprietary suites (e.g. SPEC) require suite source
- **Consistent build environment:**
  - ❖ Easy add hooks to build for profiling/instrumentation
  - ❖ Easy to get performance numbers from entire test suite
- **Entire test suite is checked every night:**
  - ❖ Hosted on Linux, Solaris, FreeBSD on X86, Sparc & PPC

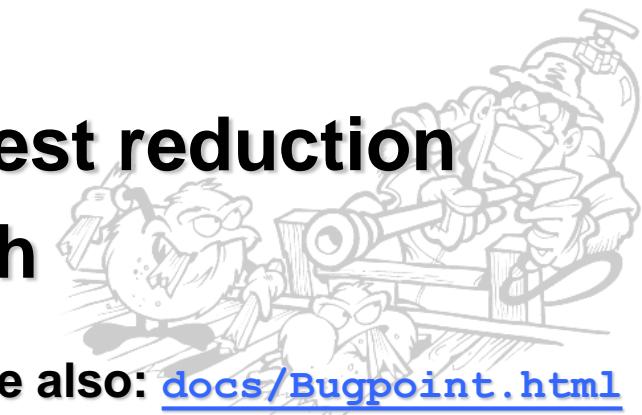
See also: [docs/TestingGuide.html](#)

# Integrated Debugging Tools

- **Extensive assertions throughout code**
  - ❖ Find problems as early as possible (close to source)
- **LLVM IR Verifier: Checks modules for validity**
  - ❖ Checks type properties, dominance properties, etc.
  - ❖ Automatically run by opt
  - ❖ Problem found?: print an error message and abort
- **LLVM IR Leak Detector**
  - ❖ Efficient and simple “garbage collector” for IR objects
  - ❖ Ensure IR objects are deallocated appropriately

# The Bugpoint automated bug finder

- **Simple idea: automate ‘binary’ search for bug**
  - ❖ Bug isolation: which passes interact to produce bug
  - ❖ *Test case reduction*: reduce input program
- **Optimizer/Codegen crashes:**
  - ❖ Throw portion of test case away, check for crash
    - If so, keep going
    - Otherwise, revert and try something else
  - ❖ Extremely effective in practice
- **Simple greedy algorithms for test reduction**
- **Completely black-box approach**



See also: [docs/Bugpoint.html](http://llvm.cs.uiuc.edu/docs/Bugpoint.html)

# Debugging Miscompilations

## ■ Optimizer miscompilation:

- ❖ Split testcase in two, optimize one. Still broken?
- ❖ Keep shrinking the portion being optimized

## ■ Codegen miscompilation:

- ❖ Split testcase in two, compile one with CBE, broken?
- ❖ Shrink portion being compiled with non CBE codegen

## ■ Code splitting granularities:

- ❖ Take out whole functions
- ❖ Take out loop nests
- ❖ Take out individual basic blocks



# How well does this thing work?

- **Extremely effective:**

- ❖ Can often reduce a 100K LOC program and 60 passes to a few basic blocks and 1 pass in 5 minutes
- ❖ Crashes are found much faster than miscompilations
  - no need to run the program to test a reduction

- **Interacts with integrated debugging tools**

- ❖ Runtime errors are detected faster

- **Limitations:**

- ❖ Program must be deterministic
  - ... or modified to be so
- ❖ Finds “a” bug, not “the” bug



# Tutorial Overview

- **Introduction to the running example**
- **LLVM C/C++ Compiler Overview**
  - ❖ High-level view of an example LLVM compiler
- **The LLVM Virtual Instruction Set**
  - ❖ IR overview and type-system
- **LLVM C++ IR and important API's**
  - ❖ Basics, PassManager, dataflow, ArgPromotion
- **Important LLVM Tools**
  - ❖ opt, code generator, JIT, test suite, bugpoint
- **Example applications of LLVM**

# Use Case 1: Edge or Path Profiling

***Goal: Profiling Research or PGO***

■ **Implementation:**

- ❖ *FunctionPass*: LLVM-to-LLVM transformation
- ❖ Instrumentation: Use CFG, intervals, dominators
- ❖ Code generation: Use C or any native back end
- ❖ Profile feedback: Use profile query interface

■ **Core extensions needed:** *none*

■ **Major LLVM Benefits**

- ❖ Language-independence, CFG, very simple IR

# Use Case 2: Alias Analysis

**Goal:** Research on new alias analysis algorithms

- **Implementation:**

- ❖ *ModulePass*: Whole-program analysis pass on LLVM
- ❖ Use type information; SSA; heap/stack/globals
- ❖ Compare *SimpleAA*, *Steensgard's*, *Andersen's*, *DSA*
- ❖ Evaluate many clients via *AliasAnalysis* interface

- **Core extensions needed:** none

- **Major LLVM Benefits**

- ❖ Language-independence, type info, SSA, DSA, IPO
- ❖ *AliasAnalysis* interface with many pre-existing clients

# Use Case 3: LDS Prefetching

***Goal:*** Prefetching linked data structures

## ■ Implementation:

- ❖ *ModulePass*: Link-time LLVM-to-LLVM transformation
- ❖ *Code transformations*: use type info, loop analysis, unrolling, prefetch insertion
- ❖ *Data transformations* (e.g., adding history pointers): use strong type info from DSA, IPO

## ■ Core extensions needed:

- ❖ Prefetch operation: add as intrinsic (in progress)

## ■ Major LLVM Benefits

- Language-independence, type info, DSA, IPO

# Use Case 4: Language Front end

***Goal:*** Use LLVM to implement a new language

- **Implementation:**

- ❖ Parser (say to AST), Semantic checking
  - ❖ AST-to-LLVM translator

- **Core extensions needed:** depends

- ❖ High-level type system is omitted by design

- **Major LLVM Benefits**

- ❖ Low-level, but powerful type system
  - ❖ Very simple IR to generate (e.g., compare GCC RTL)
  - ❖ Extensive global and IP optimization framework
  - ❖ JIT engine, native back-ends, C back-end

# Use Case 5: JIT Compiler

***Goal:*** Write JIT compiler for a bytecode language

- **Implementation:**

- ❖ Extend the LLVM JIT framework
- ❖ *Simple JIT*: Fast translation from bytecode to LLVM  
(then use LLVM JIT + GC)
- ❖ *Optimizing JIT*: Language-specific optimizations + fast translation  
(then use LLVM optimizations, JIT, GC)

- **Core extensions needed:** *none in general*

- **Major LLVM Benefits**

- ❖ Compact, typed, language-independent IR
- ❖ Existing JIT framework and GC

# Use Case 6: Architecture Research

***Goal:*** Compiler support for new architectures

- **Implementation:**

- ❖ Add new machine description (or modify one)
  - ❖ Add any new LLVM-to-LLVM transformations

- **Core extensions needed:** depends on goals

- ❖ Imminent features: ***modulo sched; vector ops***

- **Major LLVM Benefits**

- ❖ Low-level, typed, machine-independent IR
  - ❖ Explicit register/memory architecture
  - ❖ Aggressive mid-level and back-end compiler framework
  - ❖ ***Full-system evaluation:*** applications, libraries, even OS

# Five point LLVM Review

- **Extremely simple IR to learn and use**
  - ❖ 1-to-1 correspondence between .ll, .bc, and C++ IR
  - ❖ Very positive user reactions
- **Powerful and modular optimizer**
  - ❖ Easy to extend, or just use what is already there
- **Clean and modular code generator**
  - ❖ Easy to retarget, easy to replace/tweak components
- **Many “productivity tools” (bugpoint, verifier)**
  - ❖ Get more done, quicker!
- **Active dev community, good documentation**
  - ❖ Mailing lists, IRC, doxygen, extensive docs

# Get started with LLVM!

- **Download latest release or CVS:**

<http://llvm.cs.uiuc.edu/releases/>

- **Follow the “Getting Started Guide”:**

<http://llvm.cs.uiuc.edu/docs/GettingStarted.html>

- ❖ Walks you through install and setup
- ❖ Lots of other docs available in “docs” directory
- ❖ Join us on mailing lists and IRC

- **Happy hacking!**